

# Lecture 7: Tricks + Word Embeddings

Alan Ritter

(many slides from Greg Durrett)

# Recall: Feedforward NNs

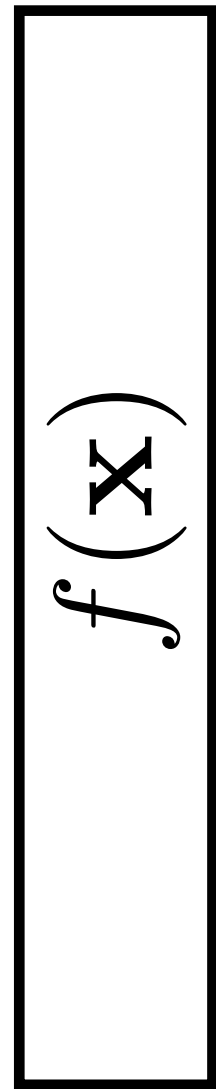
---

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$

# Recall: Feedforward NNs

---

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$

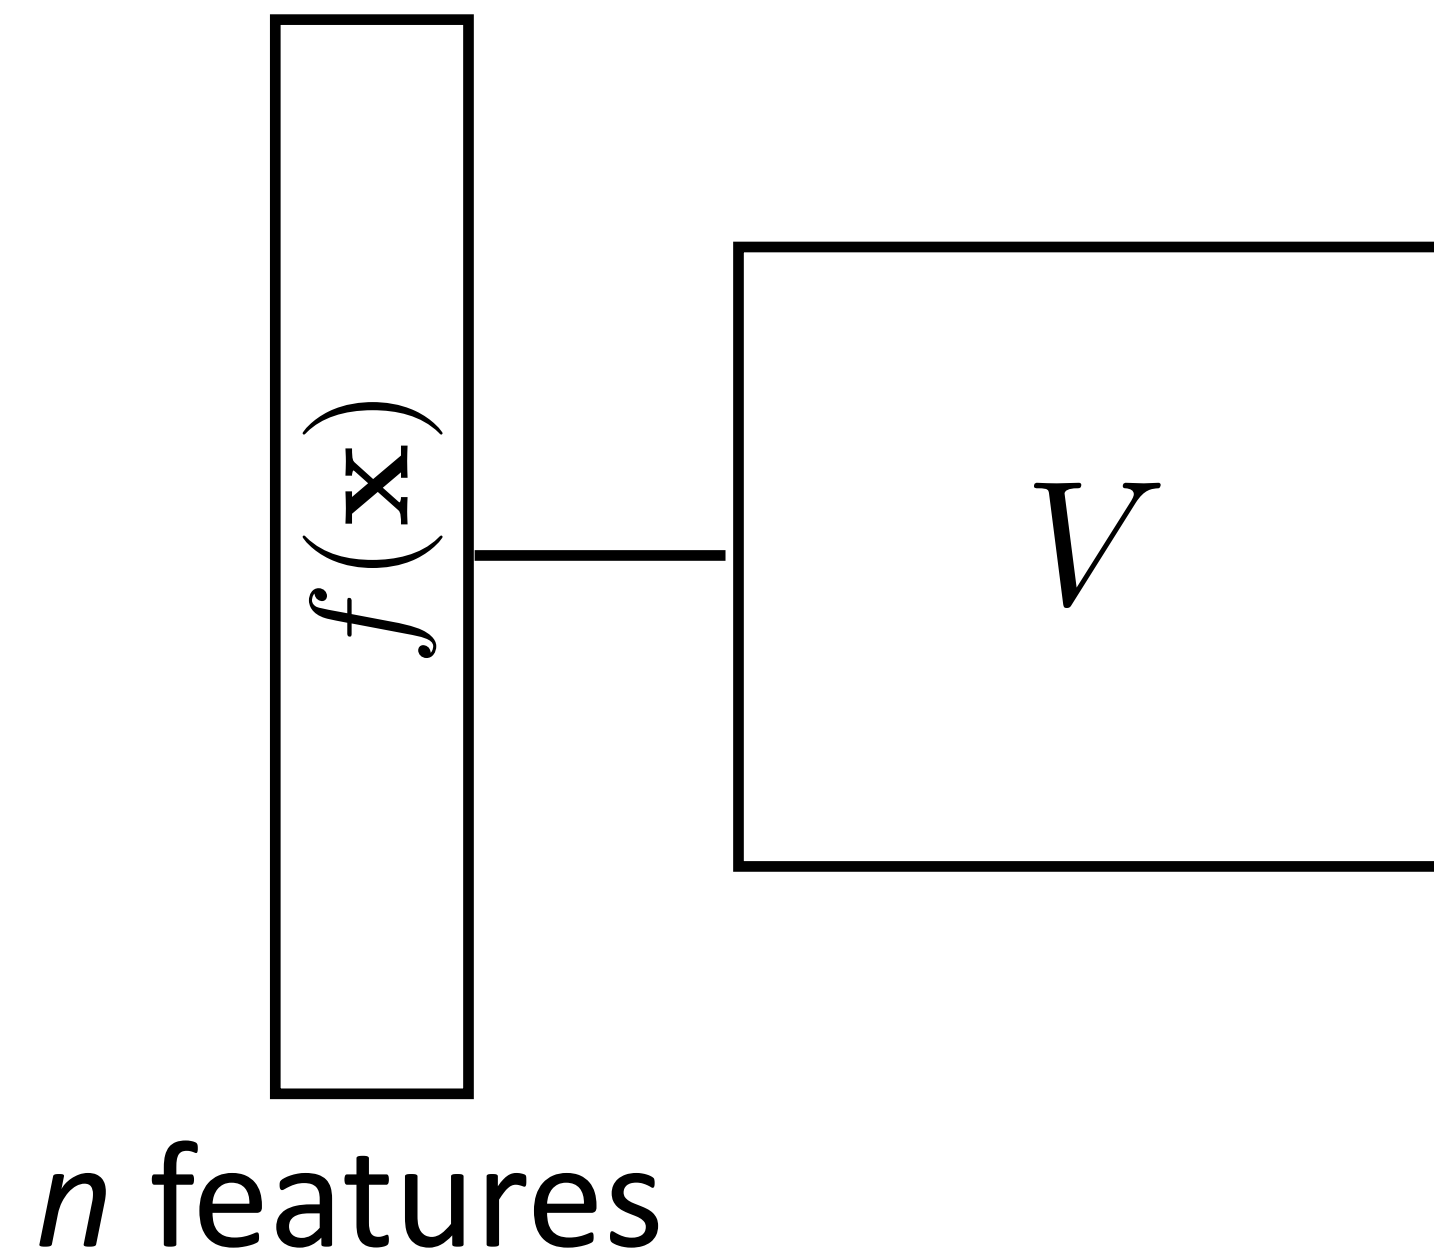


$n$  features

# Recall: Feedforward NNs

---

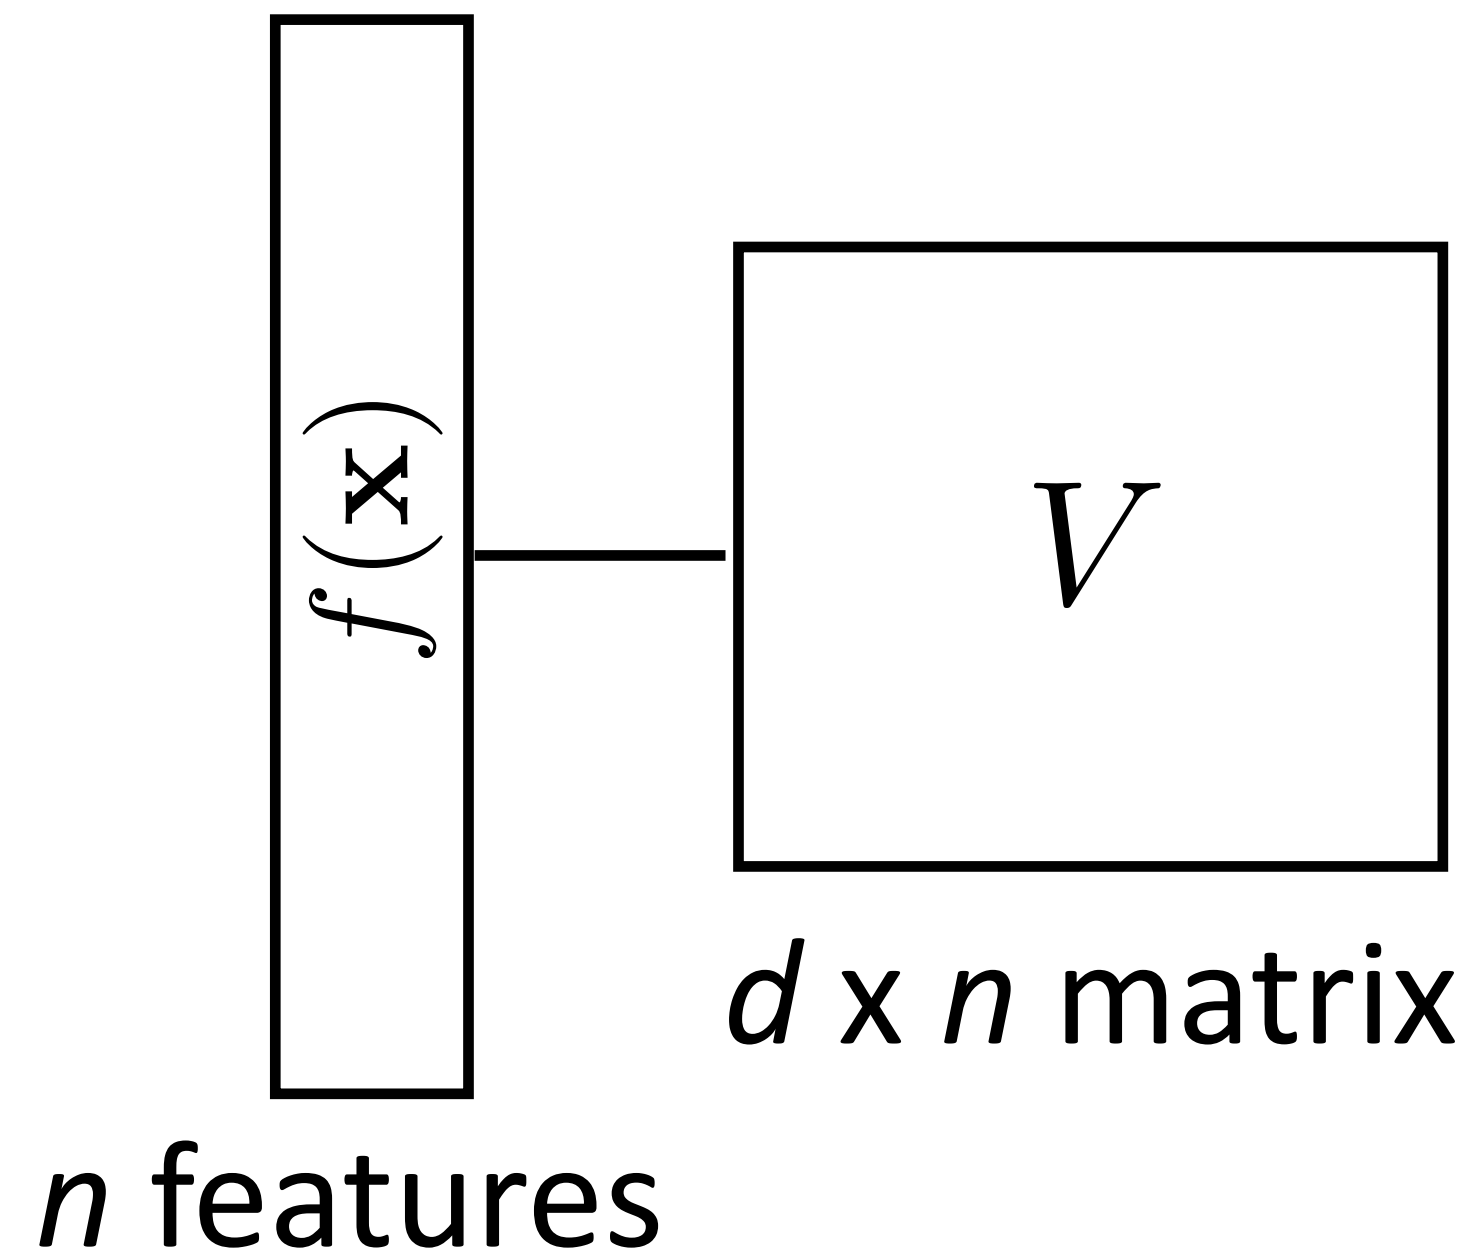
$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$



# Recall: Feedforward NNs

---

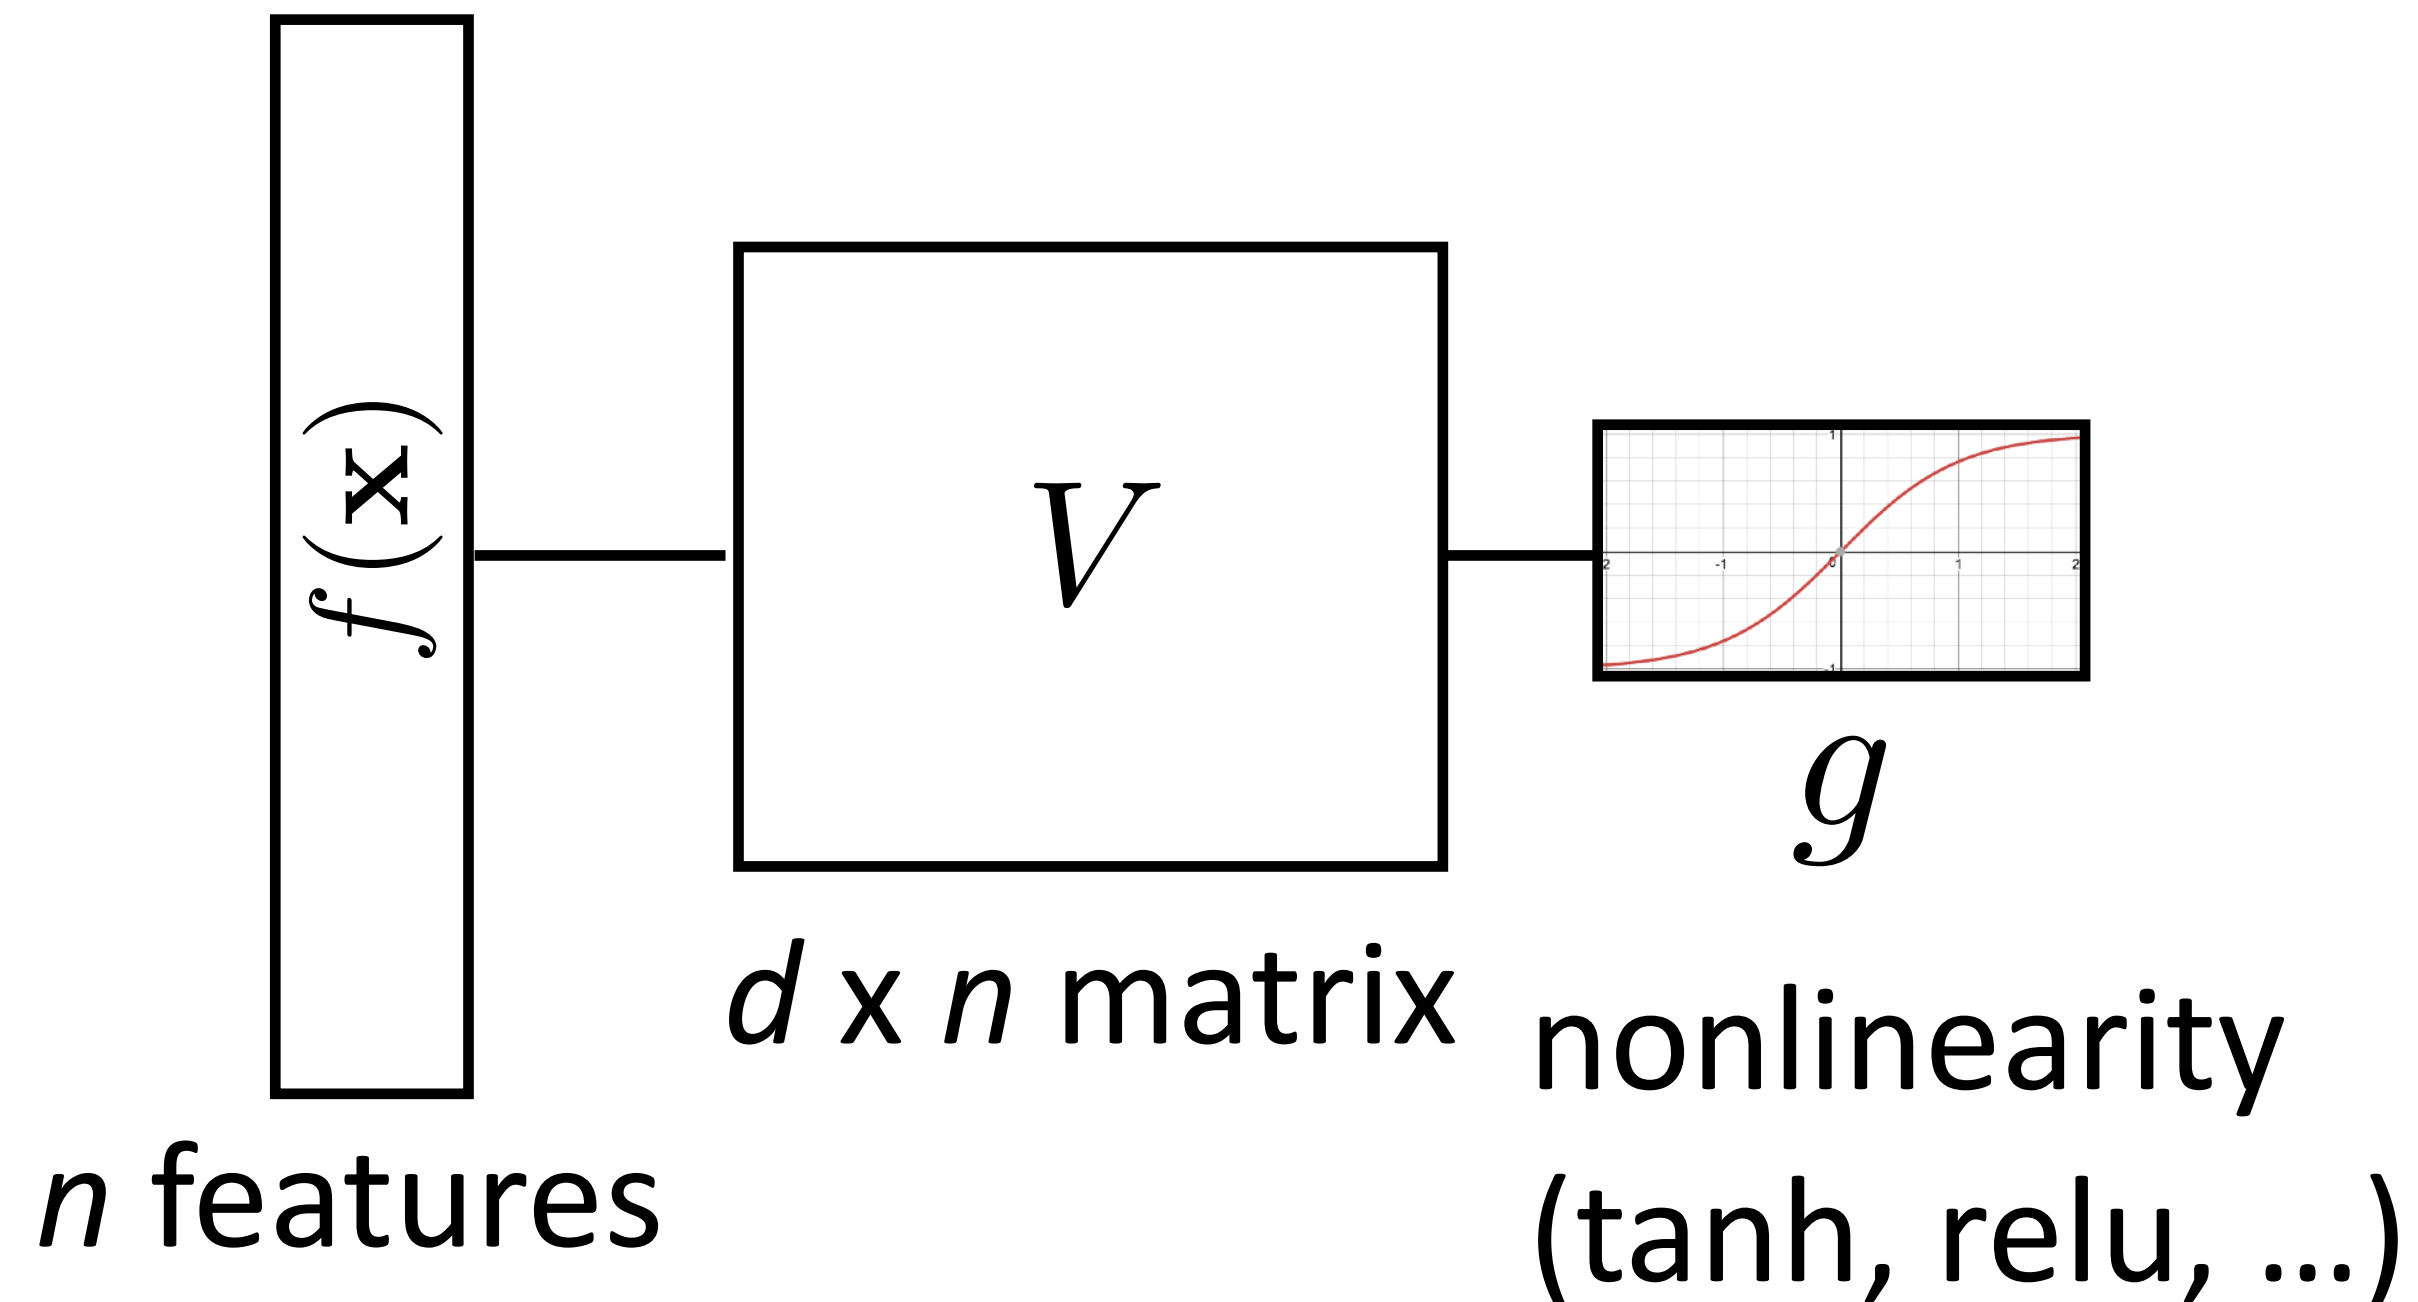
$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$



# Recall: Feedforward NNs

---

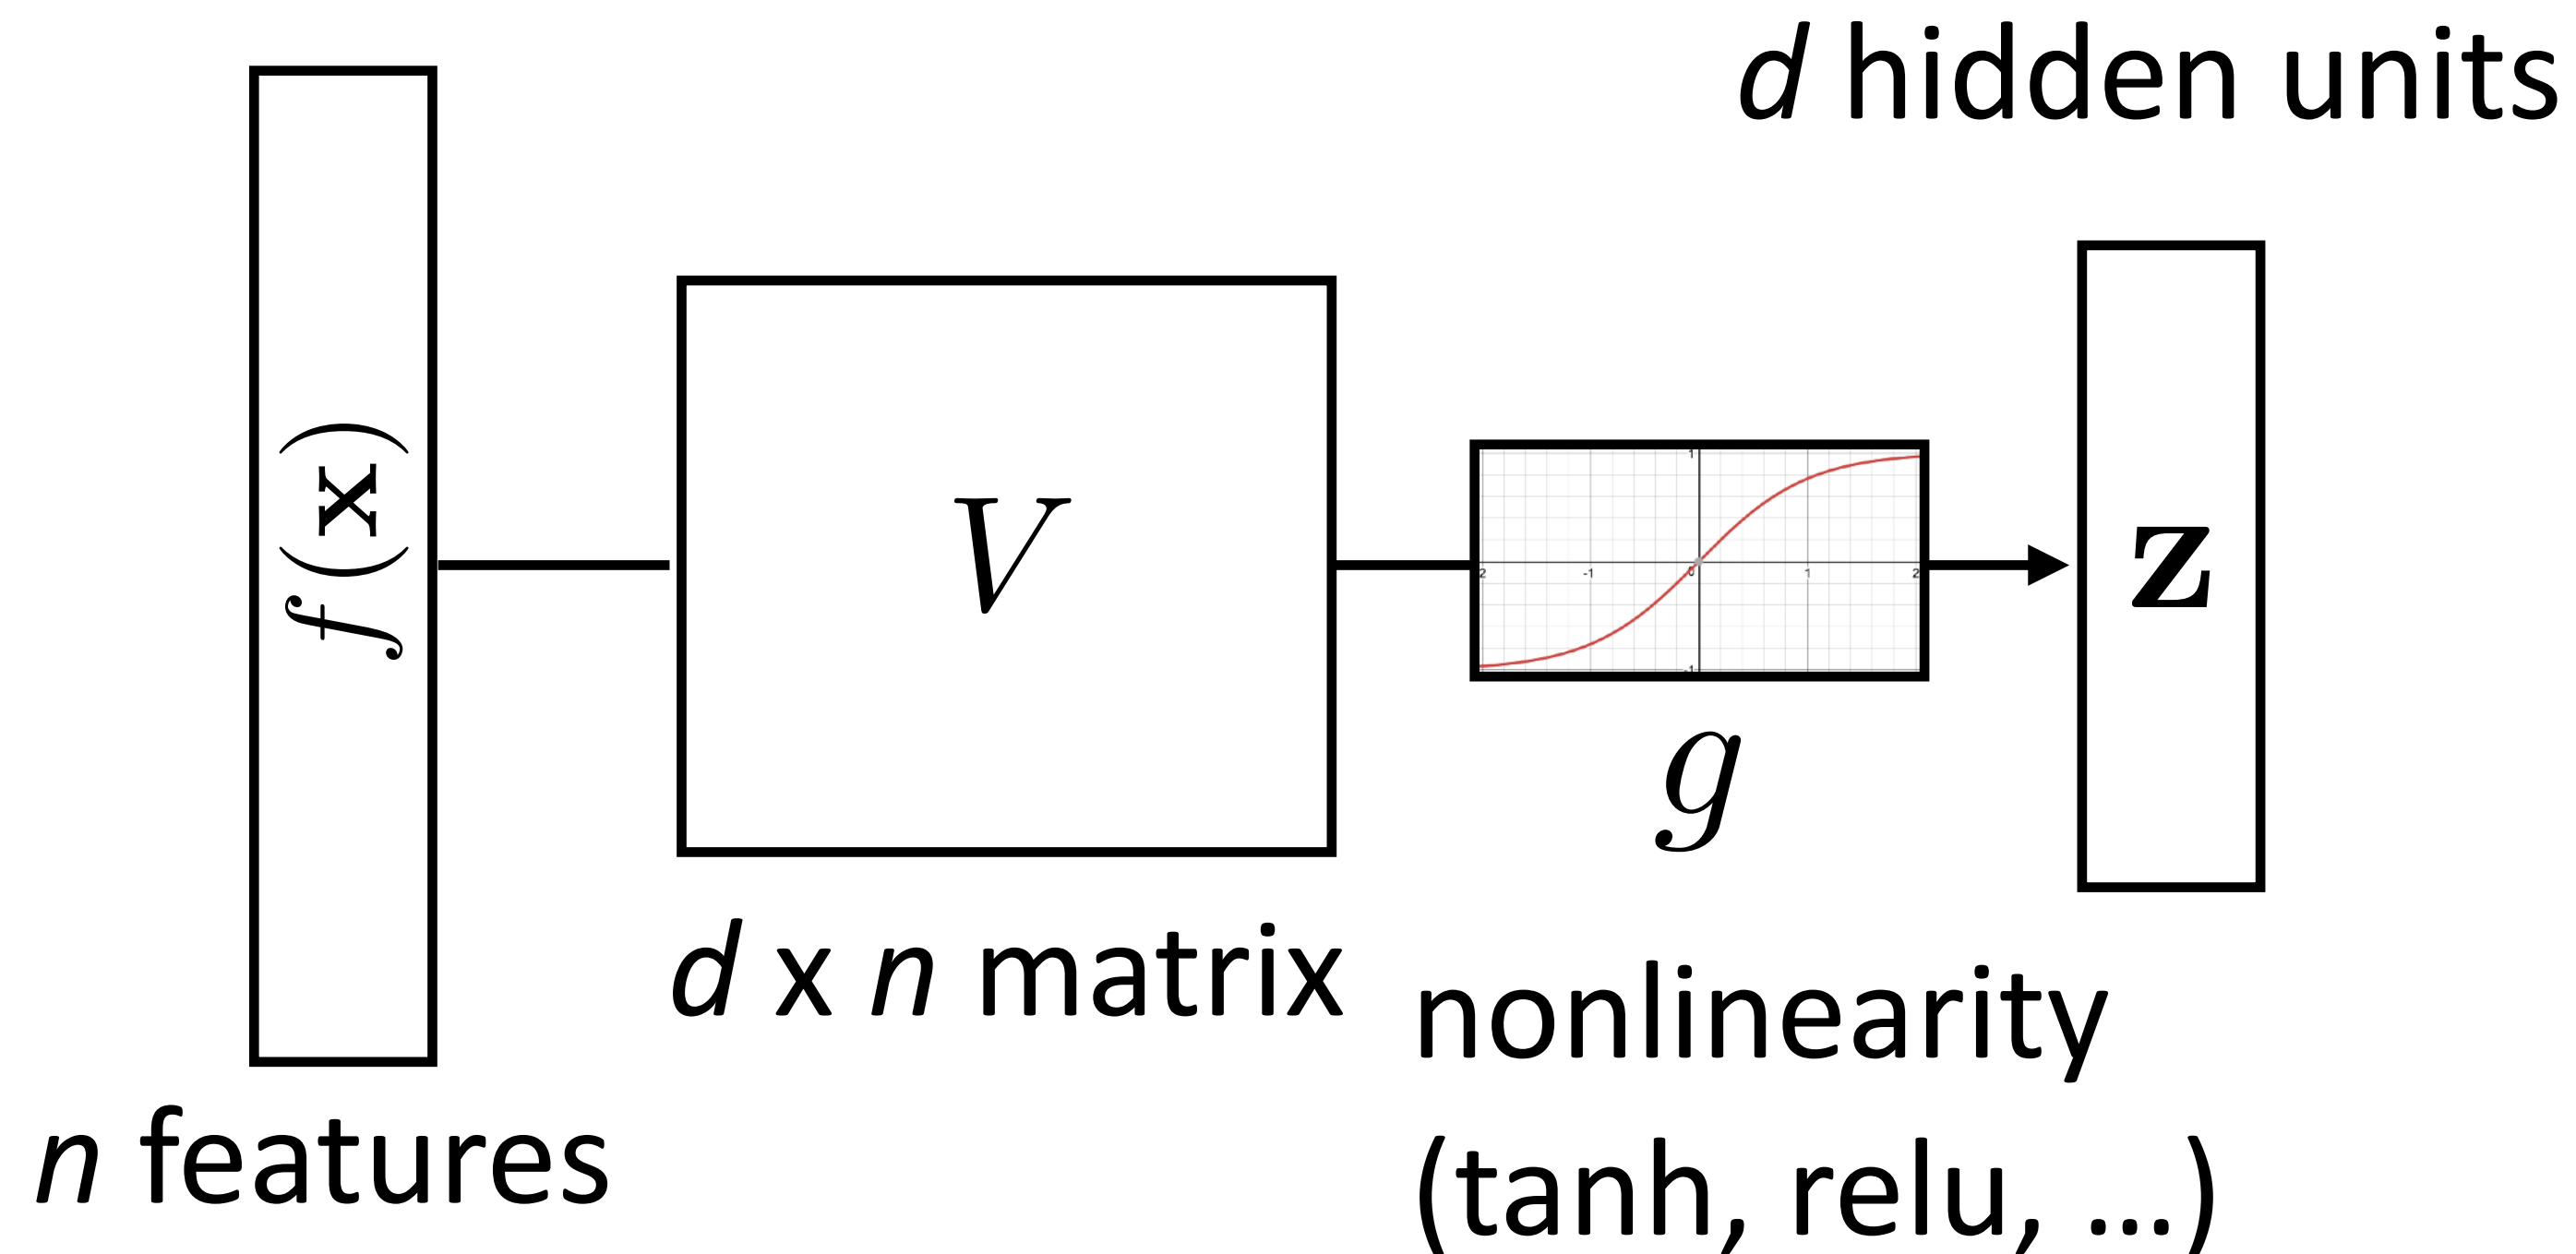
$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$



# Recall: Feedforward NNs

---

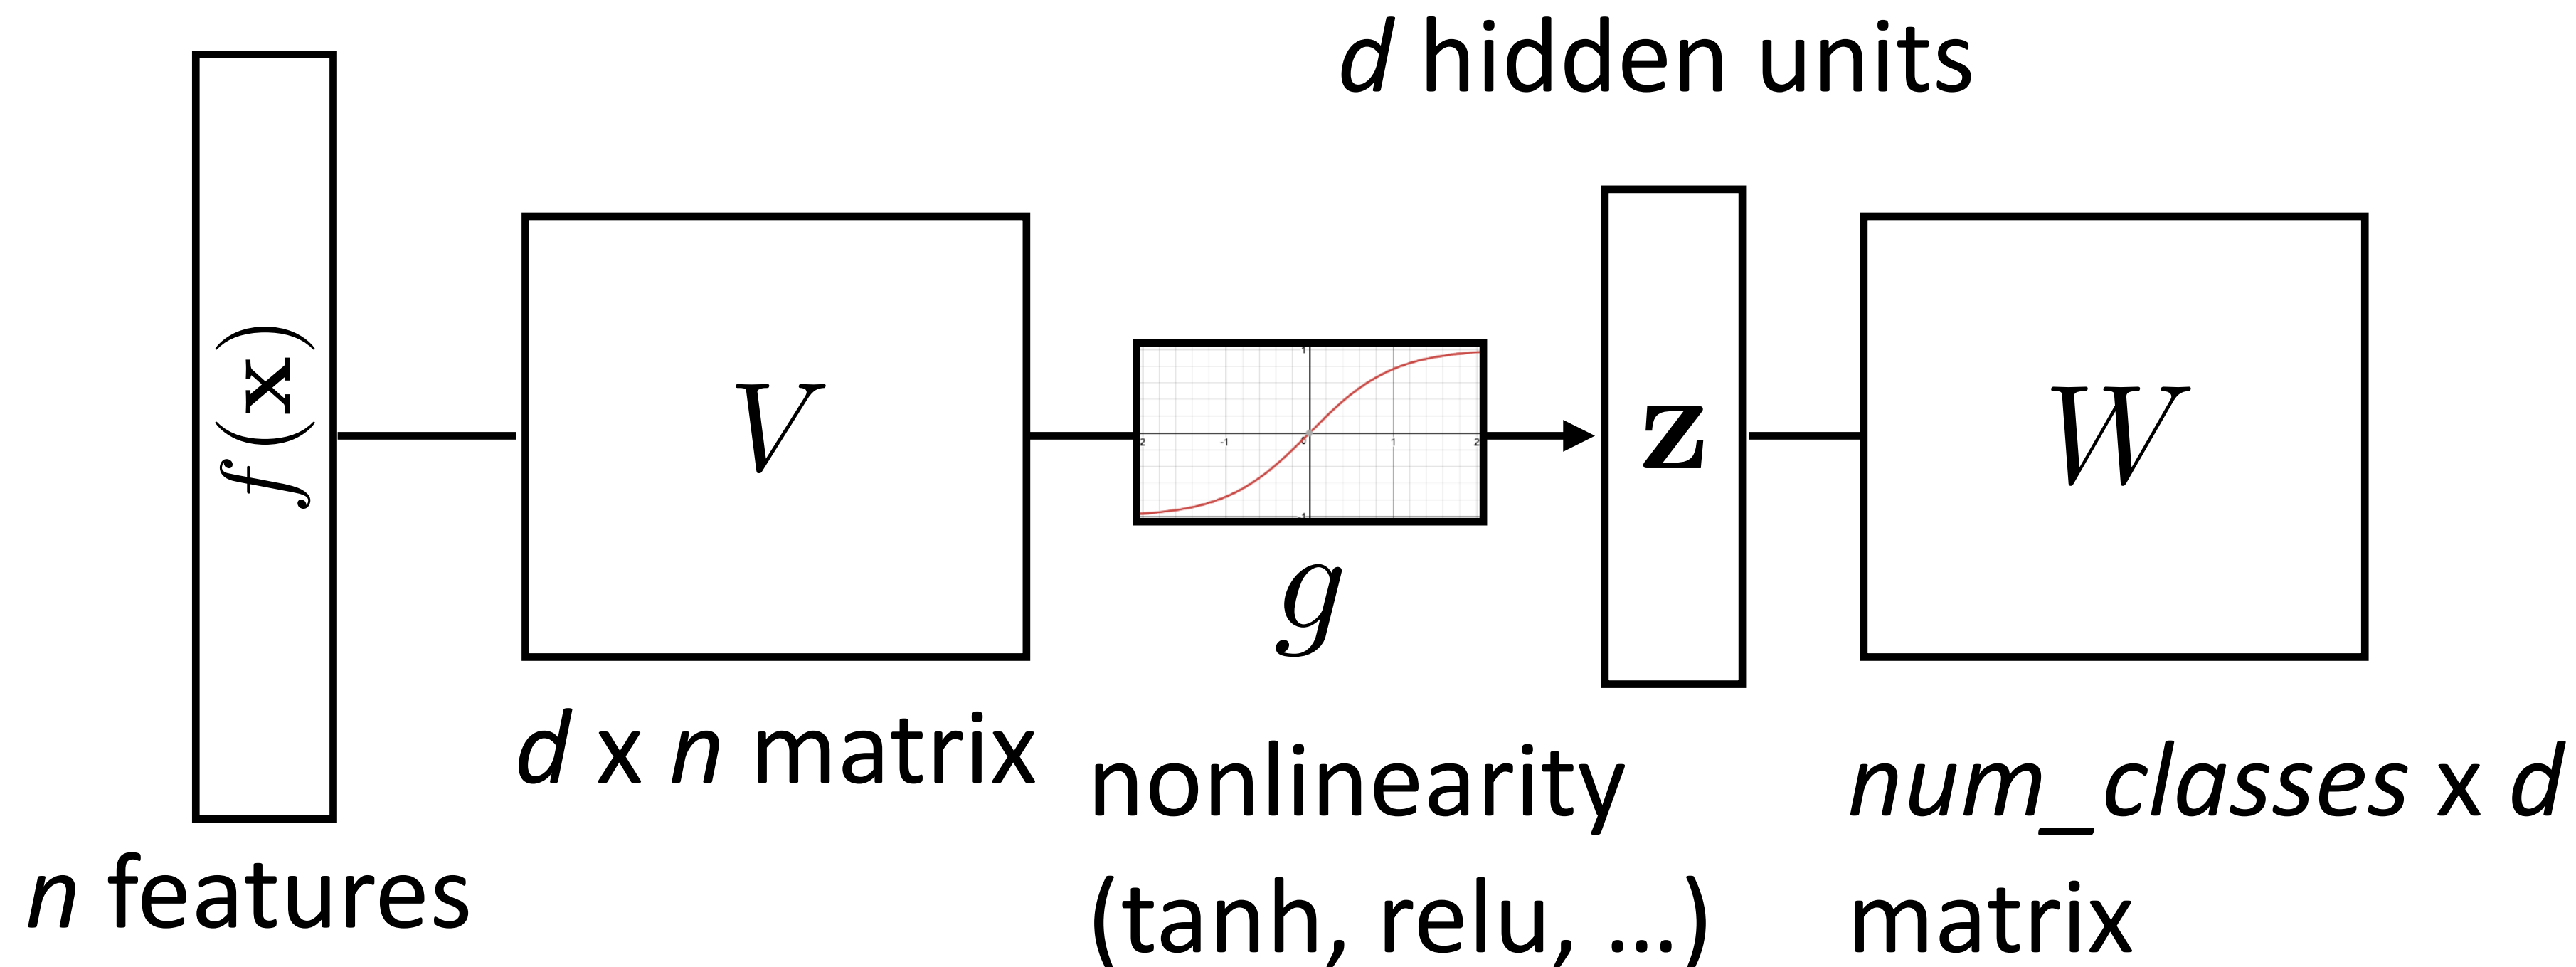
$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$



# Recall: Feedforward NNs

---

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$

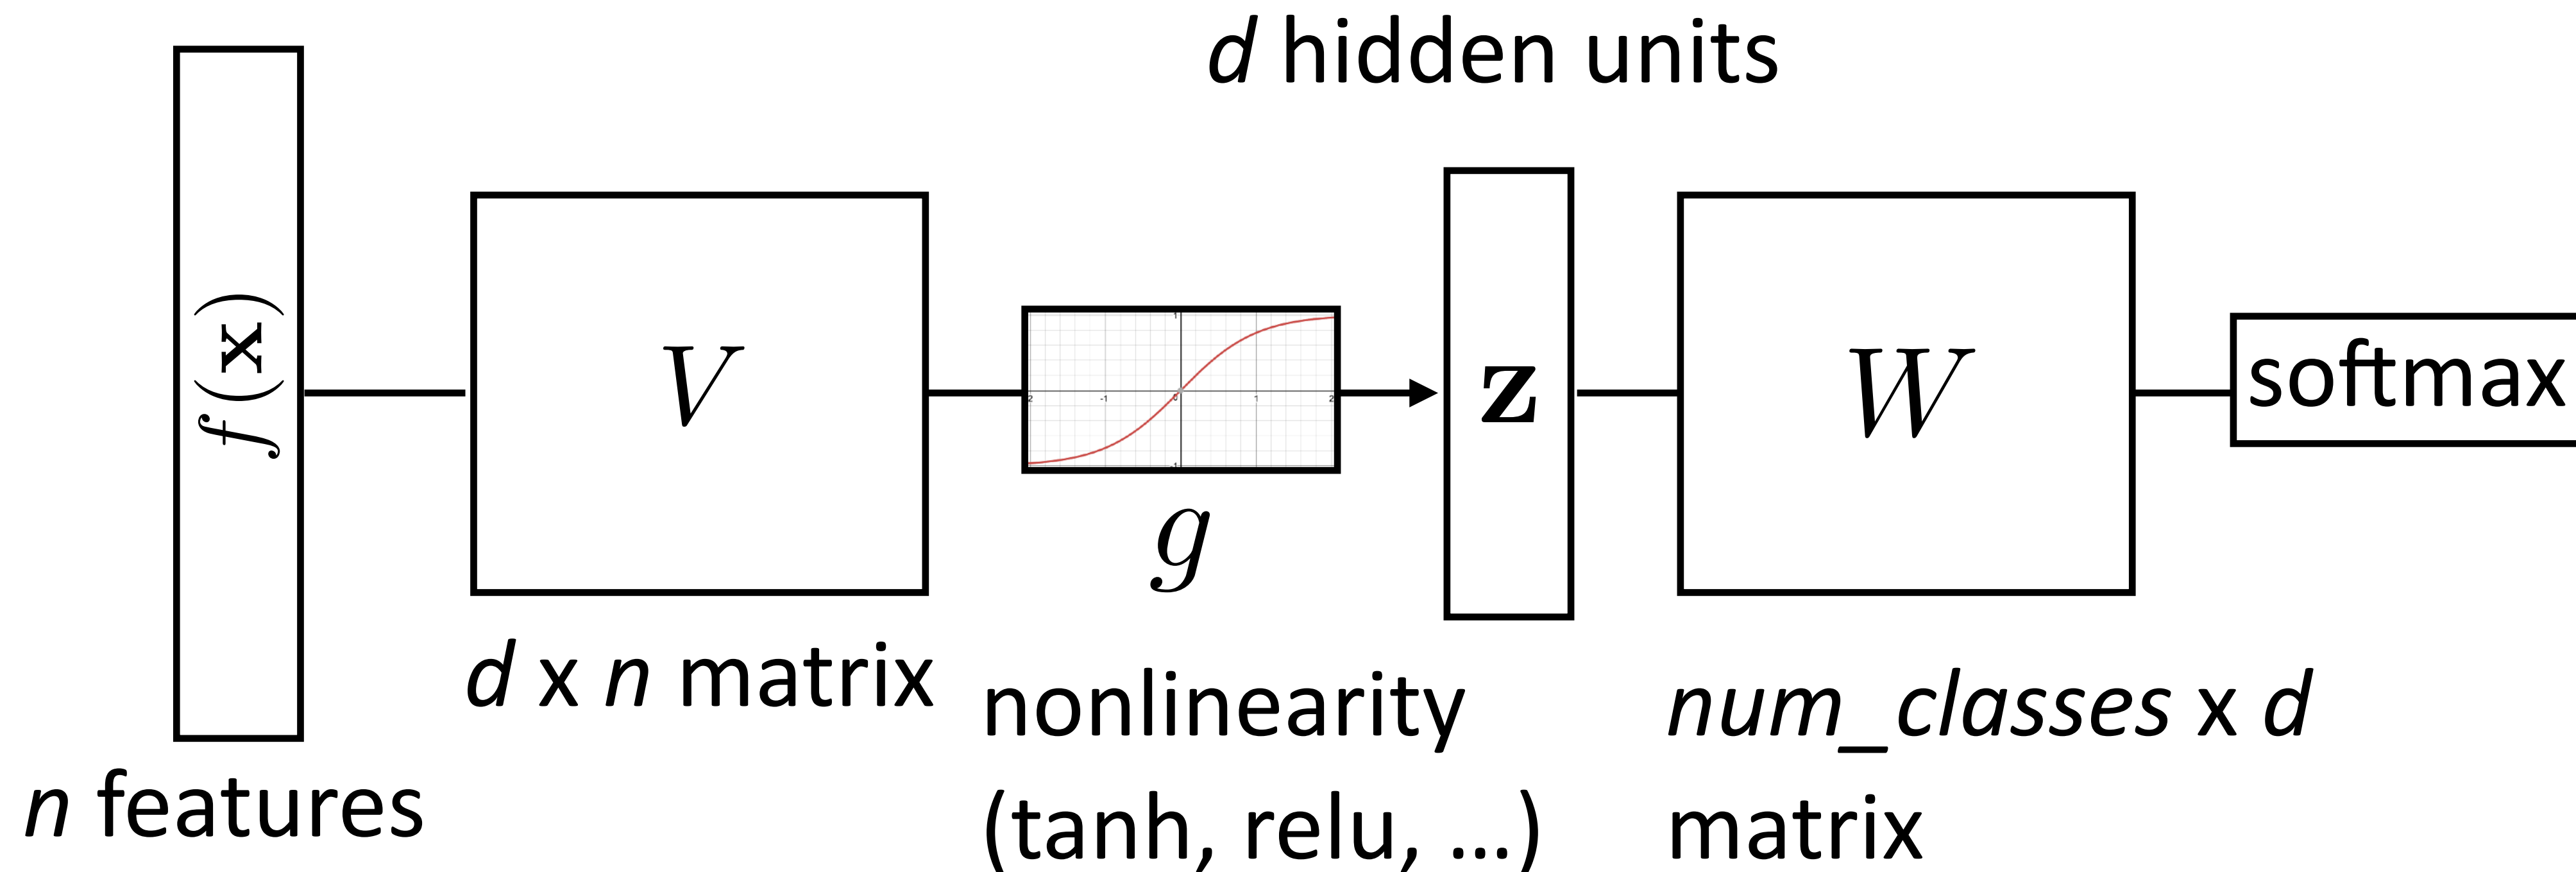




# Recall: Feedforward NNs

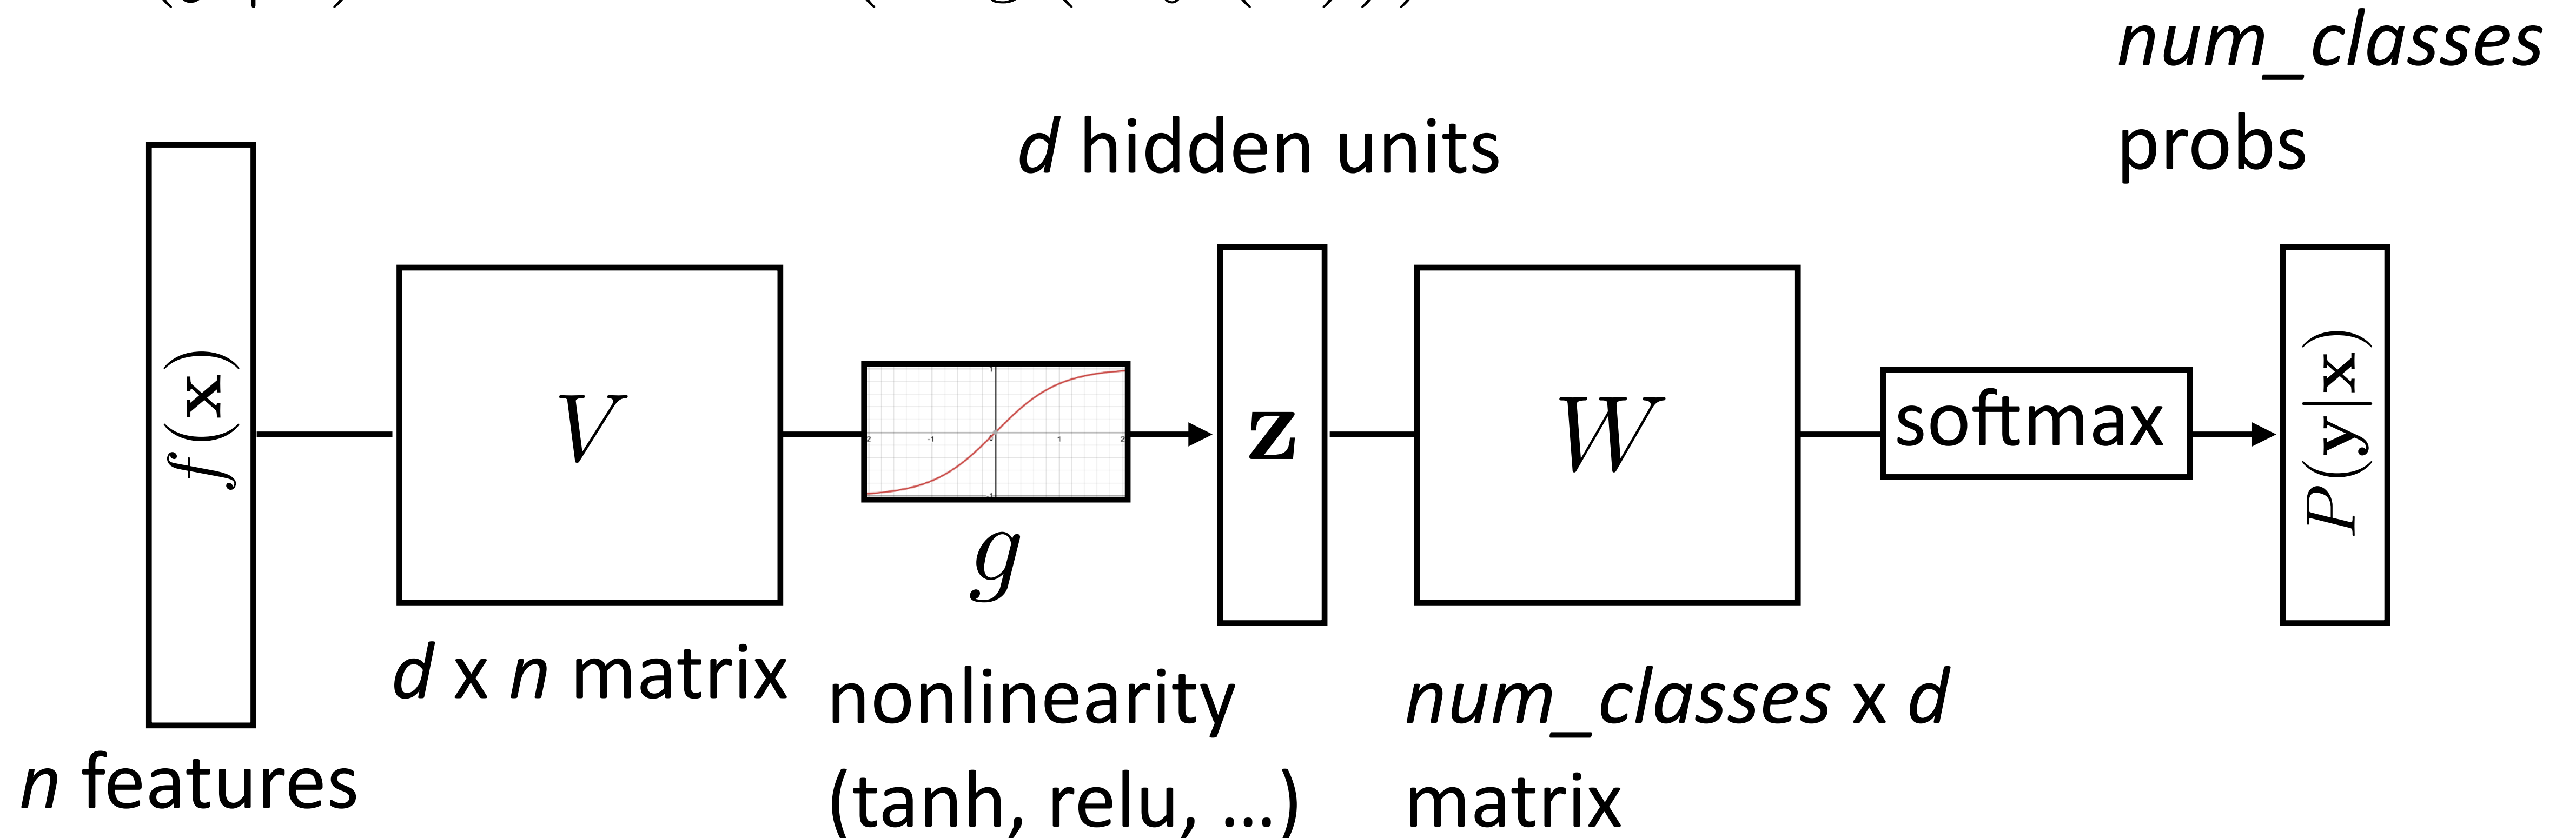
---

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(W g(V f(\mathbf{x})))$$



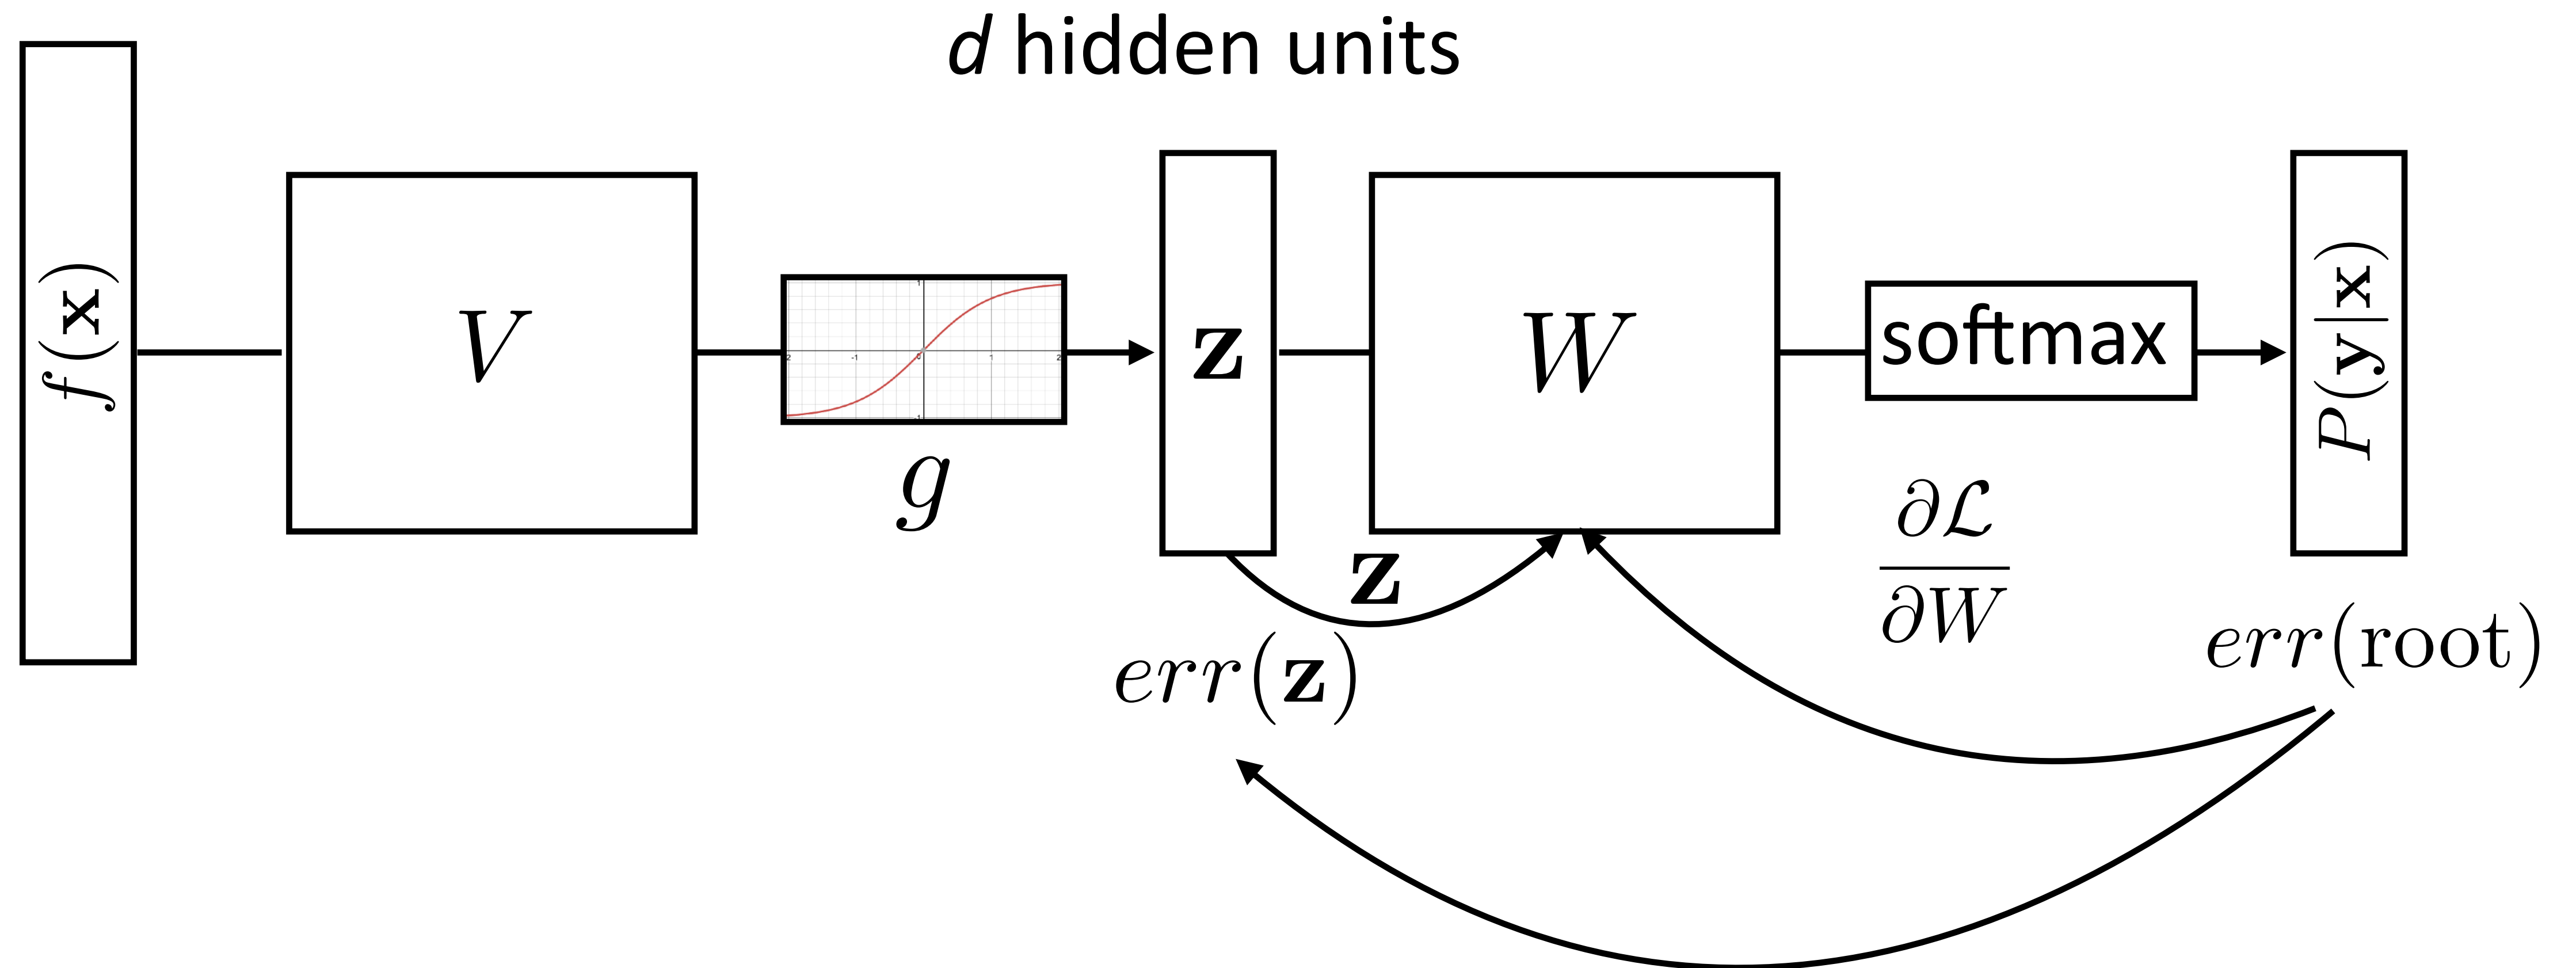
# Recall: Feedforward NNs

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(W g(V f(\mathbf{x})))$$



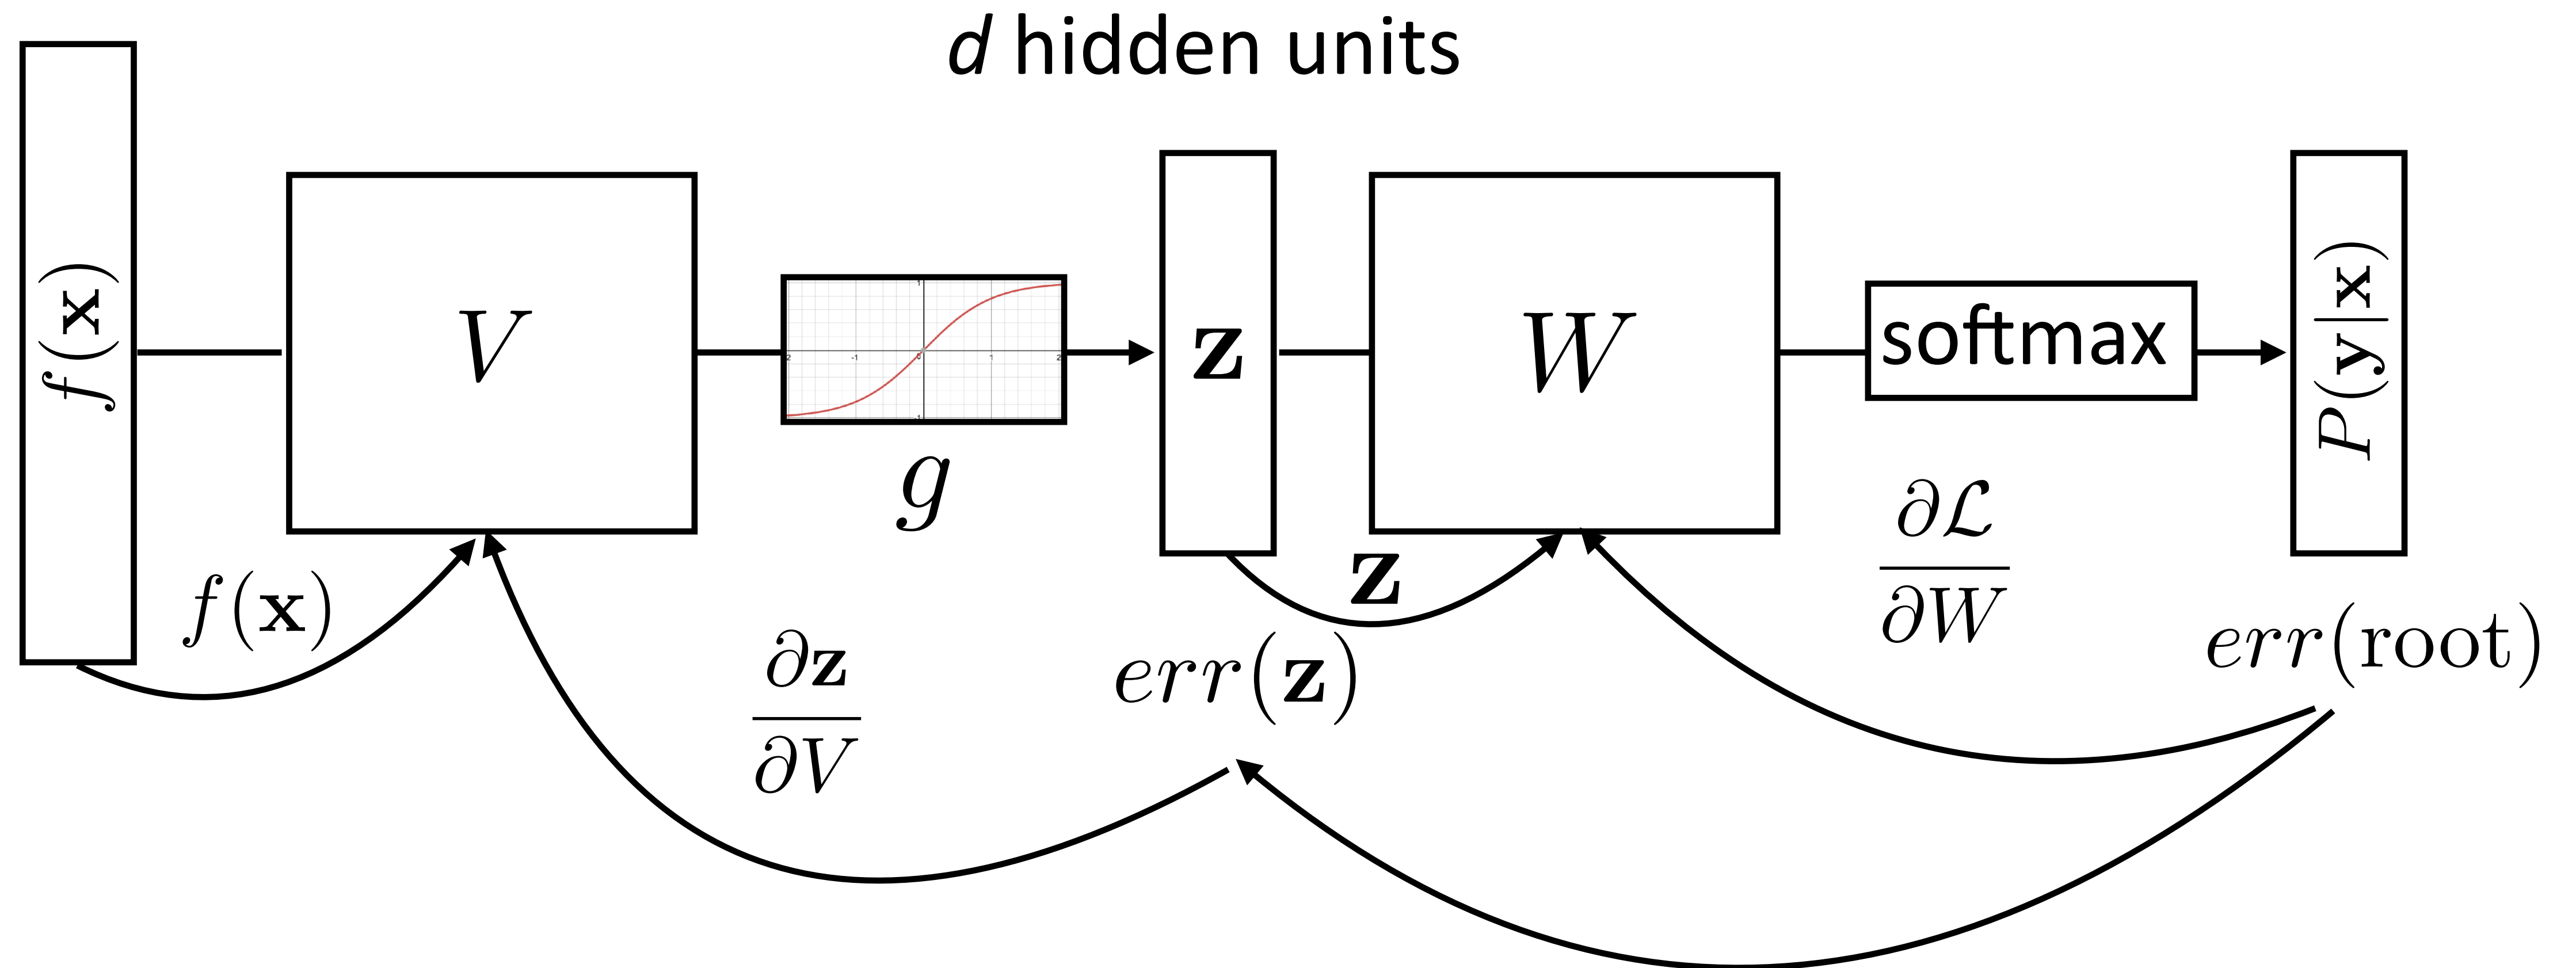
# Recall: Backpropagation

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$



# Recall: Backpropagation

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(Wg(Vf(\mathbf{x})))$$



# This Lecture

---

- ▶ Training
- ▶ Word representations
- ▶ word2vec/GloVe
- ▶ Evaluating word embeddings

# Training Tips

# Training Basics

---

# Training Basics

---

- ▶ Basic formula: compute gradients on batch, use first-order opt. method



# Training Basics

---

- ▶ Basic formula: compute gradients on batch, use first-order opt. method
- ▶ How to initialize? How to regularize? What optimizer to use?

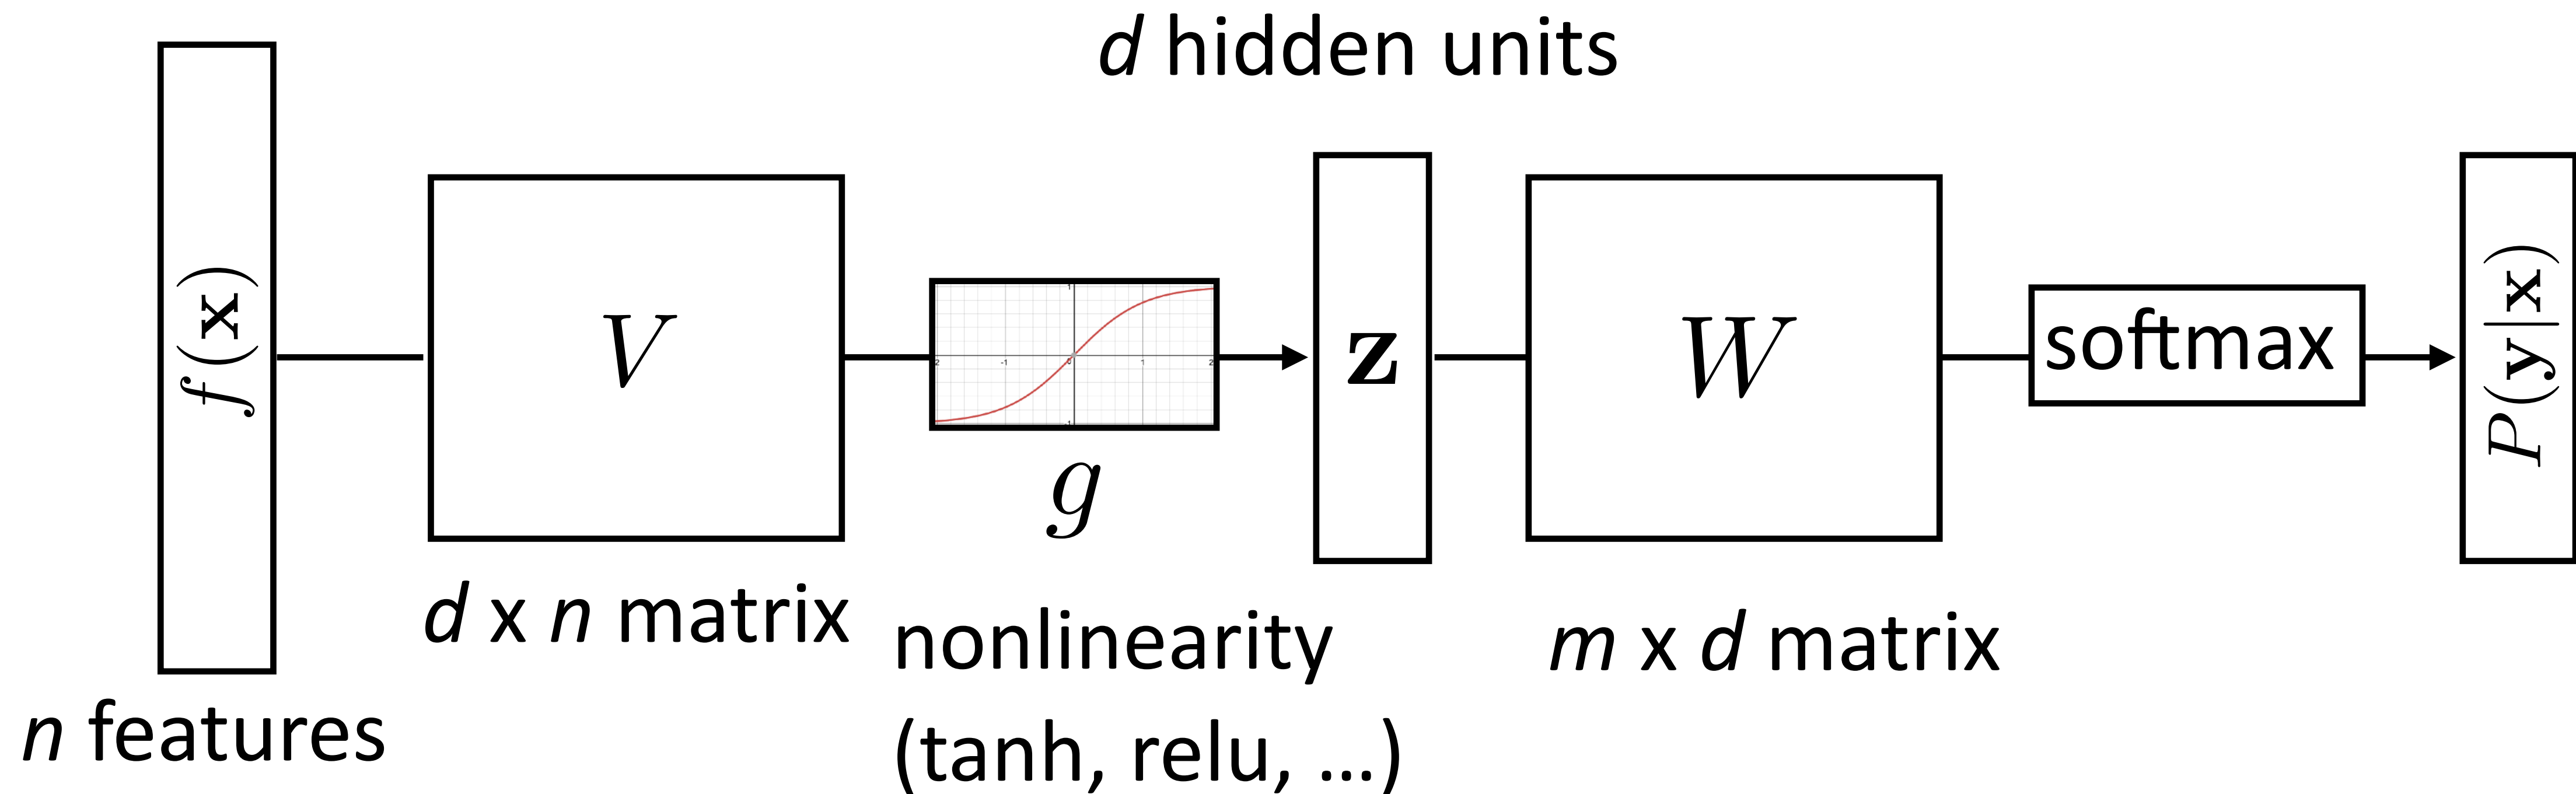
# Training Basics

---

- ▶ Basic formula: compute gradients on batch, use first-order opt. method
- ▶ How to initialize? How to regularize? What optimizer to use?
- ▶ This lecture: some practical tricks. Take deep learning or optimization courses to understand this further


# How does initialization affect learning?

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(W g(V f(\mathbf{x})))$$



# How does initialization affect learning?


$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(W g(V f(\mathbf{x})))$$



- How do we initialize  $V$  and  $W$ ? What consequences does this have?

# How does initialization affect learning?

$$P(\mathbf{y}|\mathbf{x}) = \text{softmax}(W g(V f(\mathbf{x})))$$



- ▶ How do we initialize  $V$  and  $W$ ? What consequences does this have?
- ▶ *Nonconvex* problem, so initialization matters!

# How does initialization affect learning?

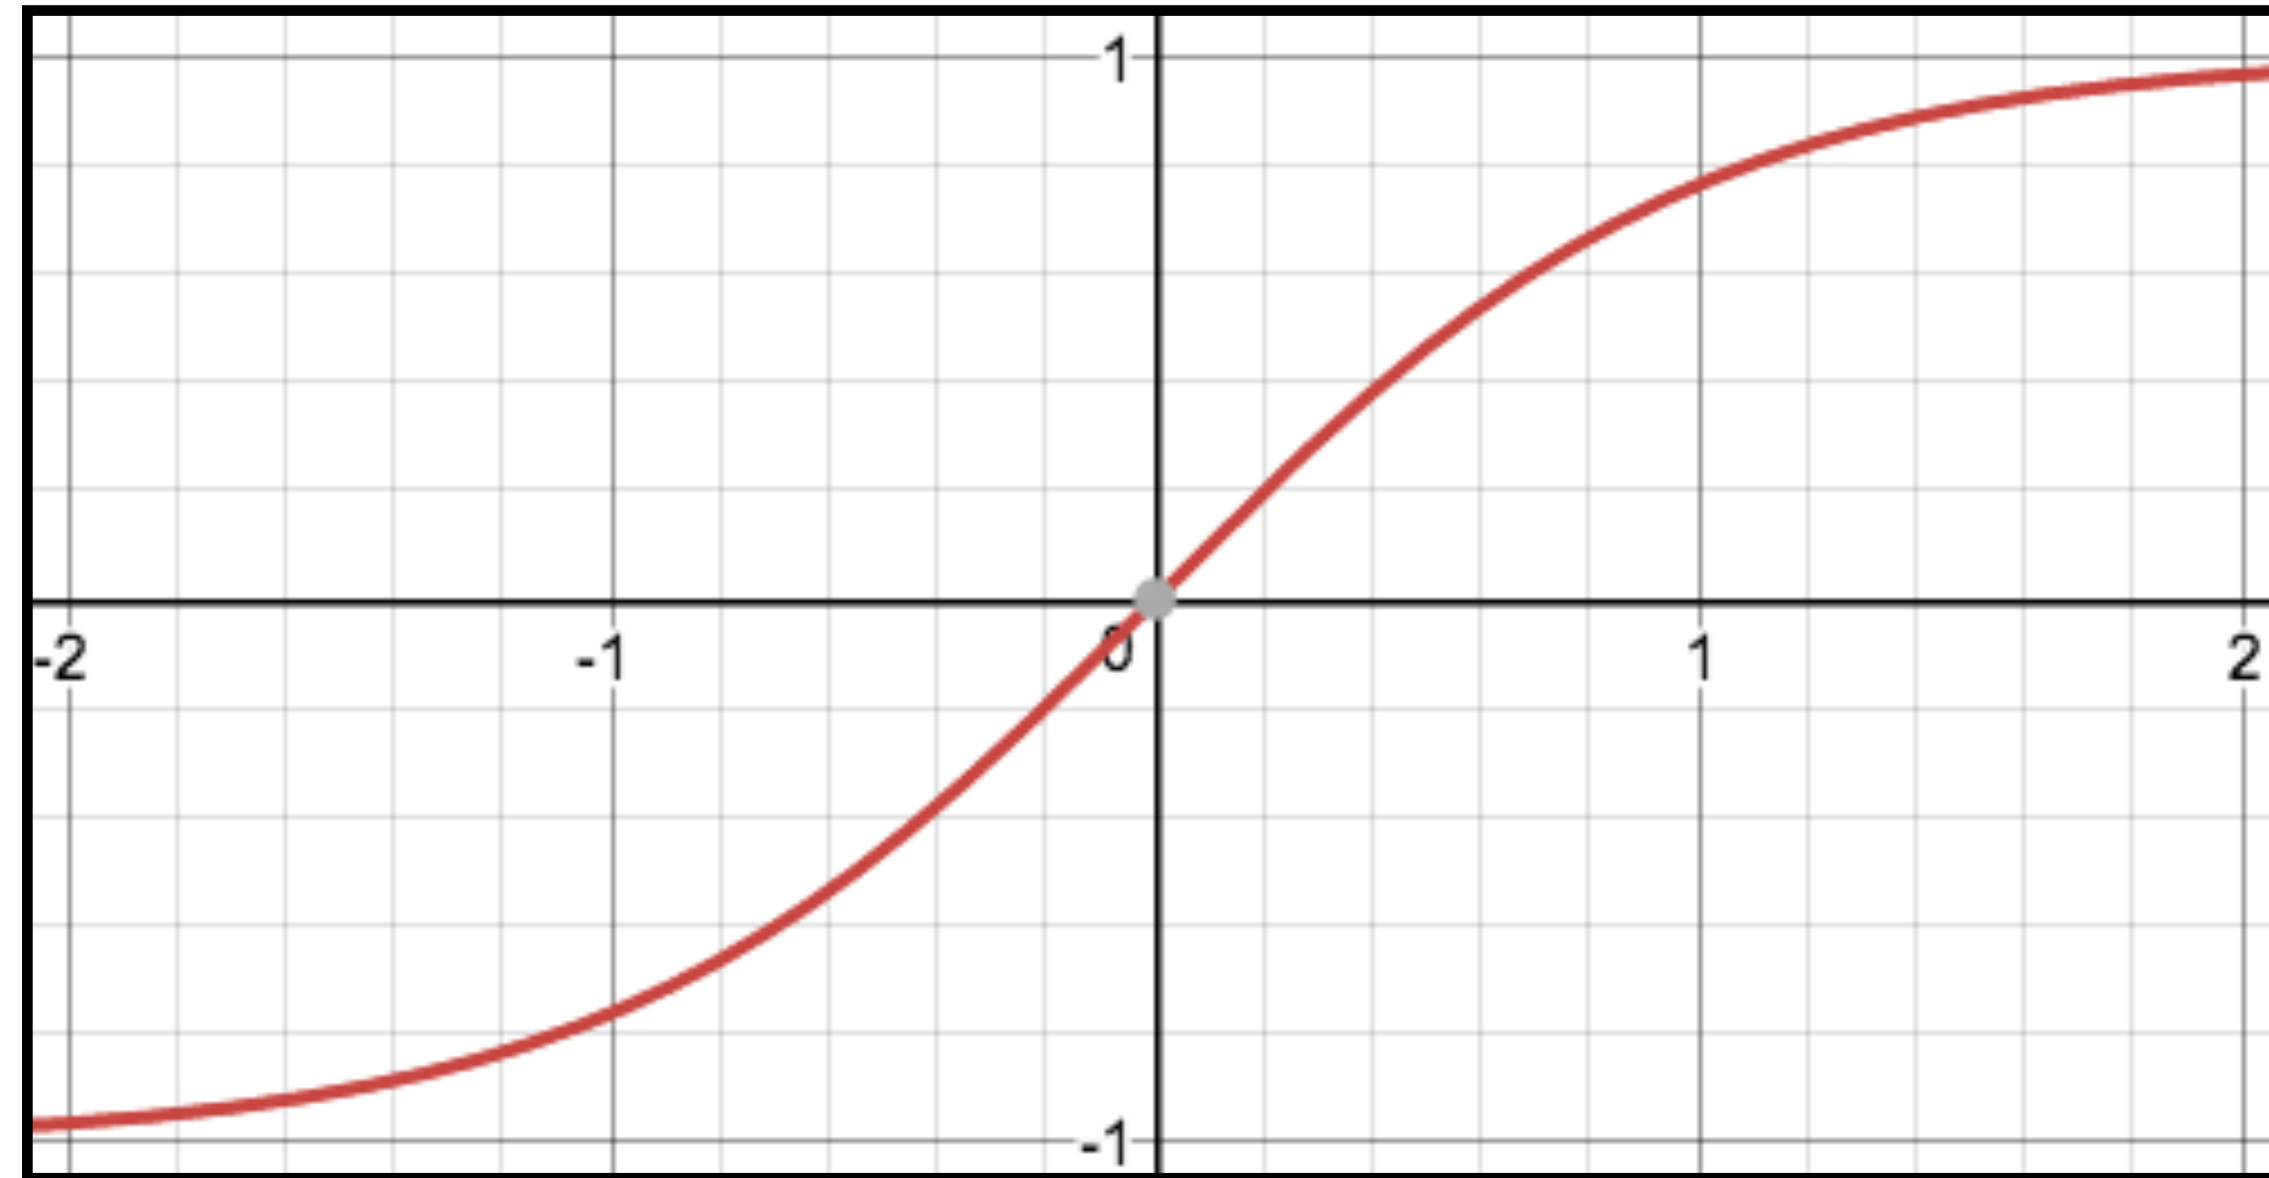
---

- ▶ Nonlinear model...how does this affect things?

# How does initialization affect learning?

---

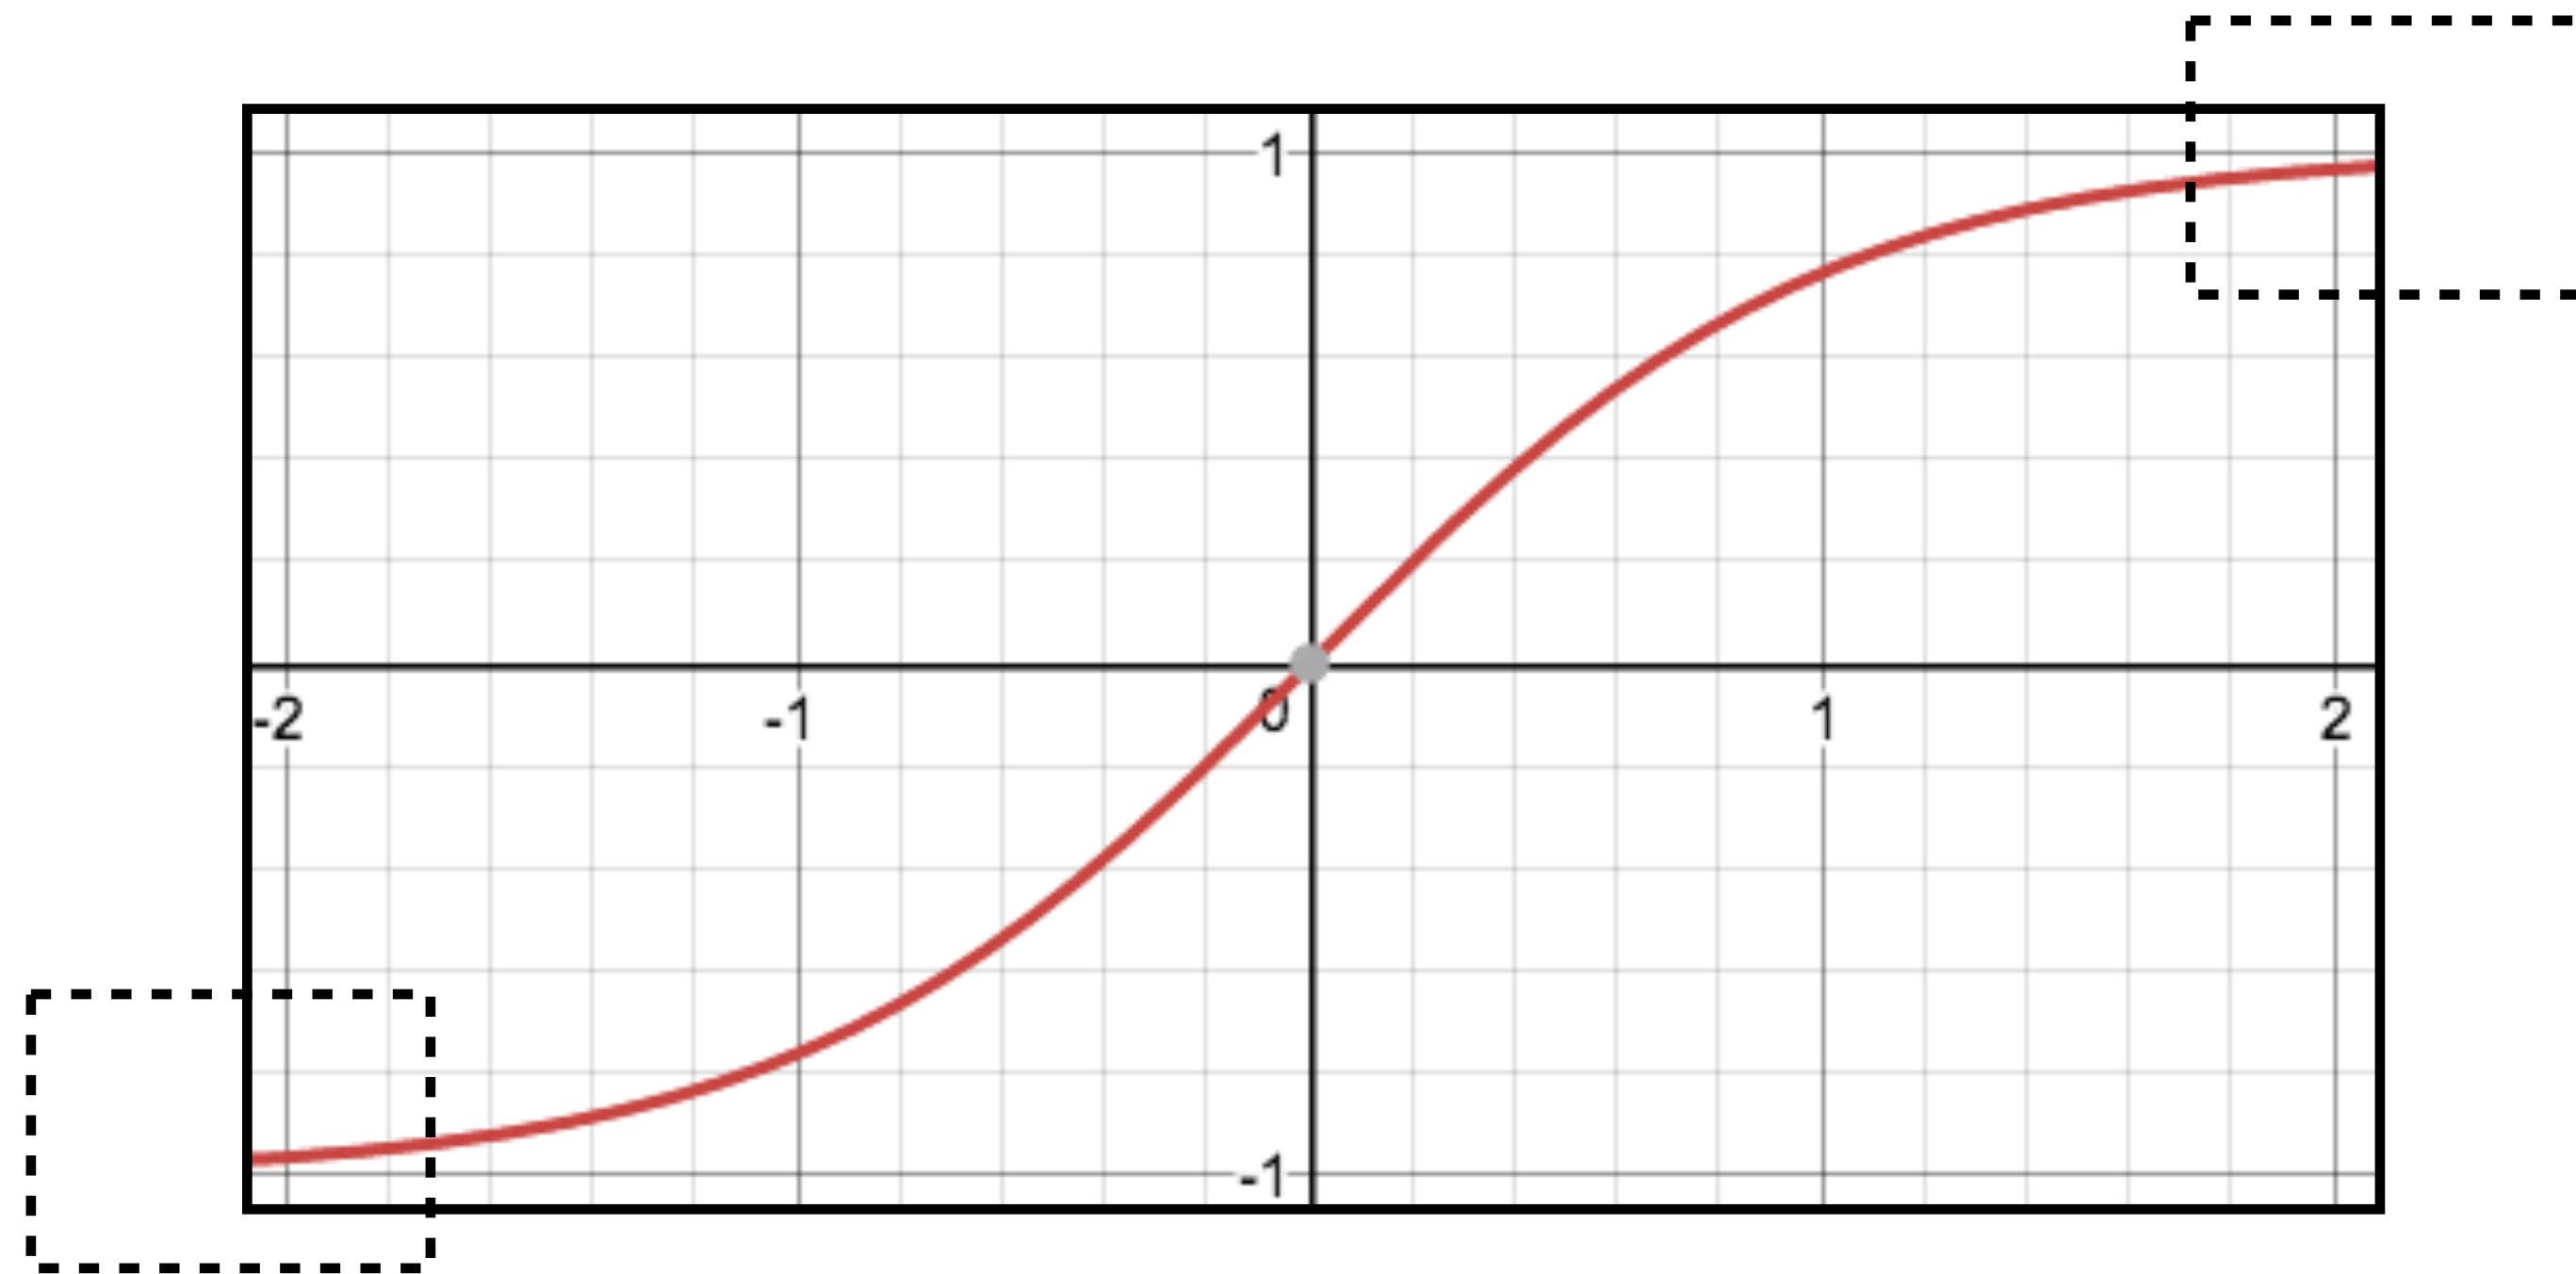
- ▶ Nonlinear model...how does this affect things?



# How does initialization affect learning?

---

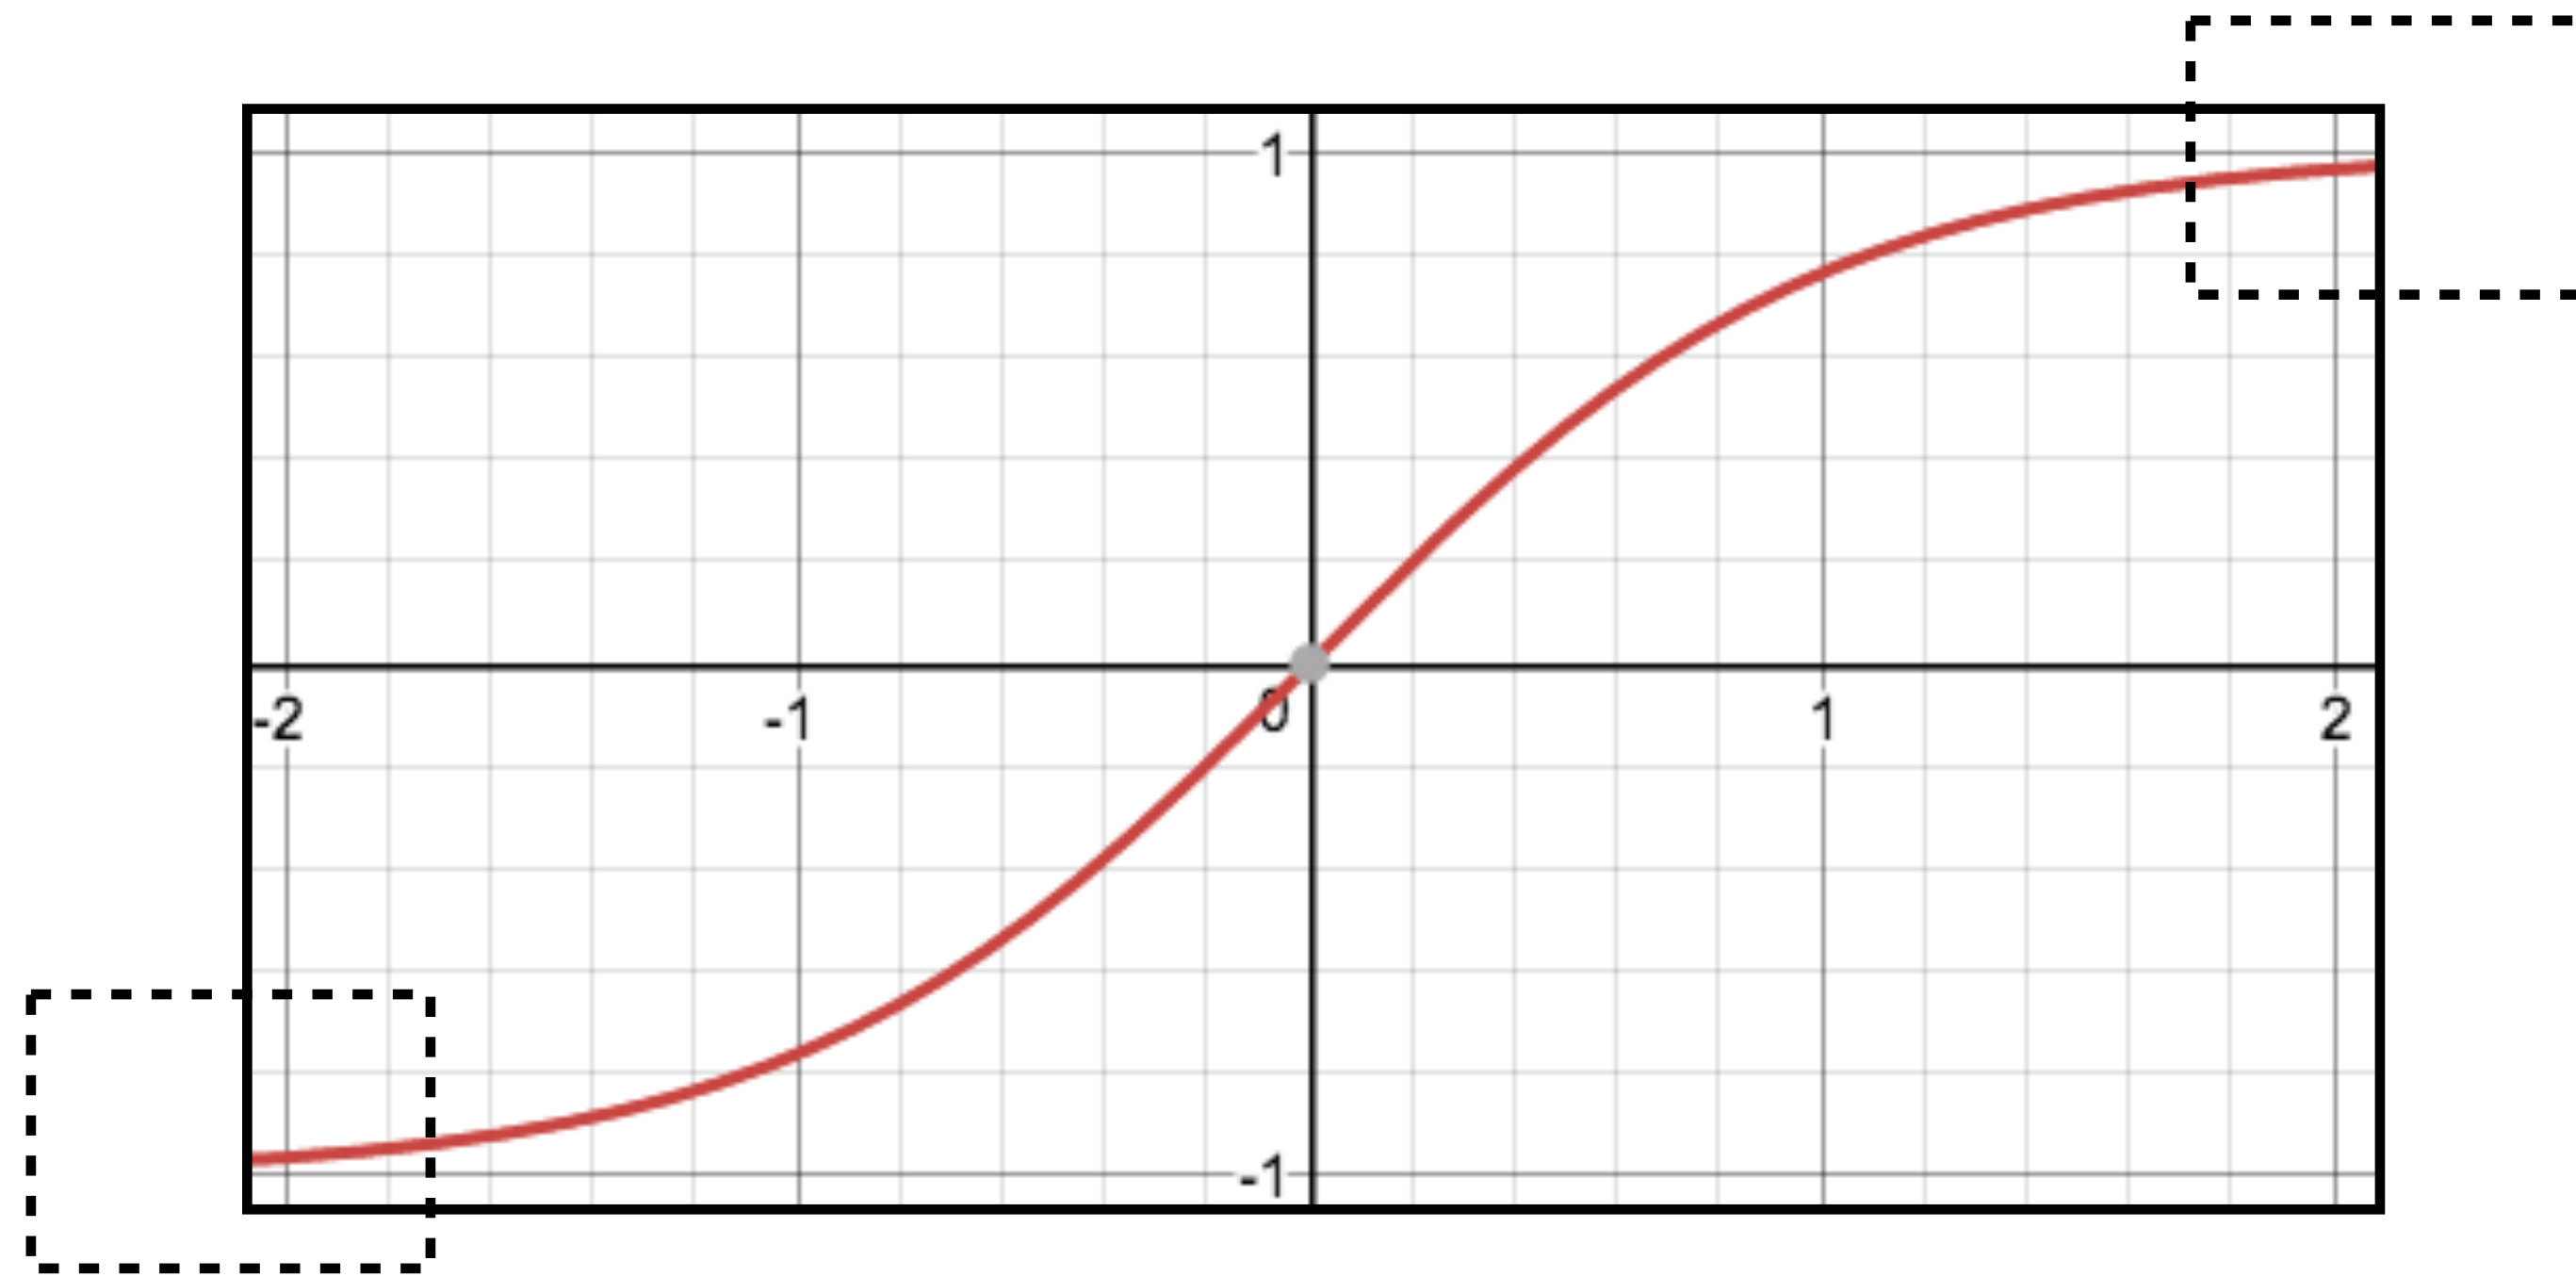
- ▶ Nonlinear model...how does this affect things?





# How does initialization affect learning?

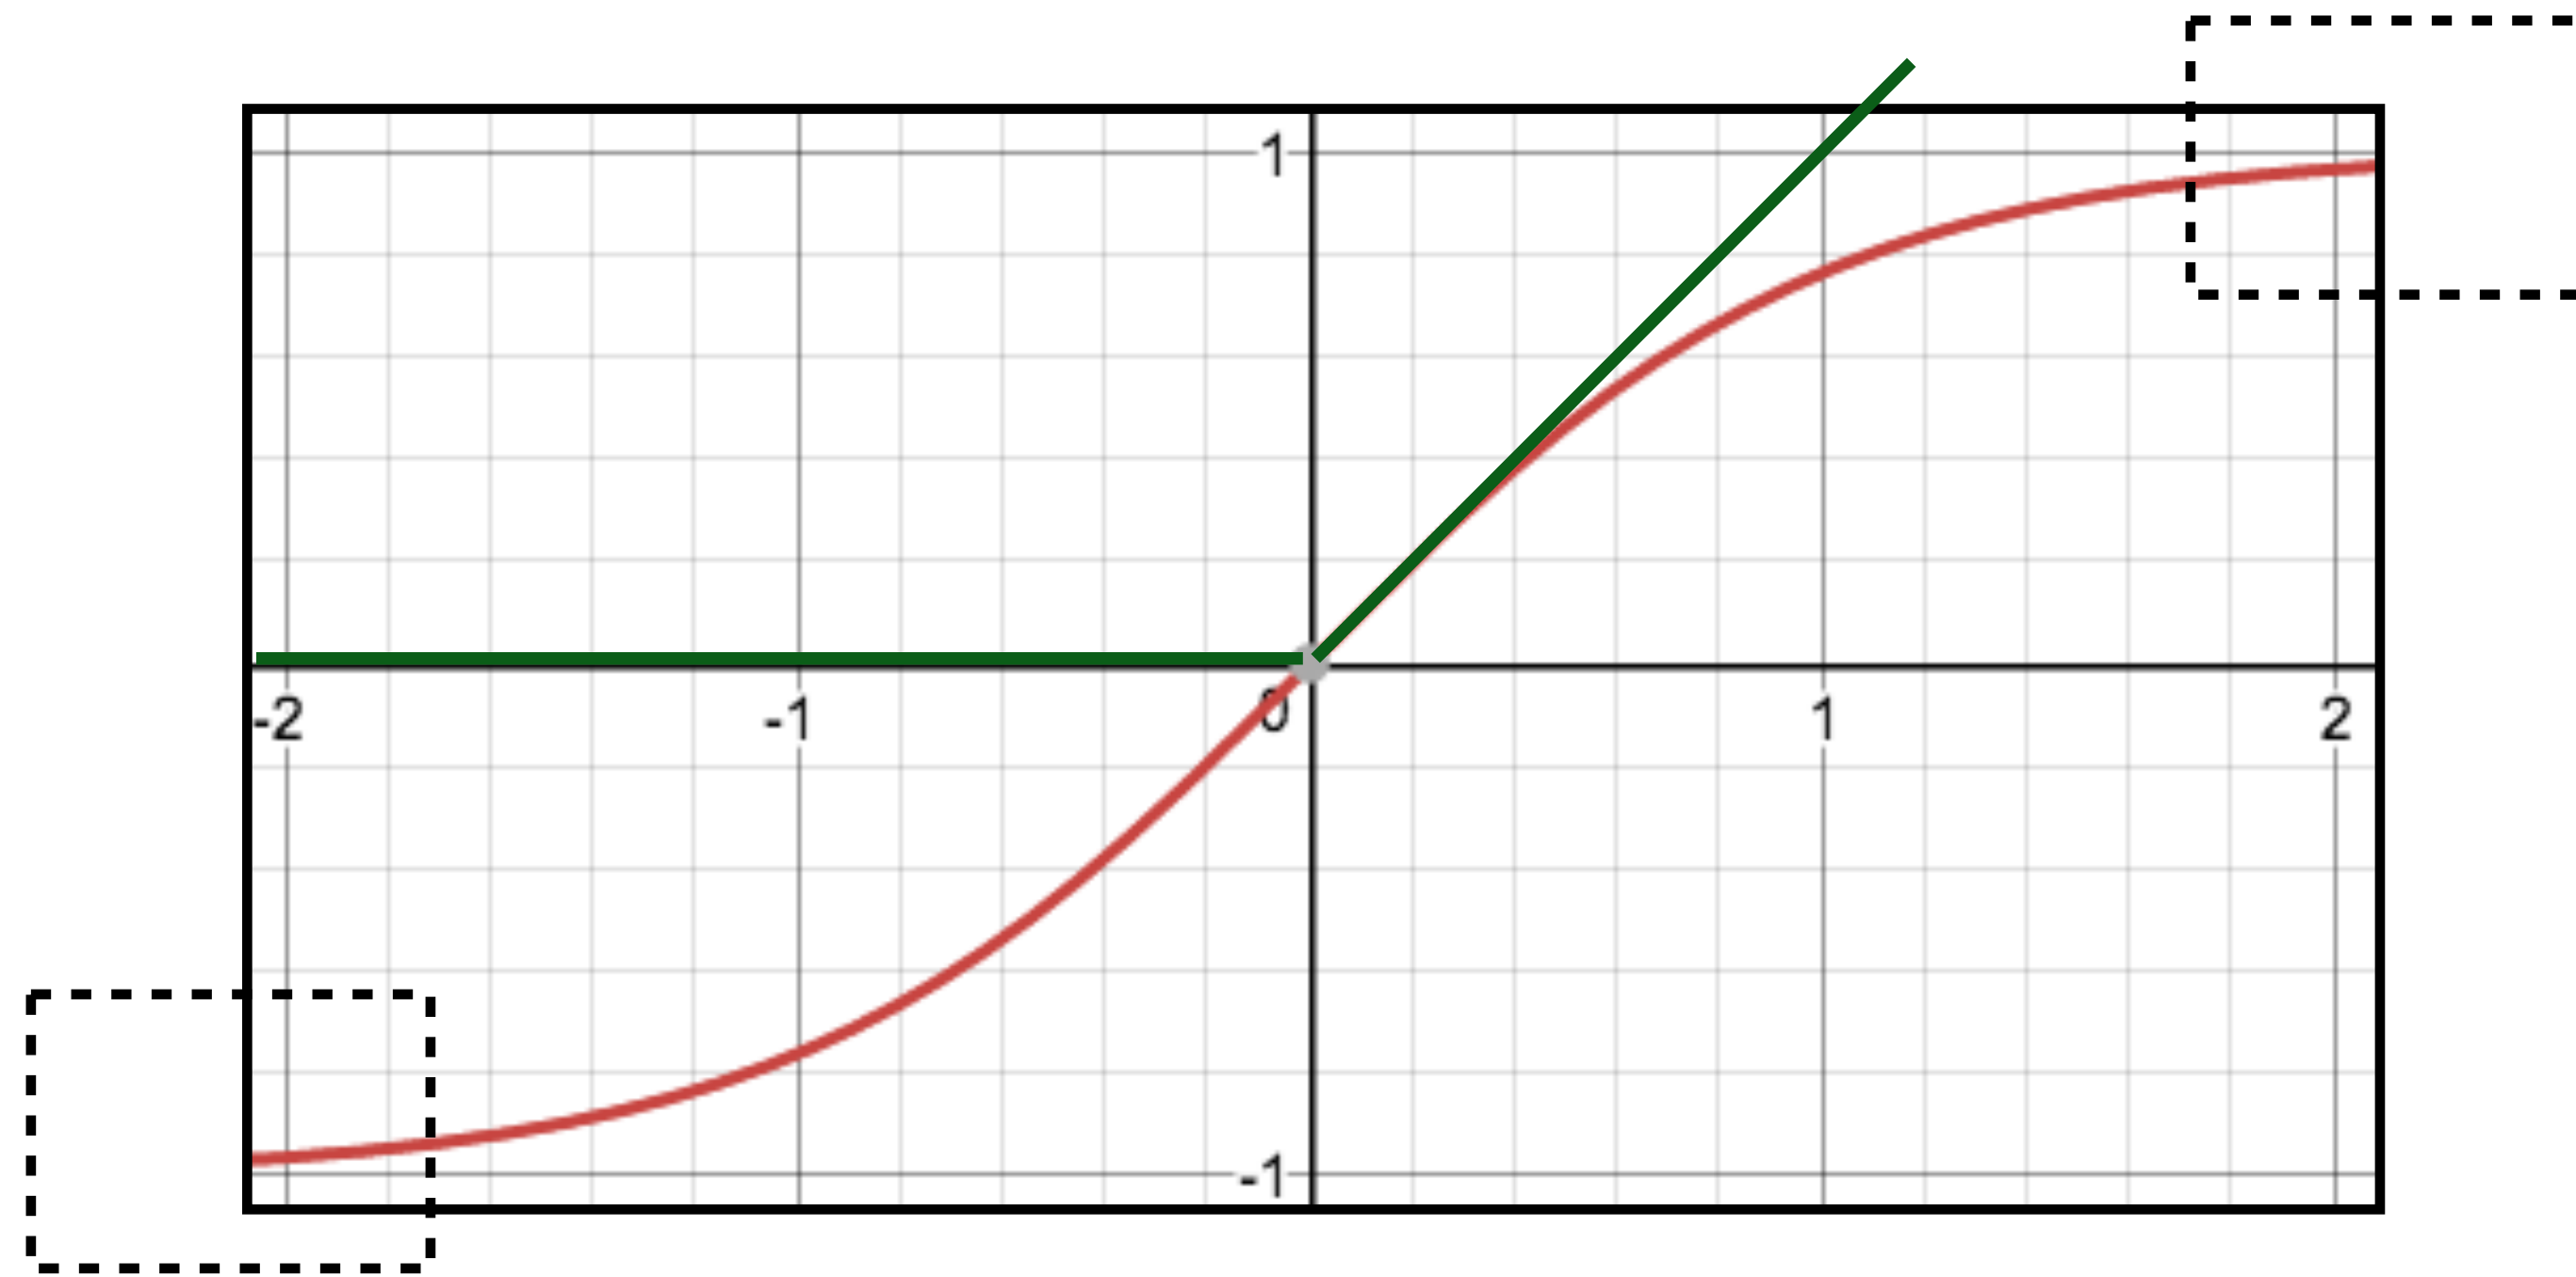
- ▶ Nonlinear model...how does this affect things?



- ▶ If cell activations are too large in absolute value, gradients are small

# How does initialization affect learning?

- ▶ Nonlinear model...how does this affect things?



- ▶ If cell activations are too large in absolute value, gradients are small
- ▶ **ReLU**: larger dynamic range (all positive numbers), but can produce big values, can break down if everything is too negative

# Initialization

---

# Initialization

---

1) Can't use zeroes for parameters to produce hidden layers: all values in that hidden layer are always 0 and have gradients of 0, never change

# Initialization

---

- 1) Can't use zeroes for parameters to produce hidden layers: all values in that hidden layer are always 0 and have gradients of 0, never change
- 2) Initialize too large and cells are saturated

# Initialization

---

- 1) Can't use zeroes for parameters to produce hidden layers: all values in that hidden layer are always 0 and have gradients of 0, never change
  - 2) Initialize too large and cells are saturated
- ▶ Can do random uniform / normal initialization with appropriate scale

# Initialization

---

1) Can't use zeroes for parameters to produce hidden layers: all values in that hidden layer are always 0 and have gradients of 0, never change

2) Initialize too large and cells are saturated

► Can do random uniform / normal initialization with appropriate scale

► Xavier initializer:  $U \left[ -\sqrt{\frac{6}{\text{fan-in} + \text{fan-out}}}, +\sqrt{\frac{6}{\text{fan-in} + \text{fan-out}}} \right]$

# Initialization

---

- 1) Can't use zeroes for parameters to produce hidden layers: all values in that hidden layer are always 0 and have gradients of 0, never change
  - 2) Initialize too large and cells are saturated
- ▶ Can do random uniform / normal initialization with appropriate scale
  - ▶ Xavier initializer:  $U \left[ -\sqrt{\frac{6}{\text{fan-in} + \text{fan-out}}}, +\sqrt{\frac{6}{\text{fan-in} + \text{fan-out}}} \right]$ 
    - ▶ Want variance of inputs and gradients for each layer to be the same



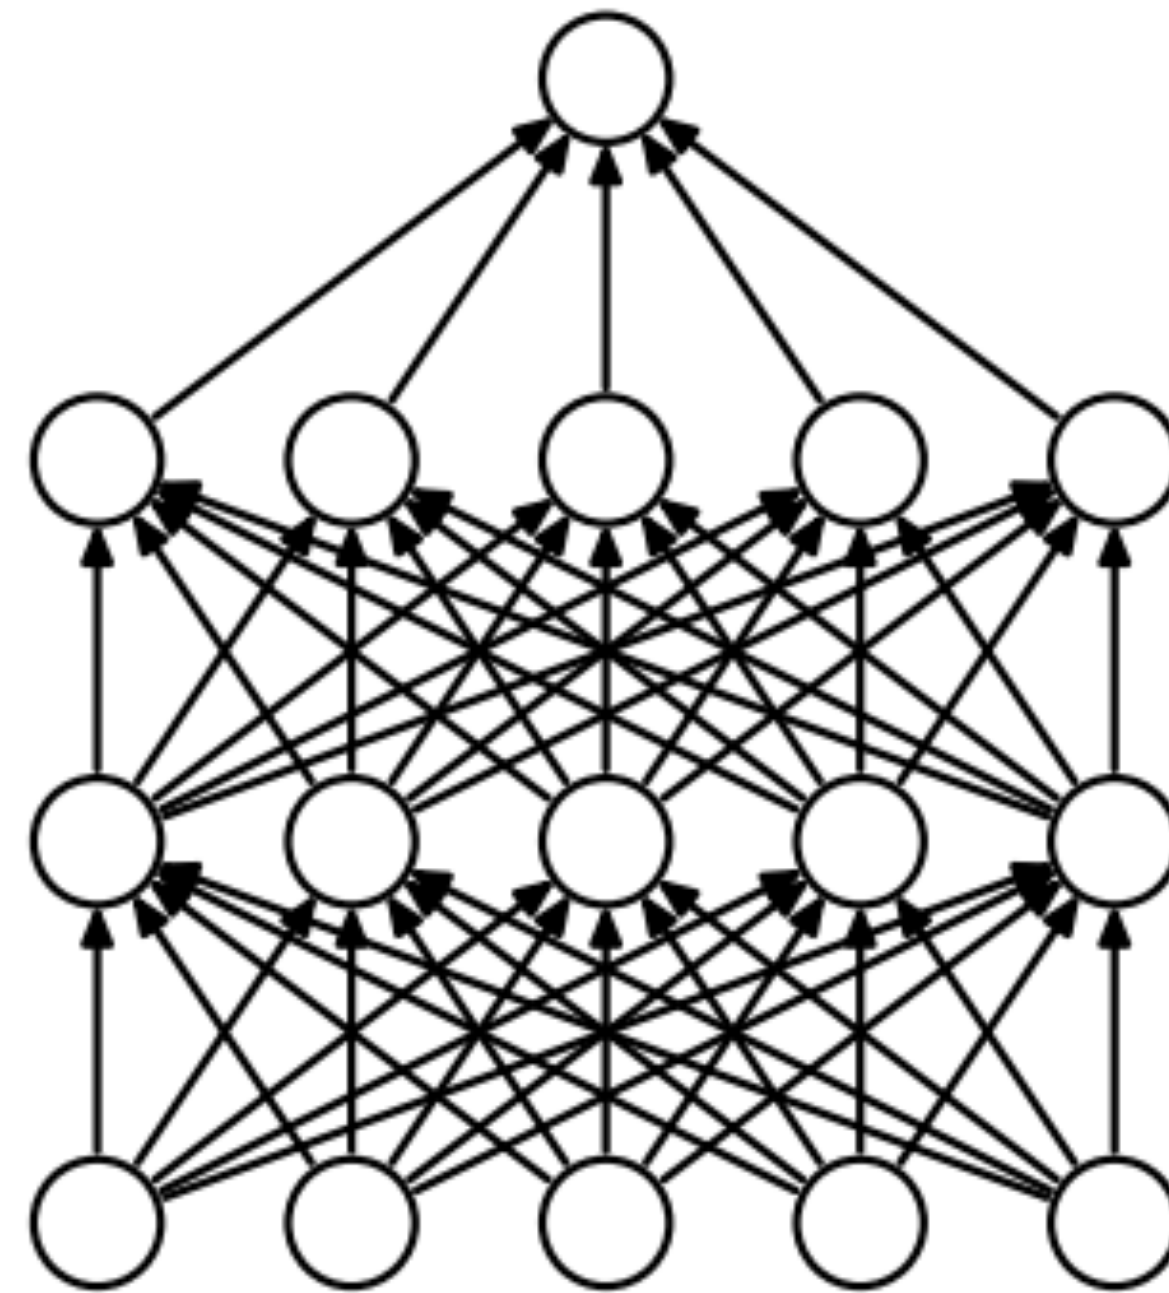
# Initialization

---

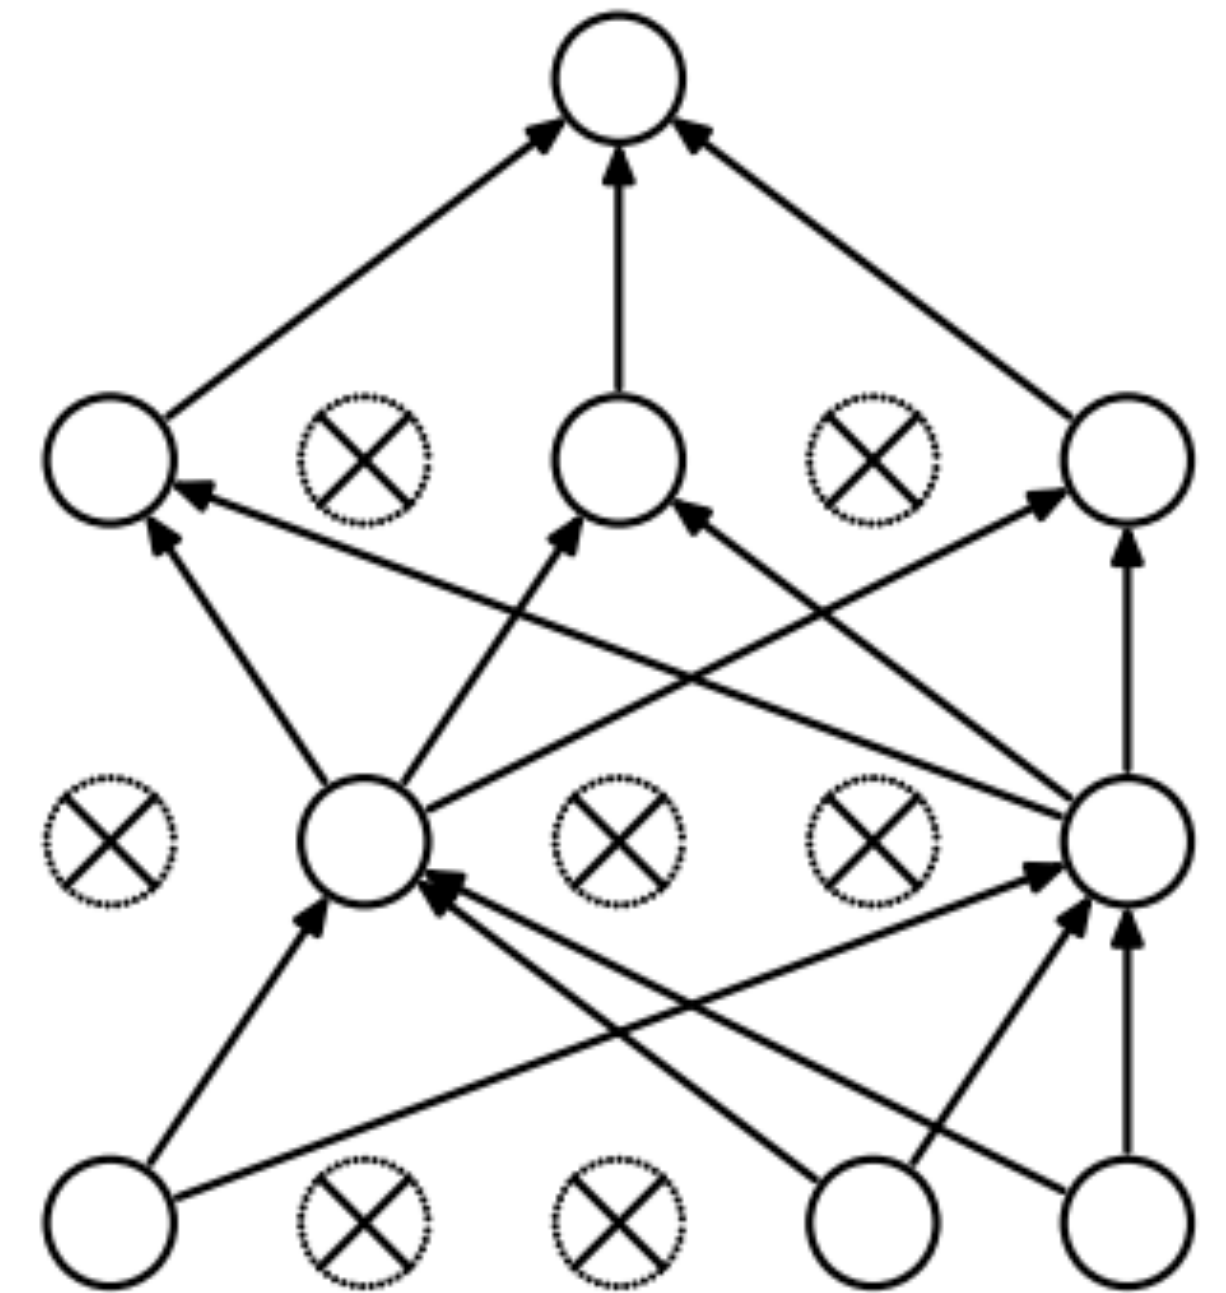
- 1) Can't use zeroes for parameters to produce hidden layers: all values in that hidden layer are always 0 and have gradients of 0, never change
  - 2) Initialize too large and cells are saturated
- ▶ Can do random uniform / normal initialization with appropriate scale
  - ▶ Xavier initializer:  $U \left[ -\sqrt{\frac{6}{\text{fan-in} + \text{fan-out}}}, +\sqrt{\frac{6}{\text{fan-in} + \text{fan-out}}} \right]$ 
    - ▶ Want variance of inputs and gradients for each layer to be the same
  - ▶ Batch normalization (Ioffe and Szegedy, 2015): periodically shift+rescale each layer to have mean 0 and variance 1 over a batch (useful if net is deep)

# Dropout

- Probabilistically zero out parts of the network during training to prevent overfitting, use whole network at test time



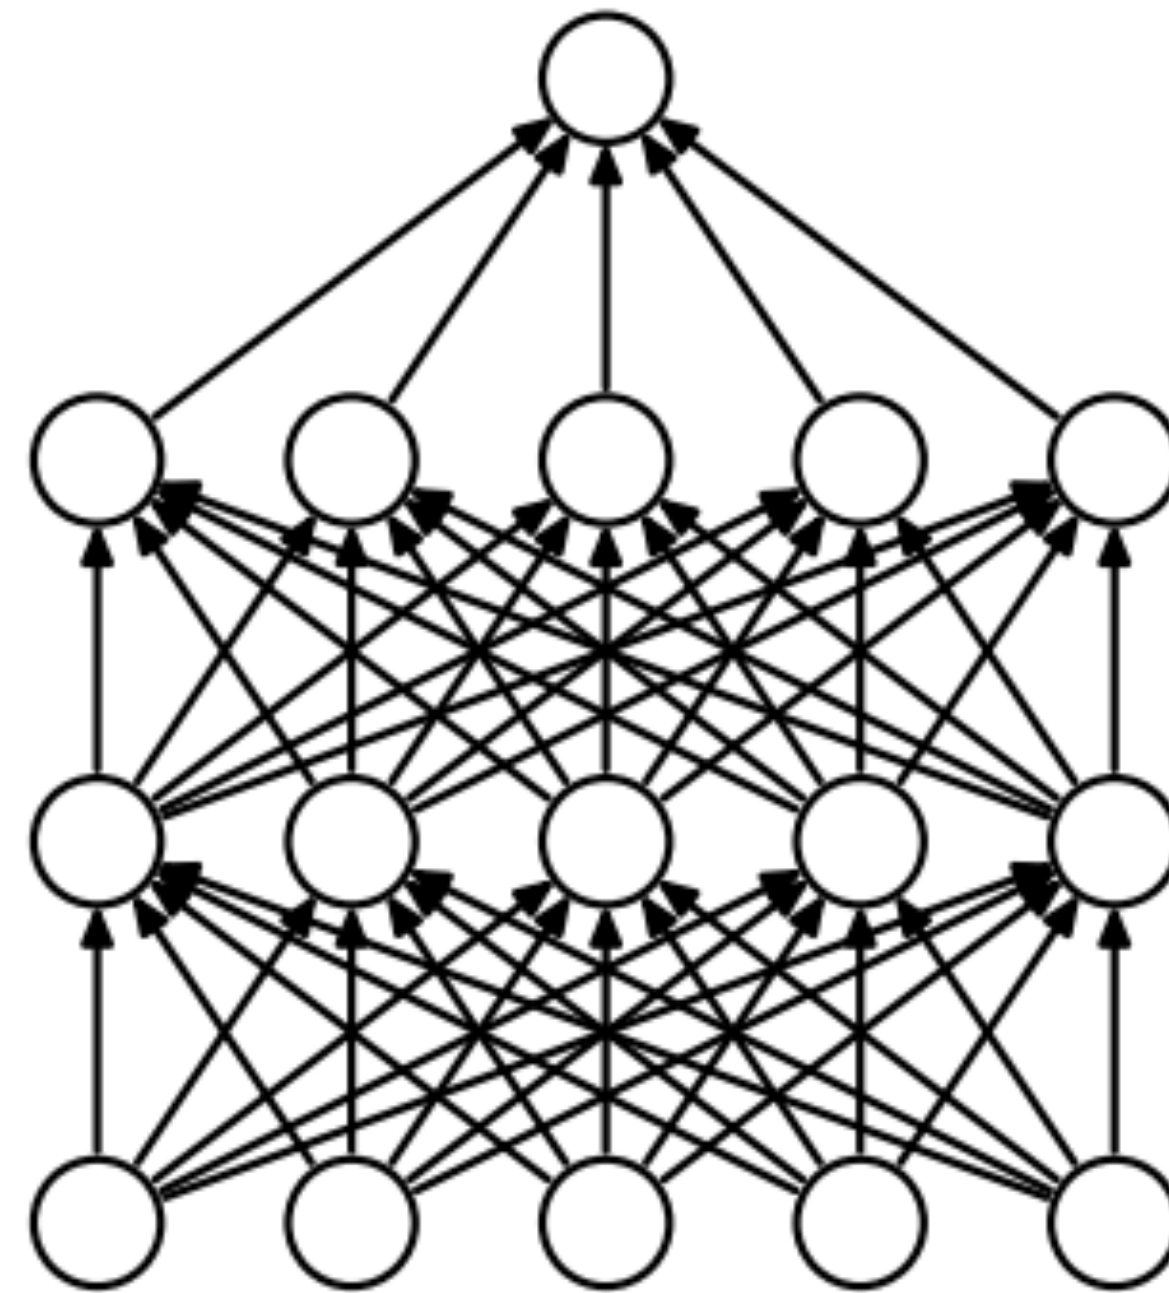
(a) Standard Neural Net



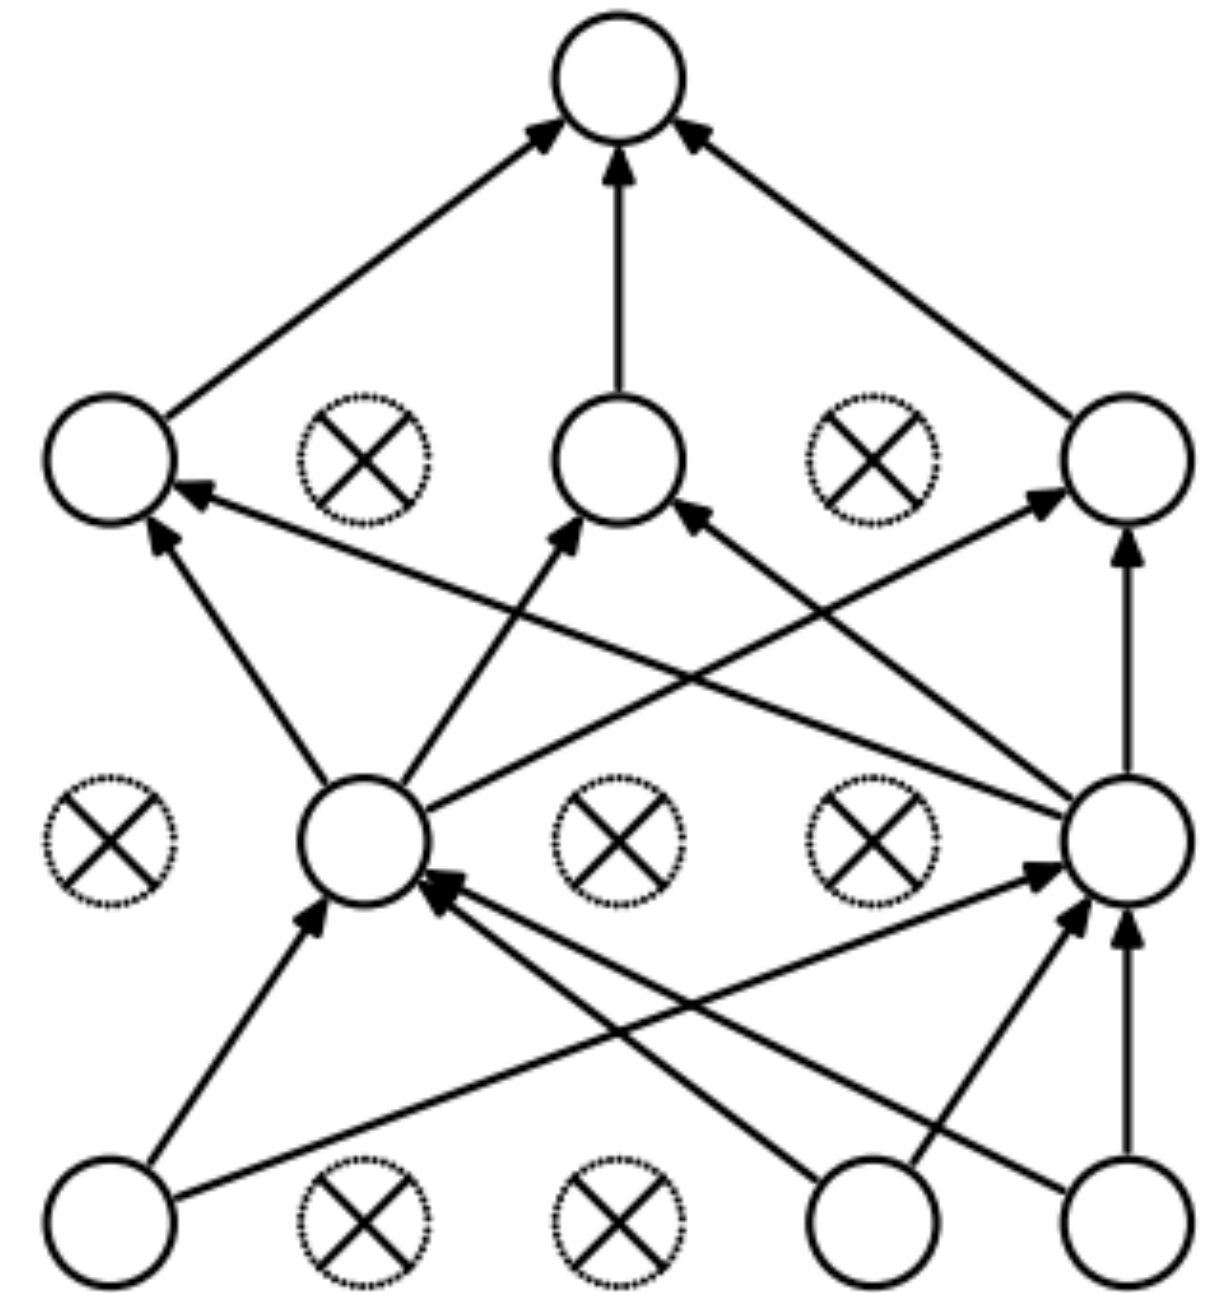
(b) After applying dropout.

# Dropout

- ▶ Probabilistically zero out parts of the network during training to prevent overfitting, use whole network at test time
- ▶ Form of stochastic regularization



(a) Standard Neural Net

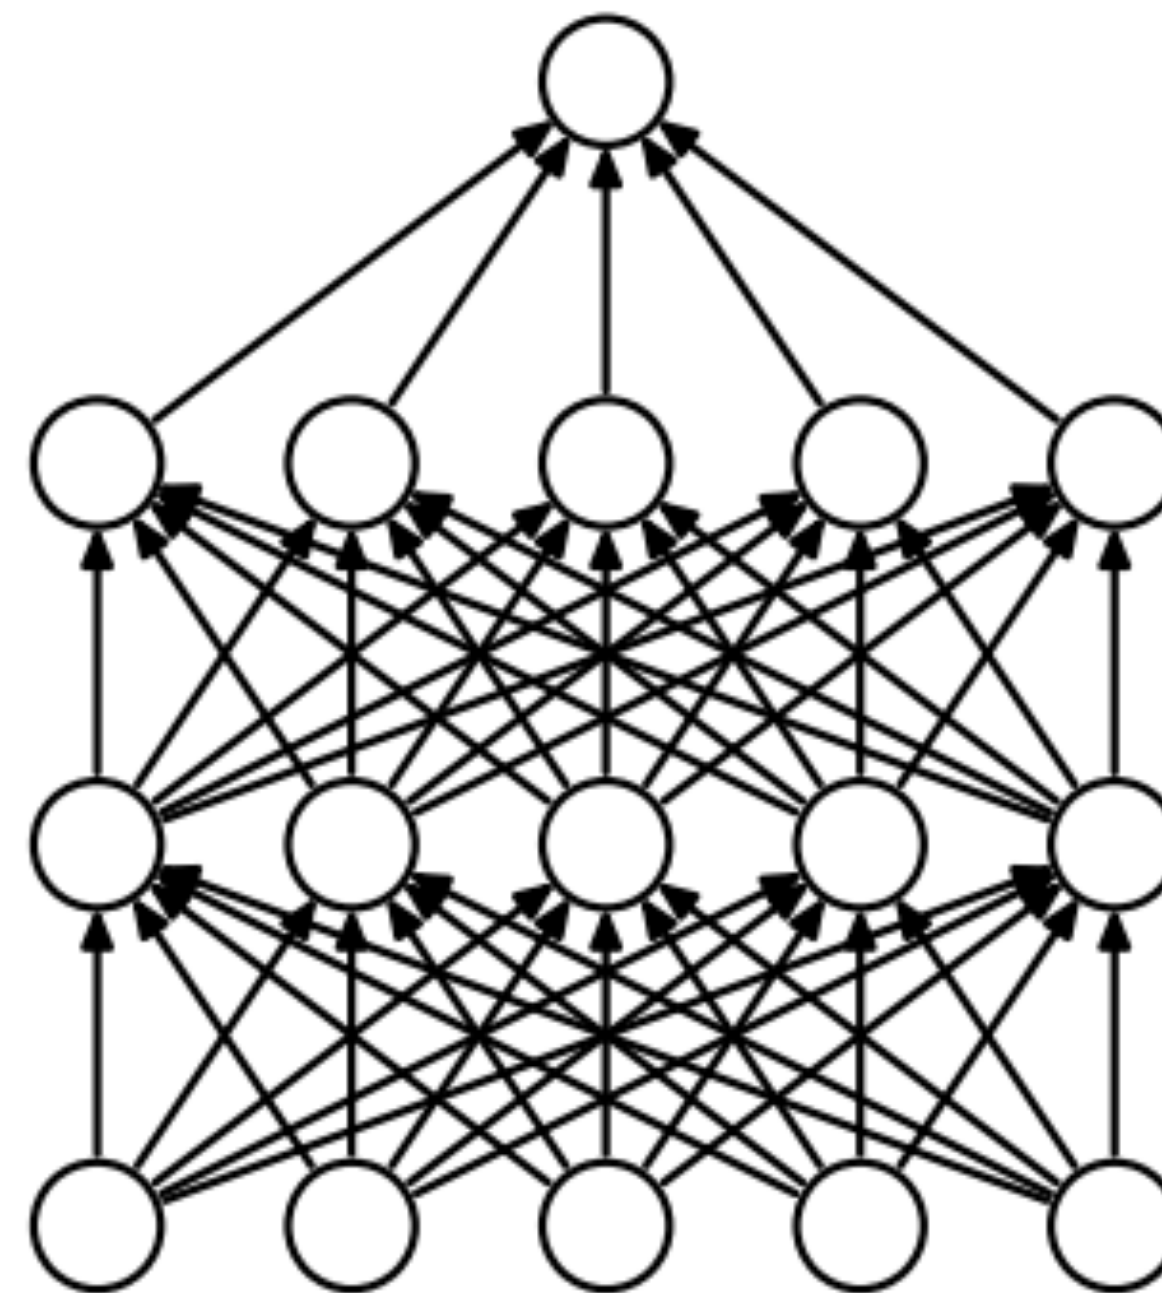


(b) After applying dropout.

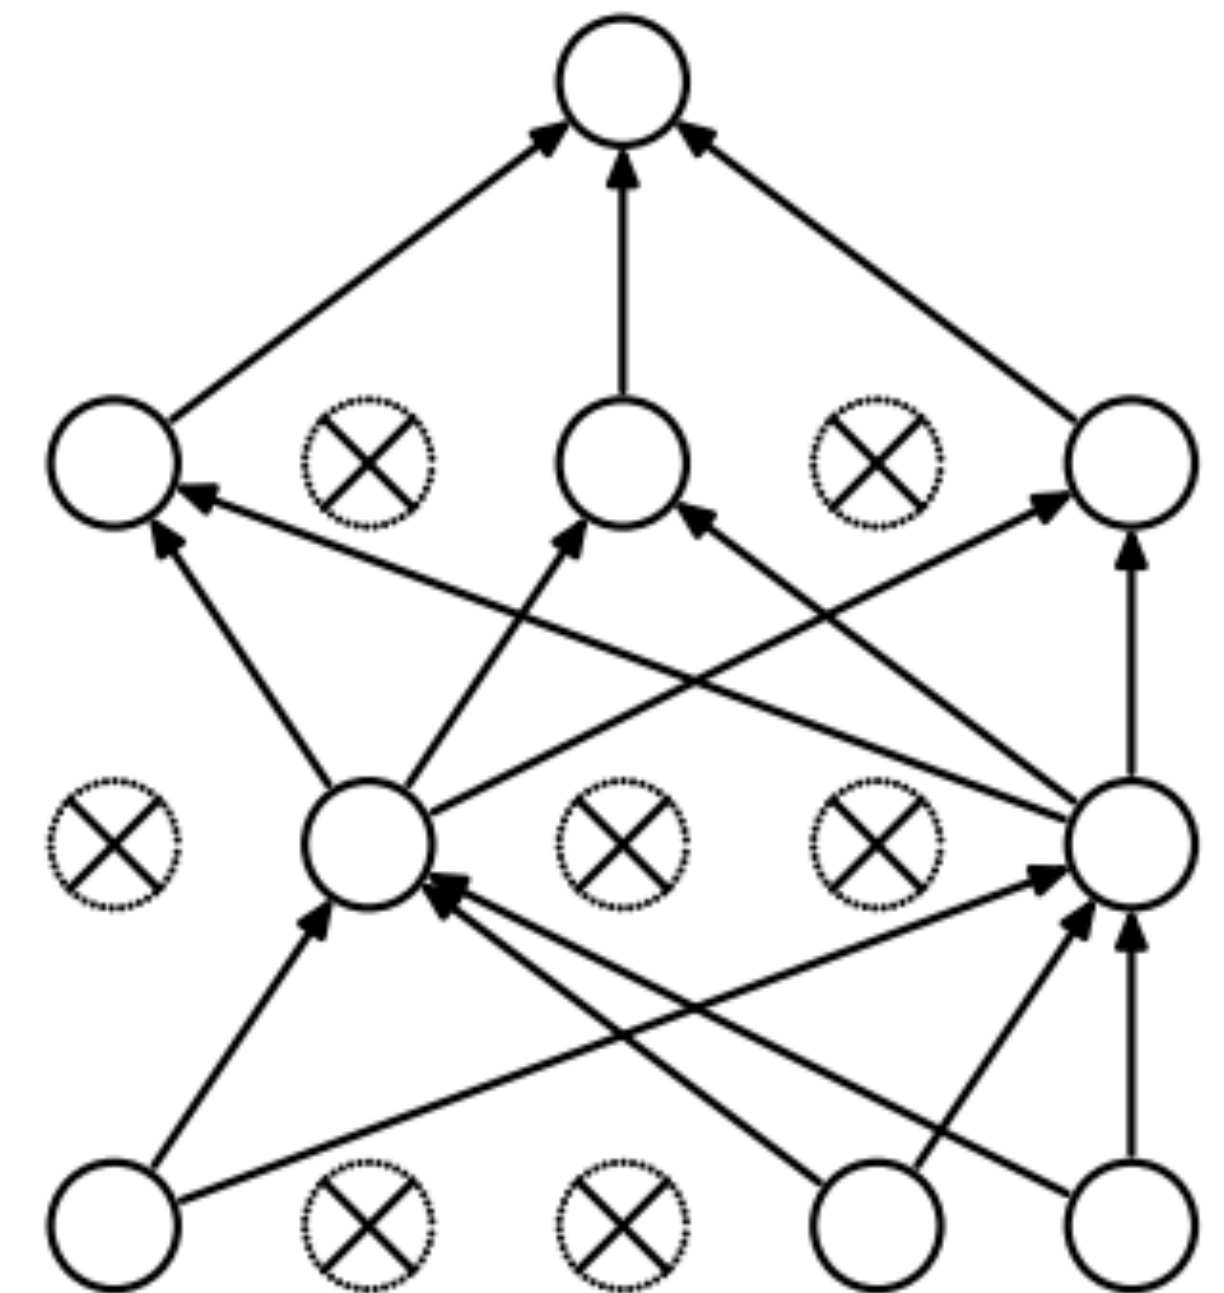


# Dropout

- ▶ Probabilistically zero out parts of the network during training to prevent overfitting, use whole network at test time
- ▶ Form of stochastic regularization
- ▶ Similar to benefits of ensembling: network needs to be robust to missing signals, so it has redundancy



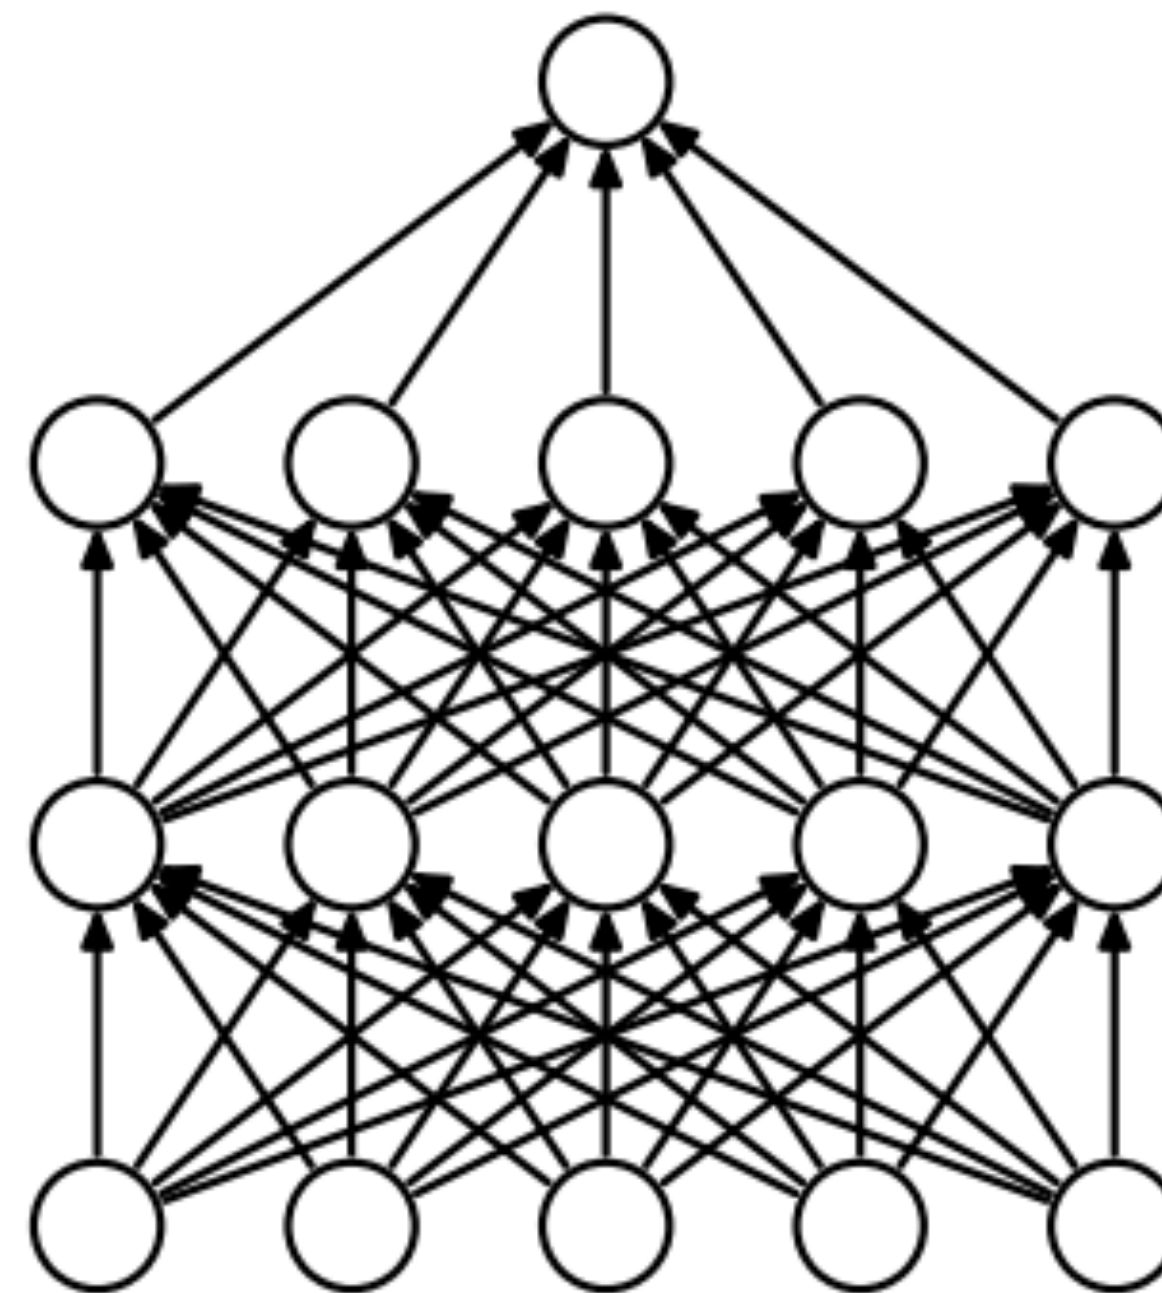
(a) Standard Neural Net




(b) After applying dropout.

# Dropout

- ▶ Probabilistically zero out parts of the network during training to prevent overfitting, use whole network at test time
- ▶ Form of stochastic regularization
- ▶ Similar to benefits of ensembling: network needs to be robust to missing signals, so it has redundancy
- ▶ One line in Pytorch/Tensorflow



(a) Standard Neural Net



(b) After applying dropout.

# Optimizer

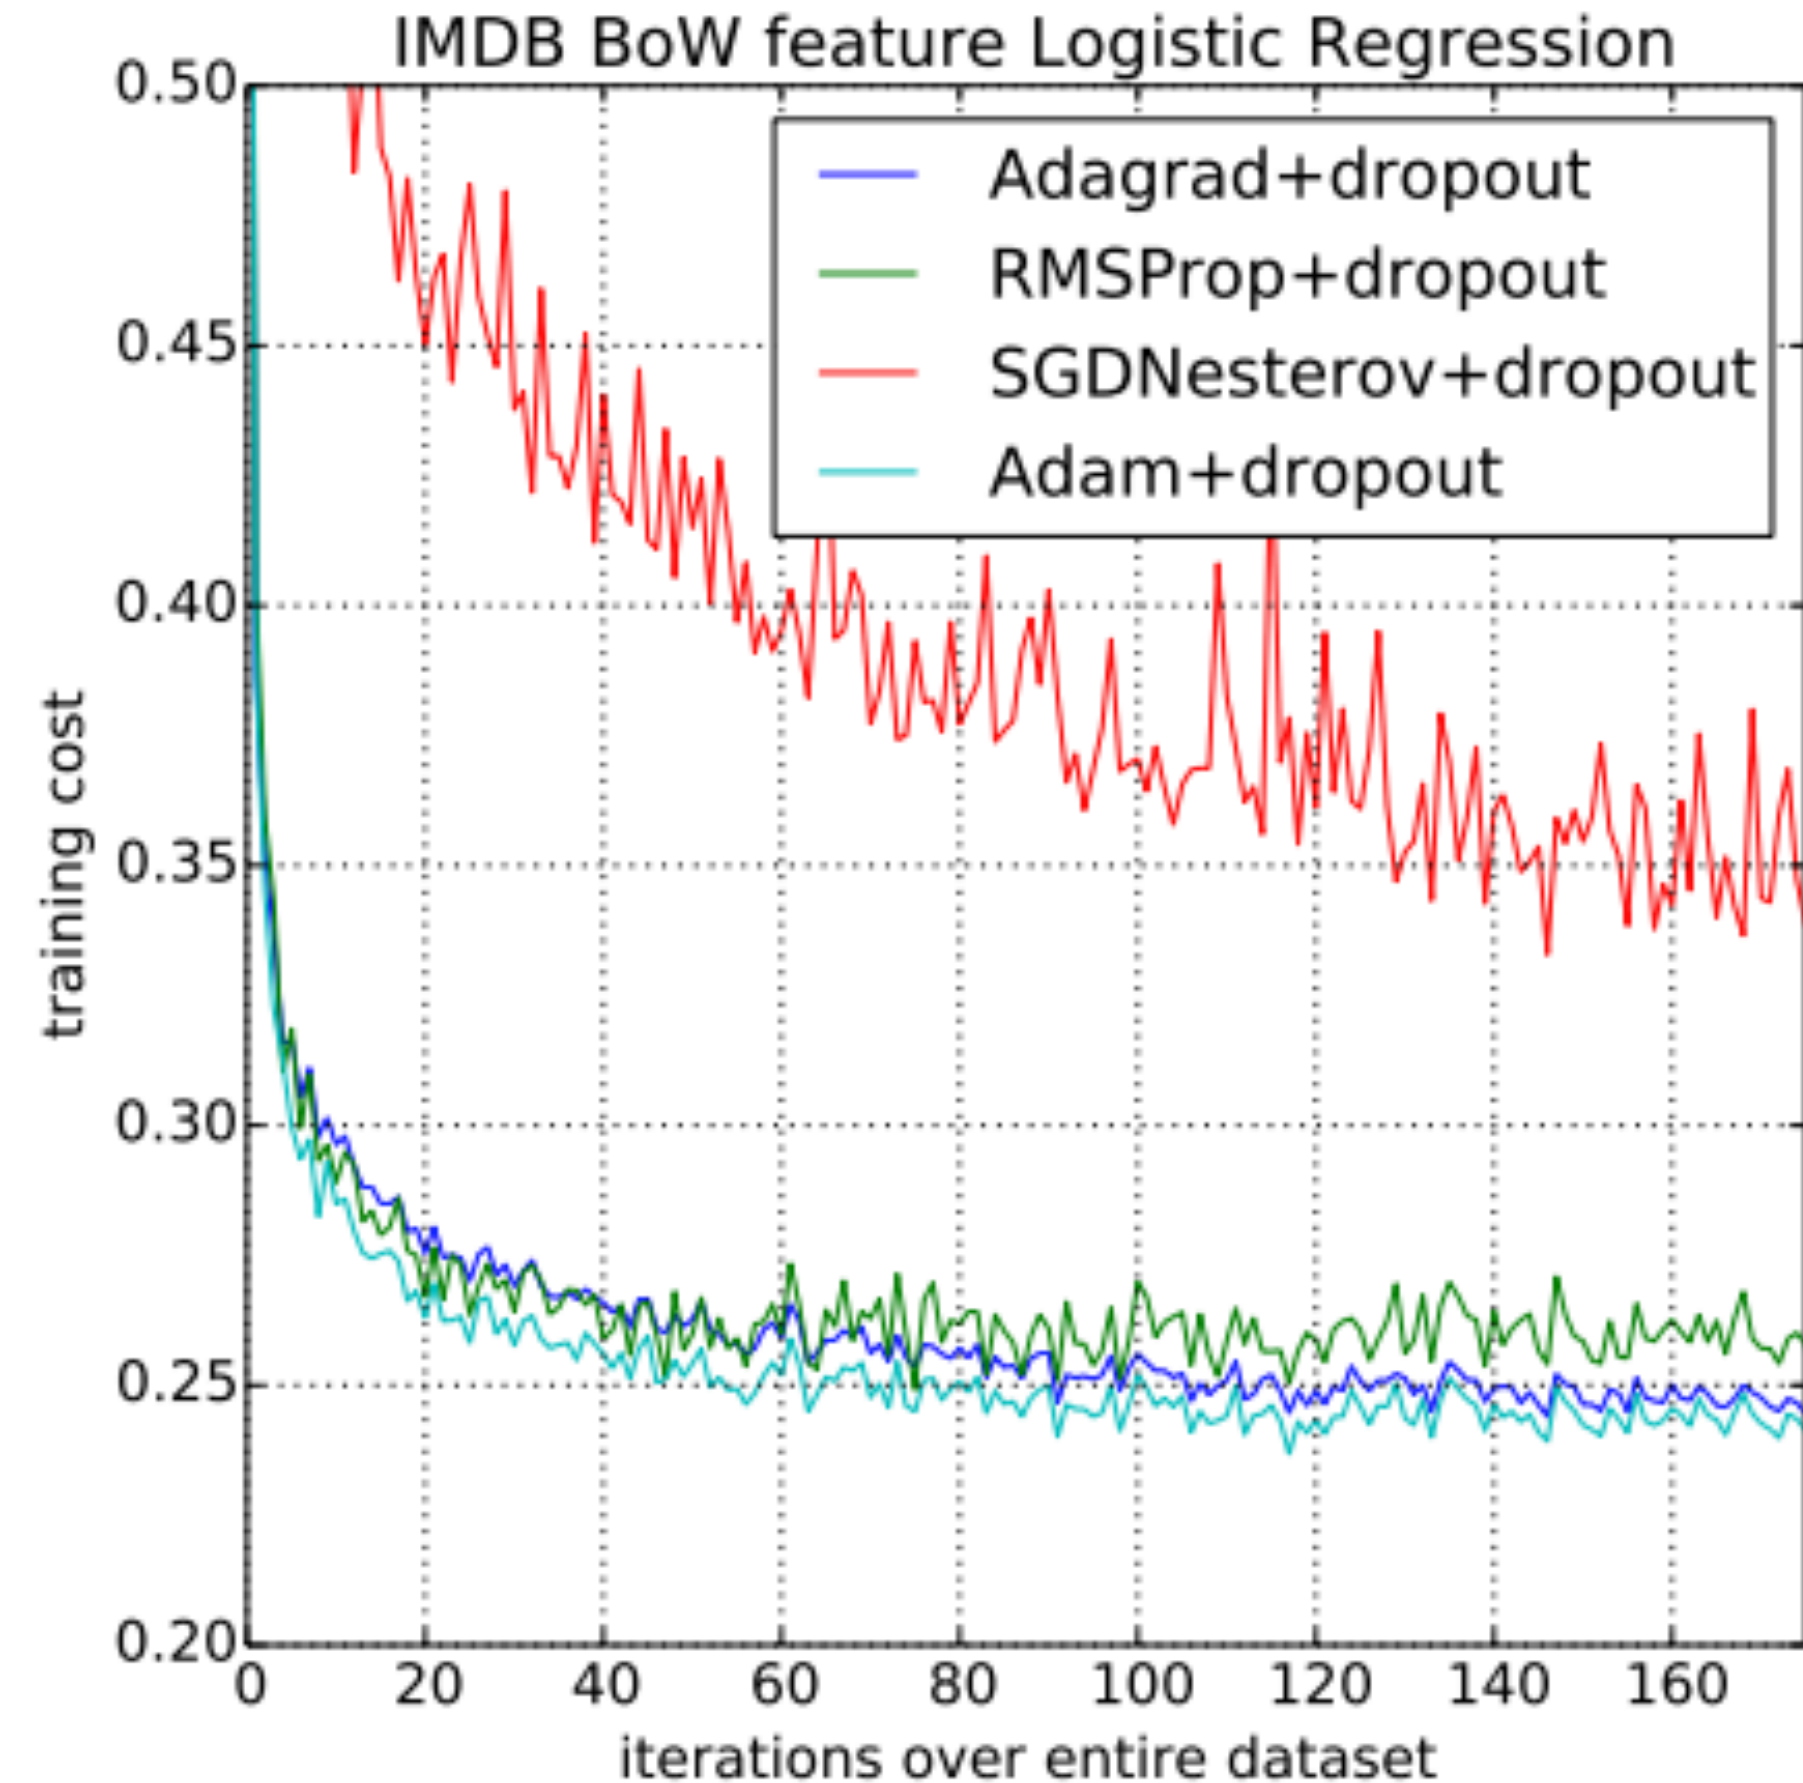
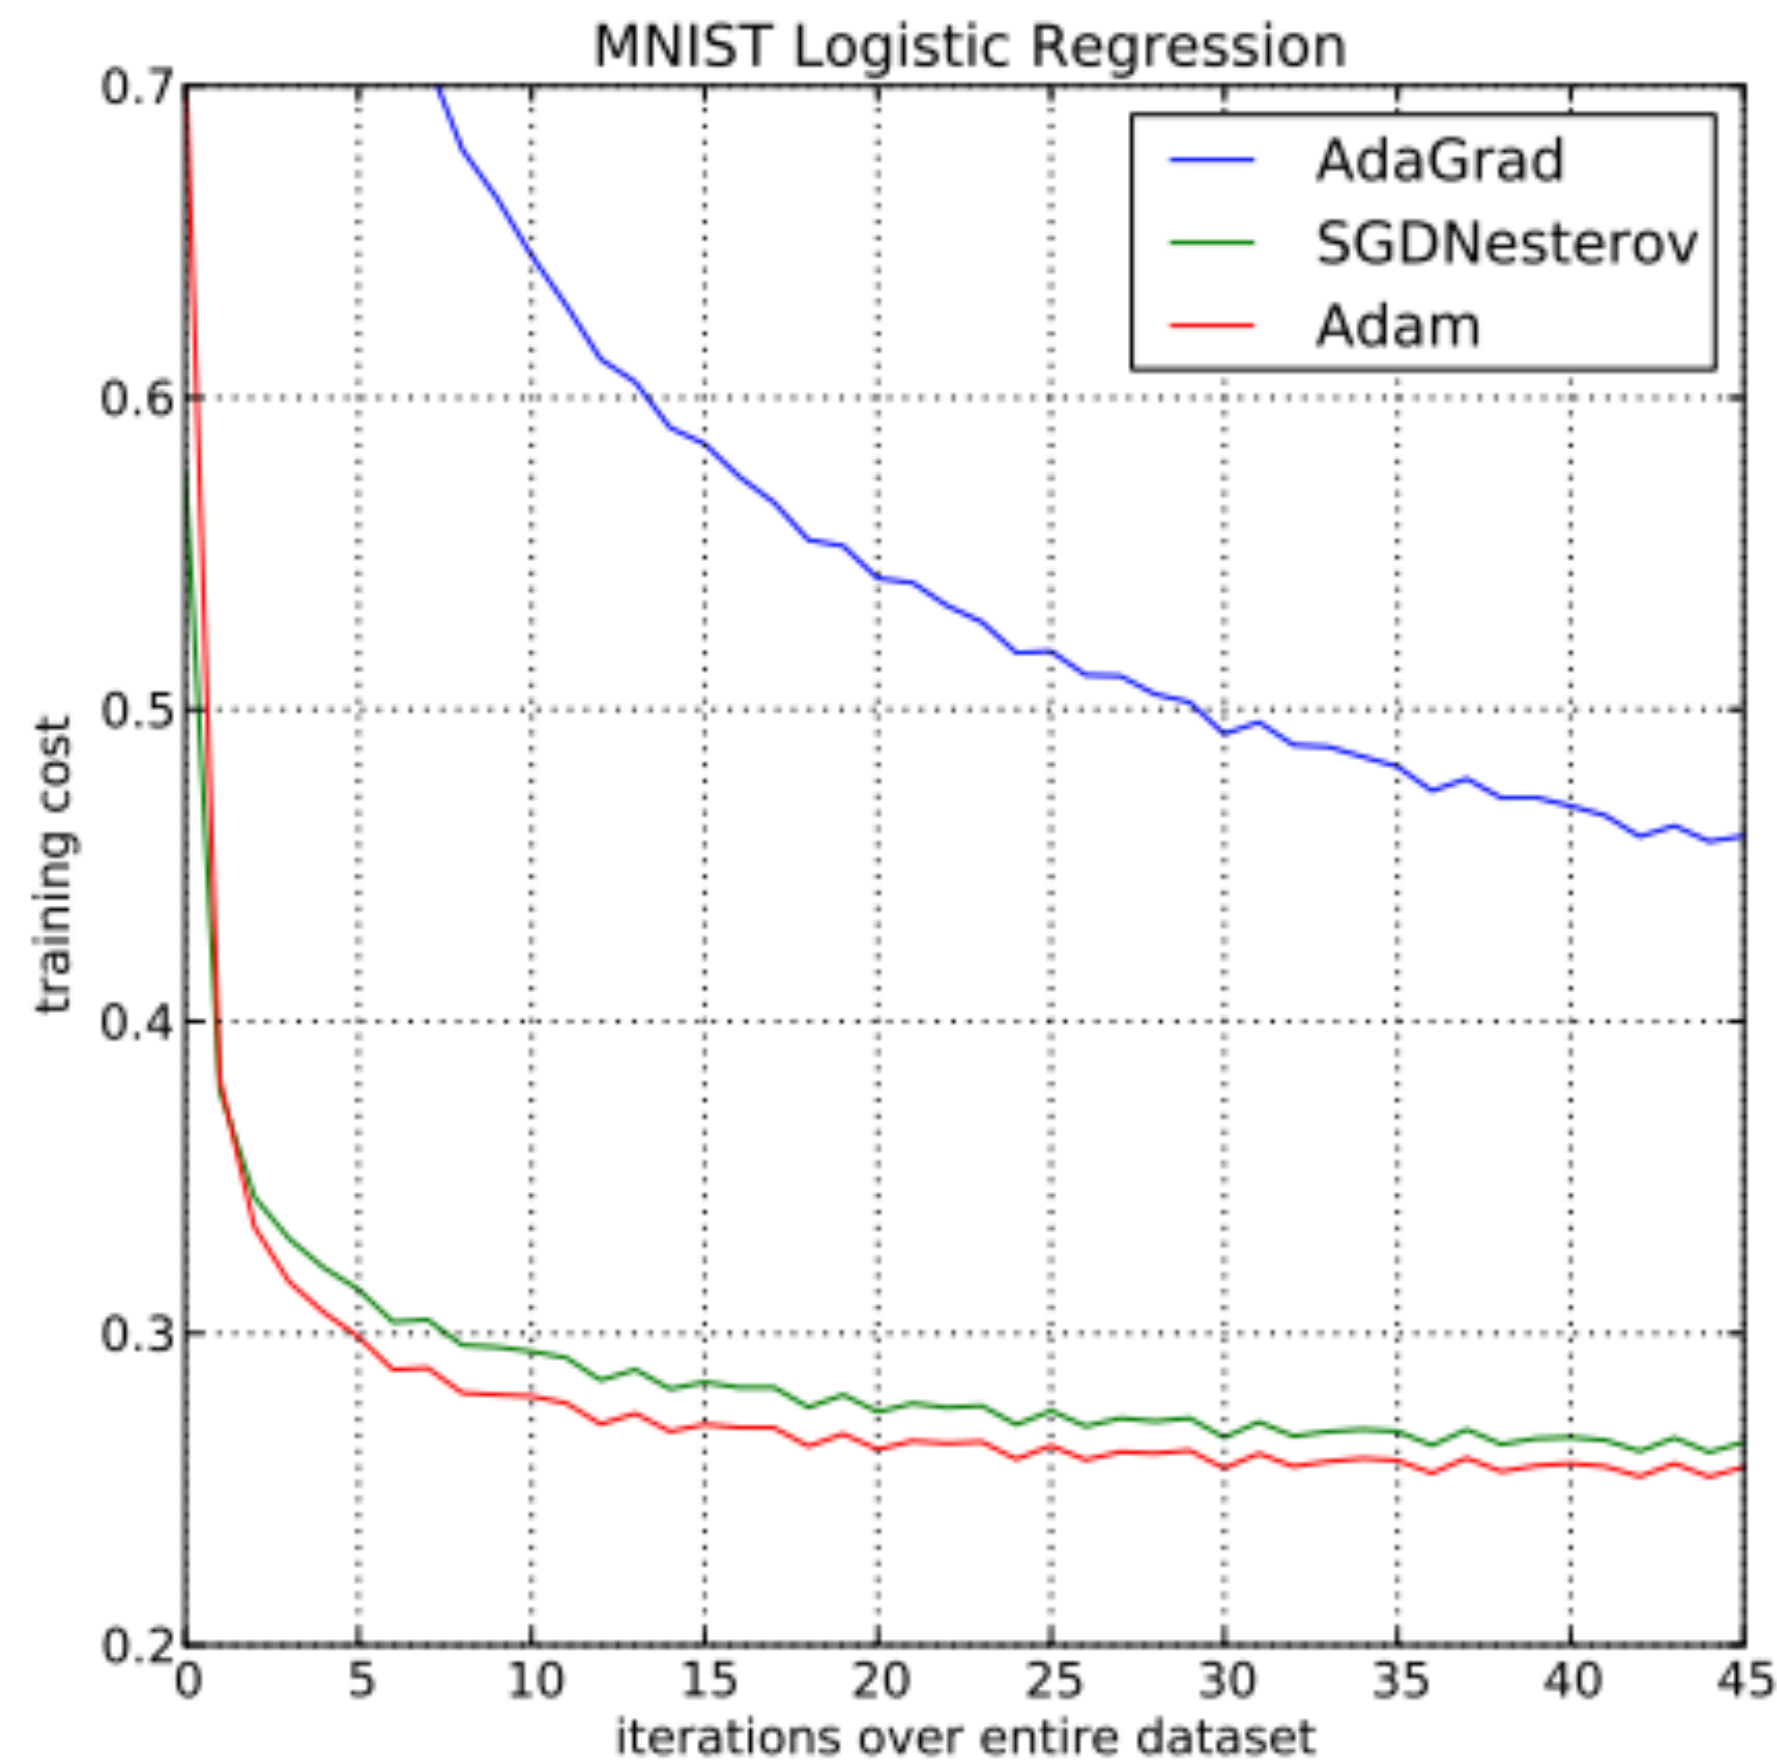
---

- ▶ Adam (Kingma and Ba, ICLR 2015) is very widely used
- ▶ Adaptive step size like Adagrad, incorporates momentum



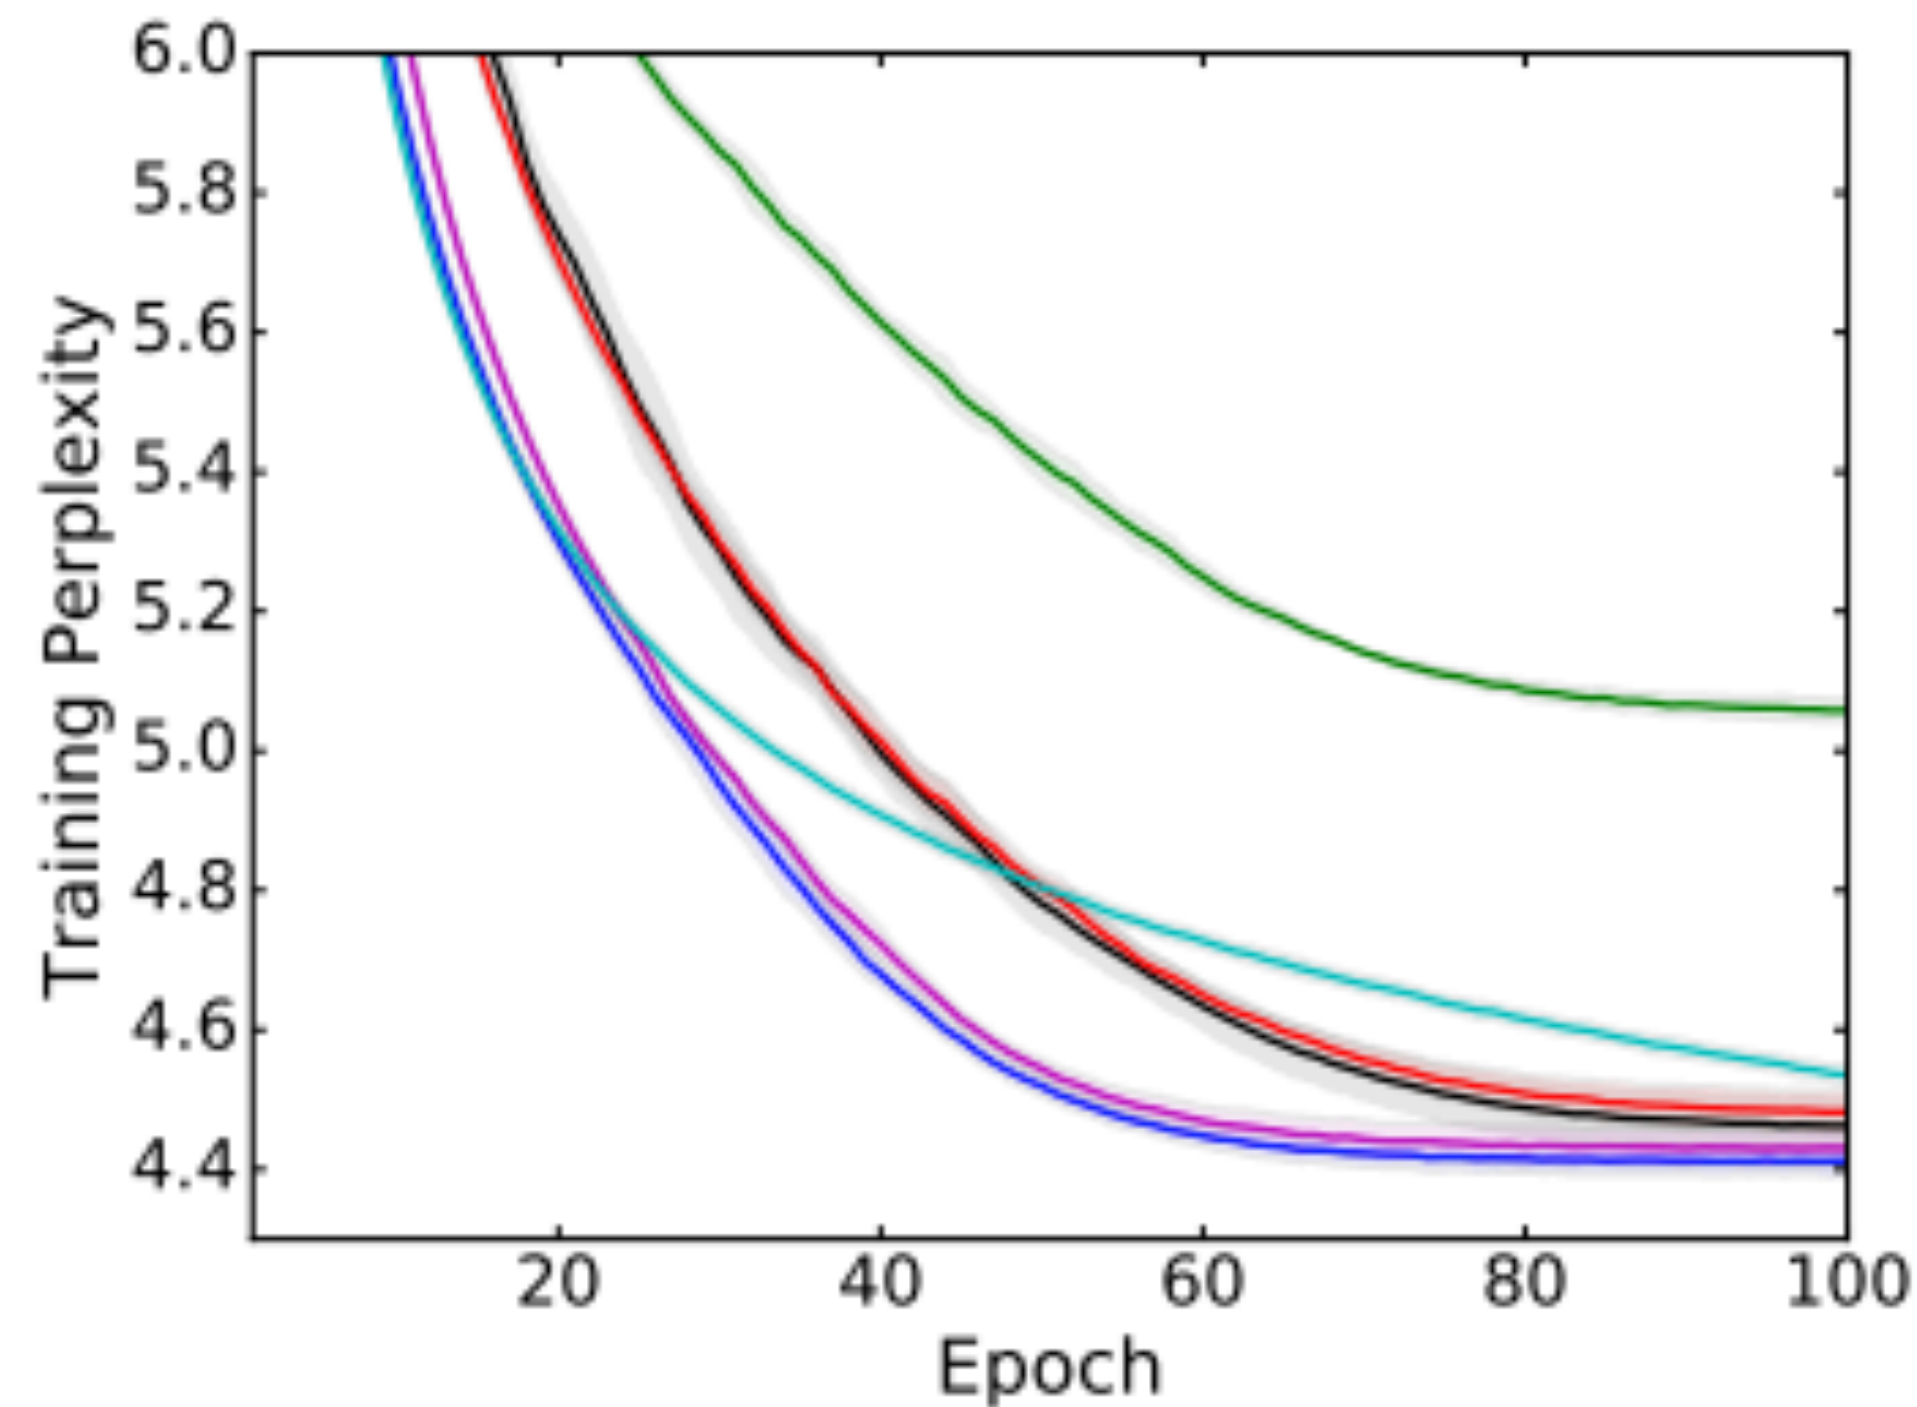
# Optimizer

- ▶ Adam (Kingma and Ba, ICLR 2015) is very widely used
- ▶ Adaptive step size like Adagrad, incorporates momentum

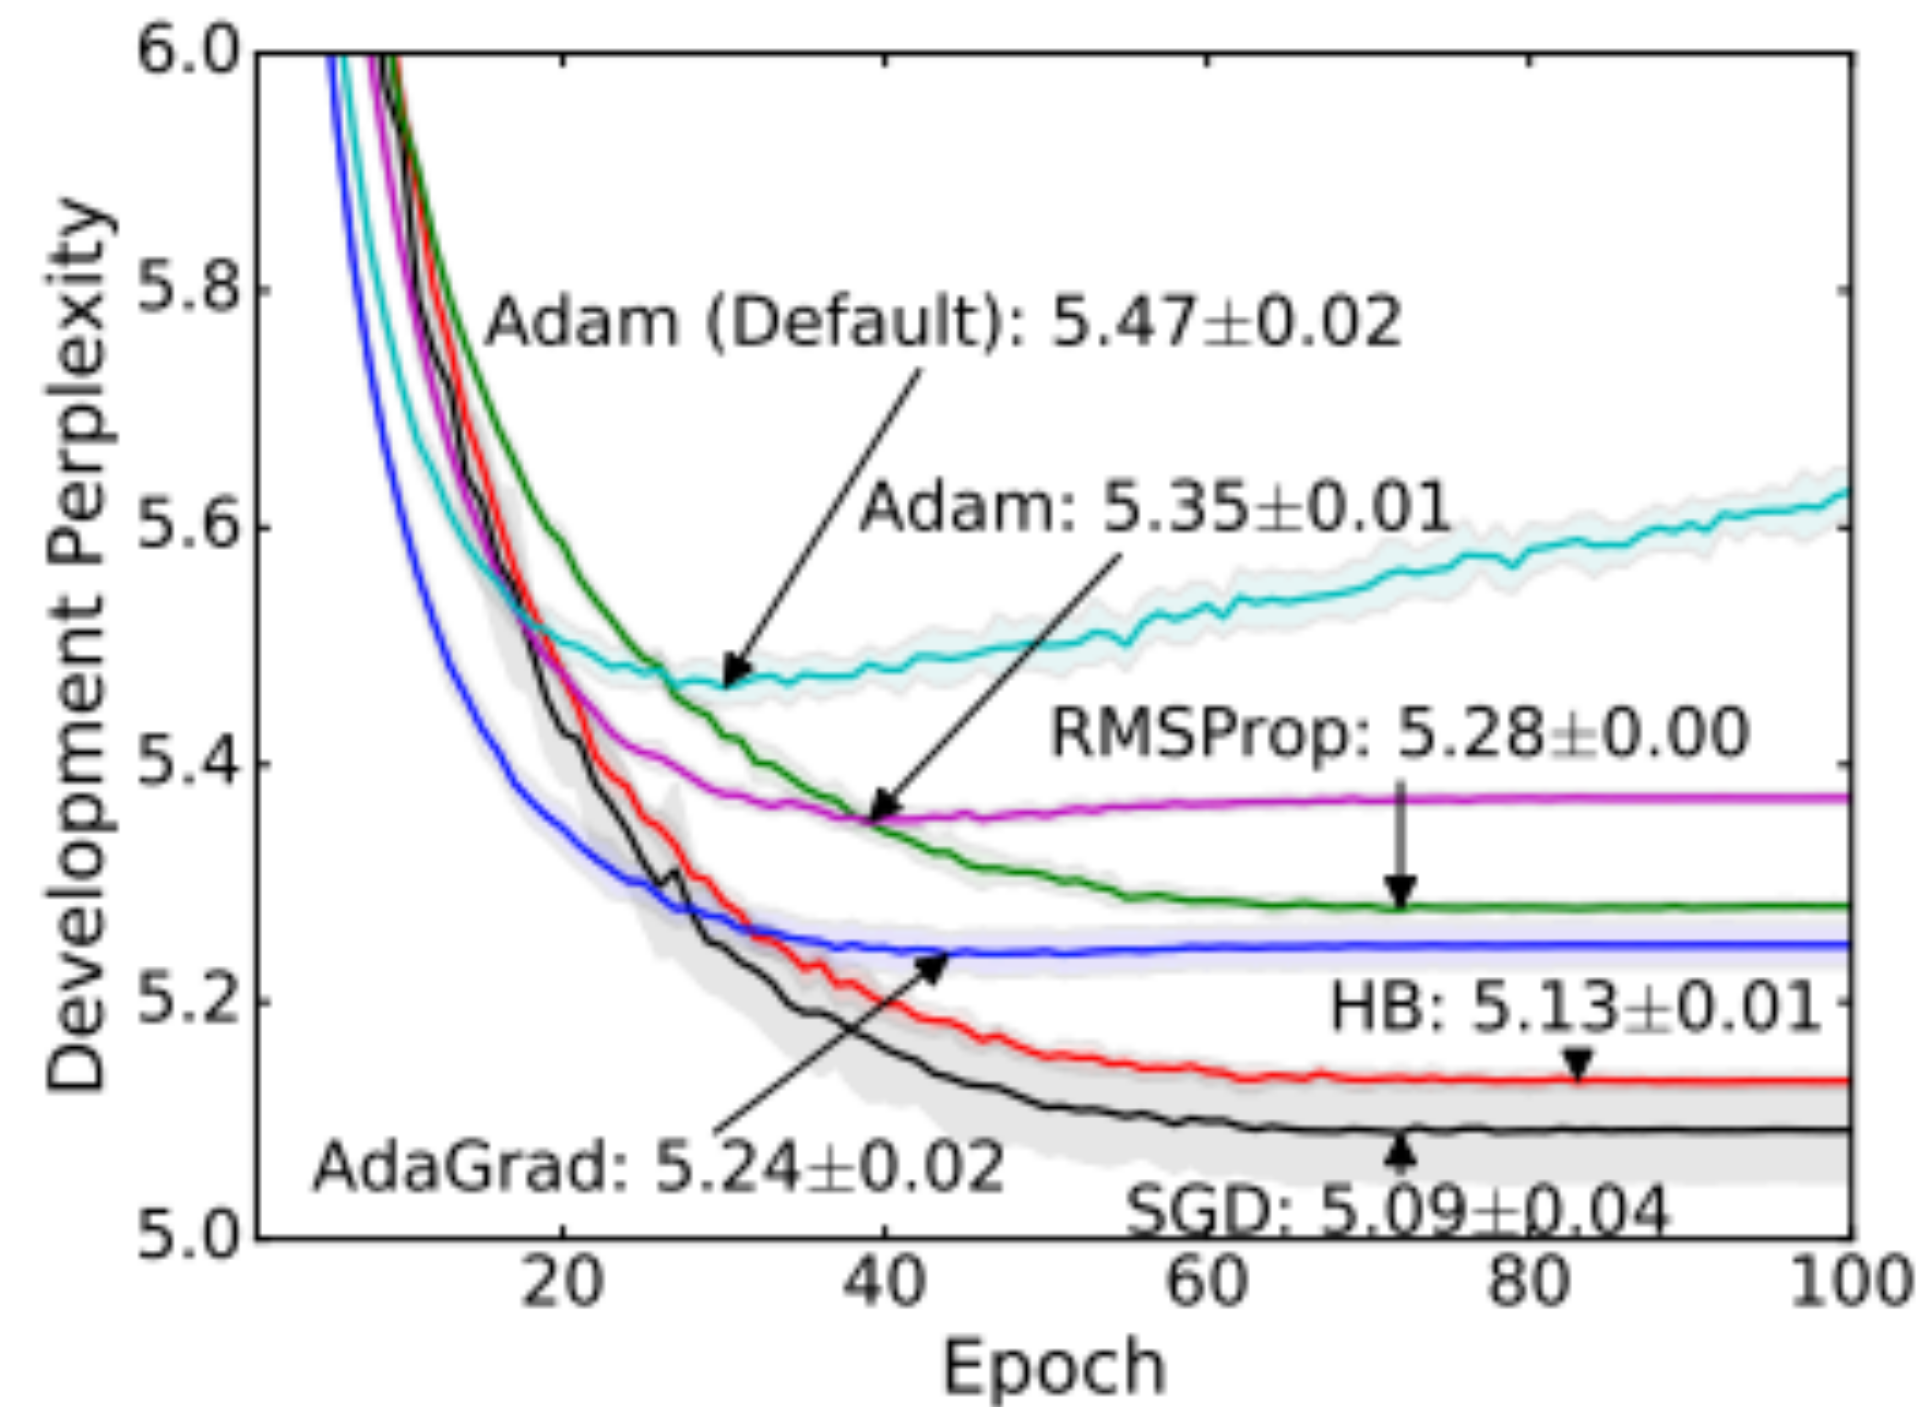


# Optimizer

- Wilson et al. NIPS 2017: adaptive methods can actually perform badly at test time (Adam is in pink, SGD in black)



(e) Generative Parsing (Training Set)

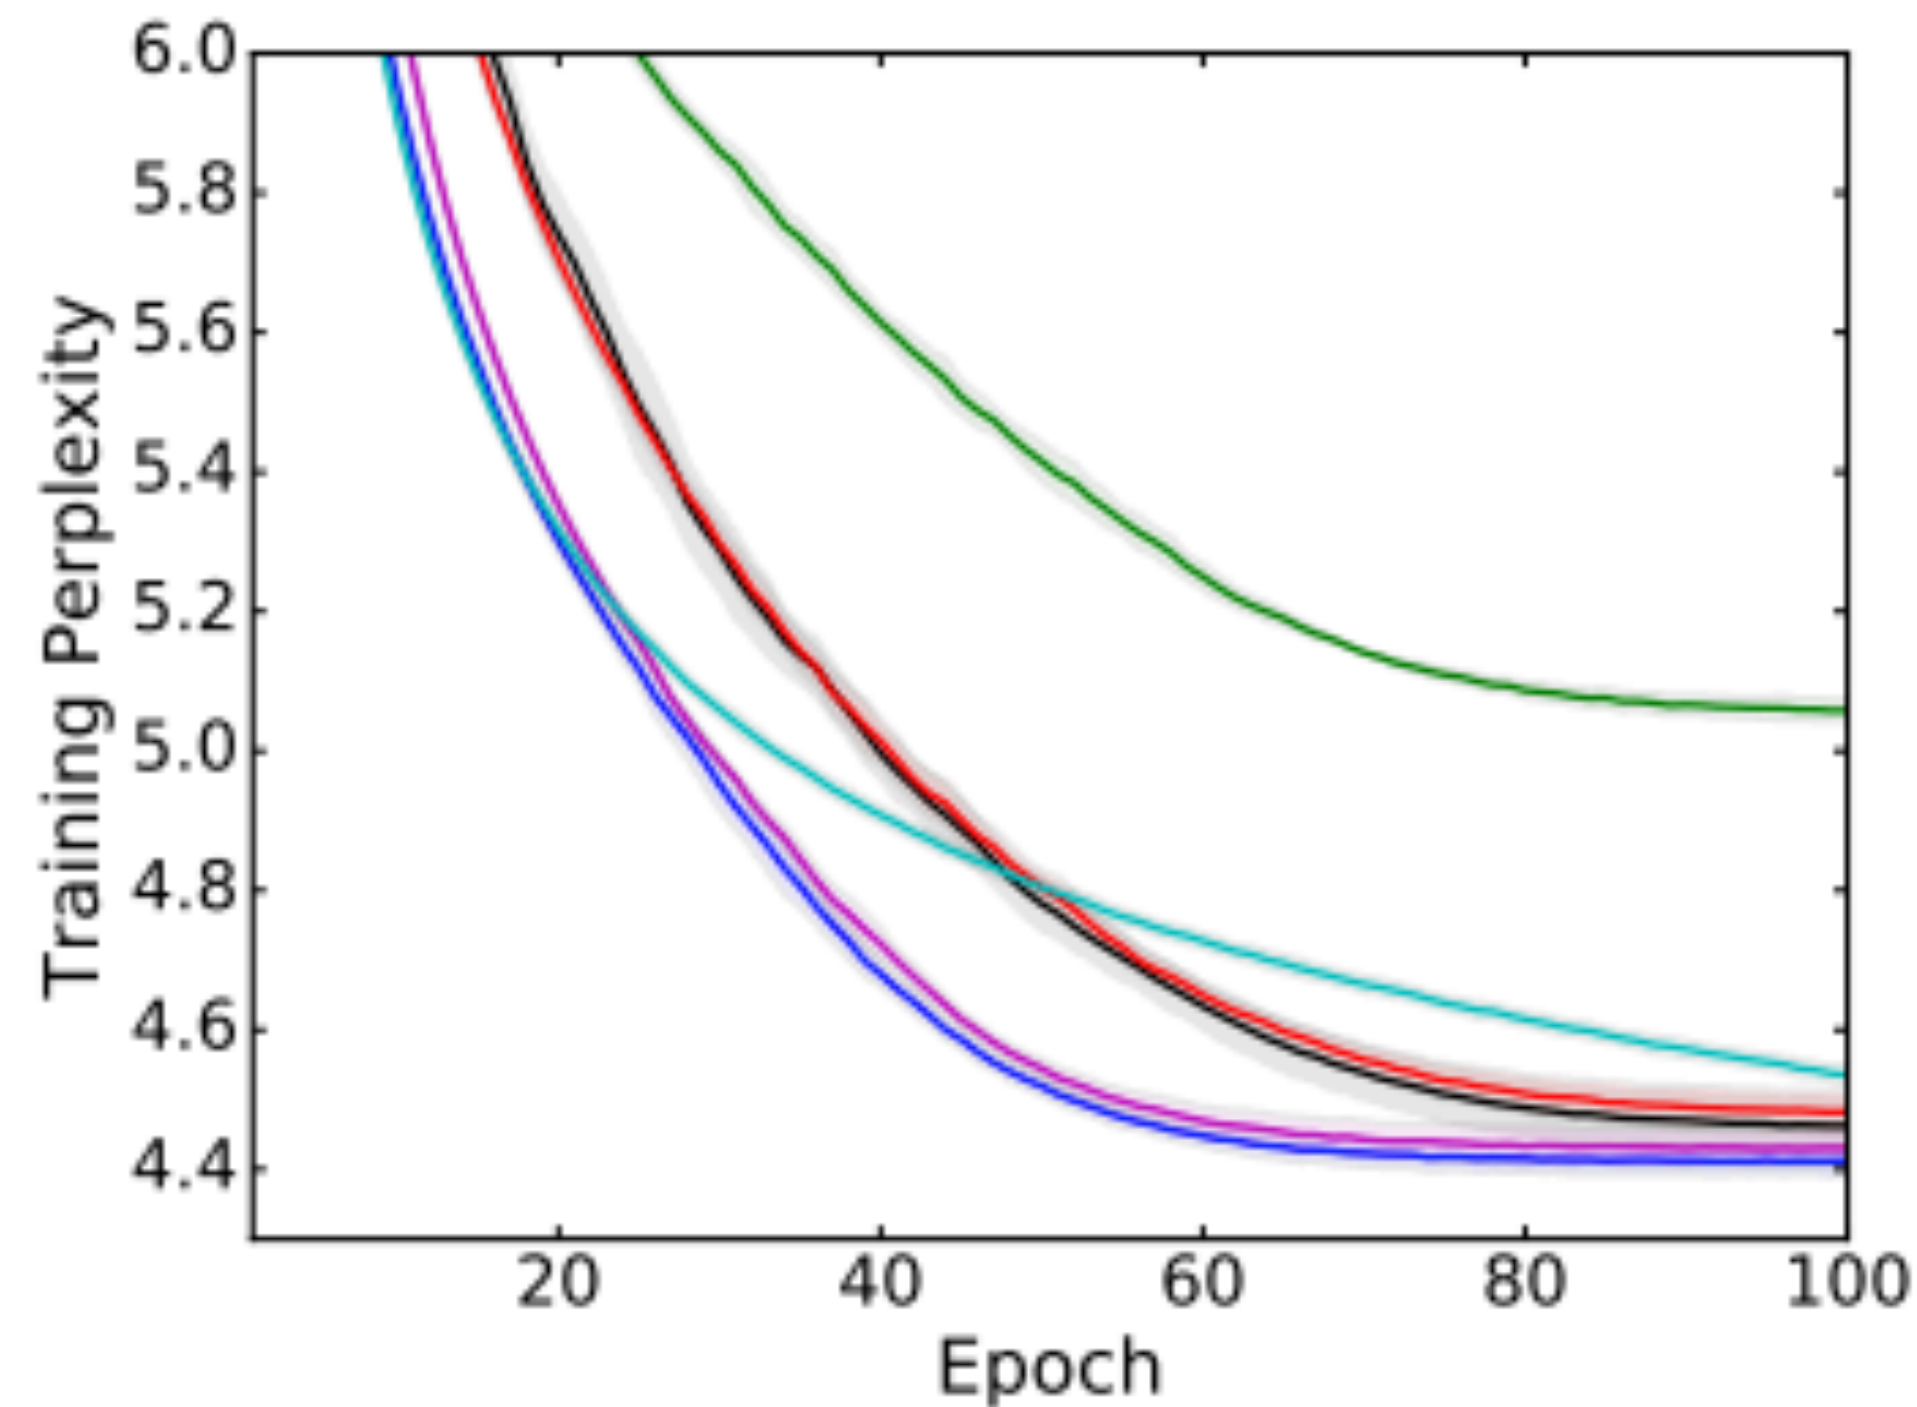


(f) Generative Parsing (Development Set)

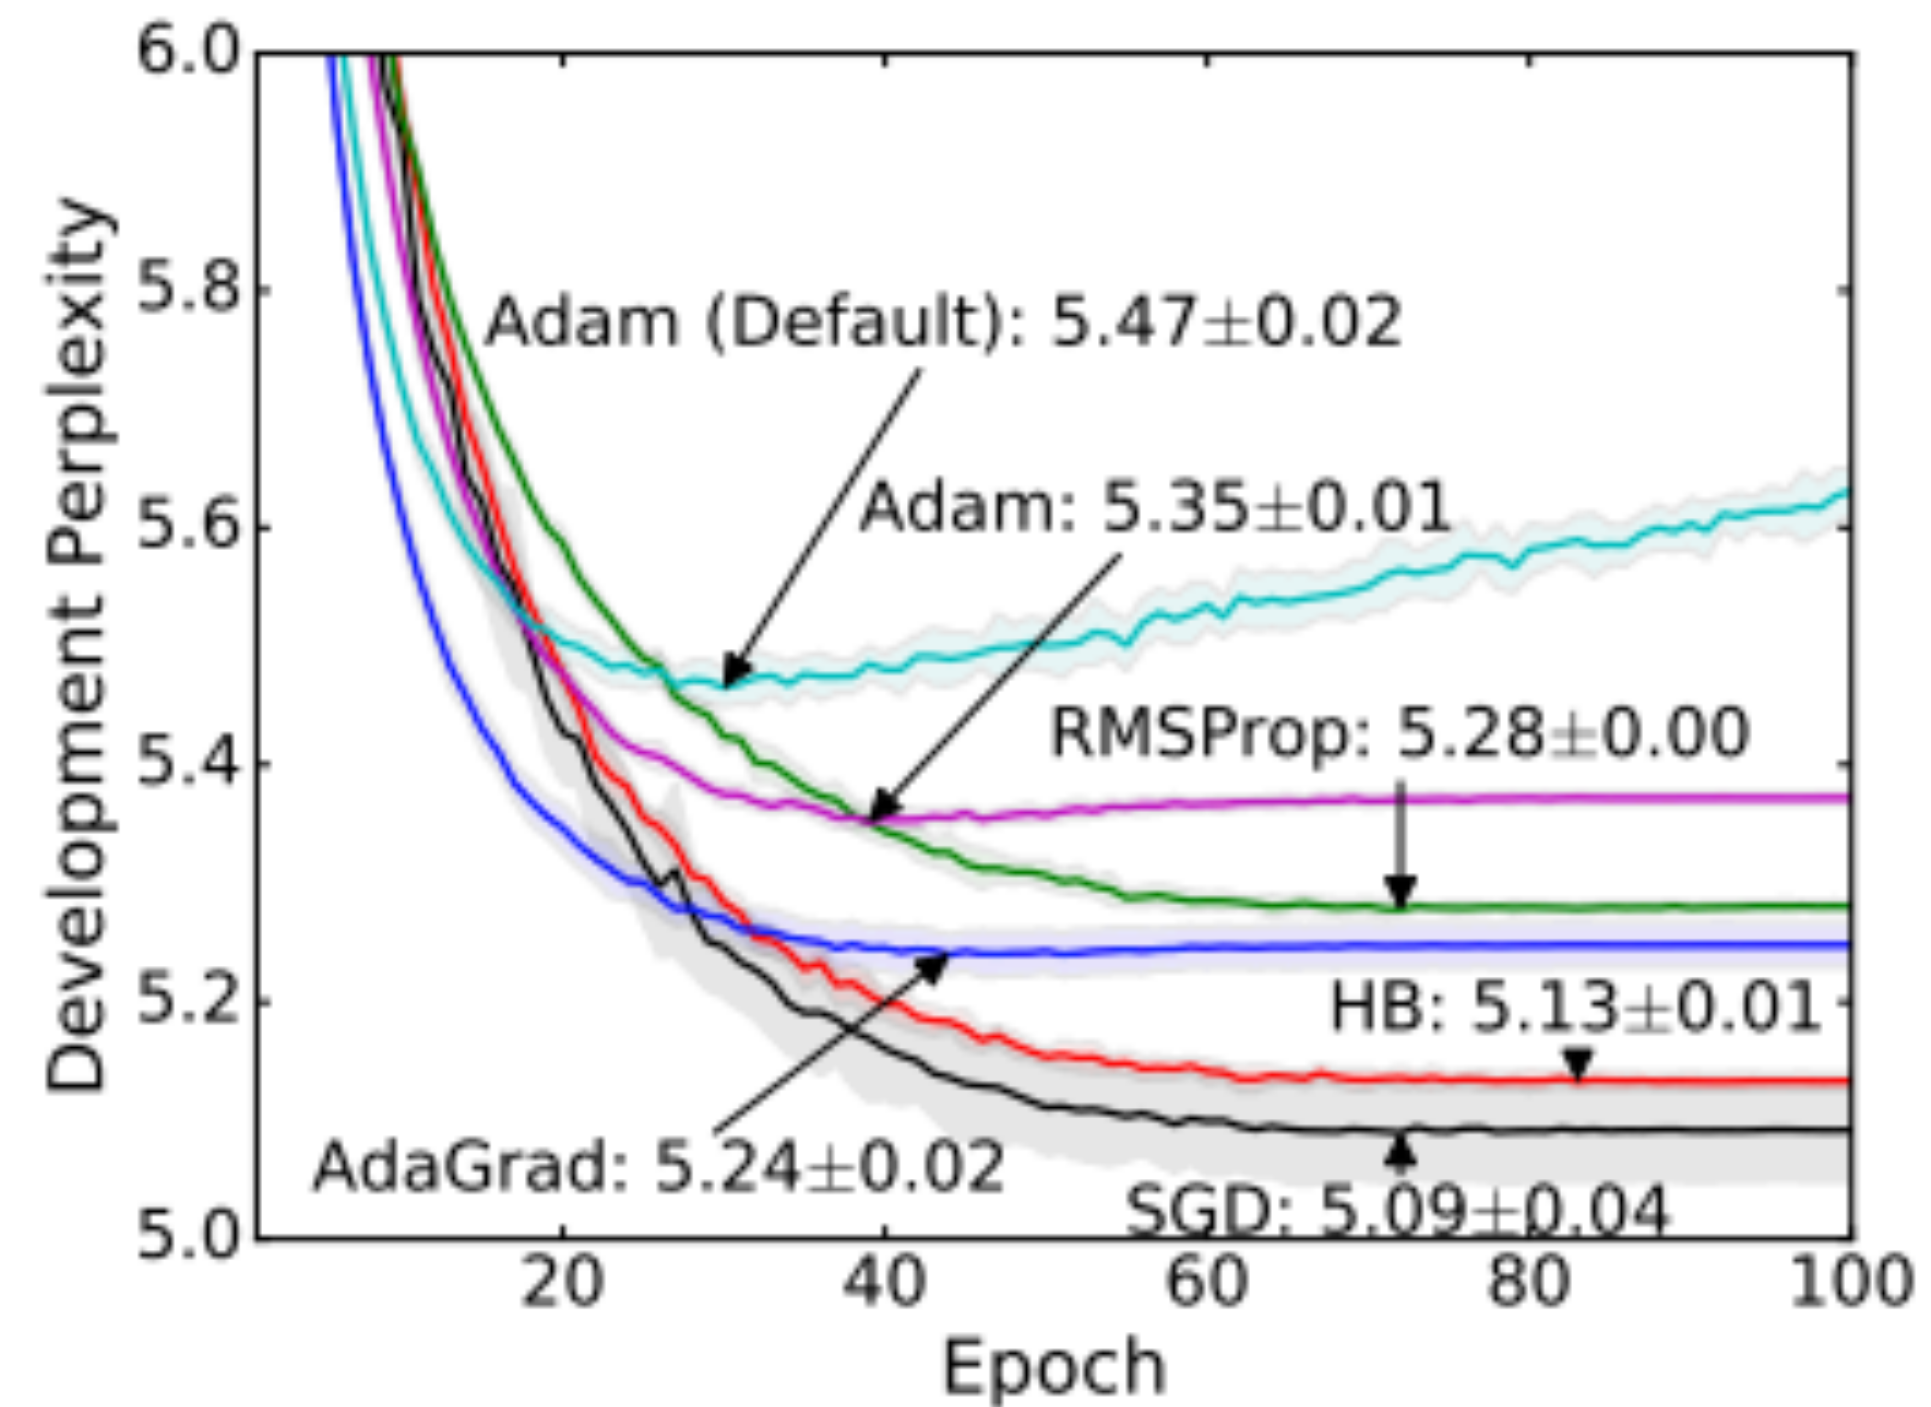


# Optimizer

- ▶ Wilson et al. NIPS 2017: adaptive methods can actually perform badly at test time (Adam is in pink, SGD in black)
- ▶ Check dev set periodically, decrease learning rate if not making progress



(e) Generative Parsing (Training Set)



(f) Generative Parsing (Development Set)

# Structured Prediction

---

- ▶ Four elements of a machine learning method:

# Structured Prediction

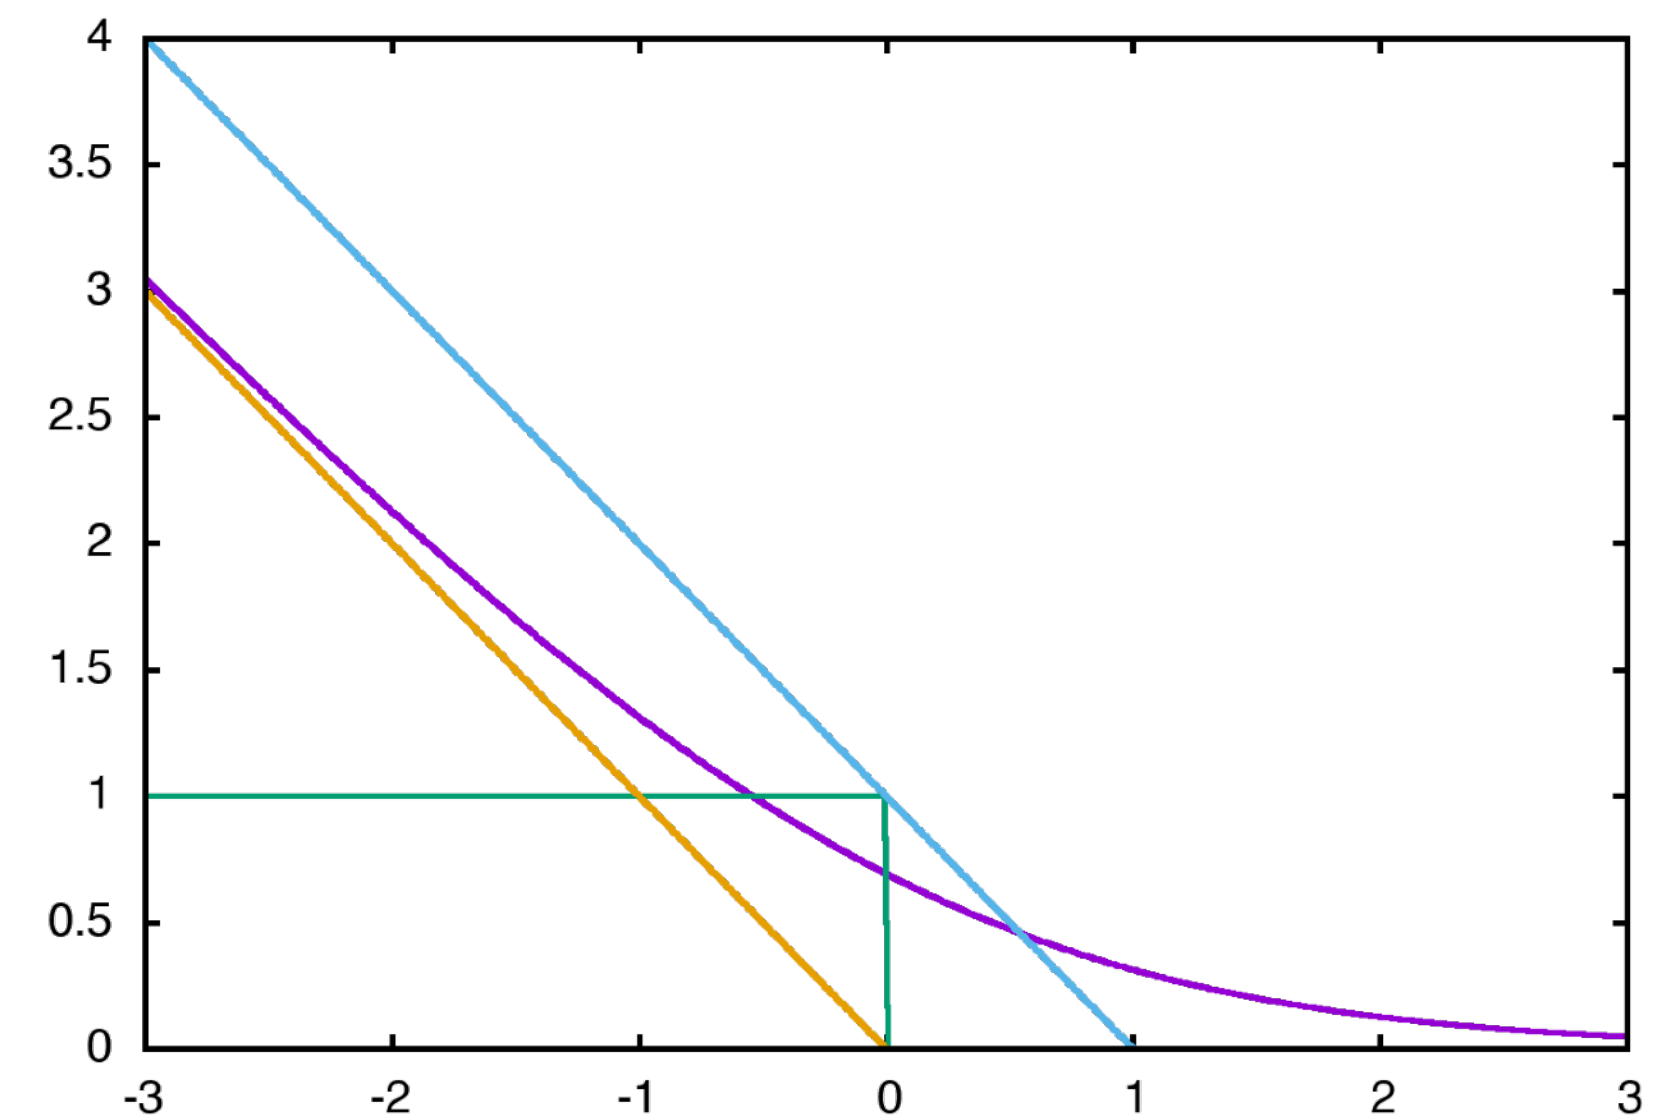
---

- ▶ Four elements of a machine learning method:
- ▶ Model: feedforward, RNNs, CNNs can be defined in a uniform framework

# Structured Prediction

---

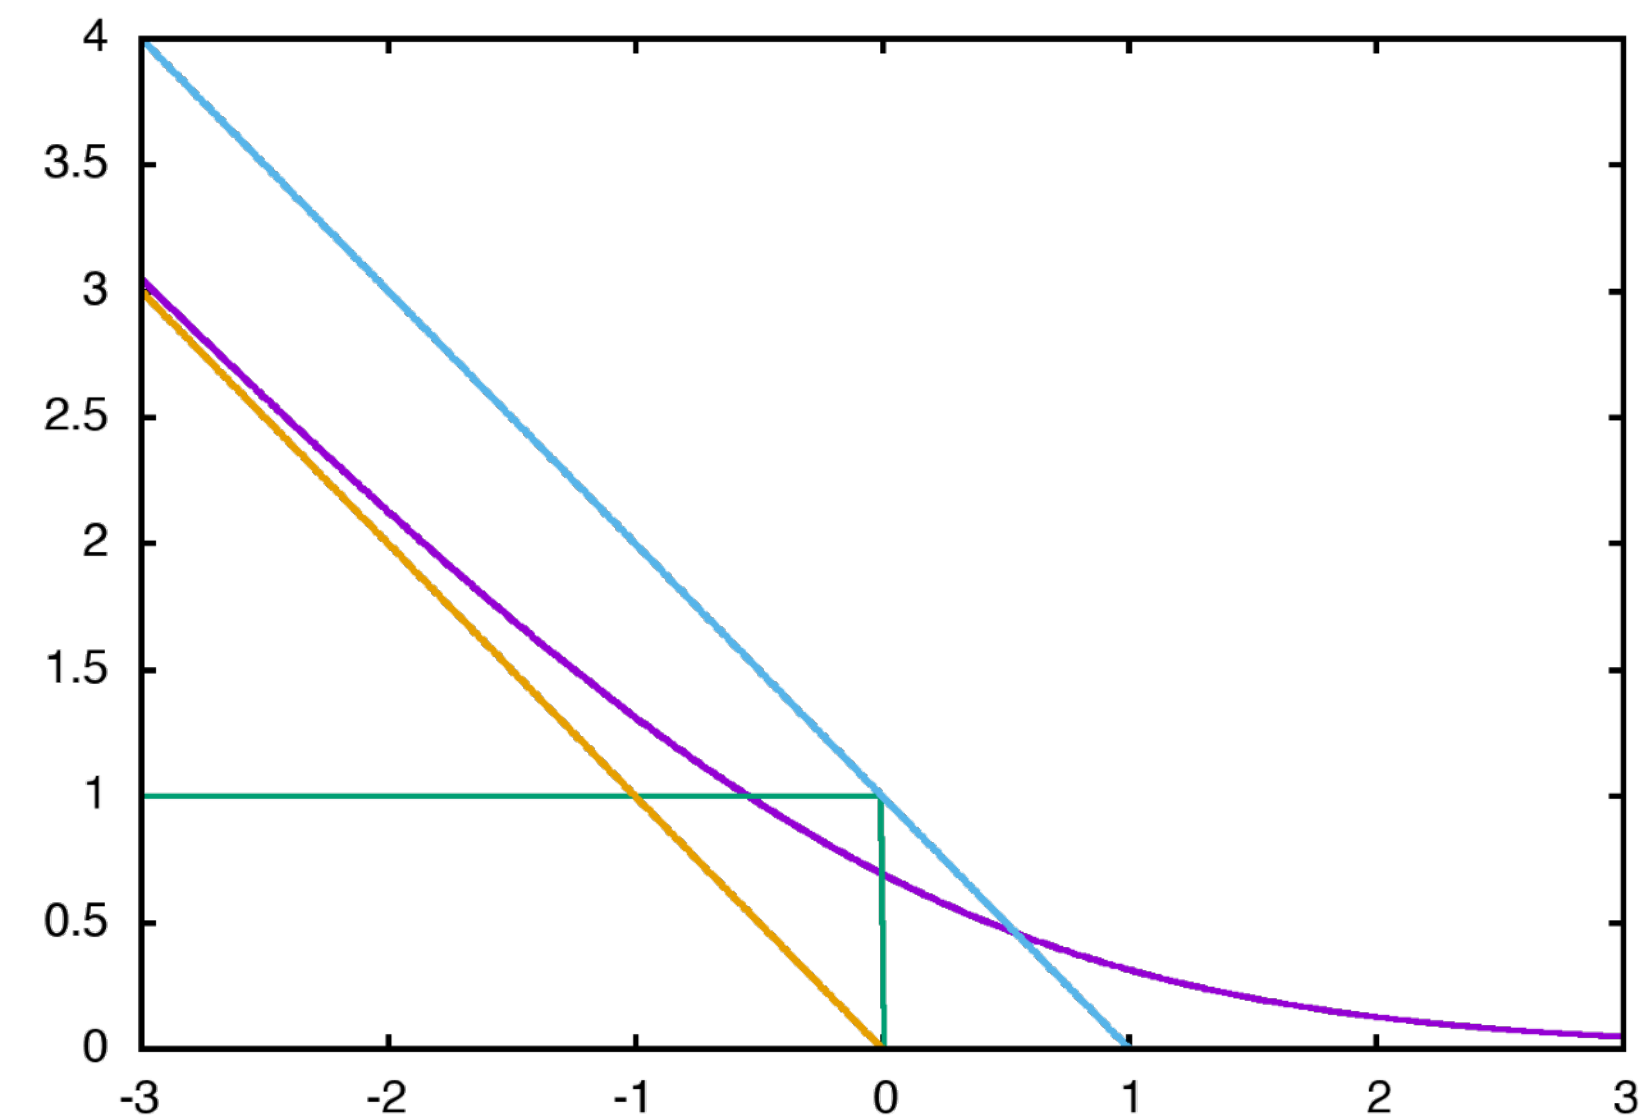
- ▶ Four elements of a machine learning method:
- ▶ Model: feedforward, RNNs, CNNs can be defined in a uniform framework
- ▶ Objective: many loss functions look similar, just changes the last layer of the neural network



# Structured Prediction

---

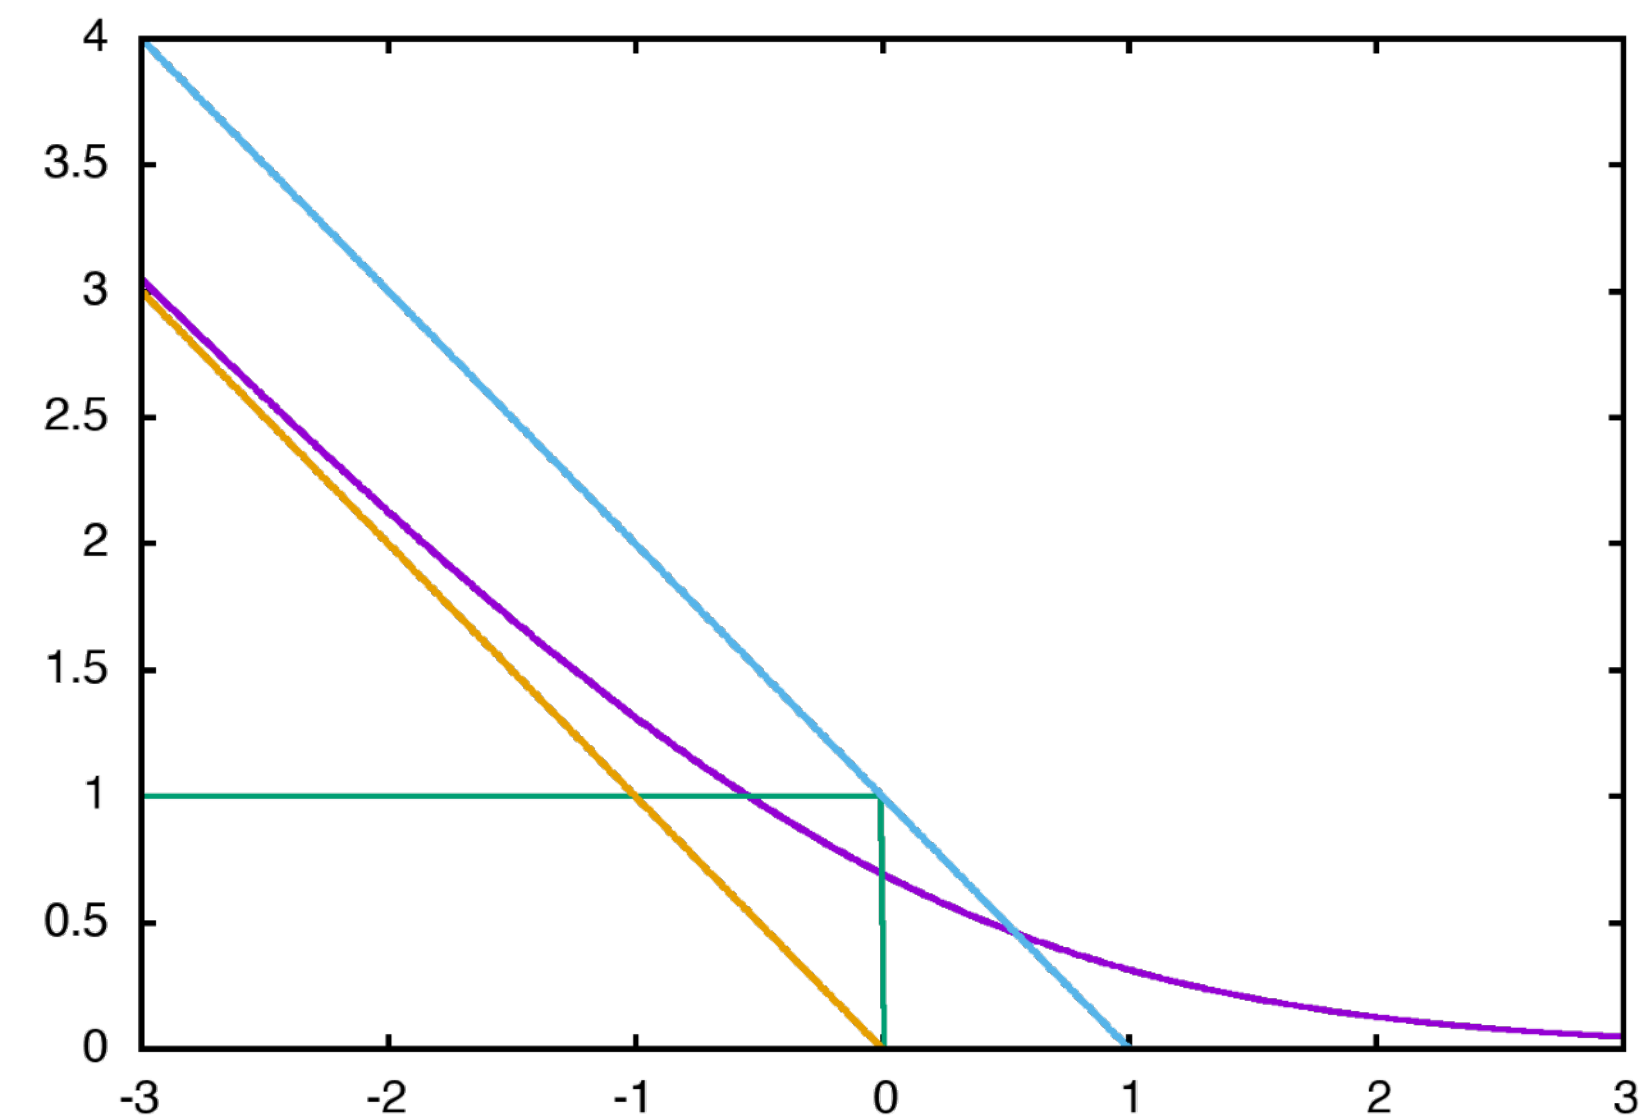
- ▶ Four elements of a machine learning method:
- ▶ Model: feedforward, RNNs, CNNs can be defined in a uniform framework
- ▶ Objective: many loss functions look similar, just changes the last layer of the neural network
- ▶ Inference: define the network, your library of choice takes care of it (mostly...)



# Structured Prediction

---

- ▶ Four elements of a machine learning method:
- ▶ Model: feedforward, RNNs, CNNs can be defined in a uniform framework
- ▶ Objective: many loss functions look similar, just changes the last layer of the neural network
- ▶ Inference: define the network, your library of choice takes care of it (mostly...)
- ▶ Training: lots of choices for optimization/hyperparameters



# Word Representations

# Word Representations

---

- ▶ Neural networks work very well at continuous data, but words are discrete



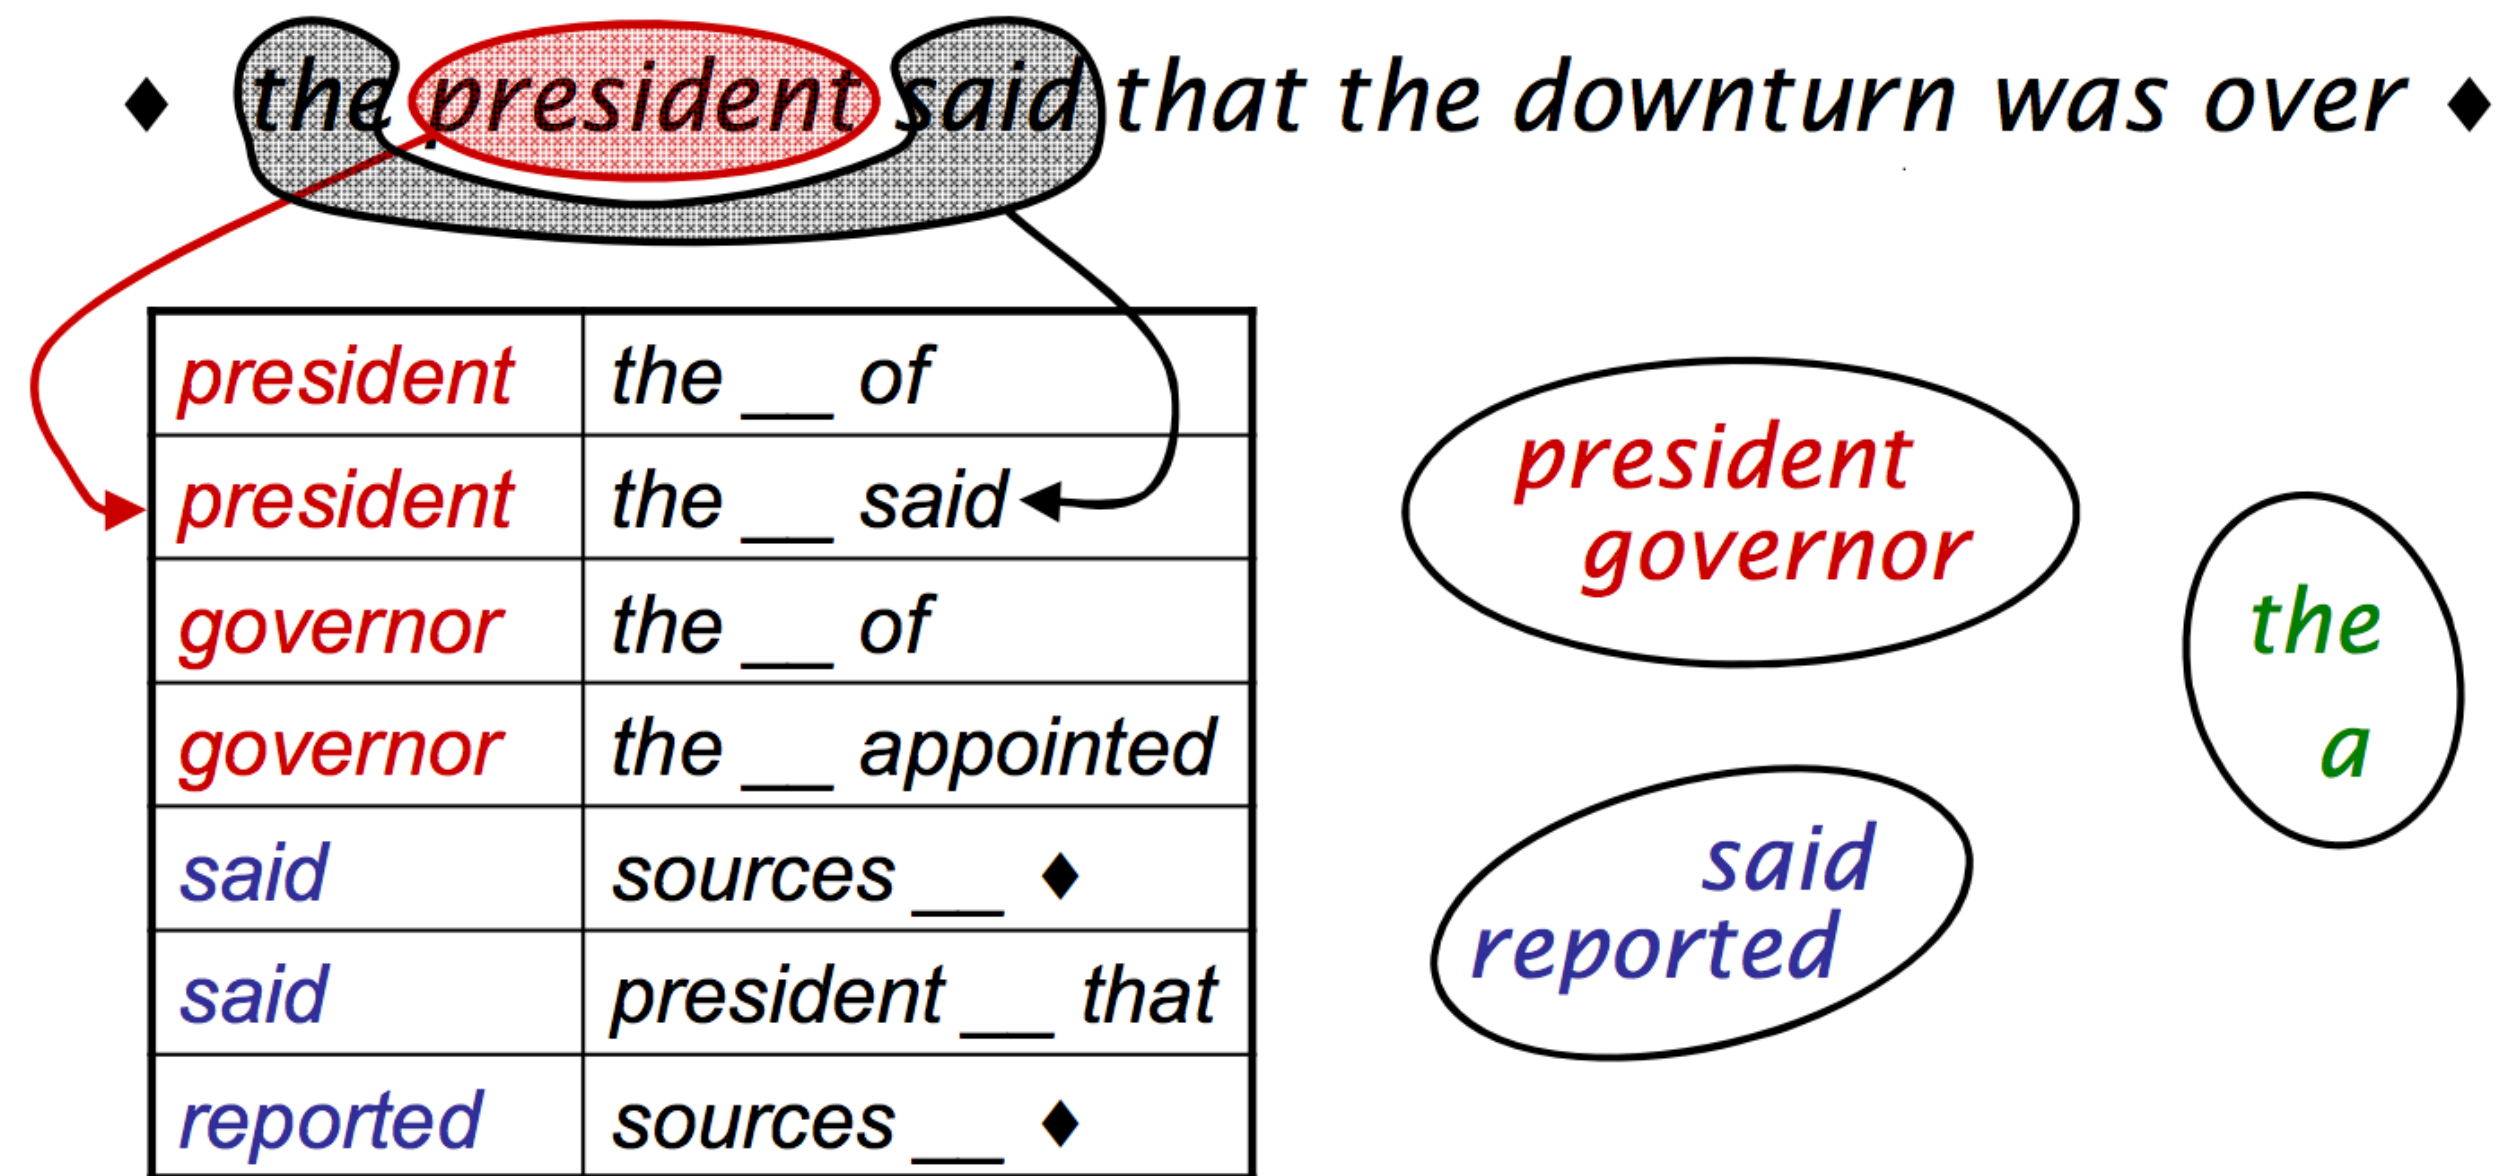
# Word Representations

---

- ▶ Neural networks work very well at continuous data, but words are discrete
- ▶ Continuous model  $\leftrightarrow$  expects continuous semantics from input

# Word Representations

- ▶ Neural networks work very well at continuous data, but words are discrete
- ▶ Continuous model  $\leftrightarrow$  expects continuous semantics from input
- ▶ “You shall know a word by the company it keeps” Firth (1957)



# Discrete Word Representations

---

# Discrete Word Representations

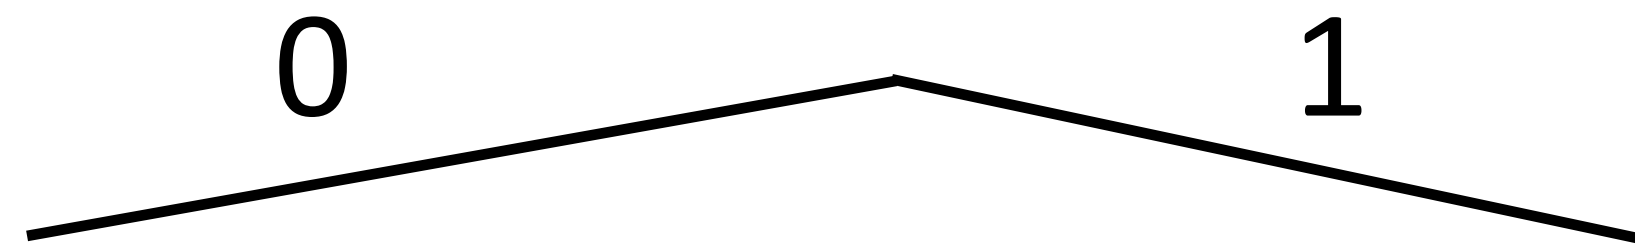
---

- ▶ Brown clusters: hierarchical agglomerative *hard* clustering (each word has one cluster, not some posterior distribution like in mixture models)

# Discrete Word Representations

---

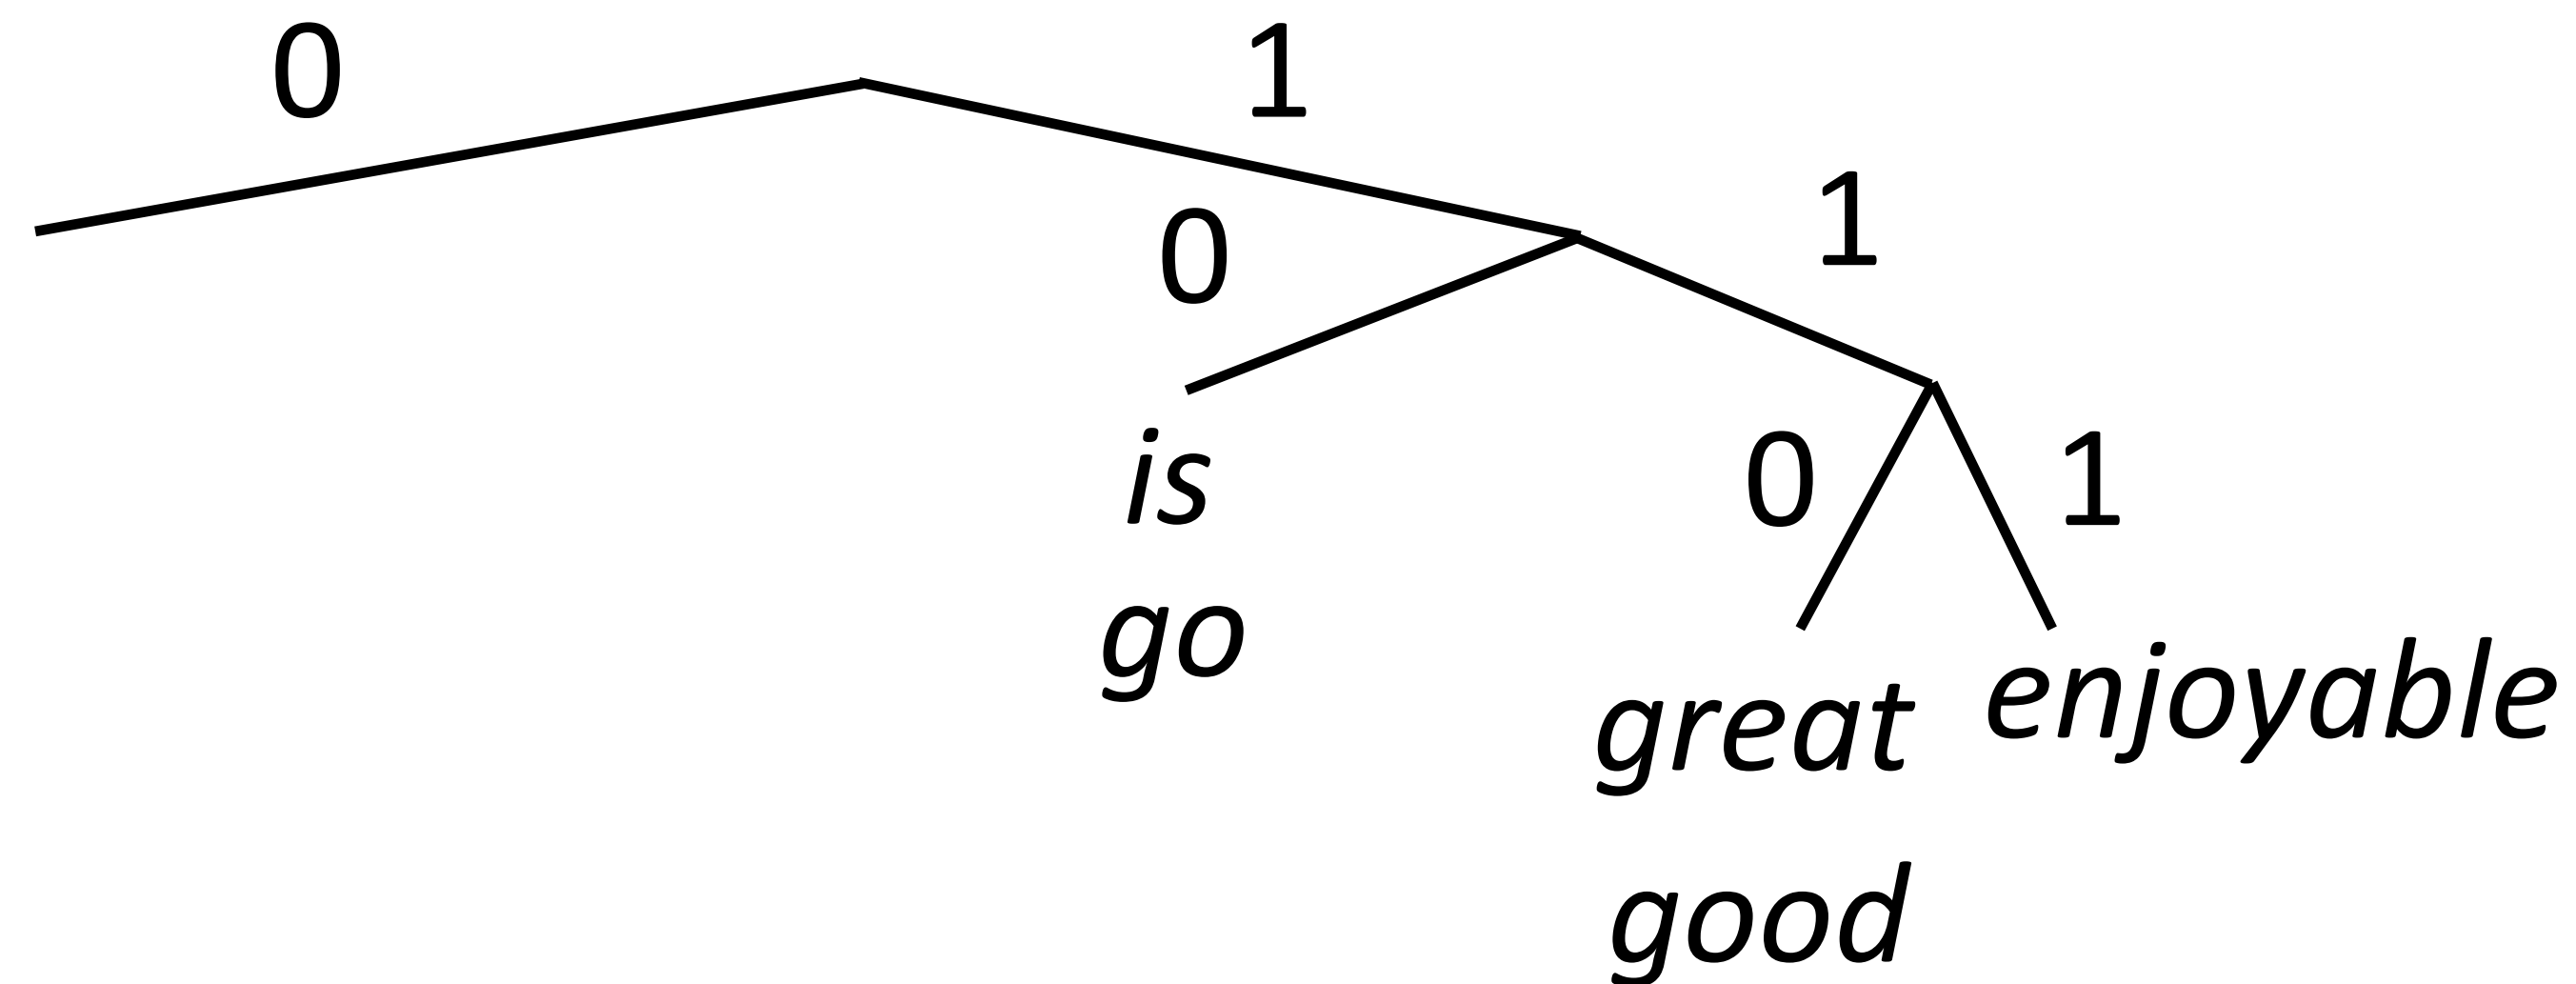
- ▶ Brown clusters: hierarchical agglomerative *hard* clustering (each word has one cluster, not some posterior distribution like in mixture models)



# Discrete Word Representations

---

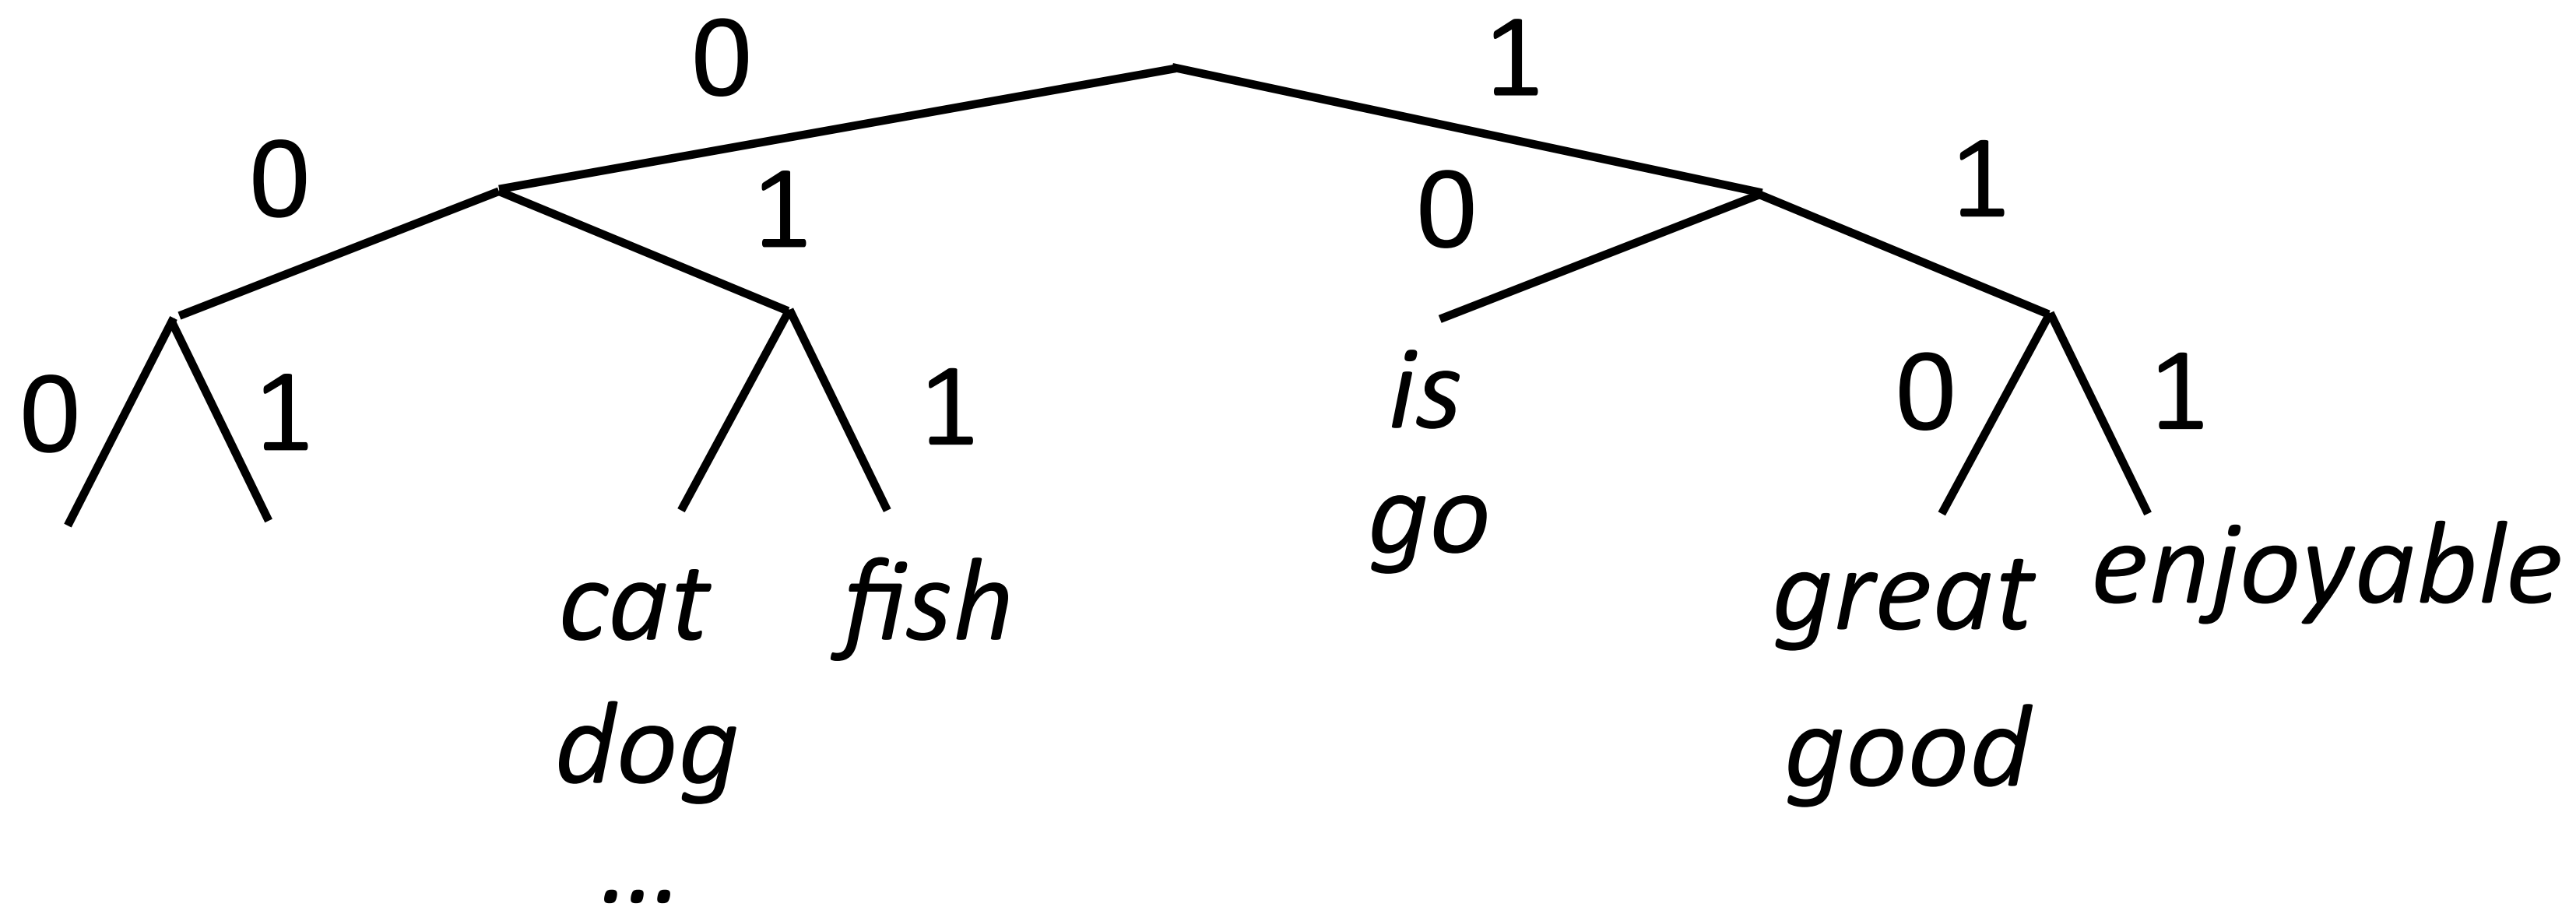
- ▶ Brown clusters: hierarchical agglomerative *hard* clustering (each word has one cluster, not some posterior distribution like in mixture models)



# Discrete Word Representations

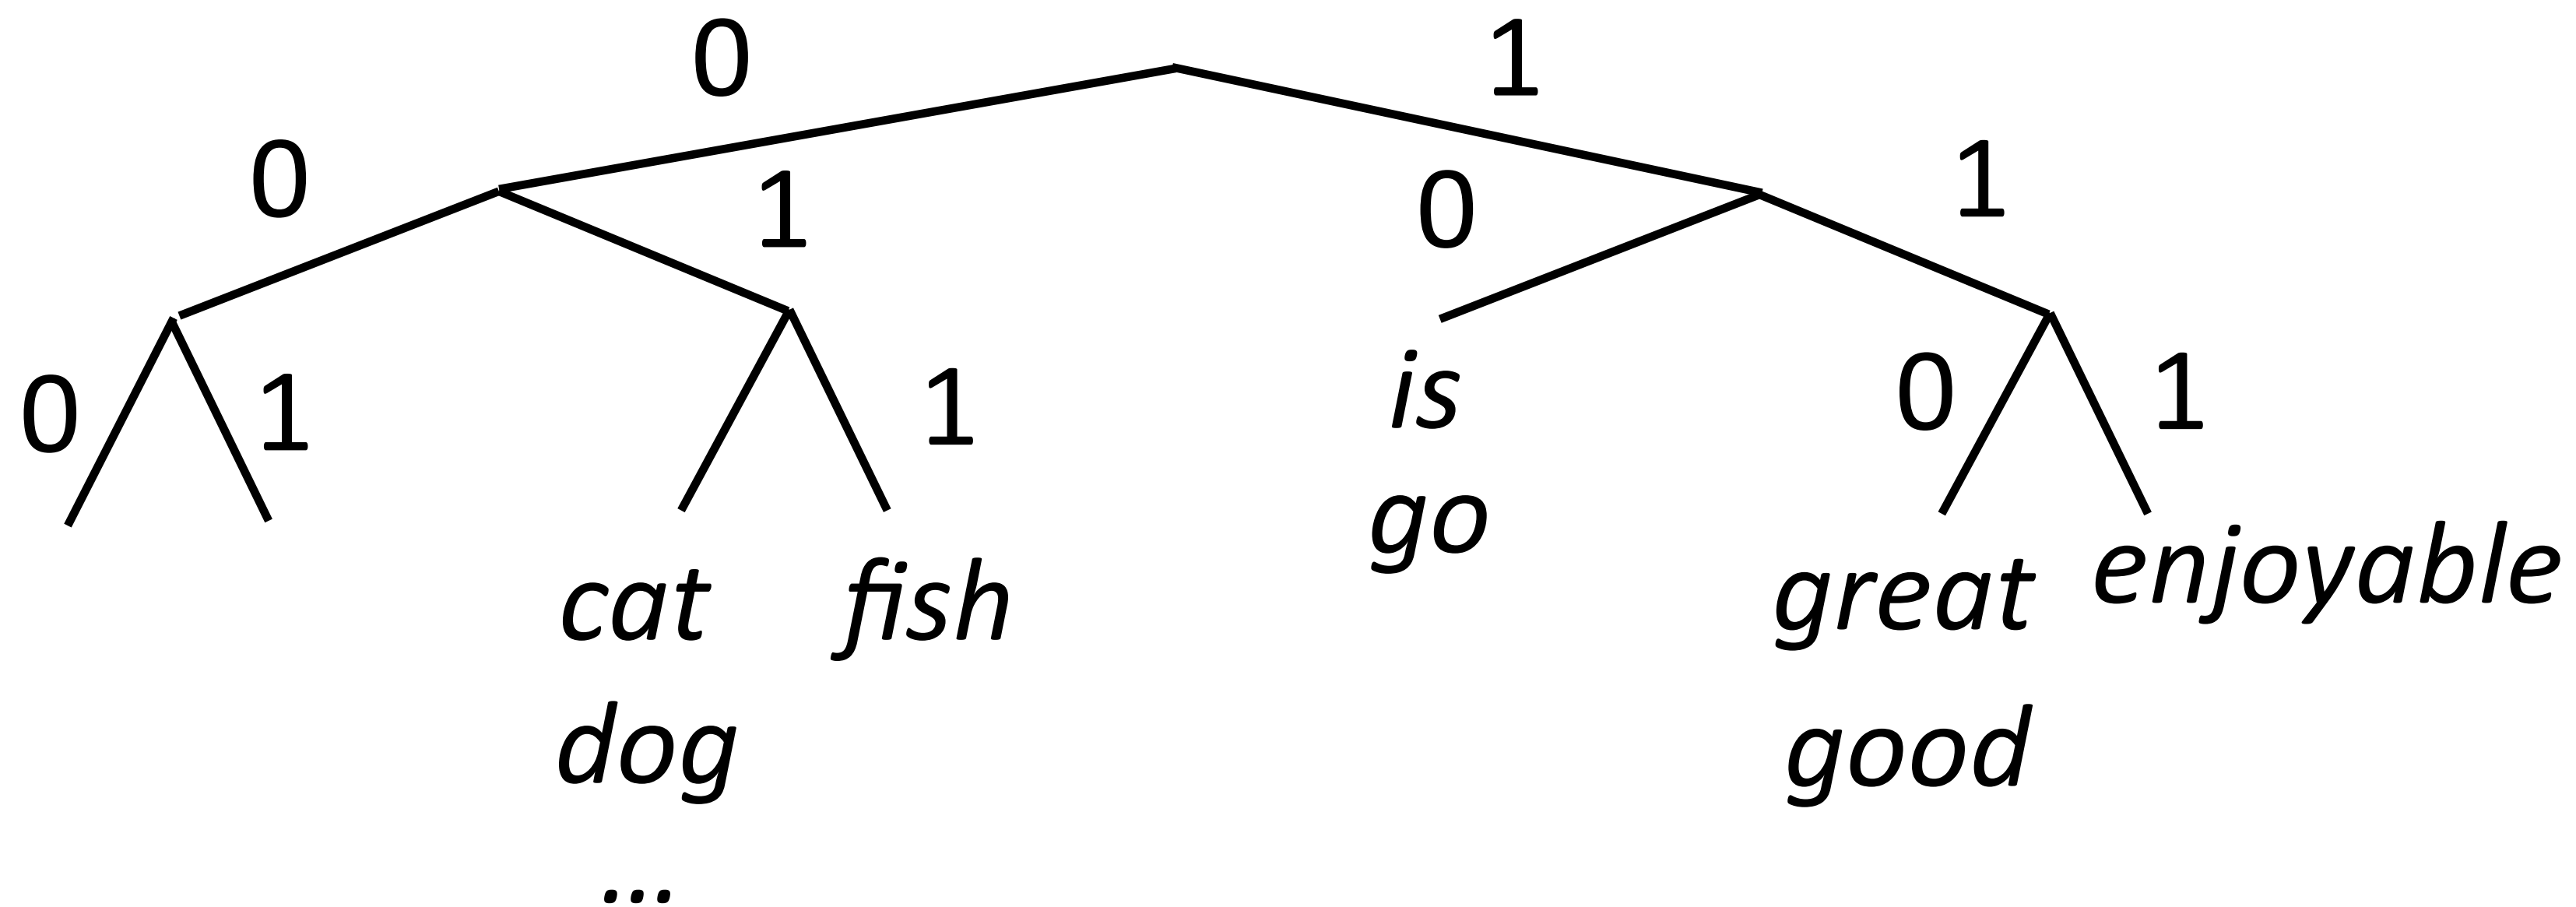
---

- ▶ Brown clusters: hierarchical agglomerative *hard* clustering (each word has one cluster, not some posterior distribution like in mixture models)



# Discrete Word Representations

- ▶ Brown clusters: hierarchical agglomerative *hard* clustering (each word has one cluster, not some posterior distribution like in mixture models)

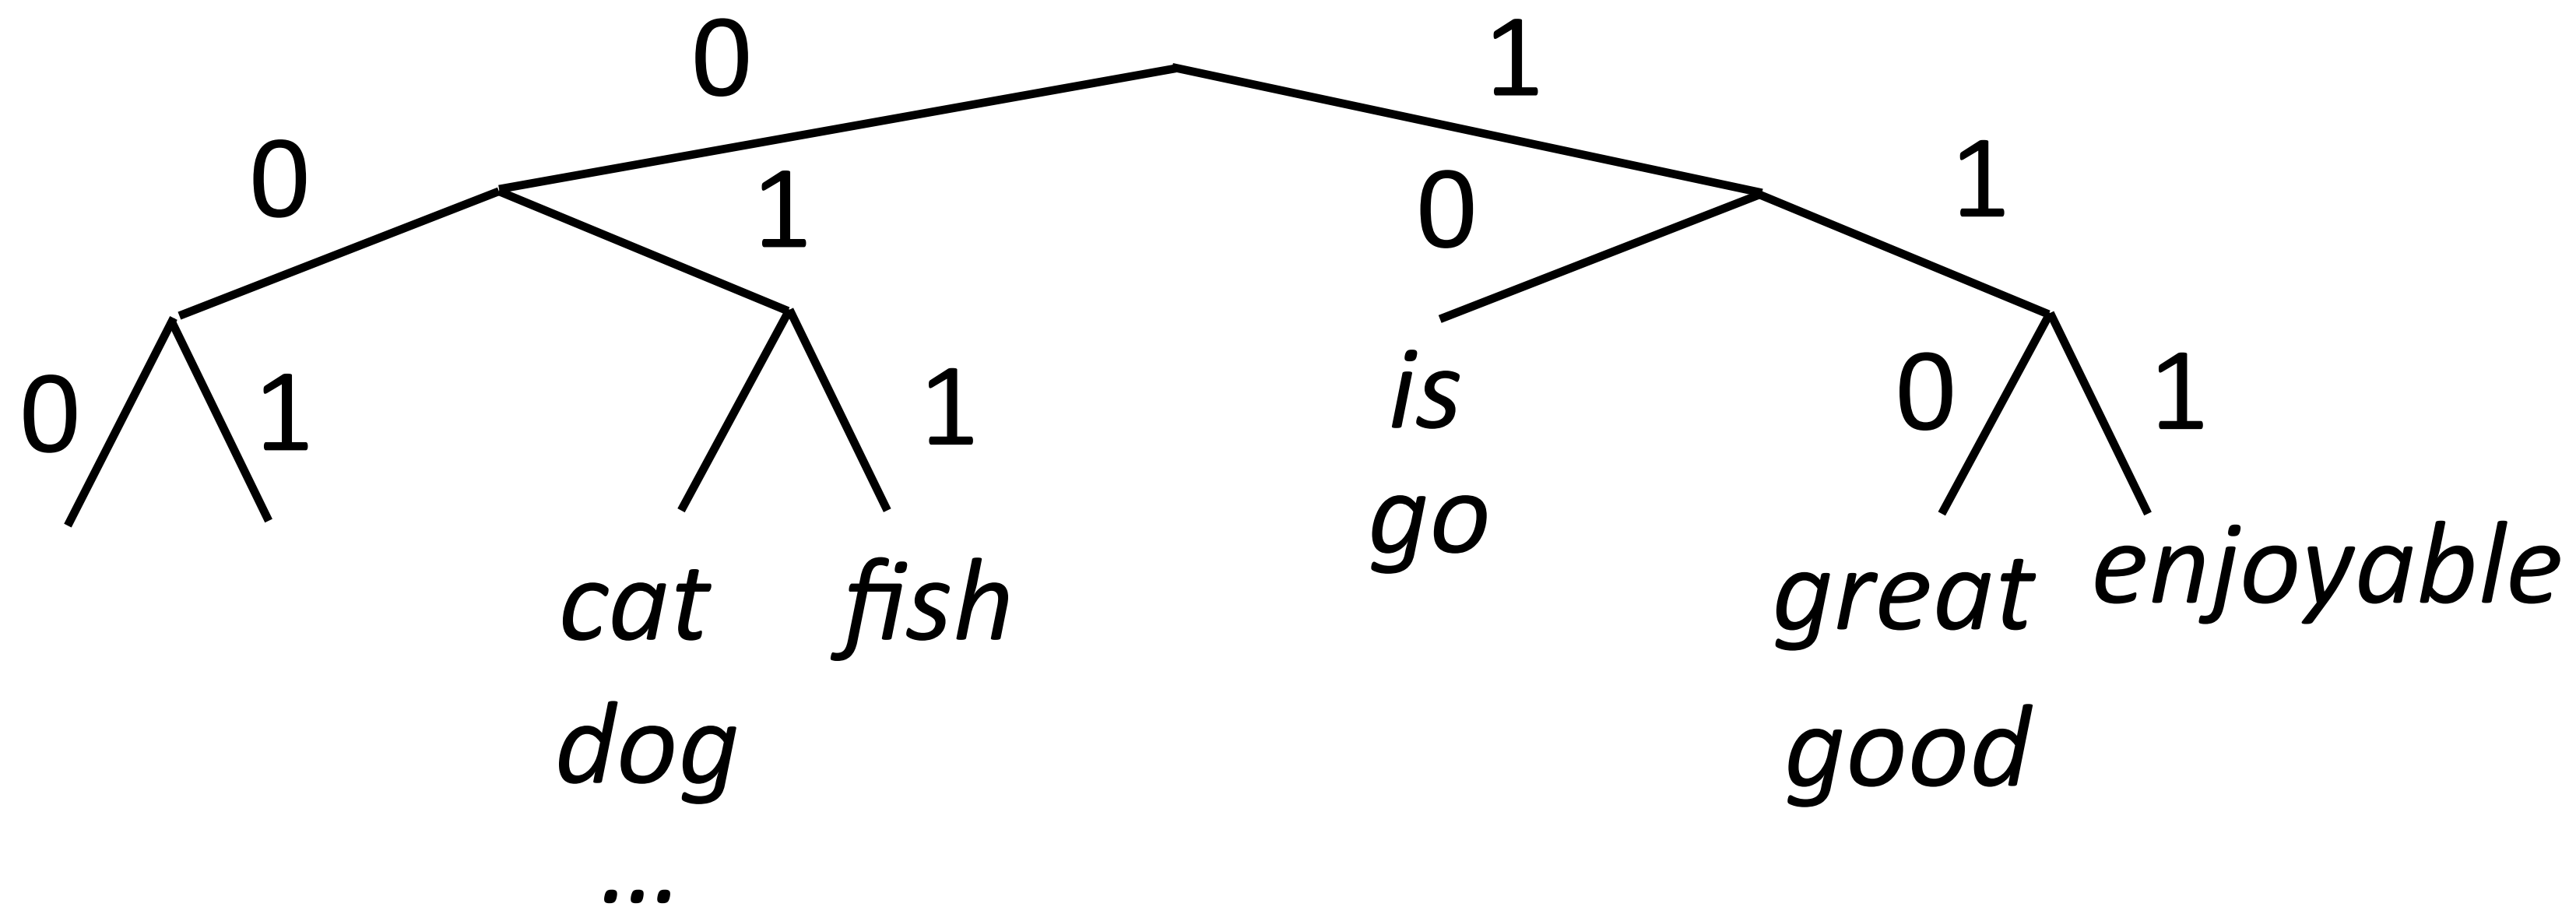


- ▶ Maximize  $P(w_i|w_{i-1}) = P(c_i|c_{i-1})P(w_i|c_i)$



# Discrete Word Representations

- ▶ Brown clusters: hierarchical agglomerative *hard* clustering (each word has one cluster, not some posterior distribution like in mixture models)



- ▶ Maximize  $P(w_i|w_{i-1}) = P(c_i|c_{i-1})P(w_i|c_i)$
  - ▶ Useful features for tasks like NER, not suitable for NNs
- Brown et al. (1992)

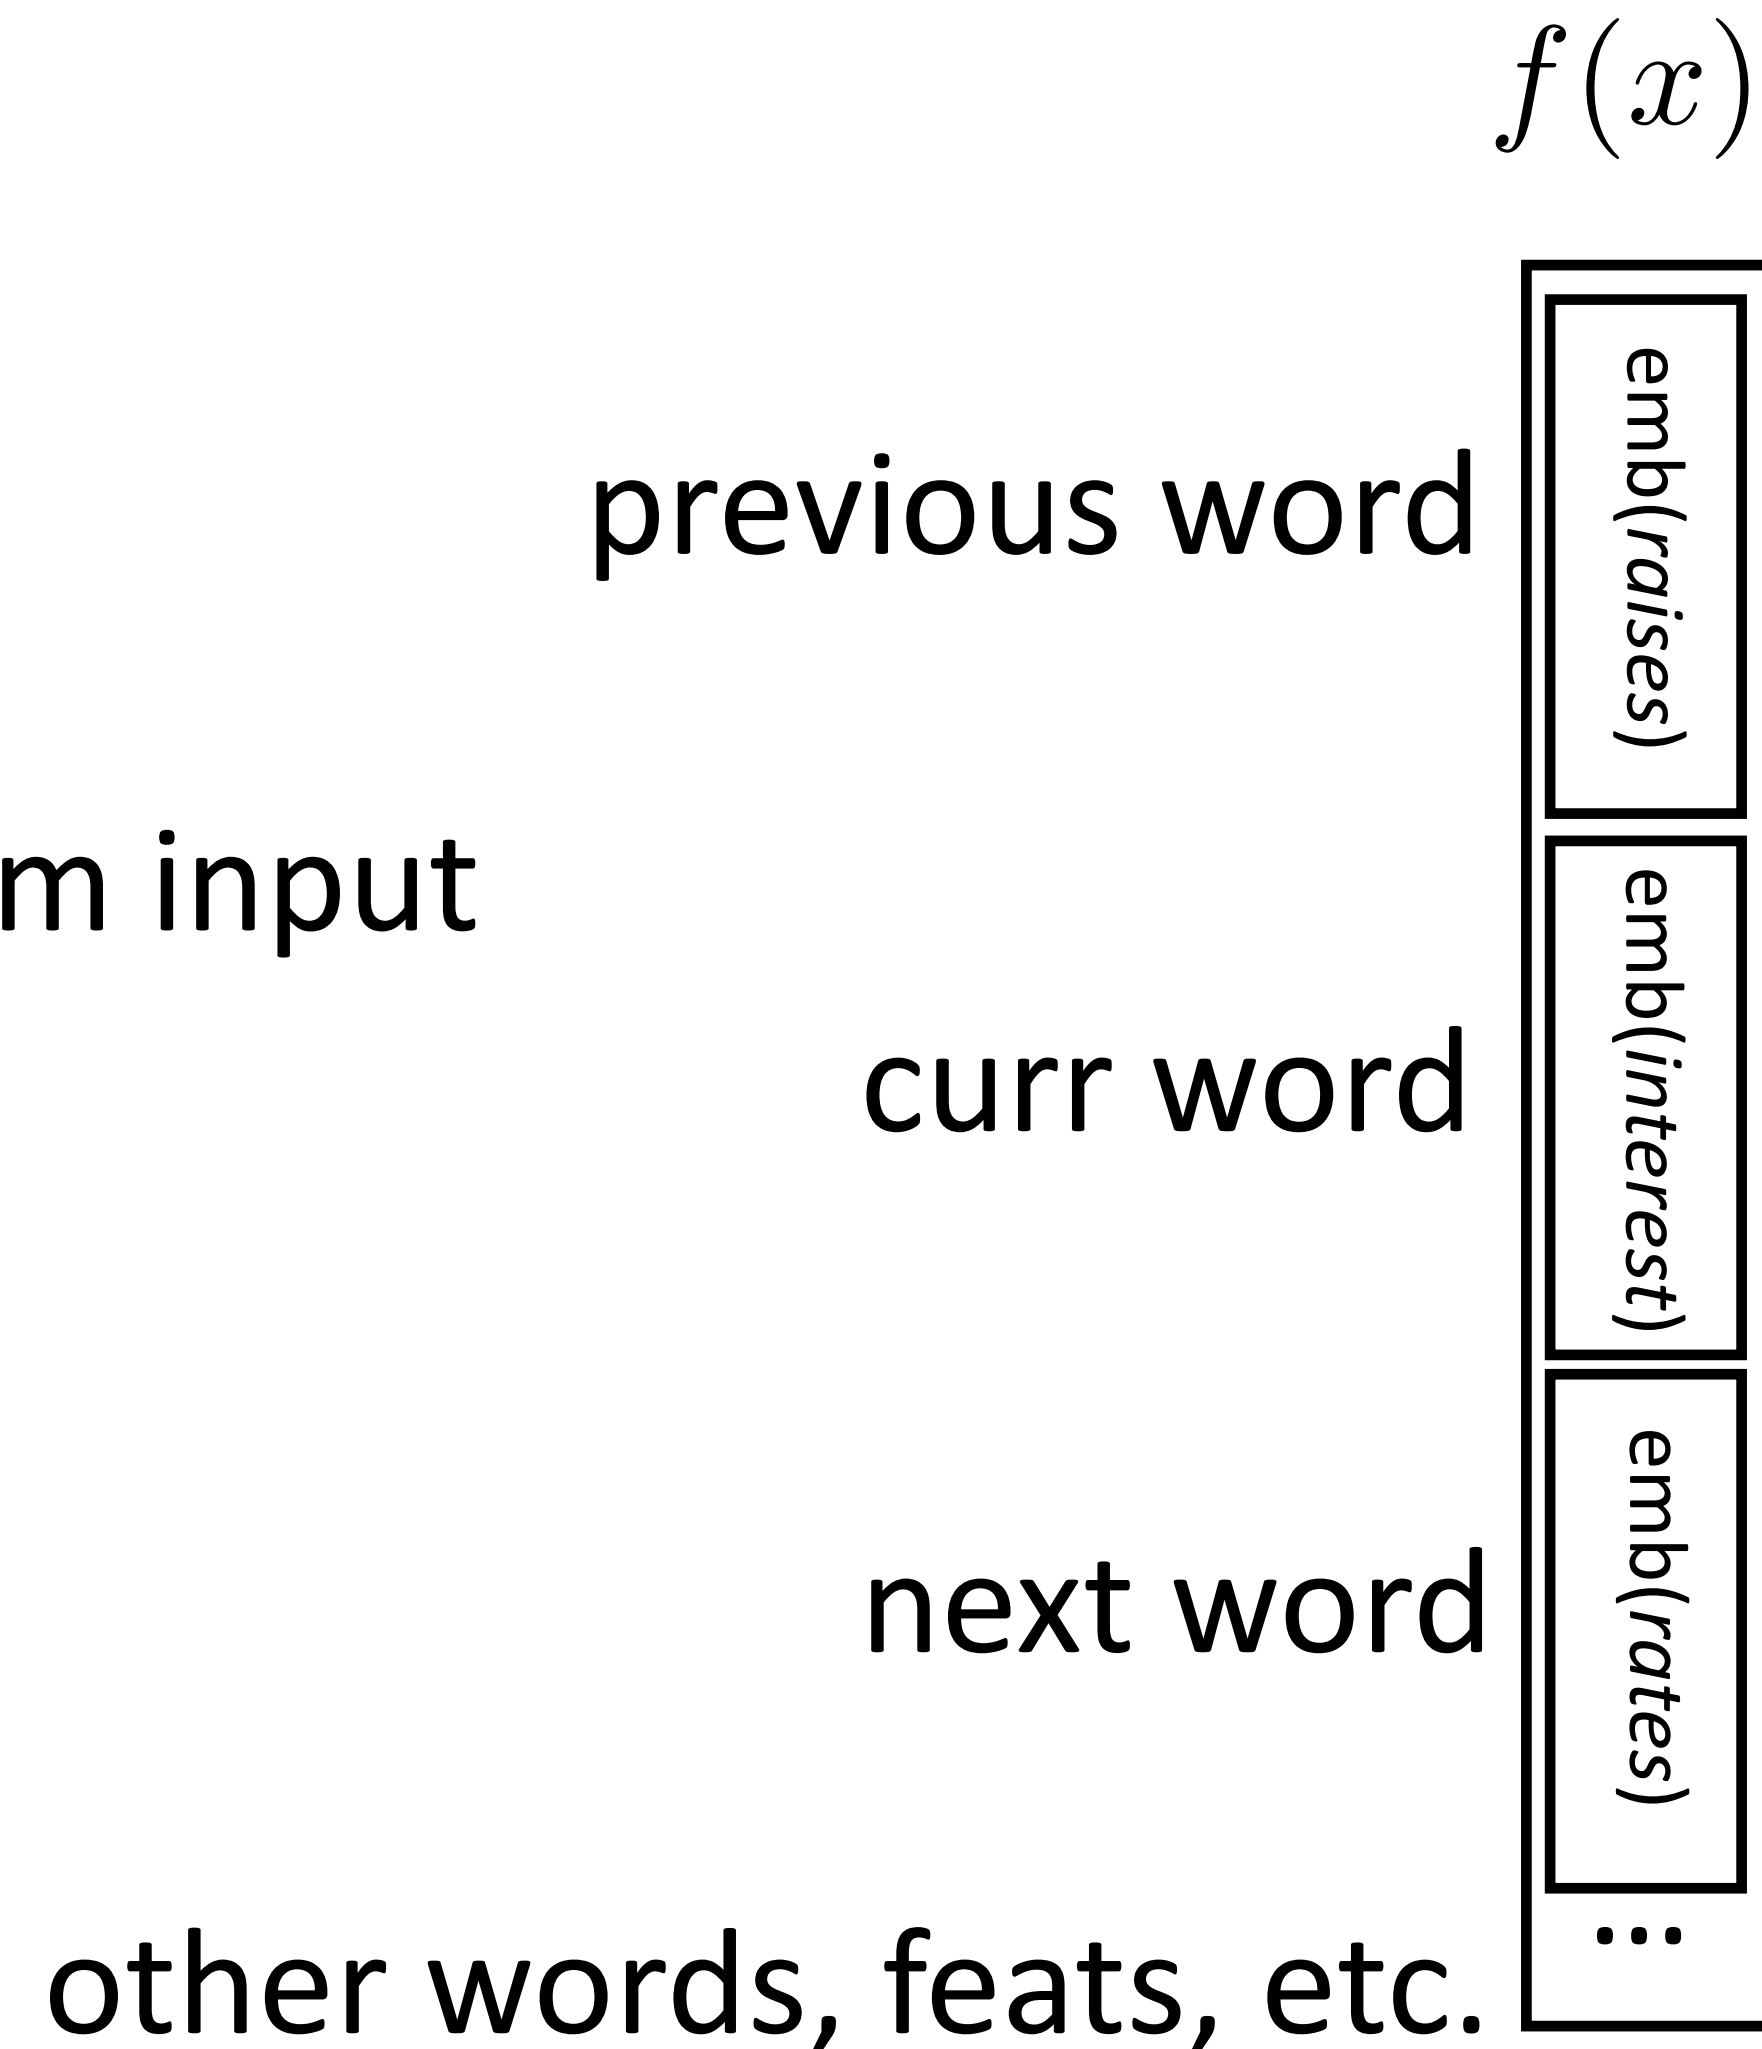
# Word Embeddings

- Part-of-speech tagging with FFNNs

??

*Fed raises **interest** rates in order to ...*

- Word embeddings for each word form input



Botha et al. (2017)

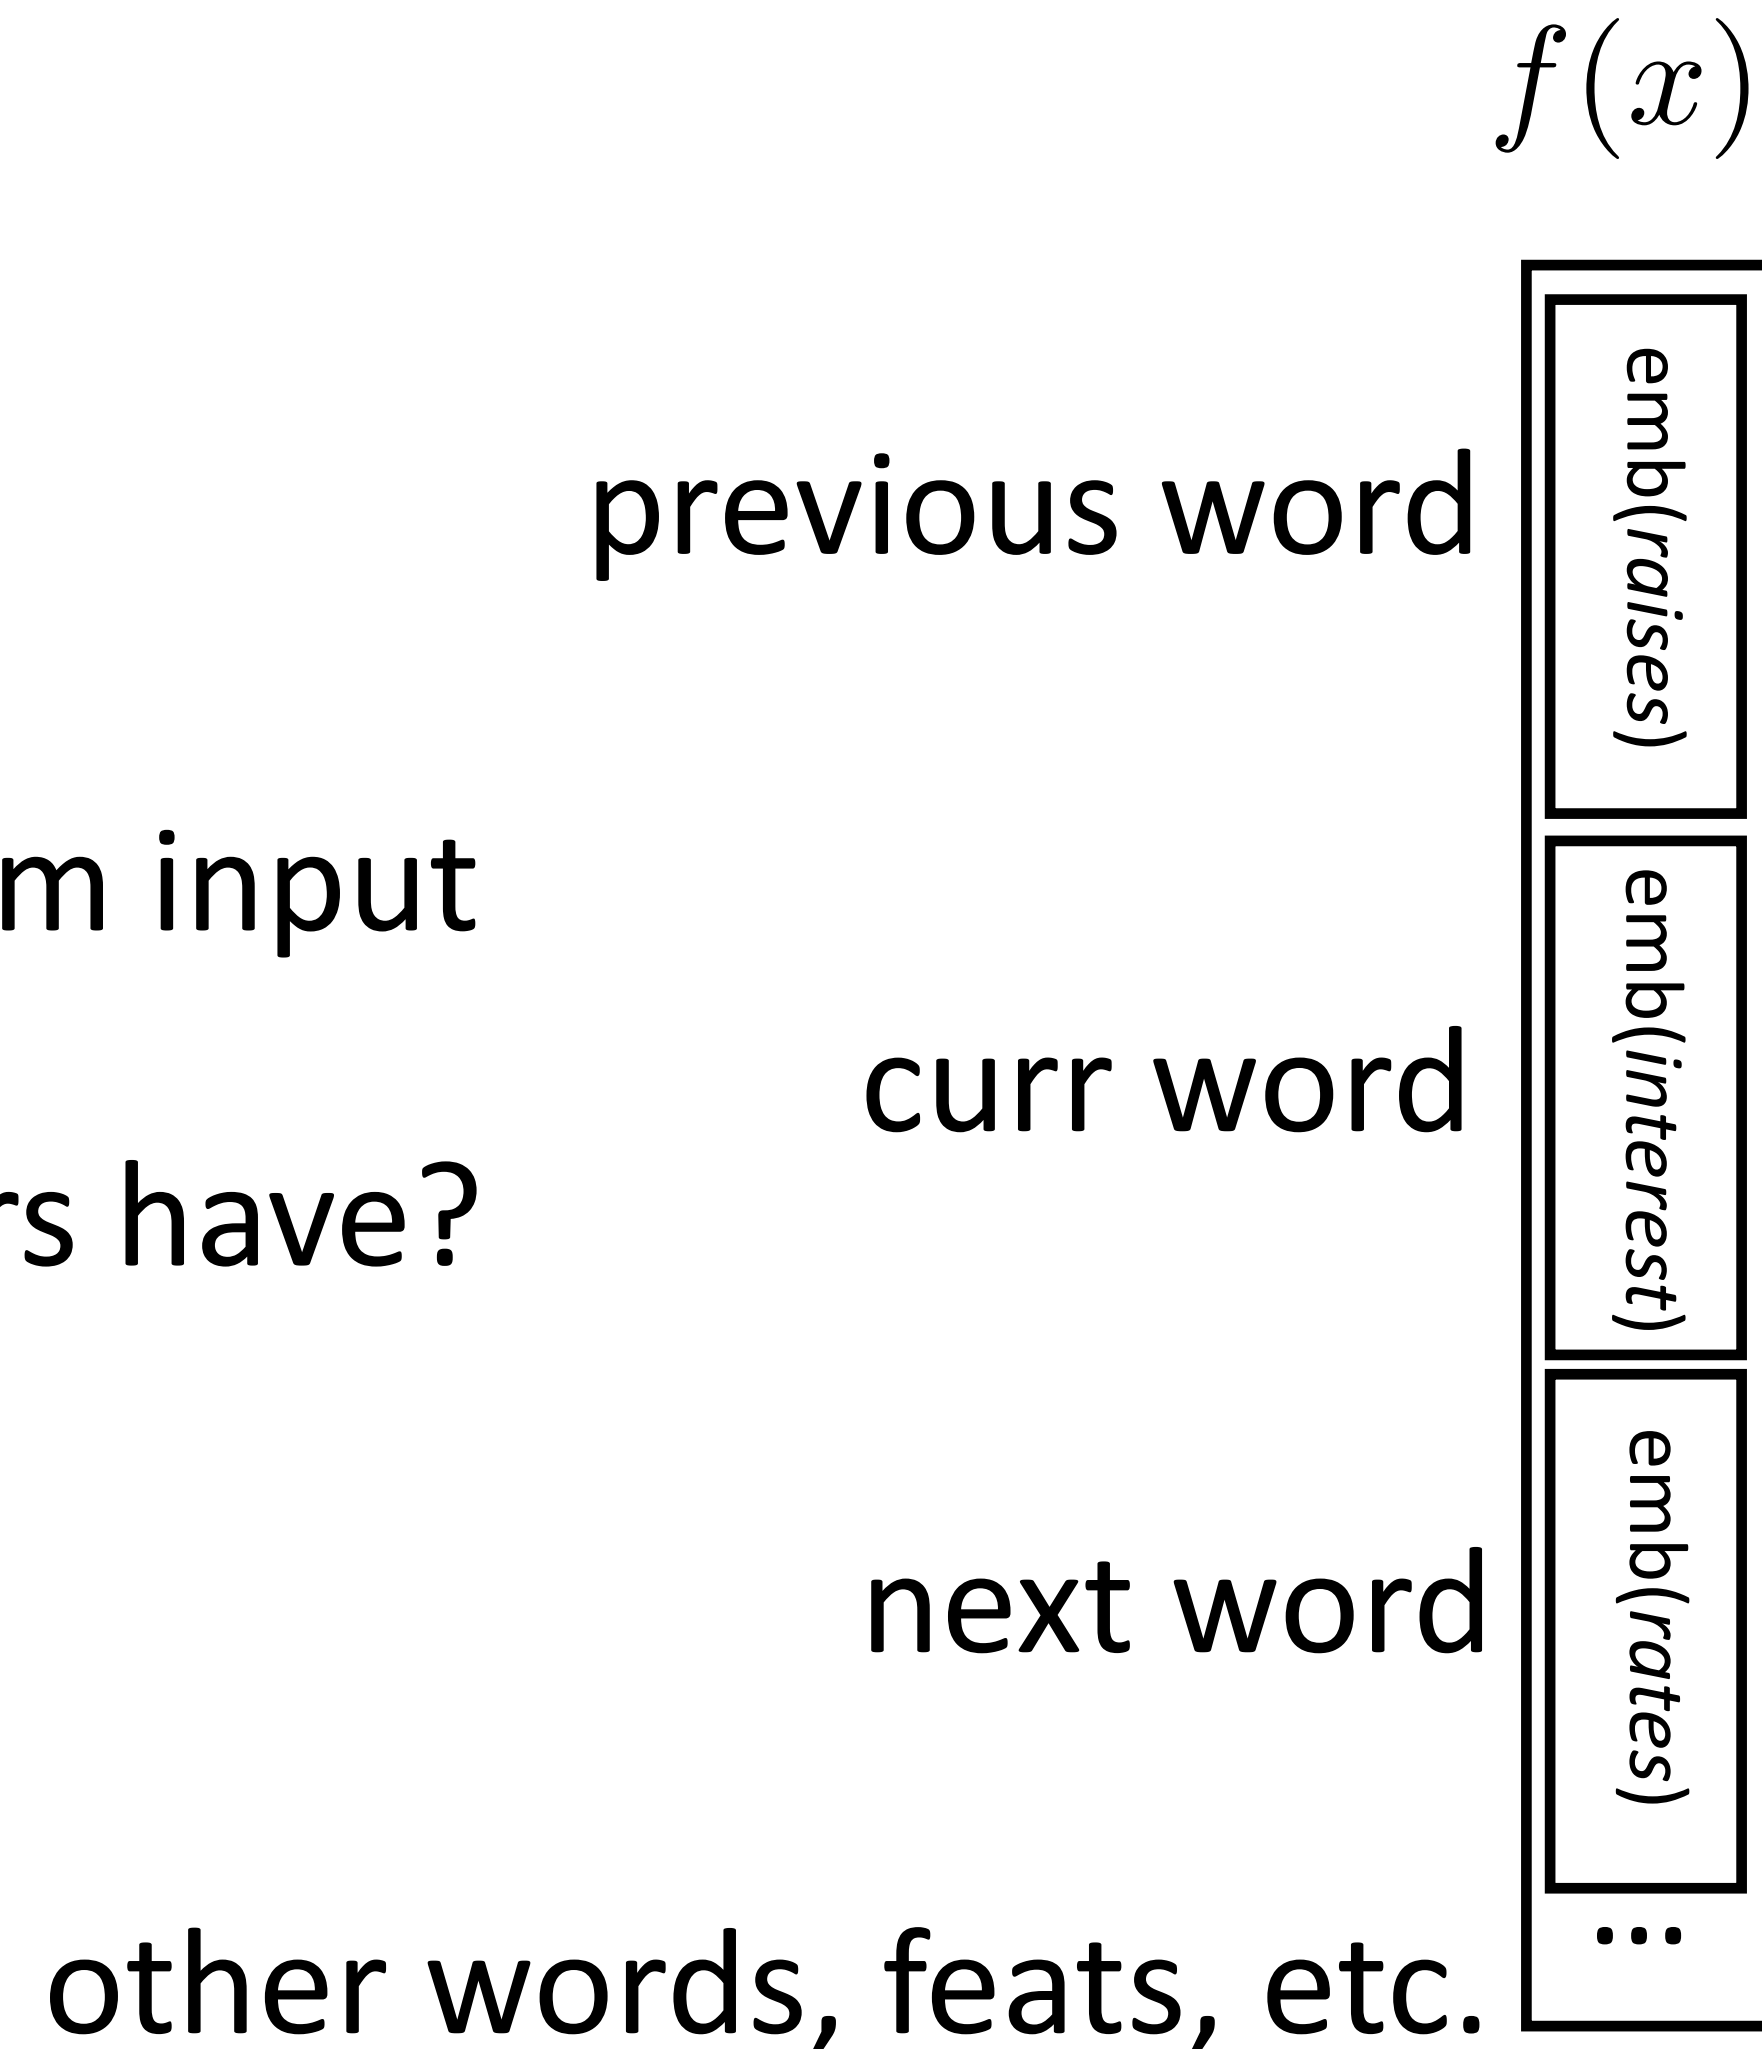
# Word Embeddings

- ▶ Part-of-speech tagging with FFNNs

??

*Fed raises **interest** rates in order to ...*

- ▶ Word embeddings for each word form input
- ▶ What properties should these vectors have?



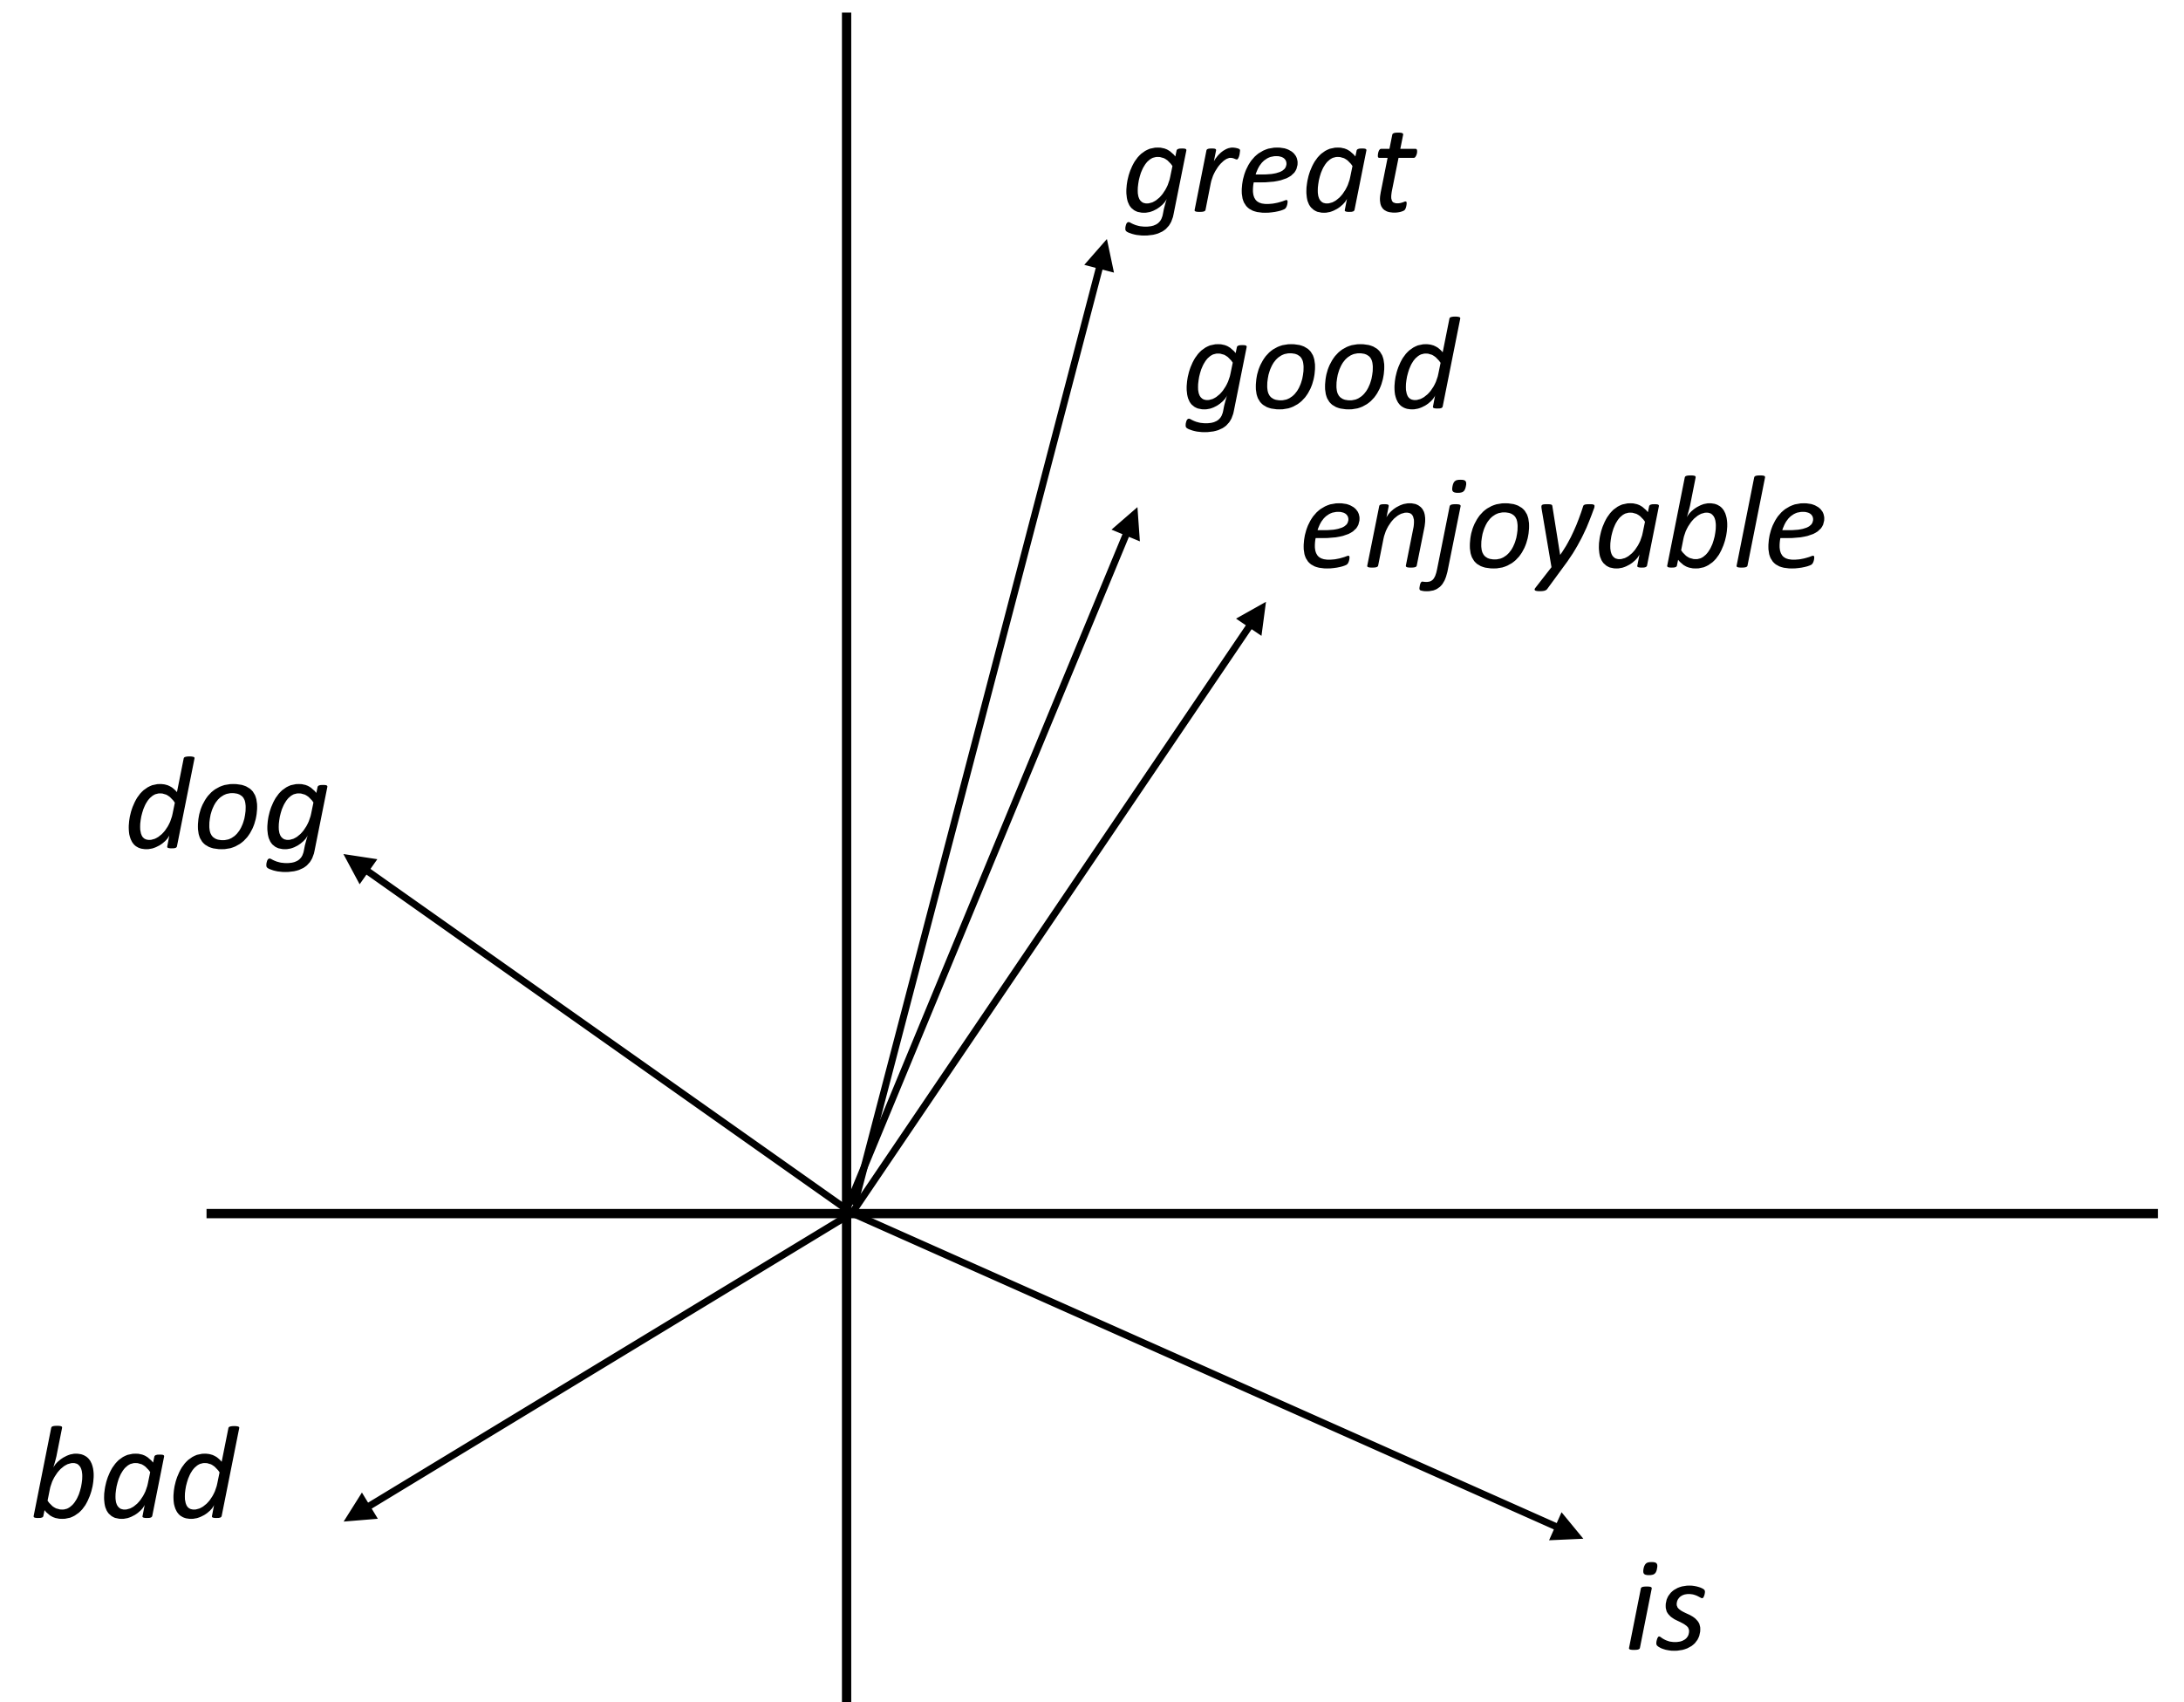
Botha et al. (2017)

# Word Embeddings

---

# Word Embeddings

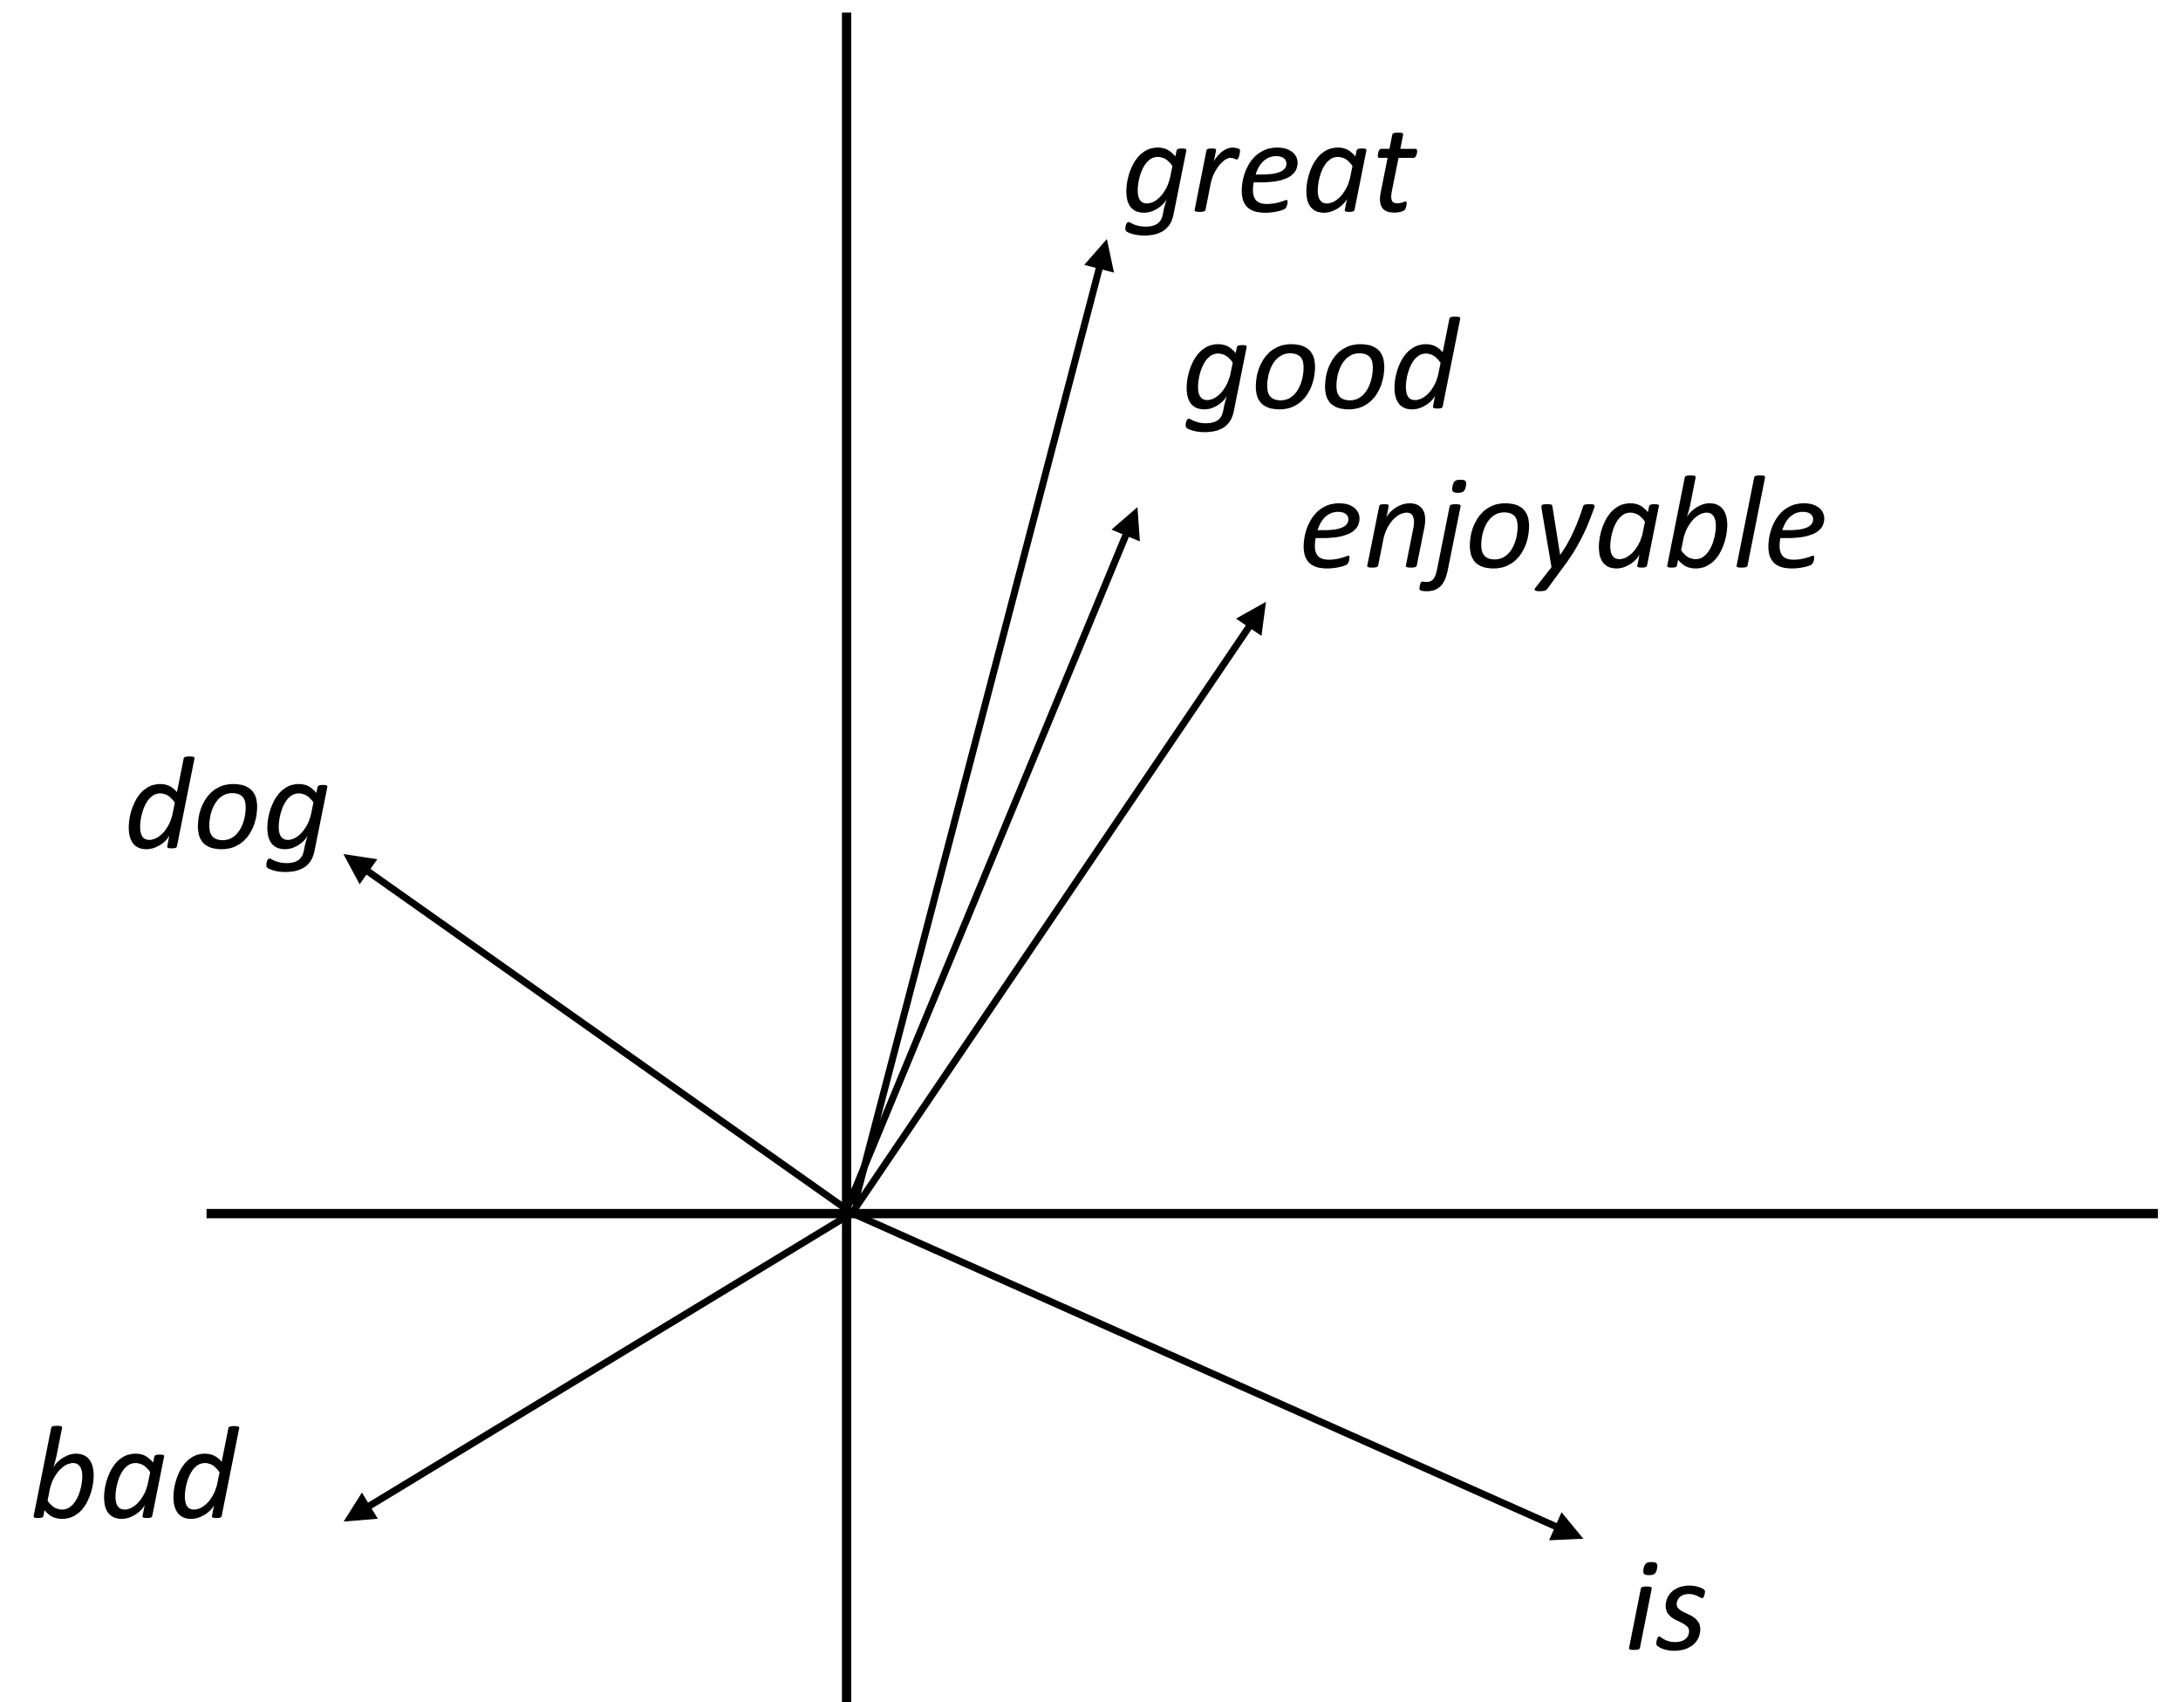
---



# Word Embeddings

---

- ▶ Want a vector space where similar words have similar embeddings



# Word Embeddings


---

- ▶ Want a vector space where similar words have similar embeddings

*the movie was great*

$\approx$

*the movie was good*



# Word Embeddings

---

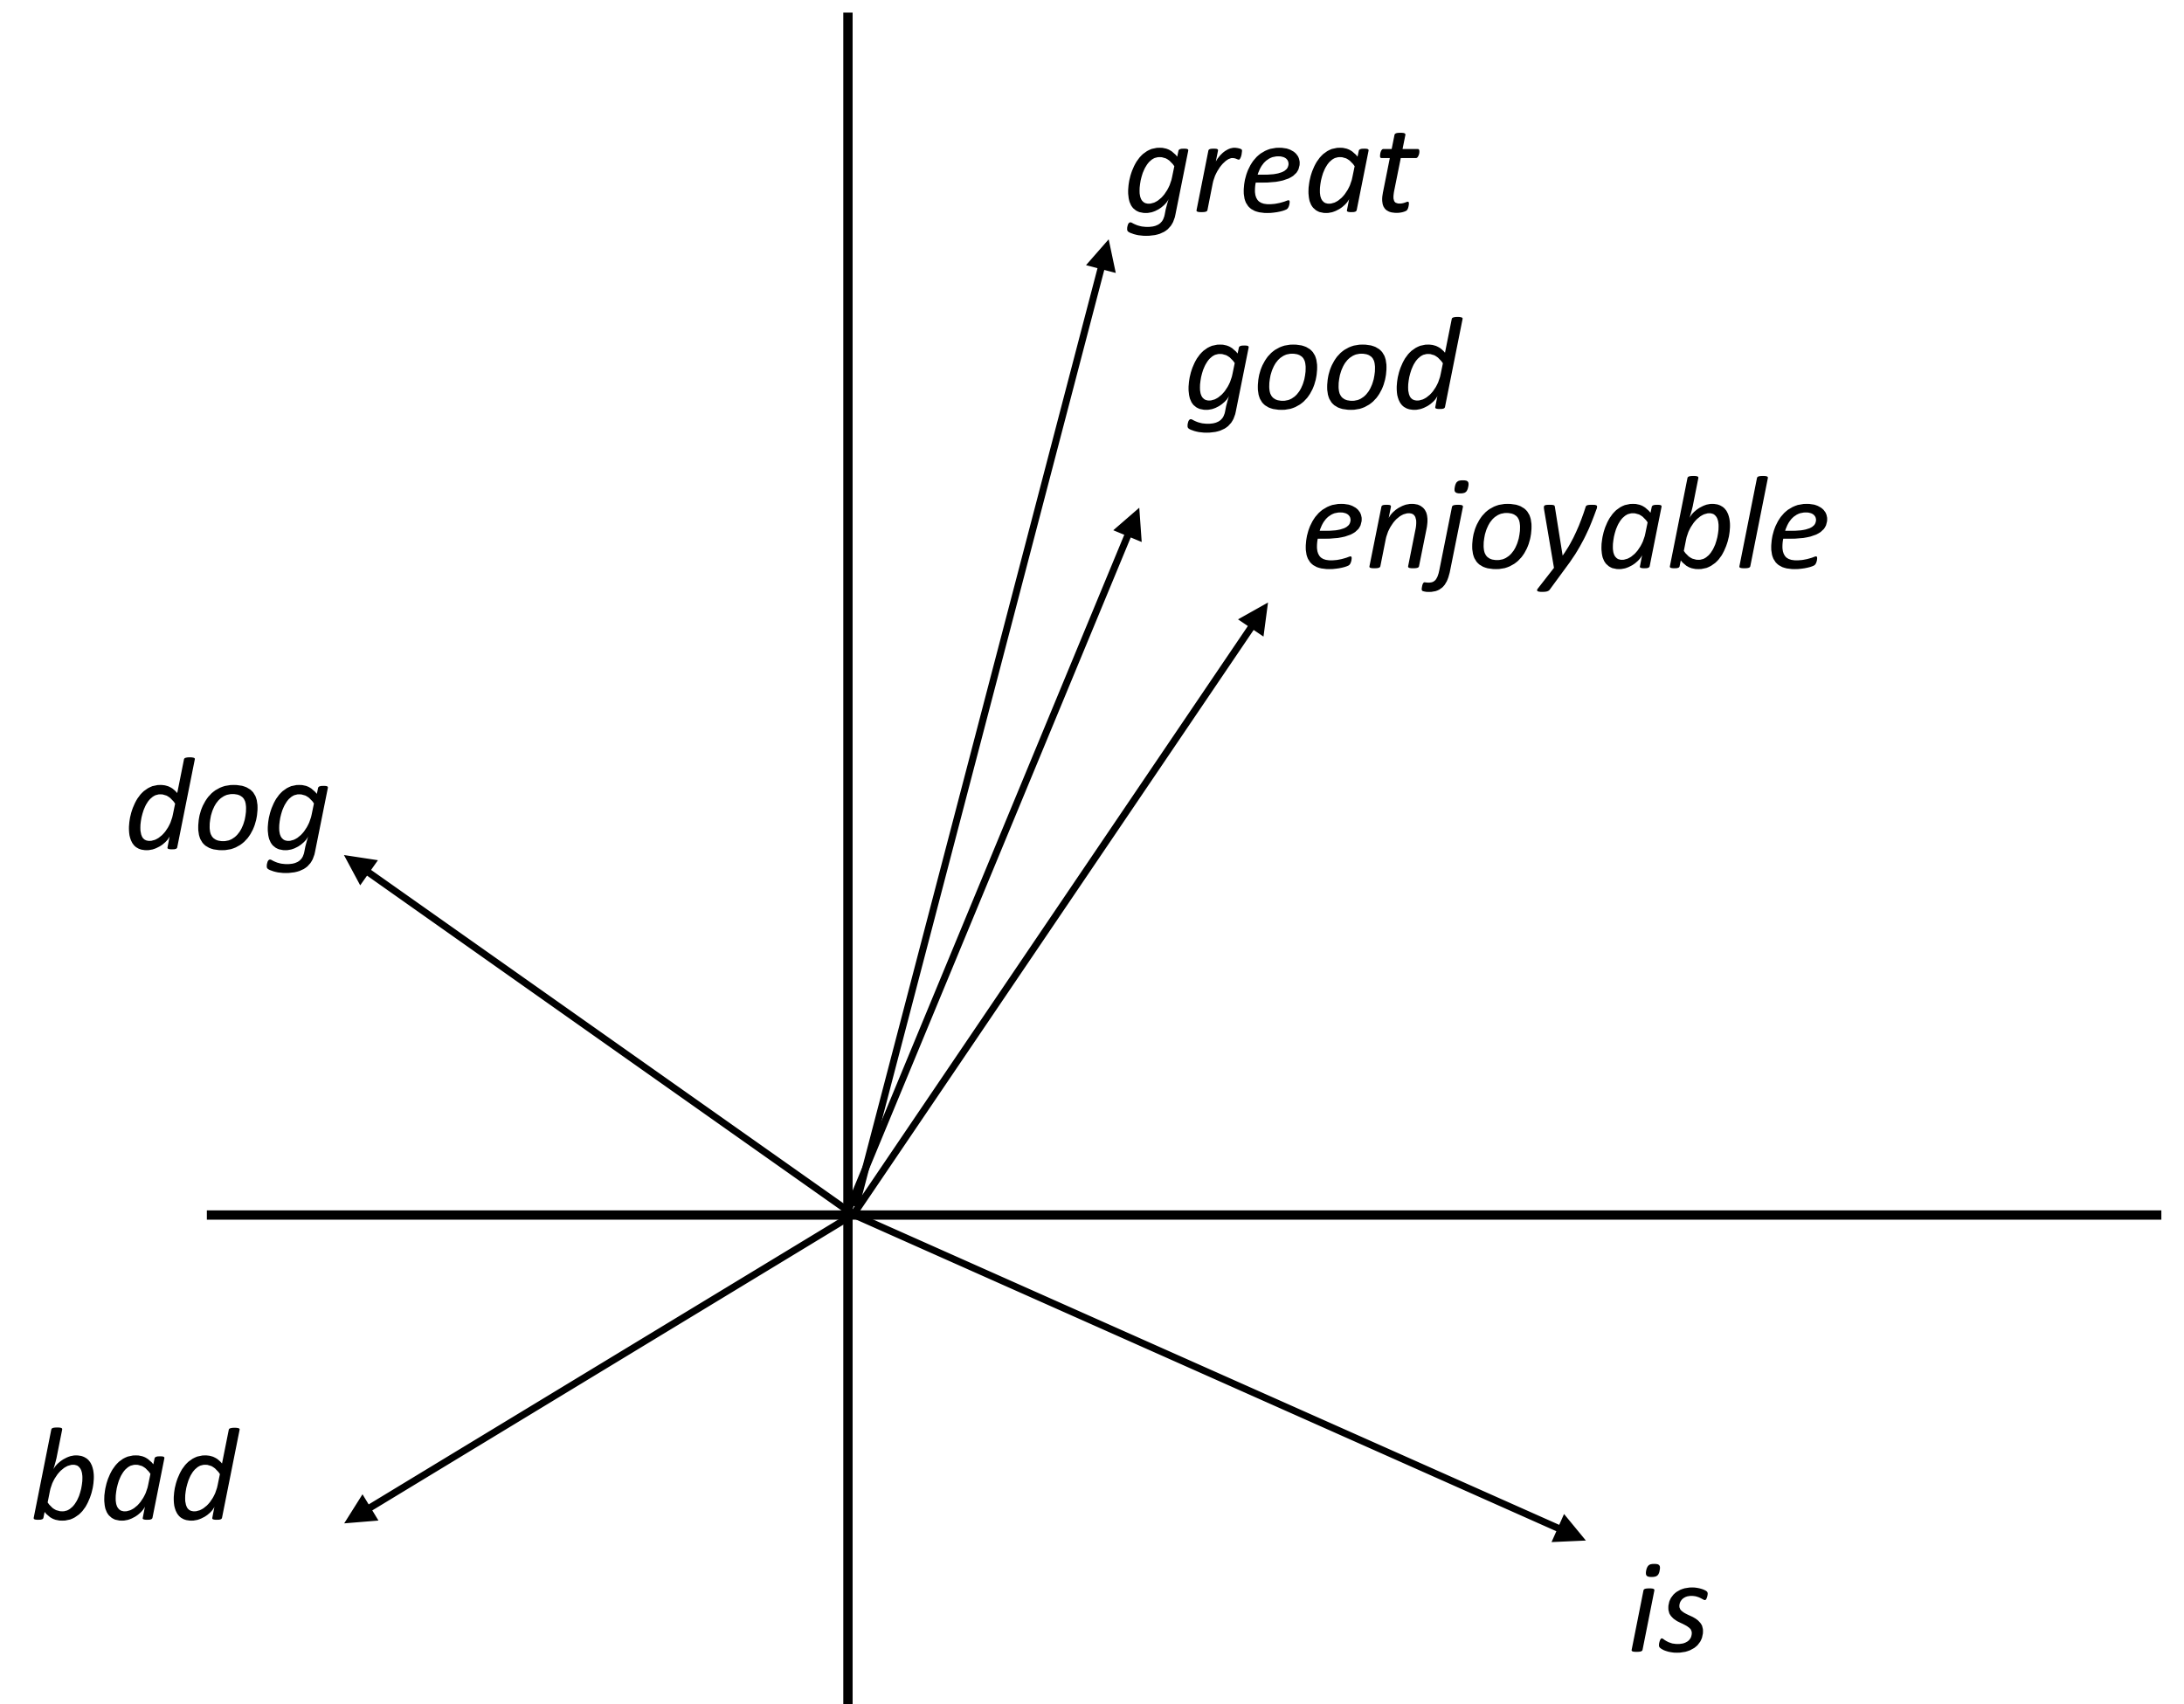
- ▶ Want a vector space where similar words have similar embeddings

*the movie was great*

$\approx$

*the movie was good*

- ▶ Goal: come up with a way to produce these embeddings





word2vec/GloVe

# Continuous Bag-of-Words

---

- ▶ Predict word from context

*the dog bit the man*




# Continuous Bag-of-Words


---

- Predict word from context

*the dog bit the man*



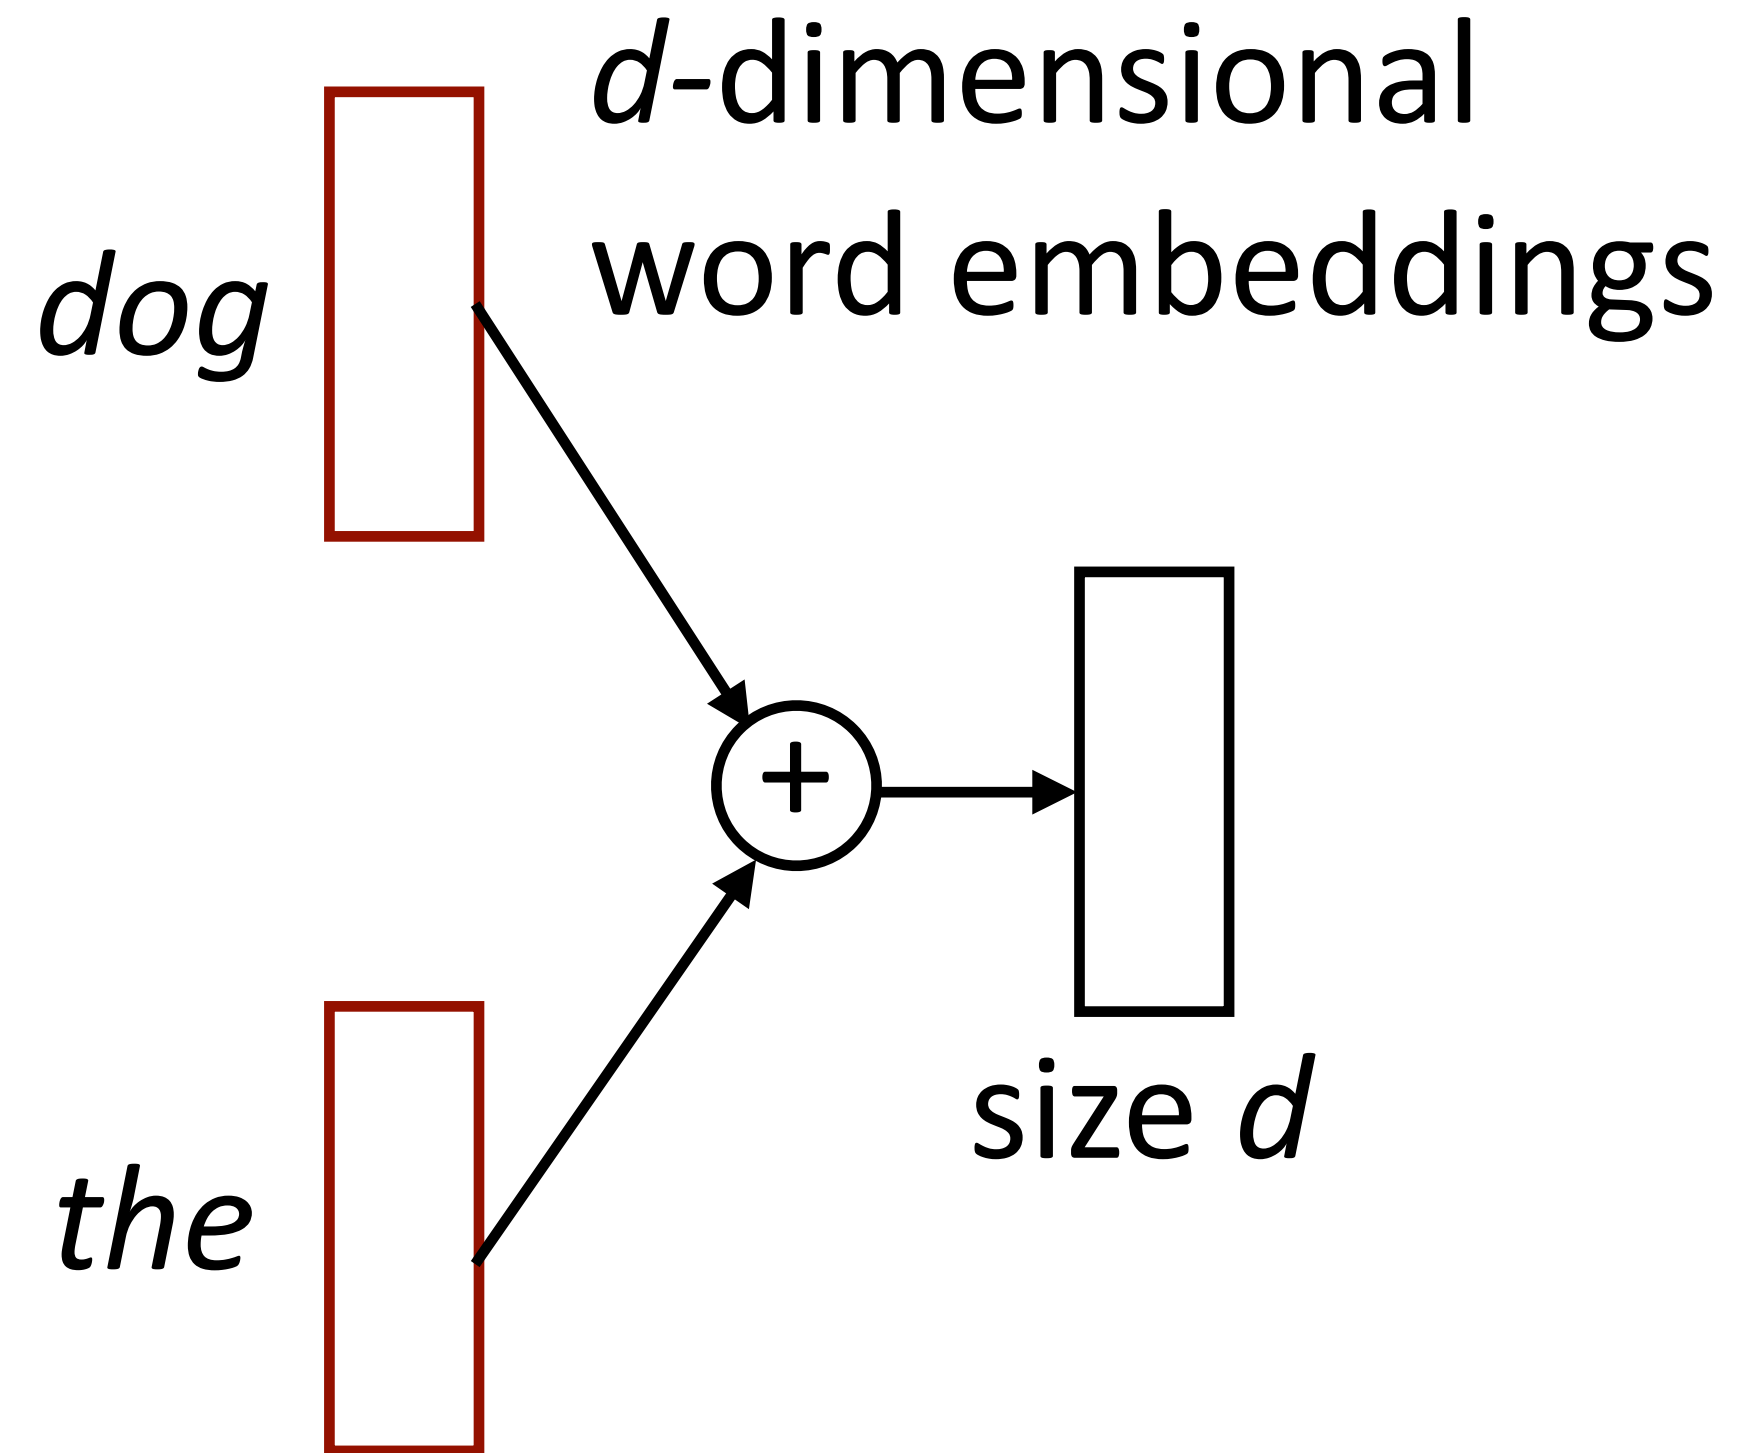
*dog*   $d$ -dimensional  
word embeddings

*the* 

# Continuous Bag-of-Words

- Predict word from context

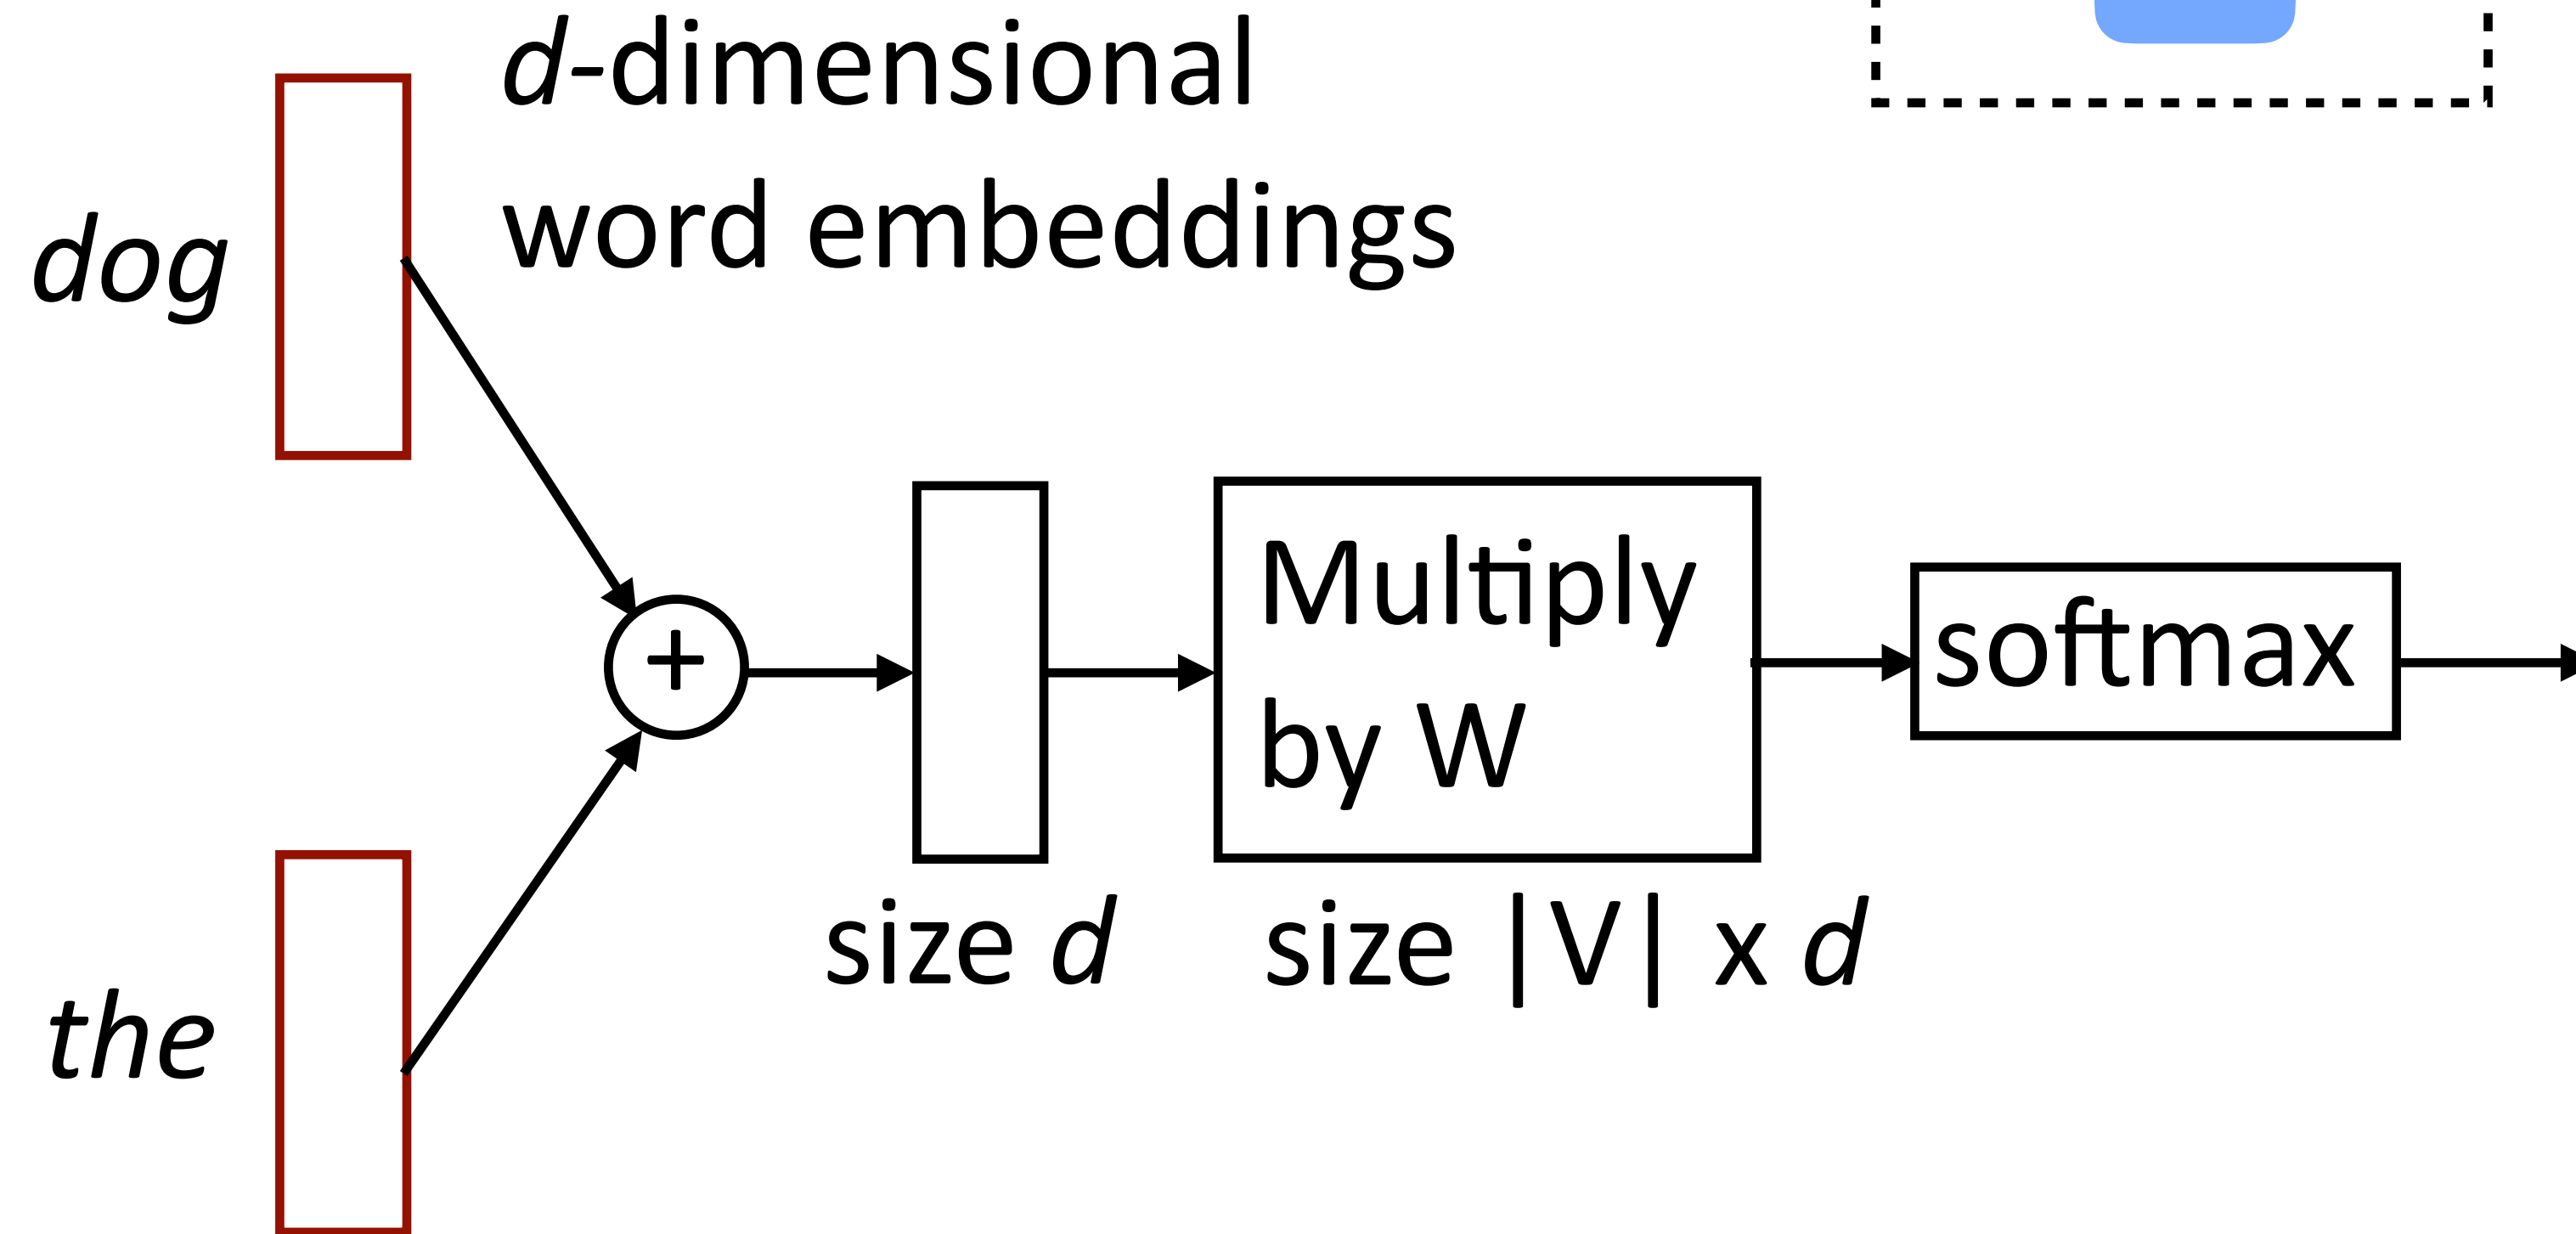
*the dog bit the man*



# Continuous Bag-of-Words

- Predict word from context

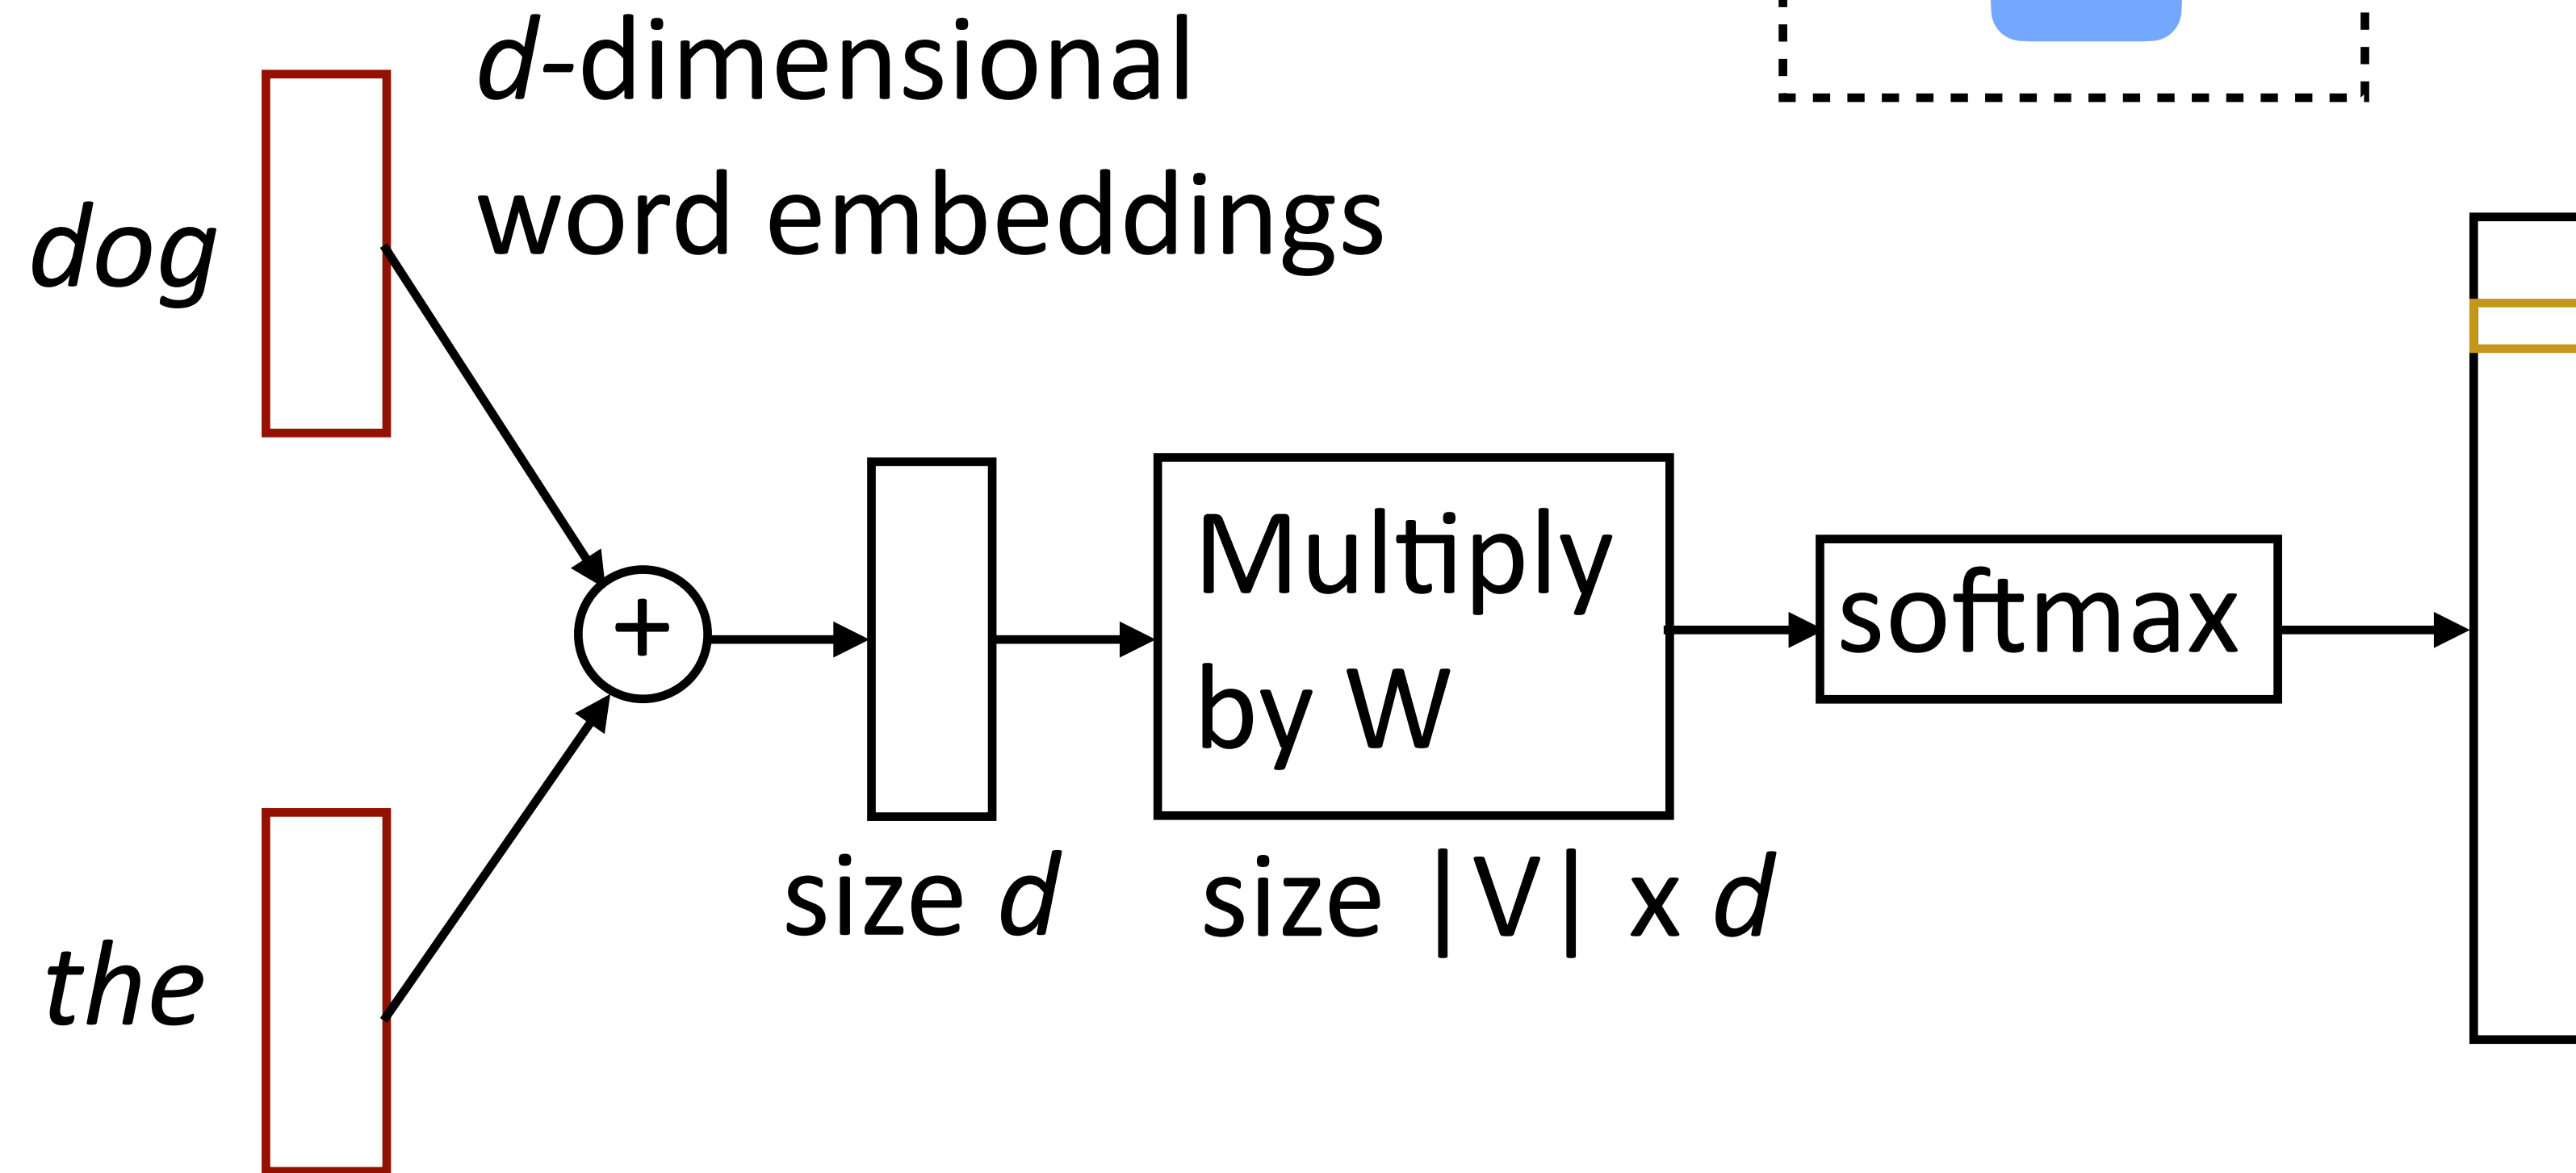
*the dog bit the man*



# Continuous Bag-of-Words

- Predict word from context

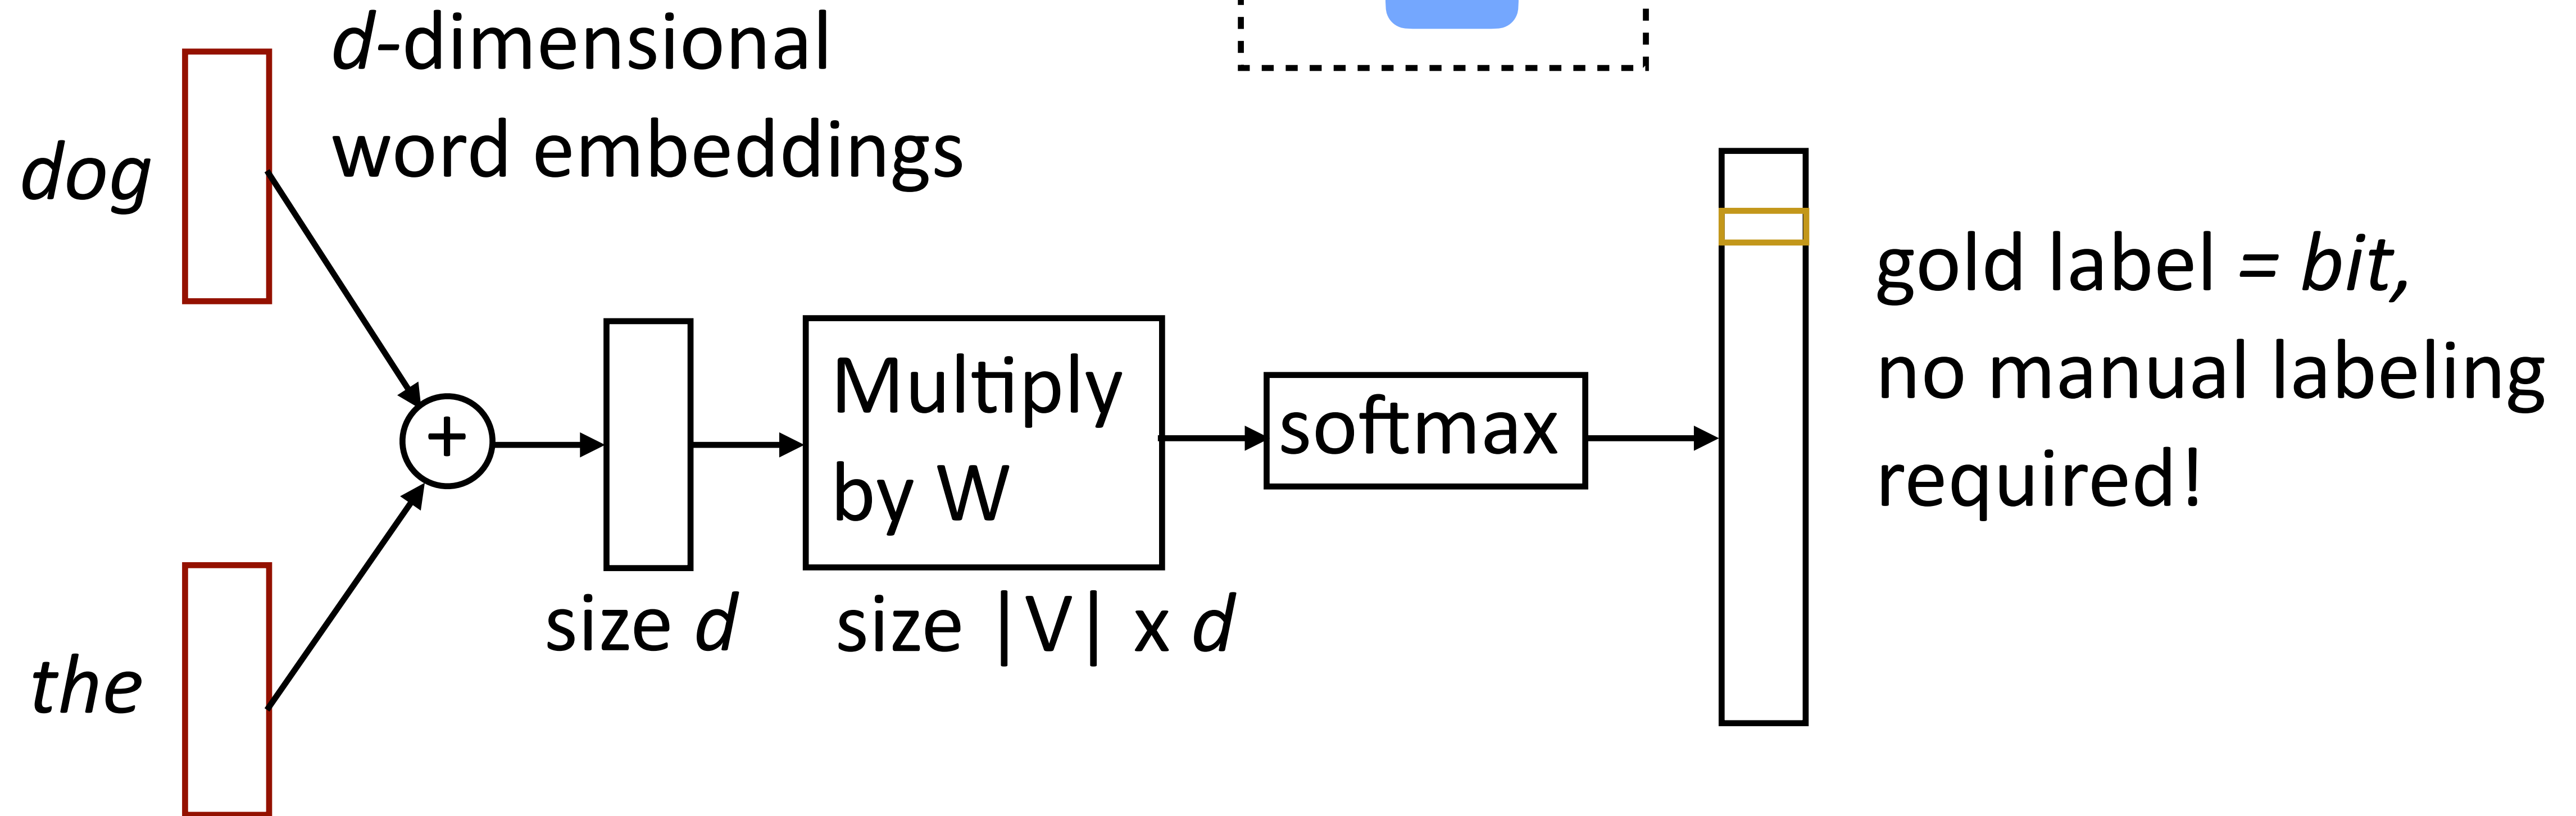
*the dog bit the man*



# Continuous Bag-of-Words

- Predict word from context

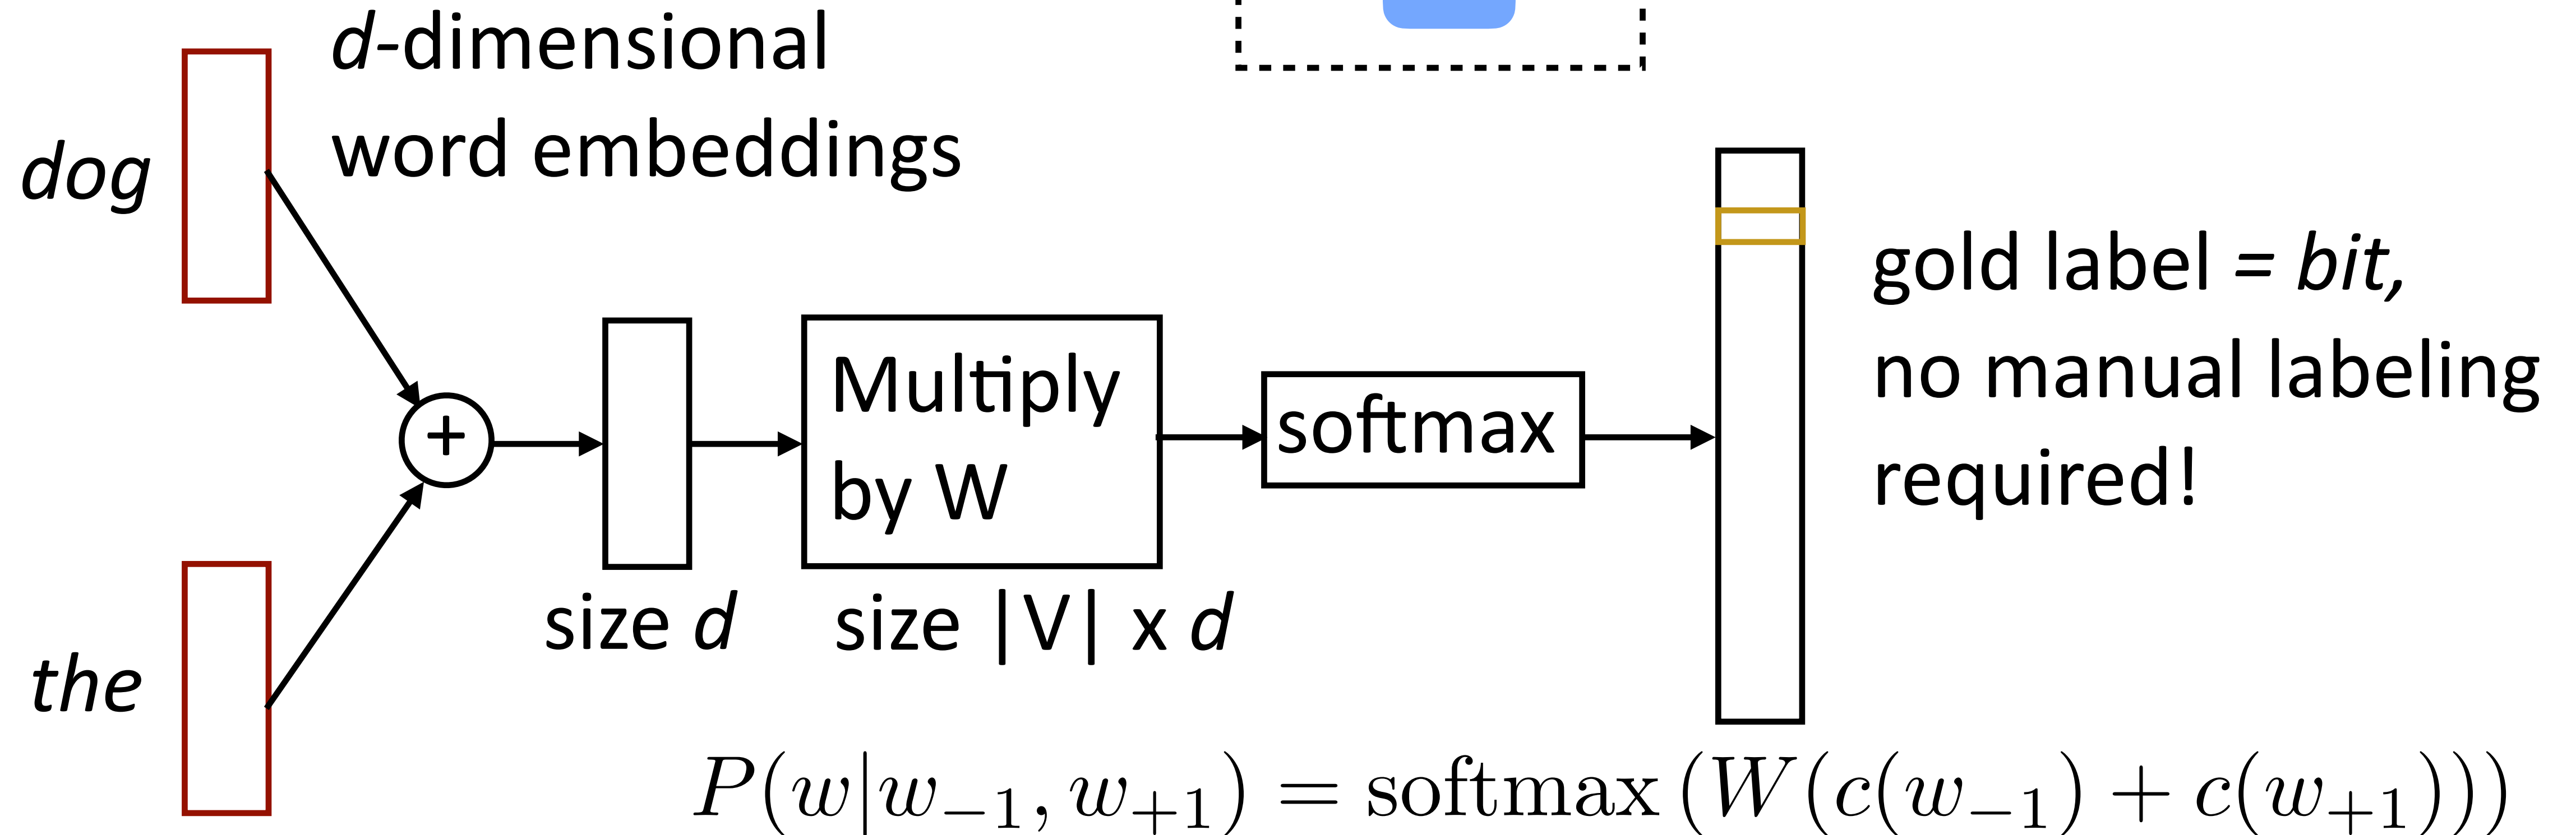
*the dog bit the man*



# Continuous Bag-of-Words

- Predict word from context

*the dog **bit** the man*

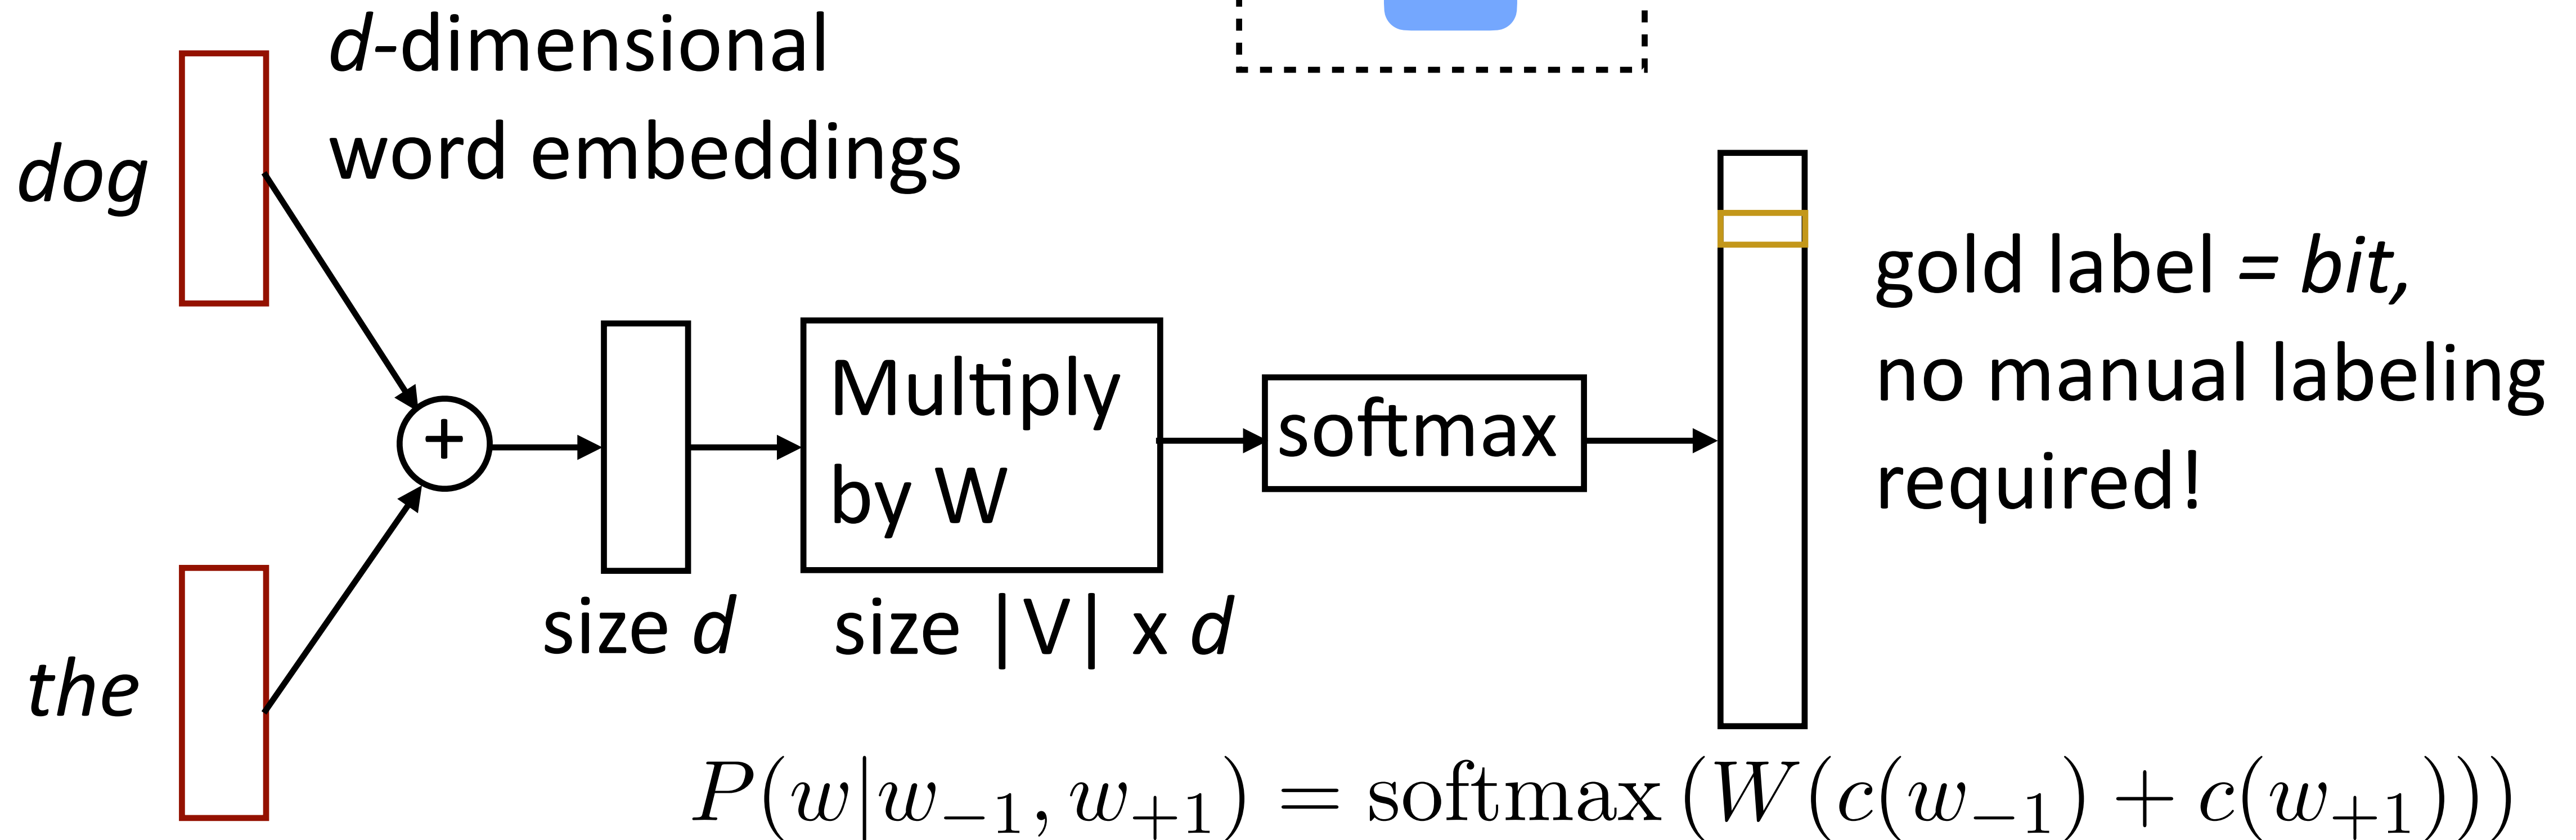




# Continuous Bag-of-Words

- Predict word from context

*the dog **bit** the man*



- Parameters:  $d \times |V|$  (one  $d$ -length **vector per voc word**),  
 $|V| \times d$  output parameters ( $W$ )

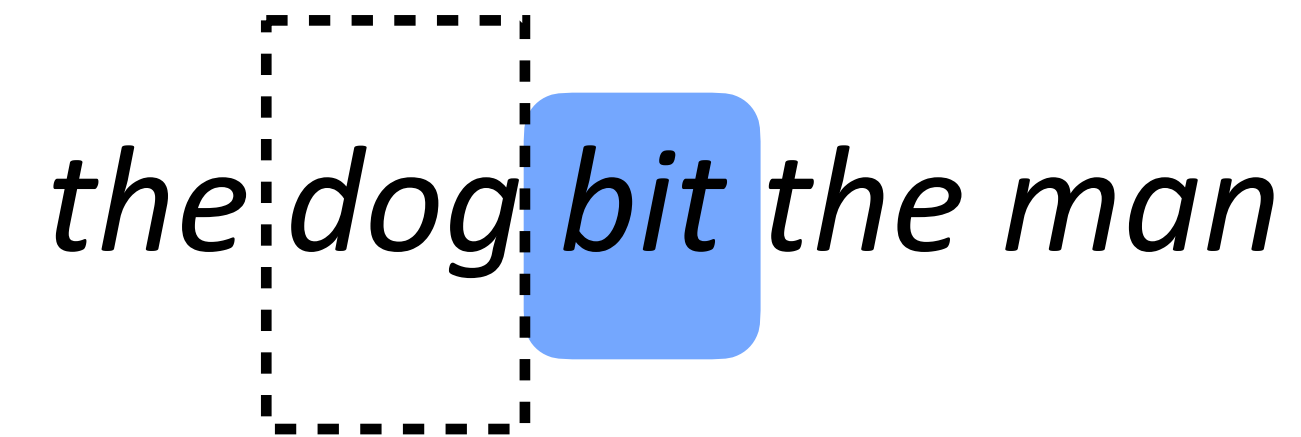
Mikolov et al. (2013)

# Skip-Gram

---

- Predict one word of context from word

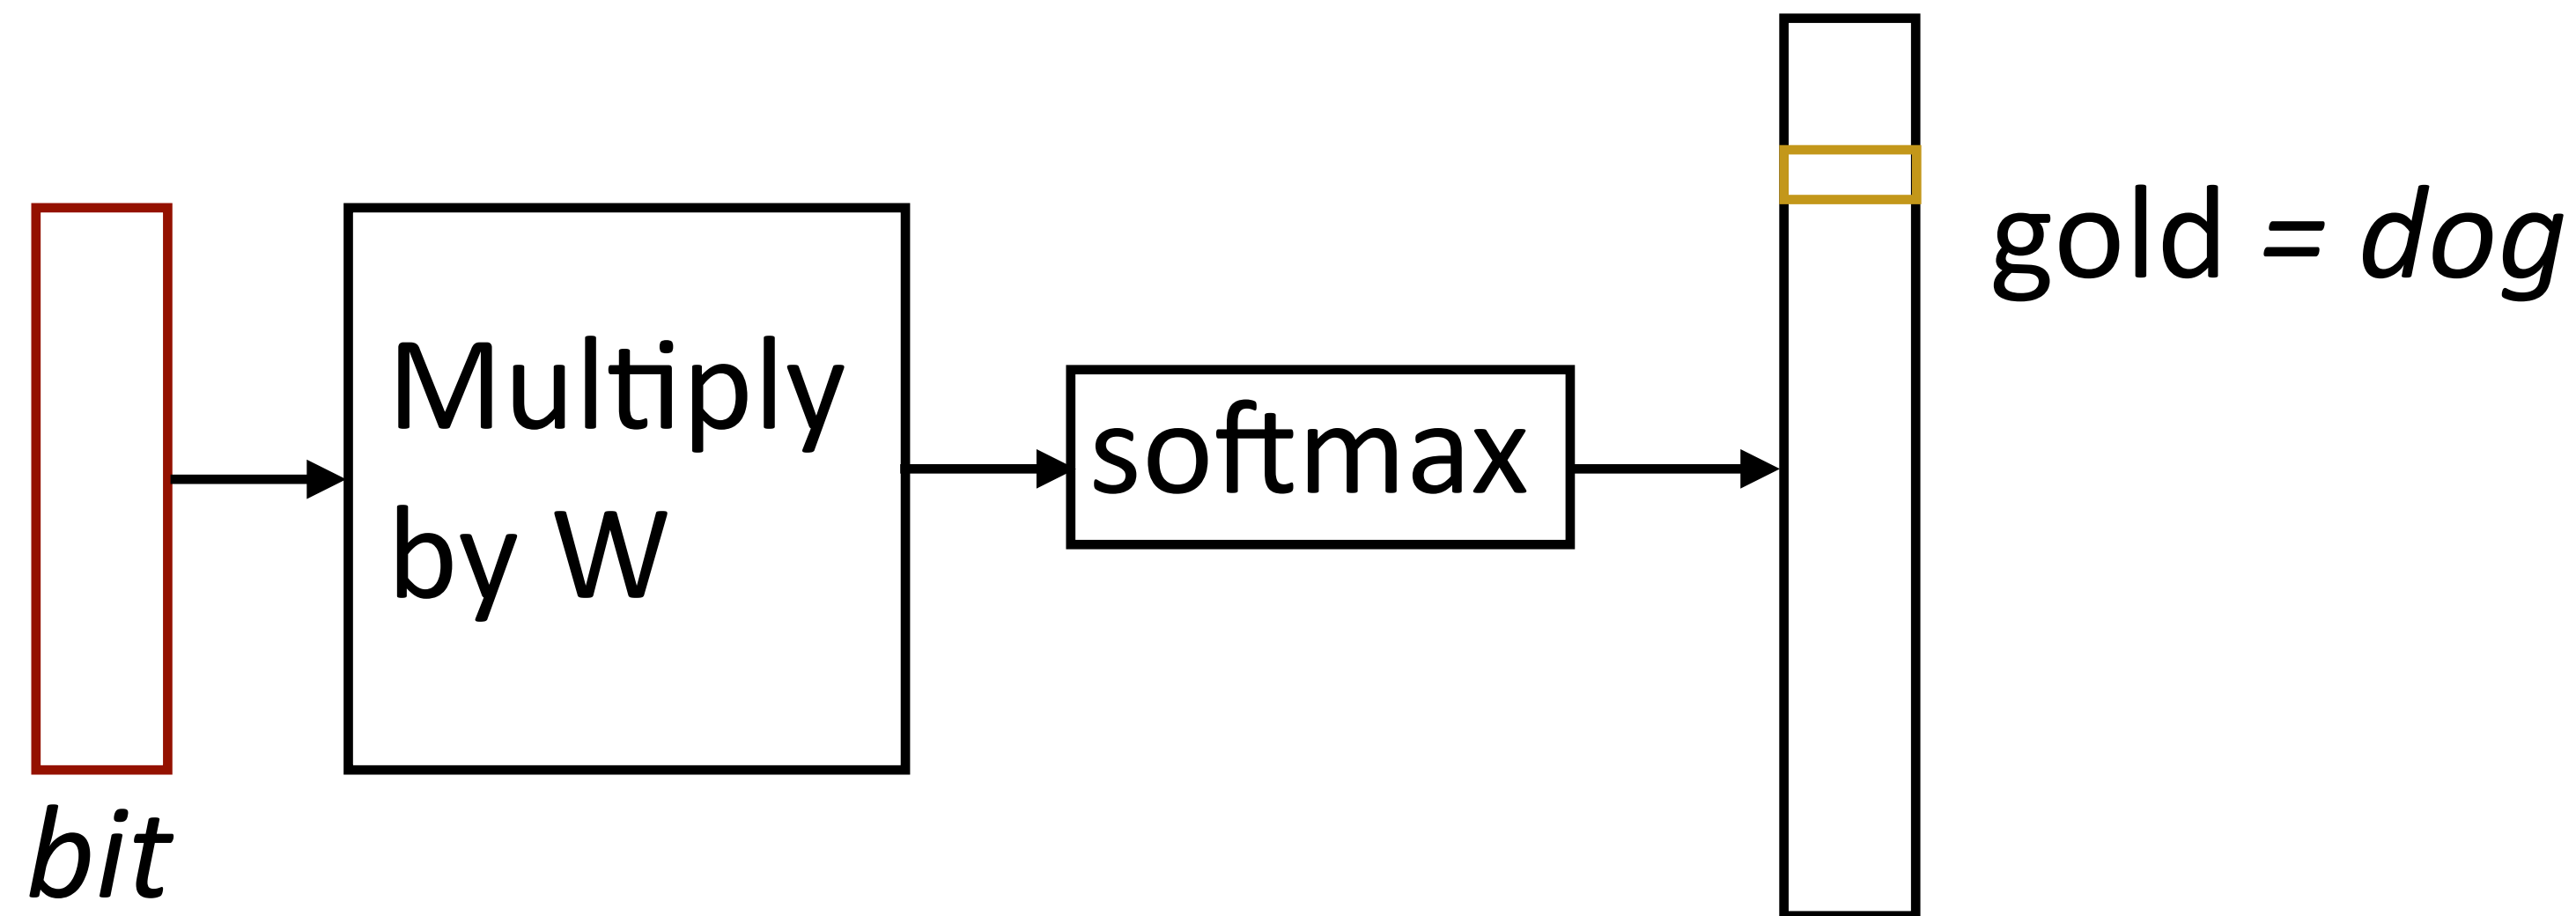
*the dog bit the man*



# Skip-Gram

- Predict one word of context from word

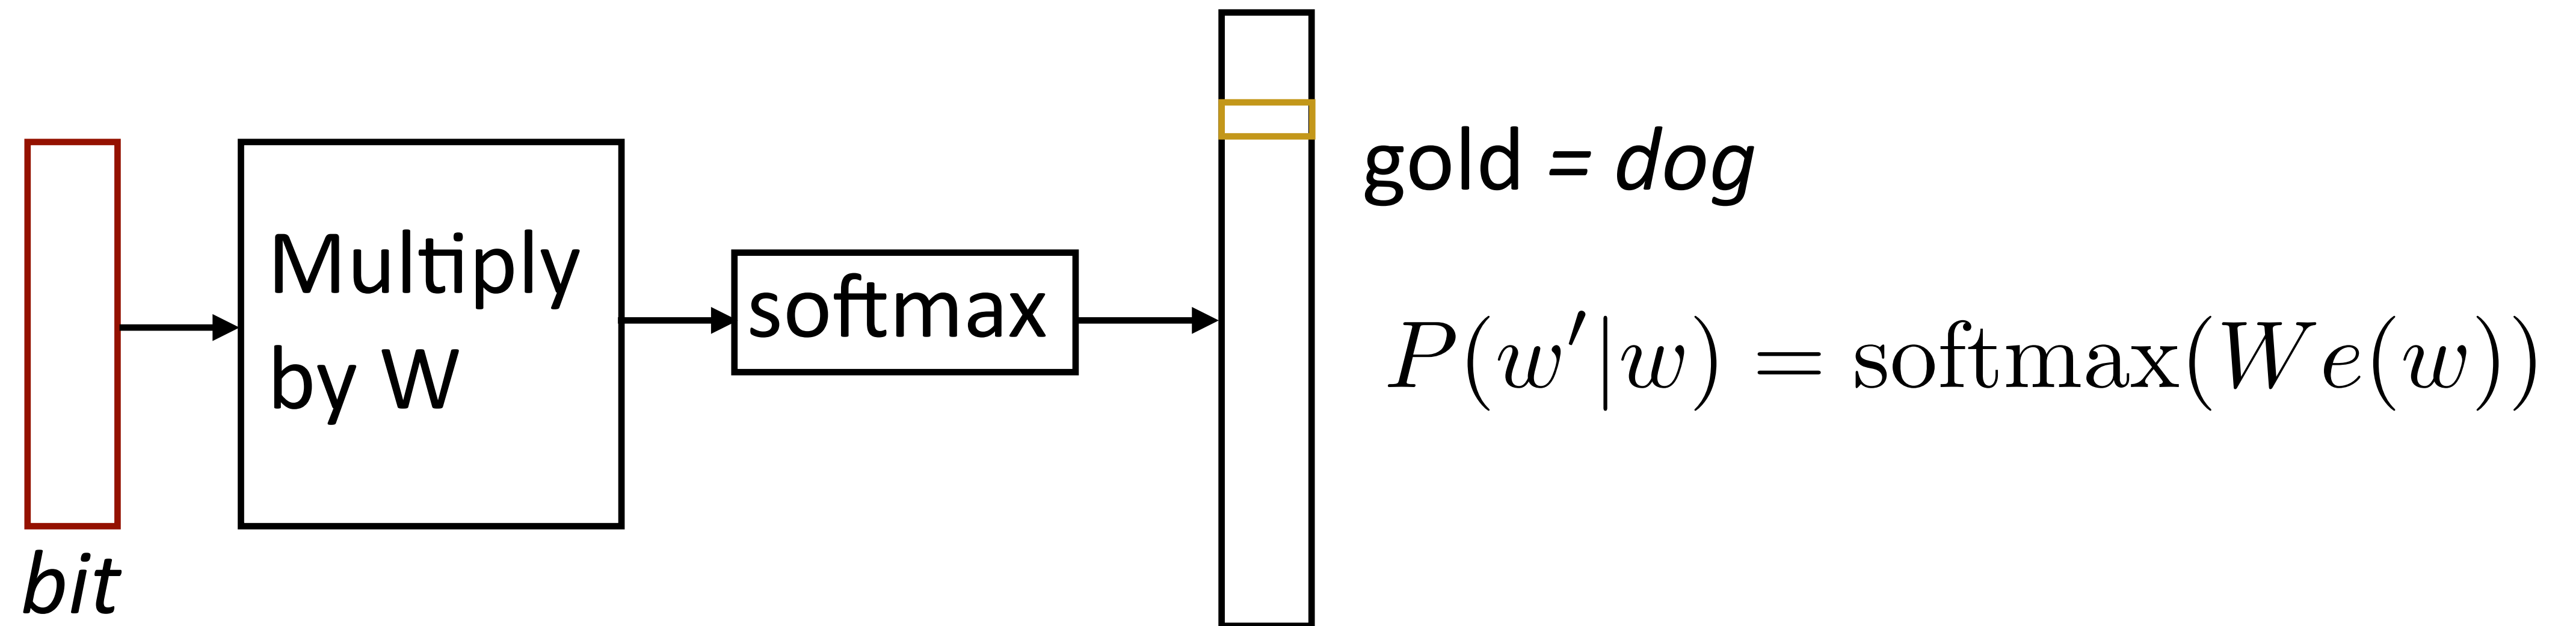
*the* *dog* *bit* *the* *man*



# Skip-Gram

- Predict one word of context from word

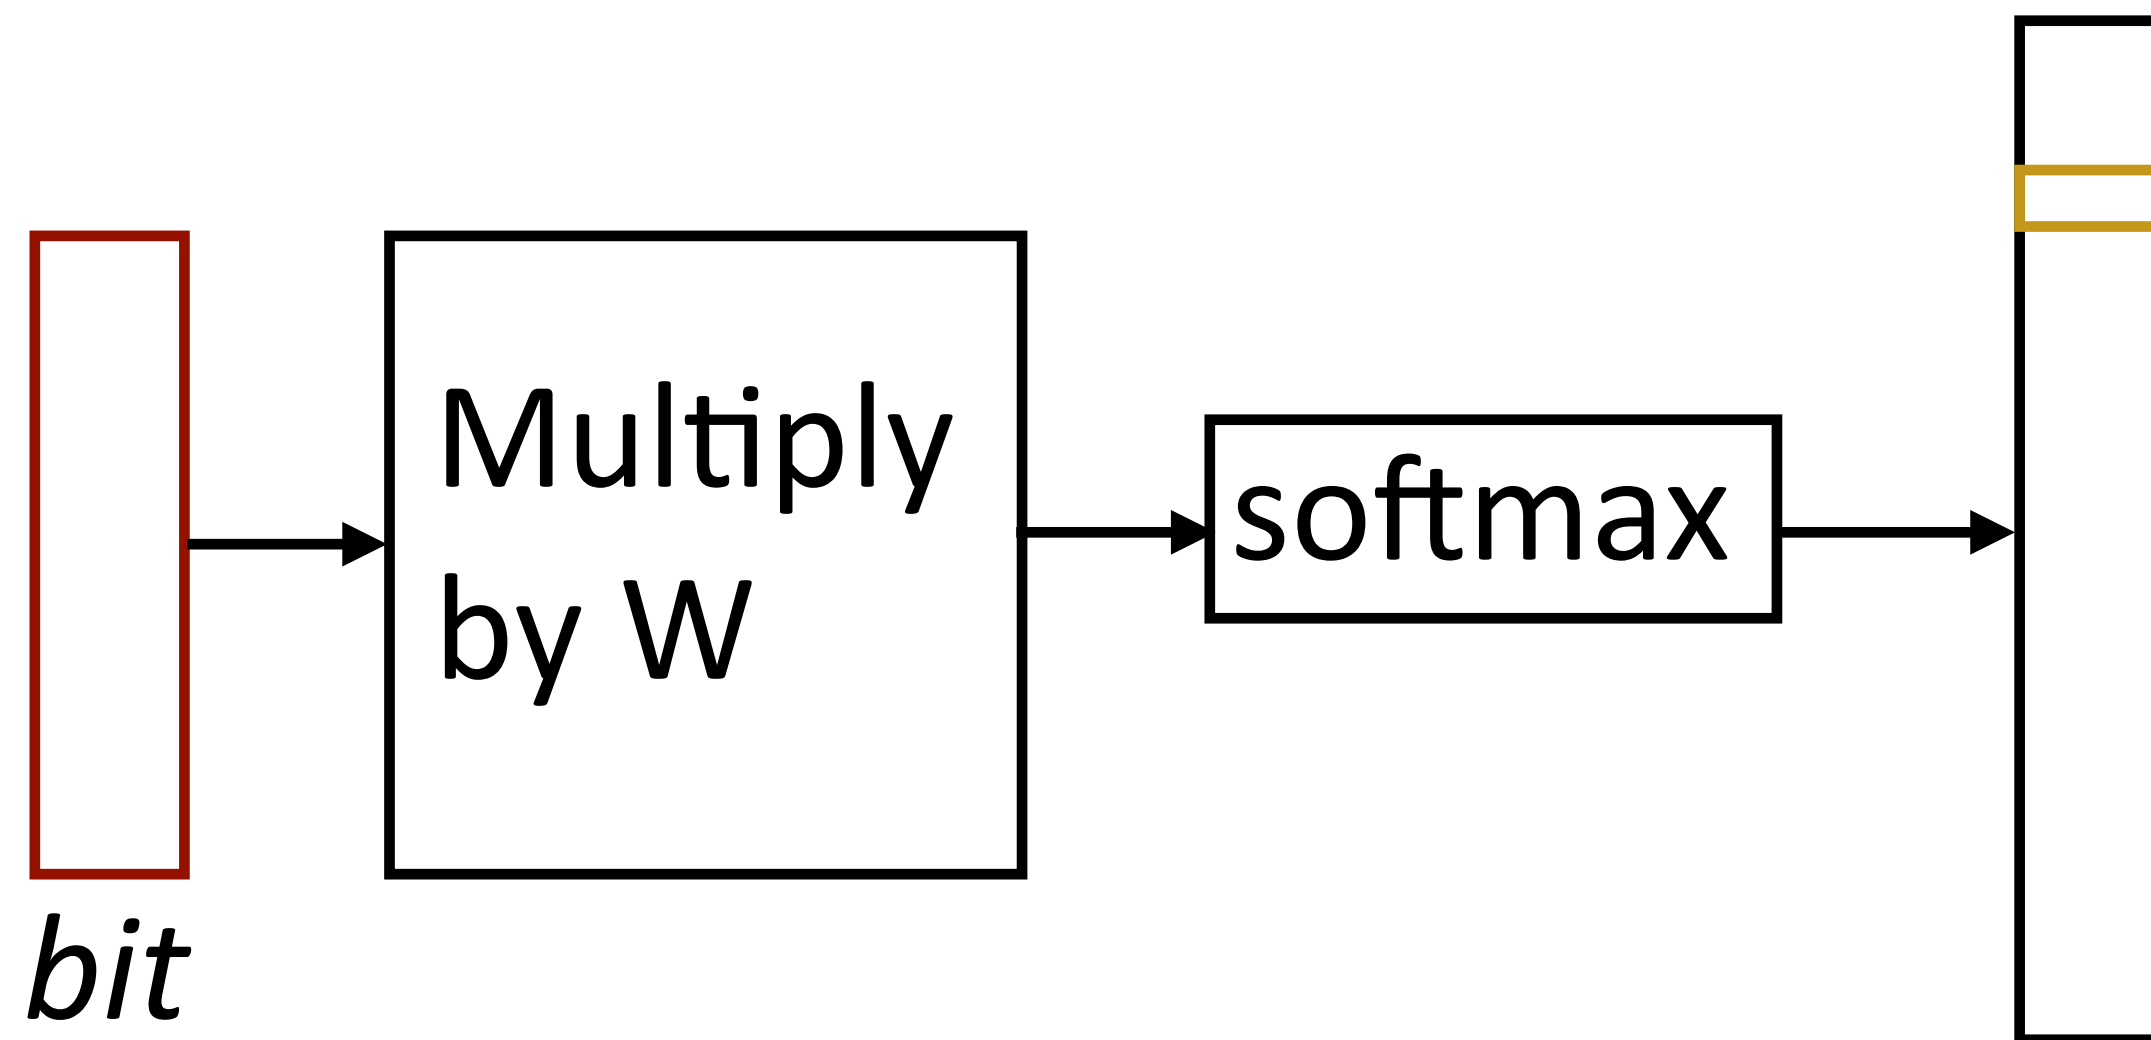
*the* *dog* *bit* *the* *man*



# Skip-Gram

- Predict one word of context from word

*the* *dog* *bit* *the* *man*



gold = *dog*

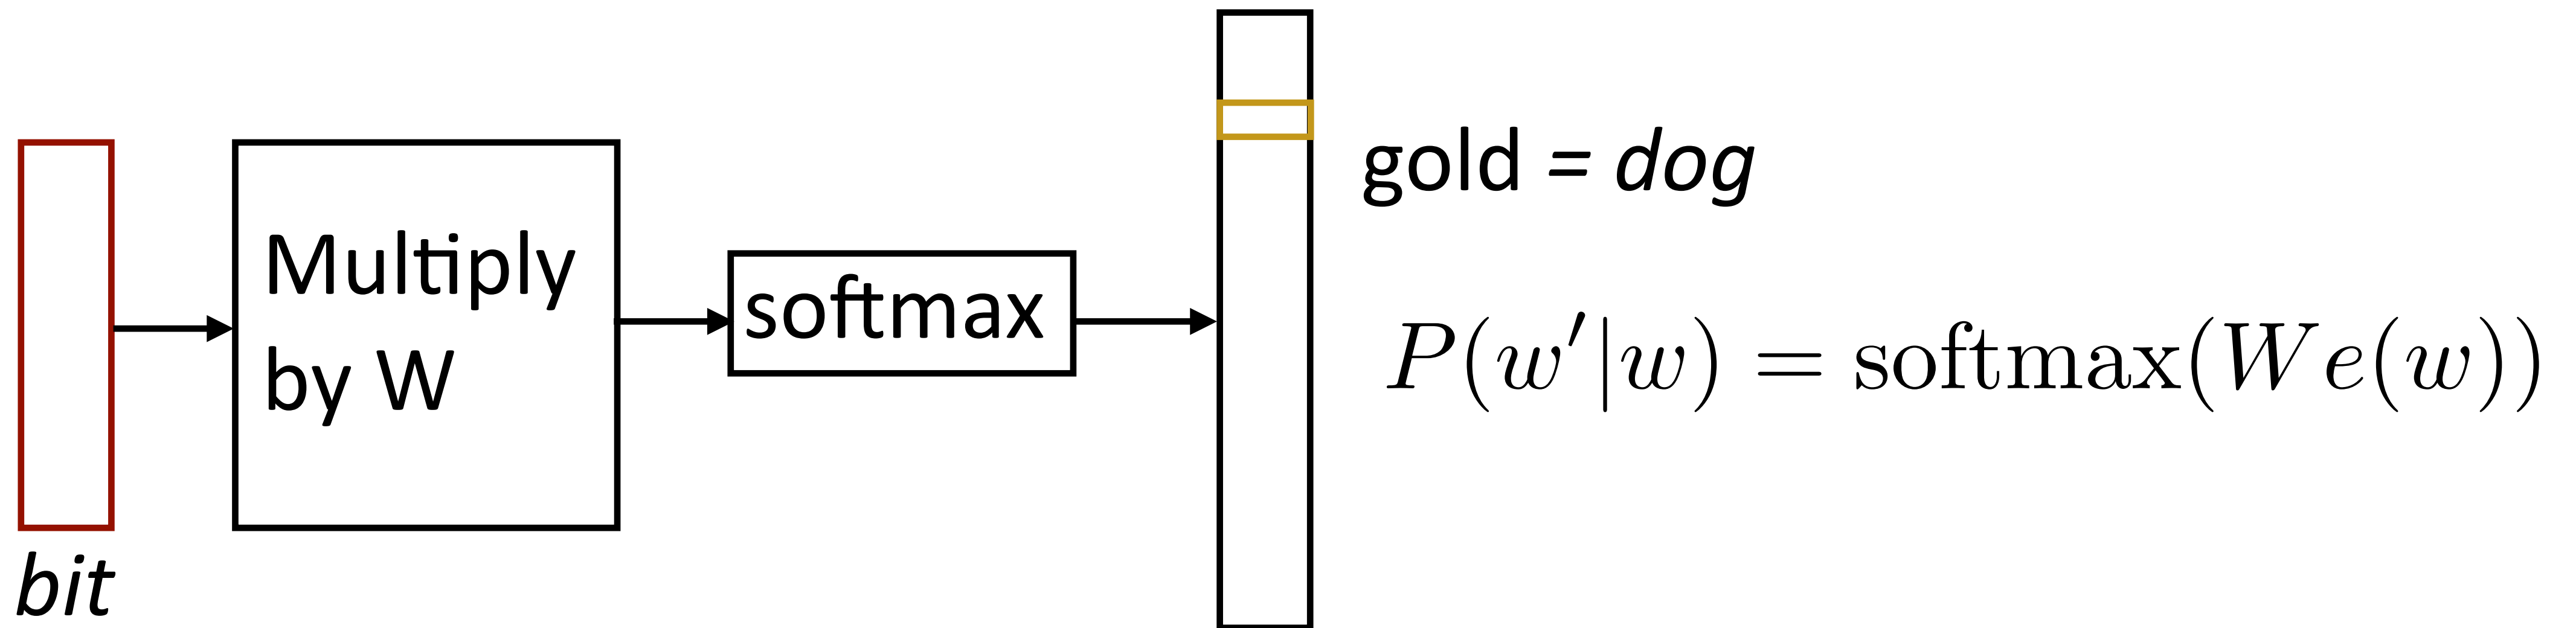
$$P(w'|w) = \text{softmax}(We(w))$$

- Another training example: *bit* -> *the*

# Skip-Gram

- Predict one word of context from word

*the* *dog* *bit* *the* *man*



- Another training example: *bit* -> *the*
- Parameters:  $d \times |V|$  **vectors**,  $|V| \times d$  output parameters ( $W$ ) (also usable as vectors!)

# Hierarchical Softmax

---

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

# Hierarchical Softmax

---

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG

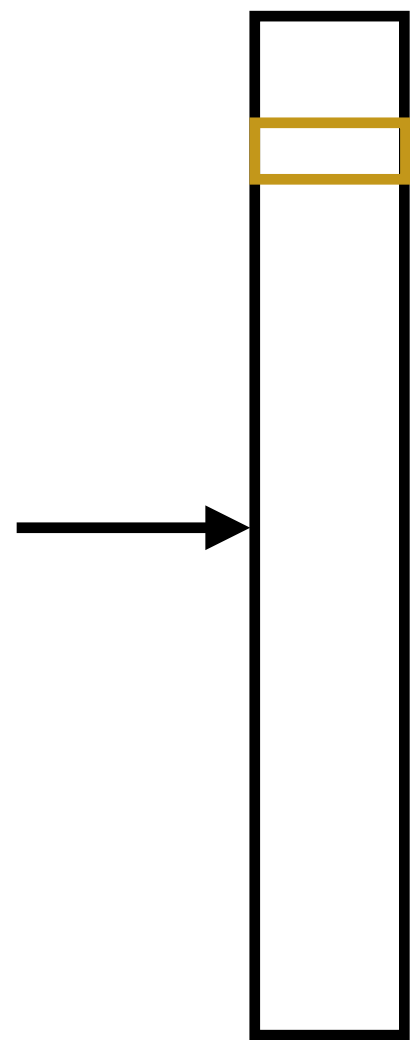


# Hierarchical Softmax

---

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG



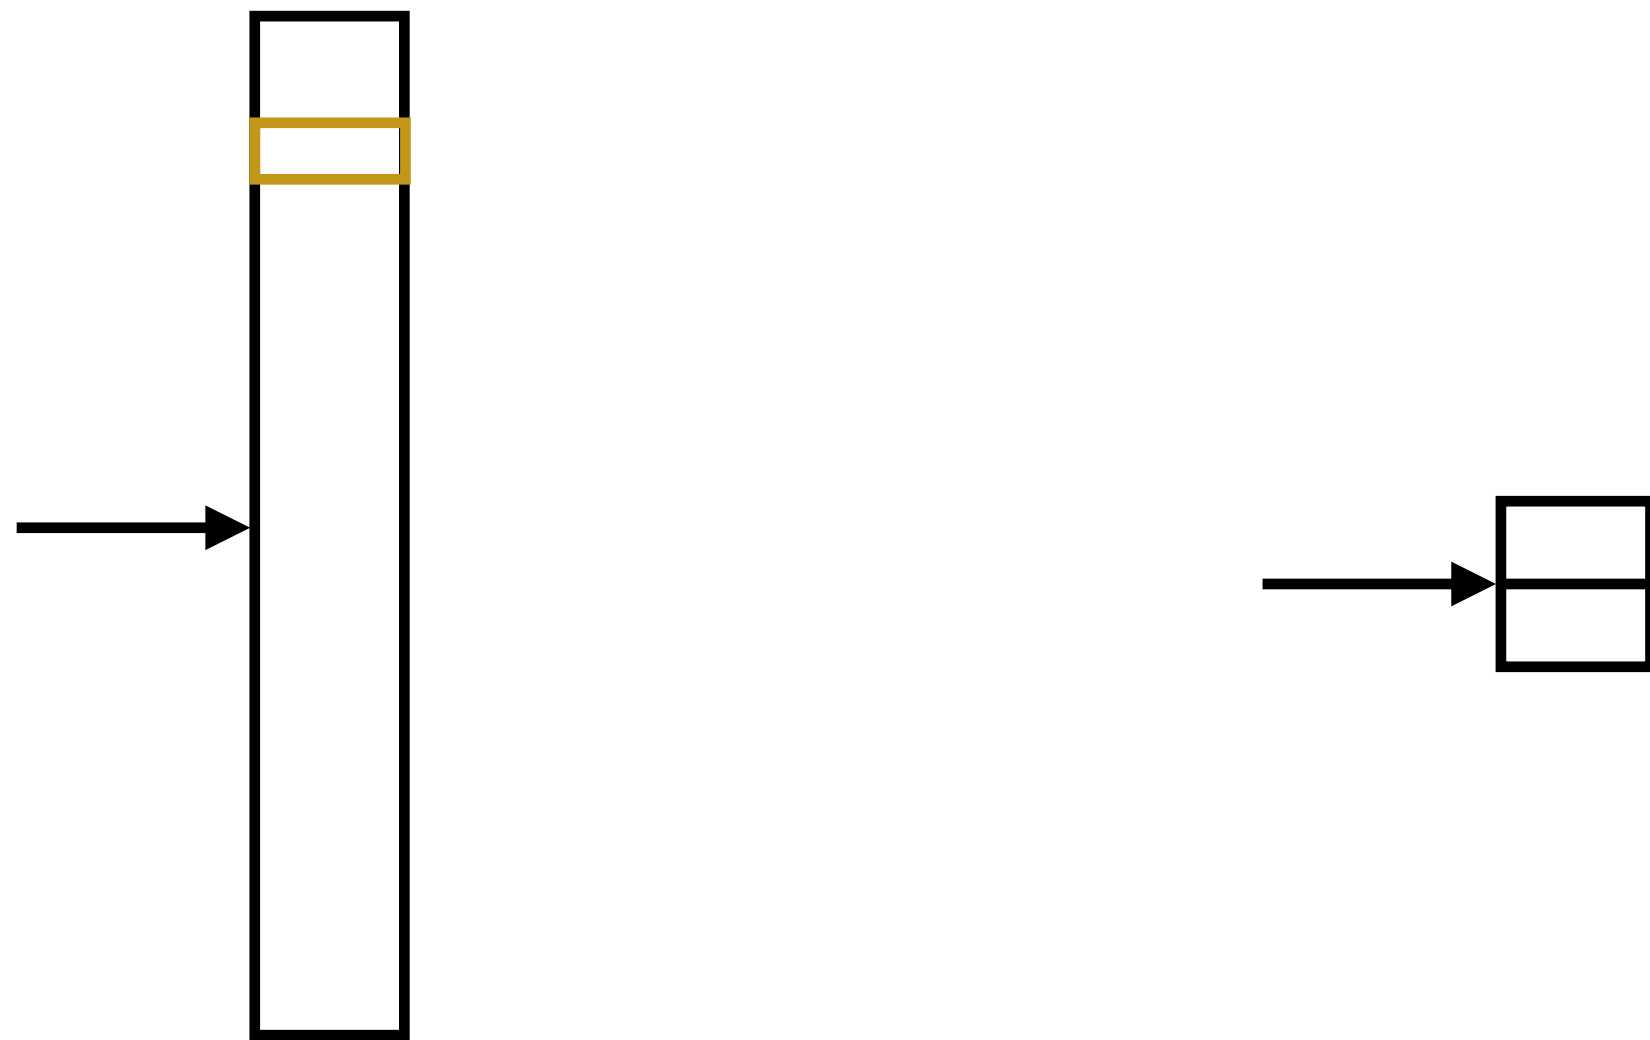
- ▶ Standard softmax:  
 $[|V| \times d] \times d$

# Hierarchical Softmax

---

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG




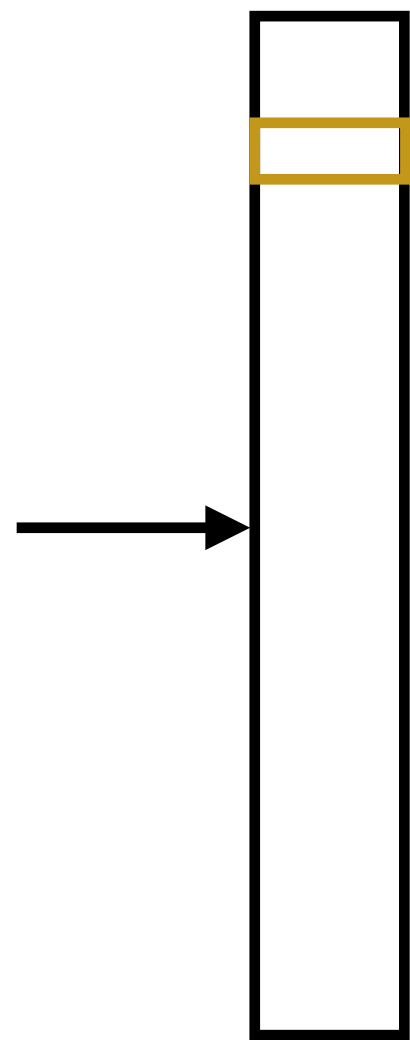
- ▶ Standard softmax:  
 $[|V| \times d] \times d$

# Hierarchical Softmax

---

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG

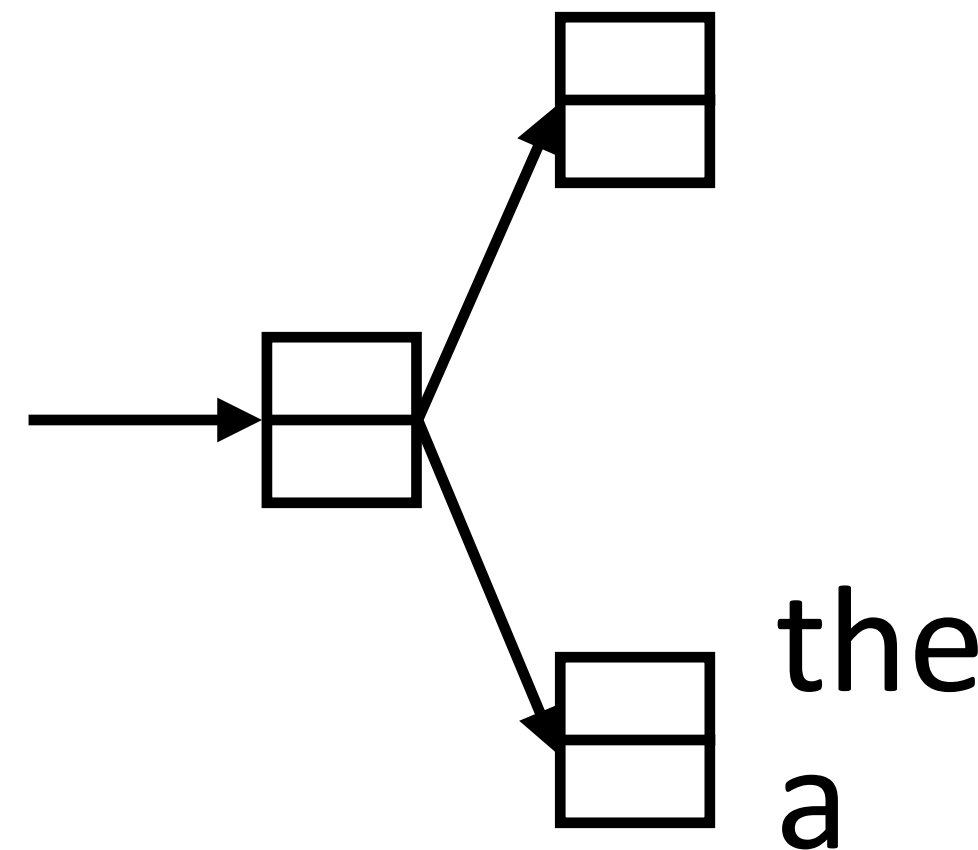
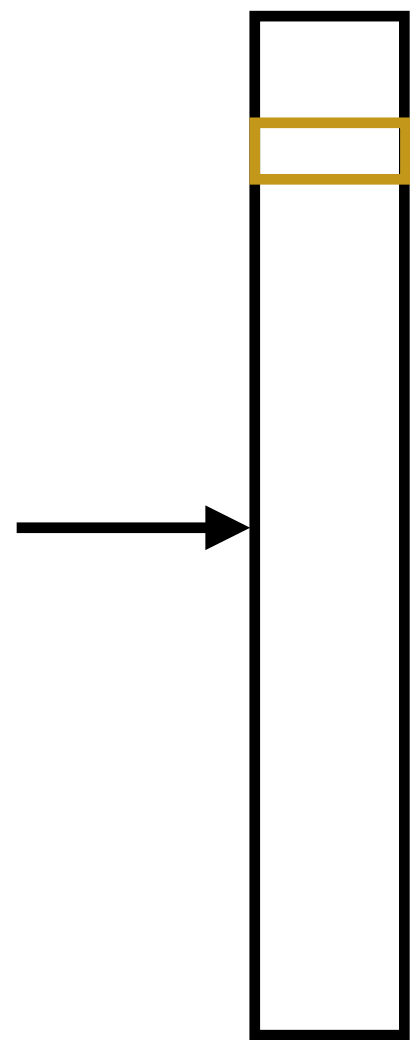


- ▶ Standard softmax:  
 $[|V| \times d] \times d$

# Hierarchical Softmax

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG

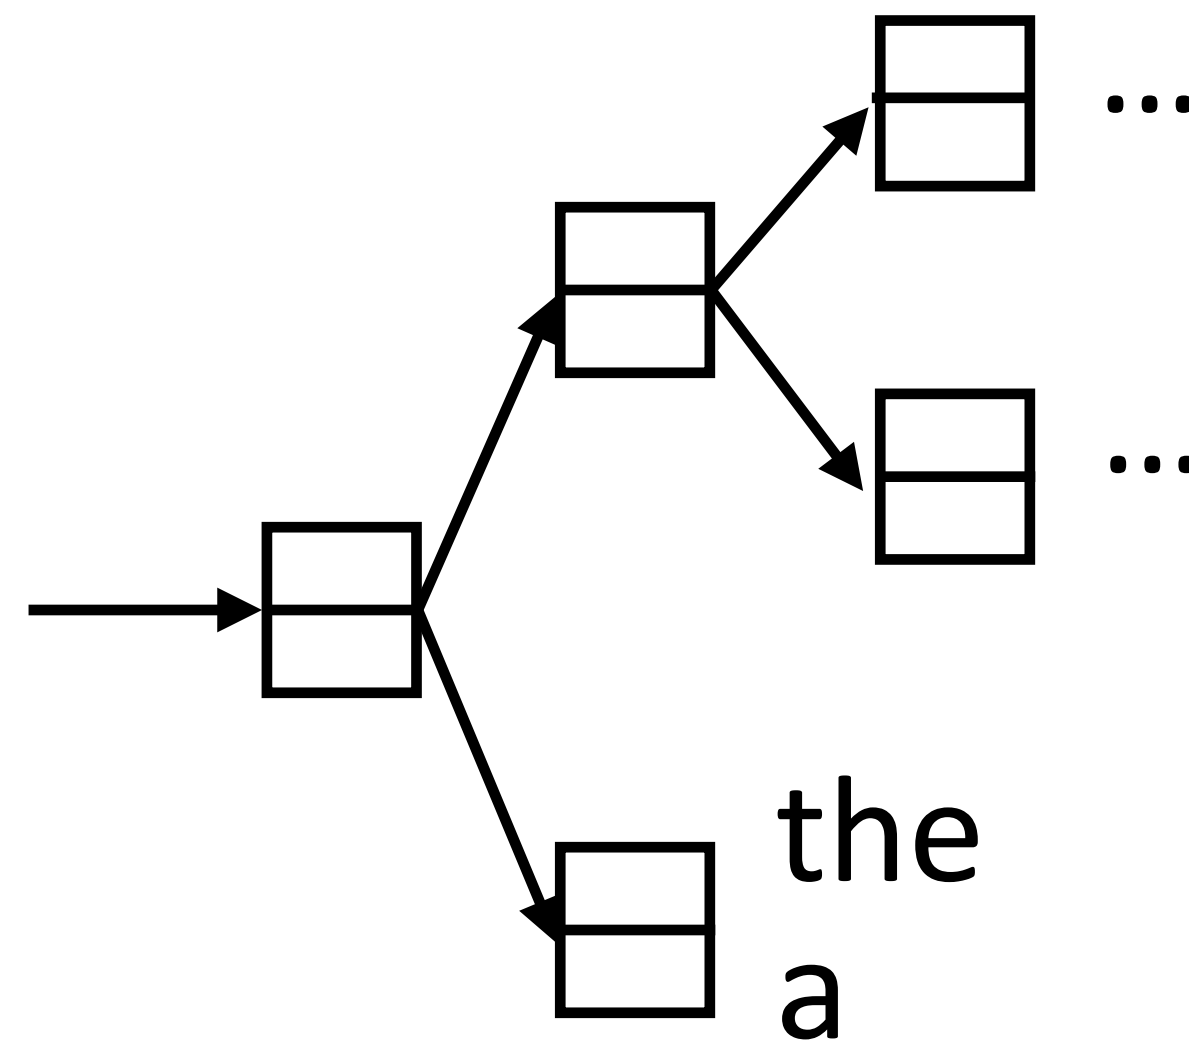



- ▶ Standard softmax:  
 $[|V| \times d] \times d$

# Hierarchical Softmax

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG

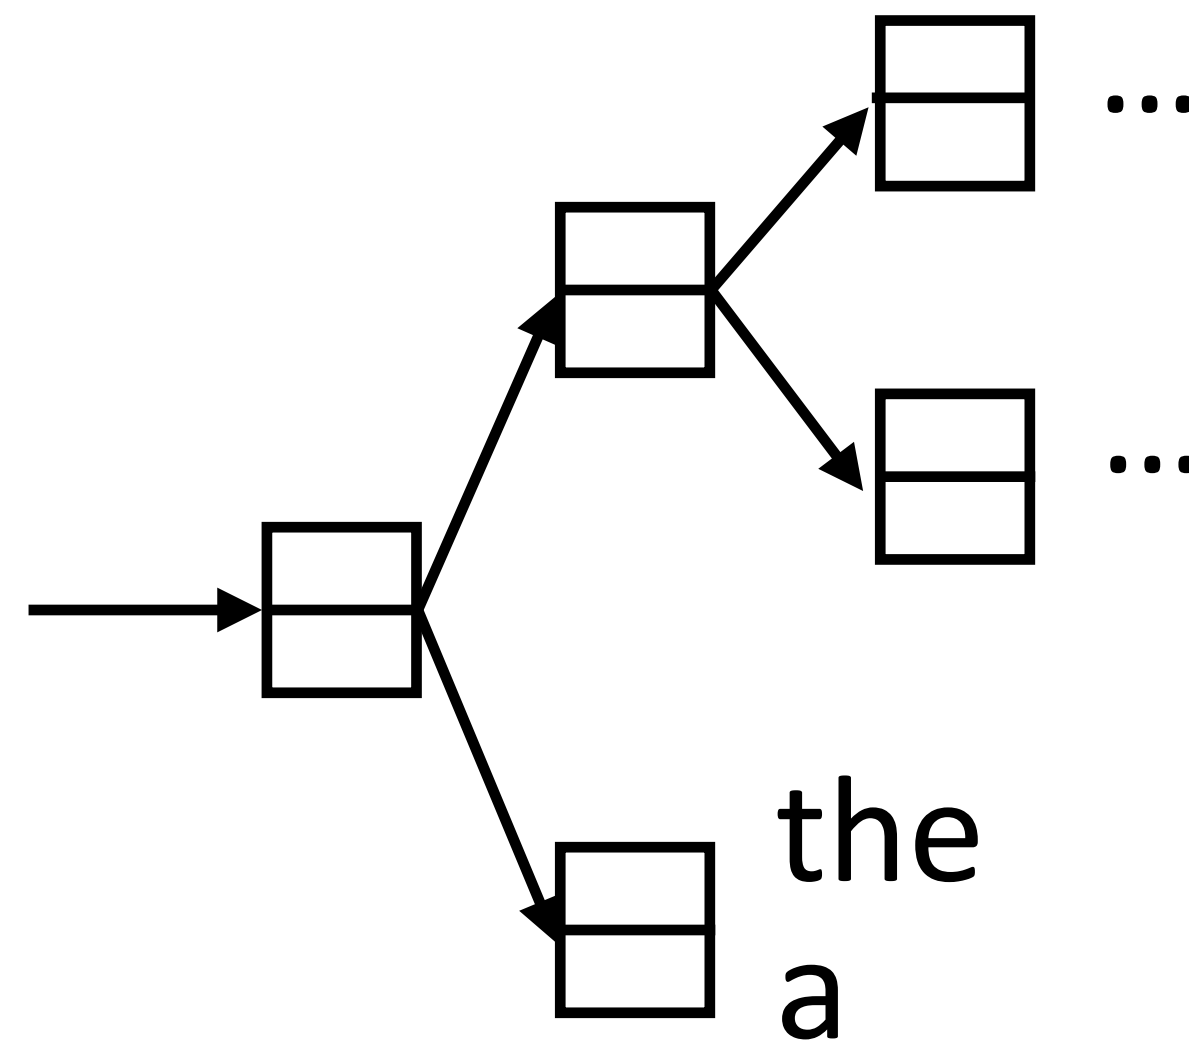
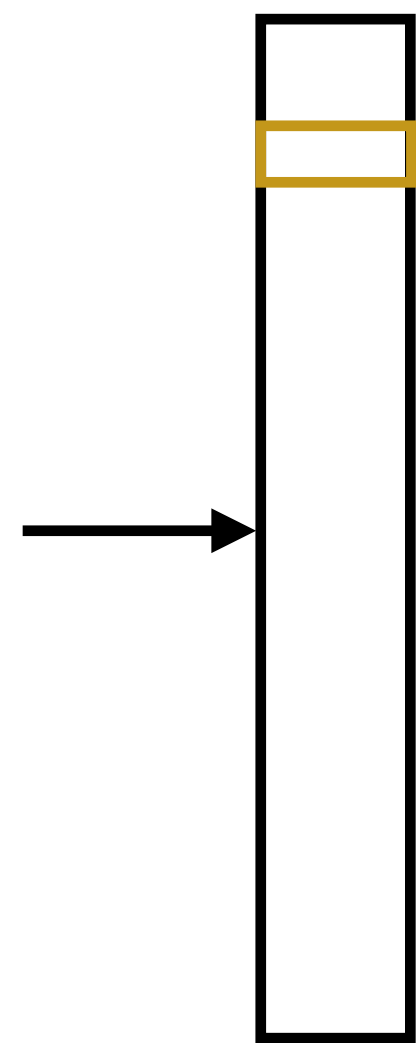


- ▶ Standard softmax:  
 $[|V| \times d] \times d$

# Hierarchical Softmax

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG



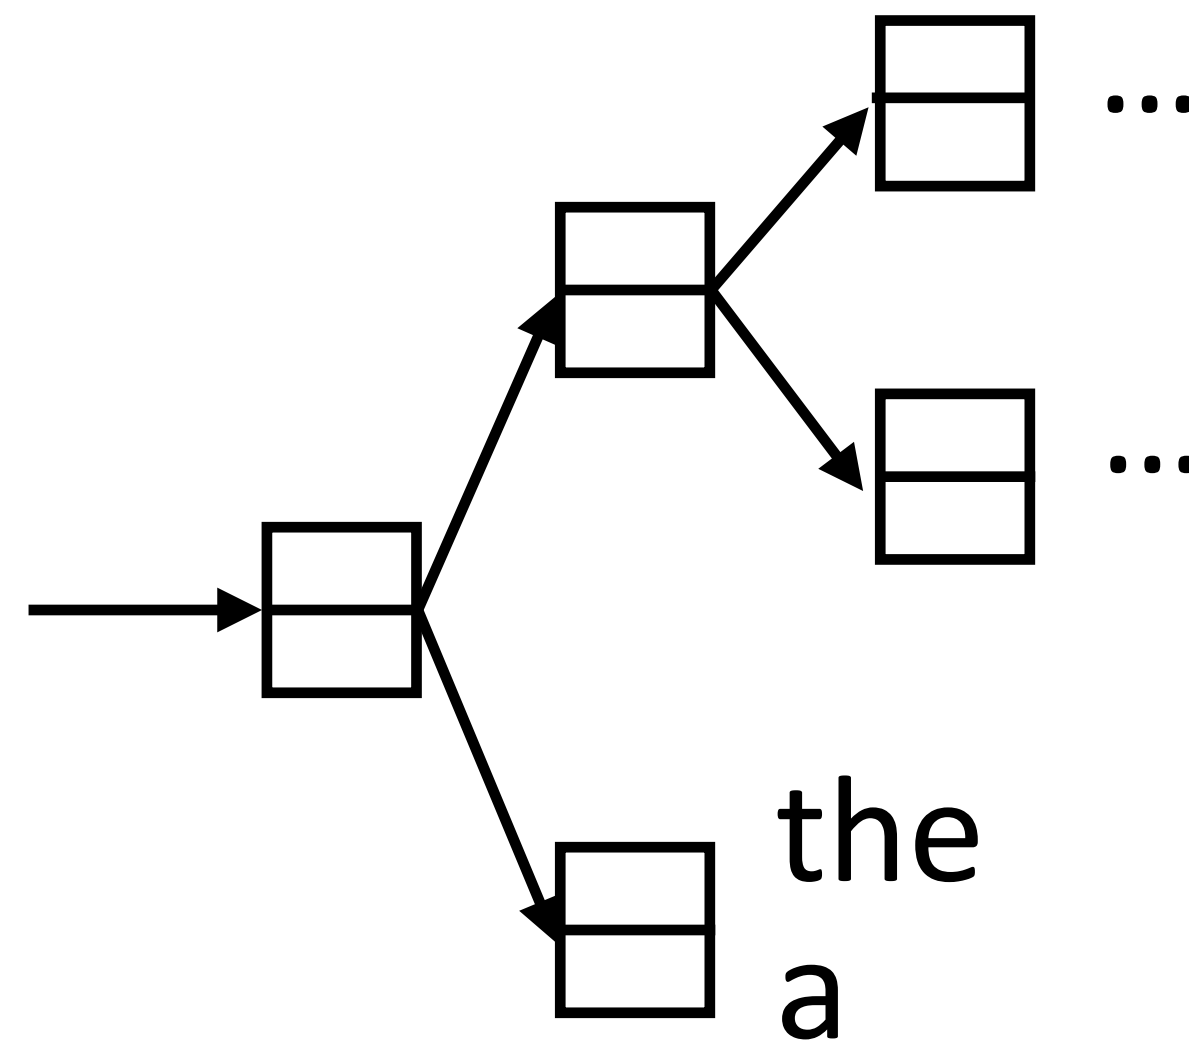
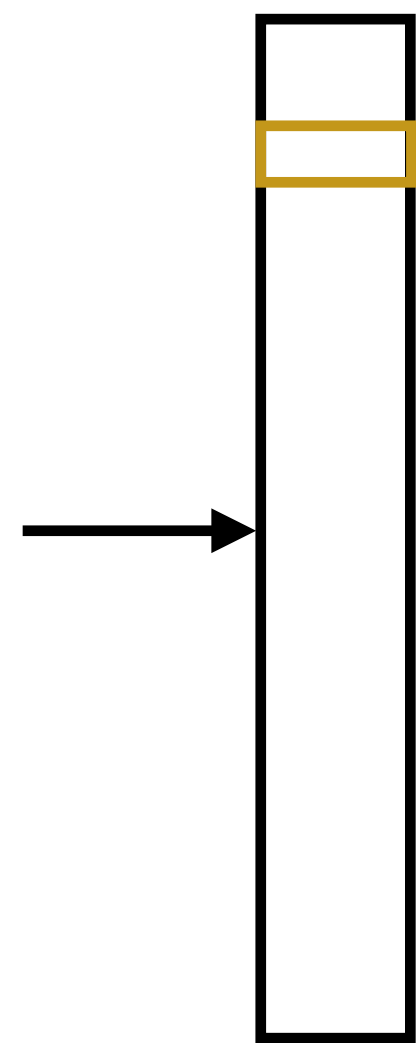
- ▶ Huffman encode vocabulary, use binary classifiers to decide which branch to take

- ▶ Standard softmax:  
 $[|V| \times d] \times d$

# Hierarchical Softmax

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG



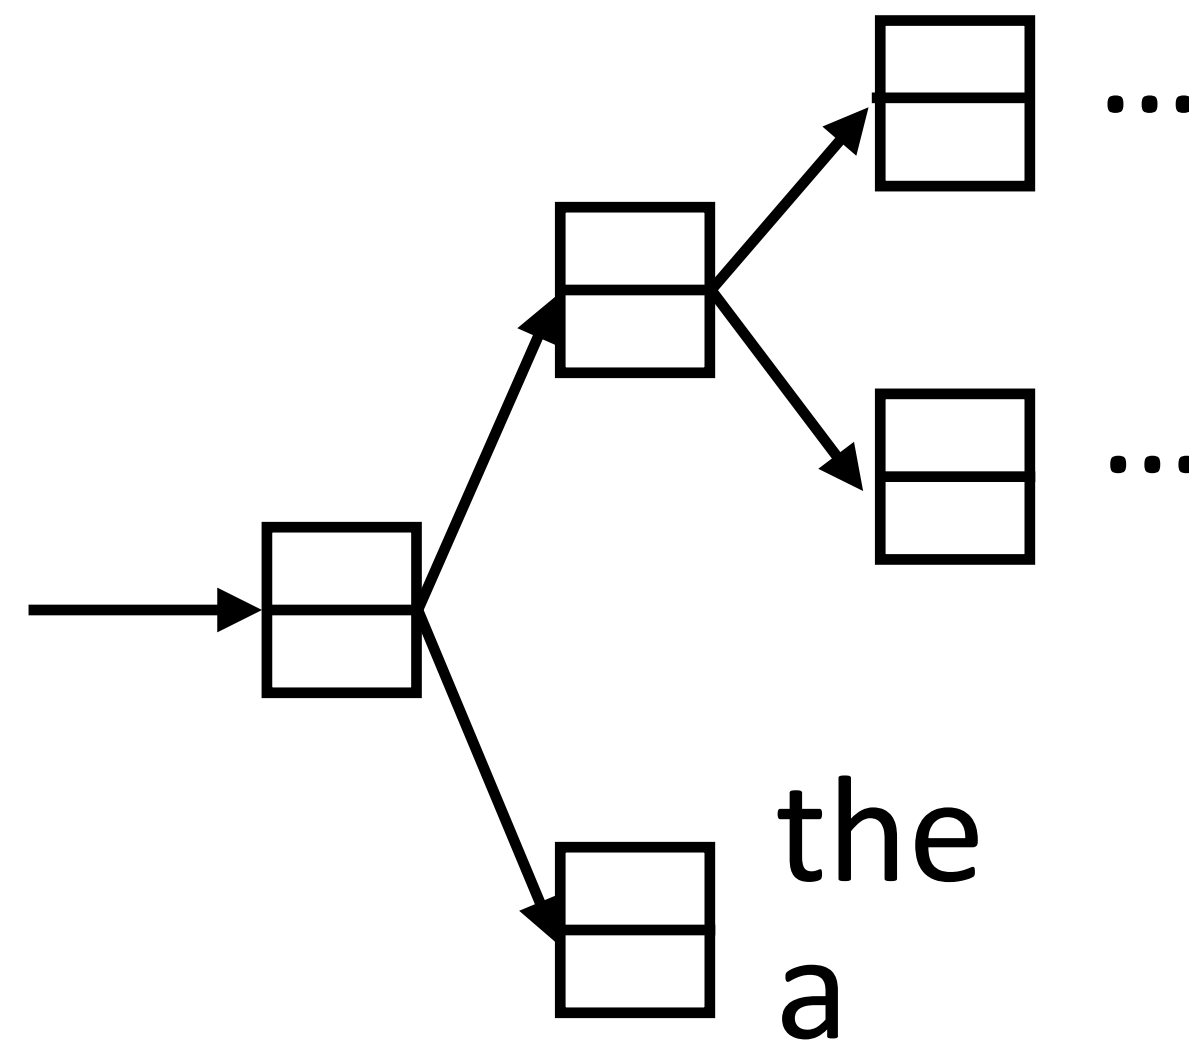
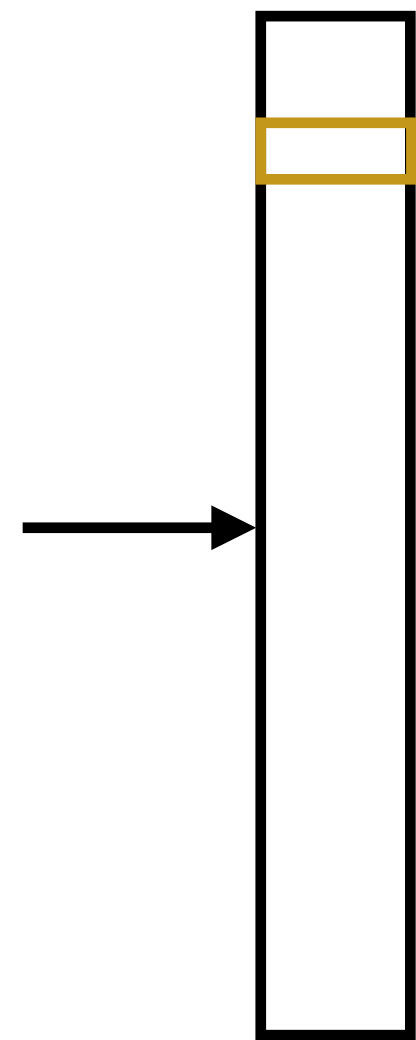
- ▶ Huffman encode vocabulary, use binary classifiers to decide which branch to take
- ▶  $\log(|V|)$  binary decisions

- ▶ Standard softmax:  
 $[|V| \times d] \times d$

# Hierarchical Softmax

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG



- ▶ Standard softmax:  
 $[|V| \times d] \times d$

- ▶ Hierarchical softmax:

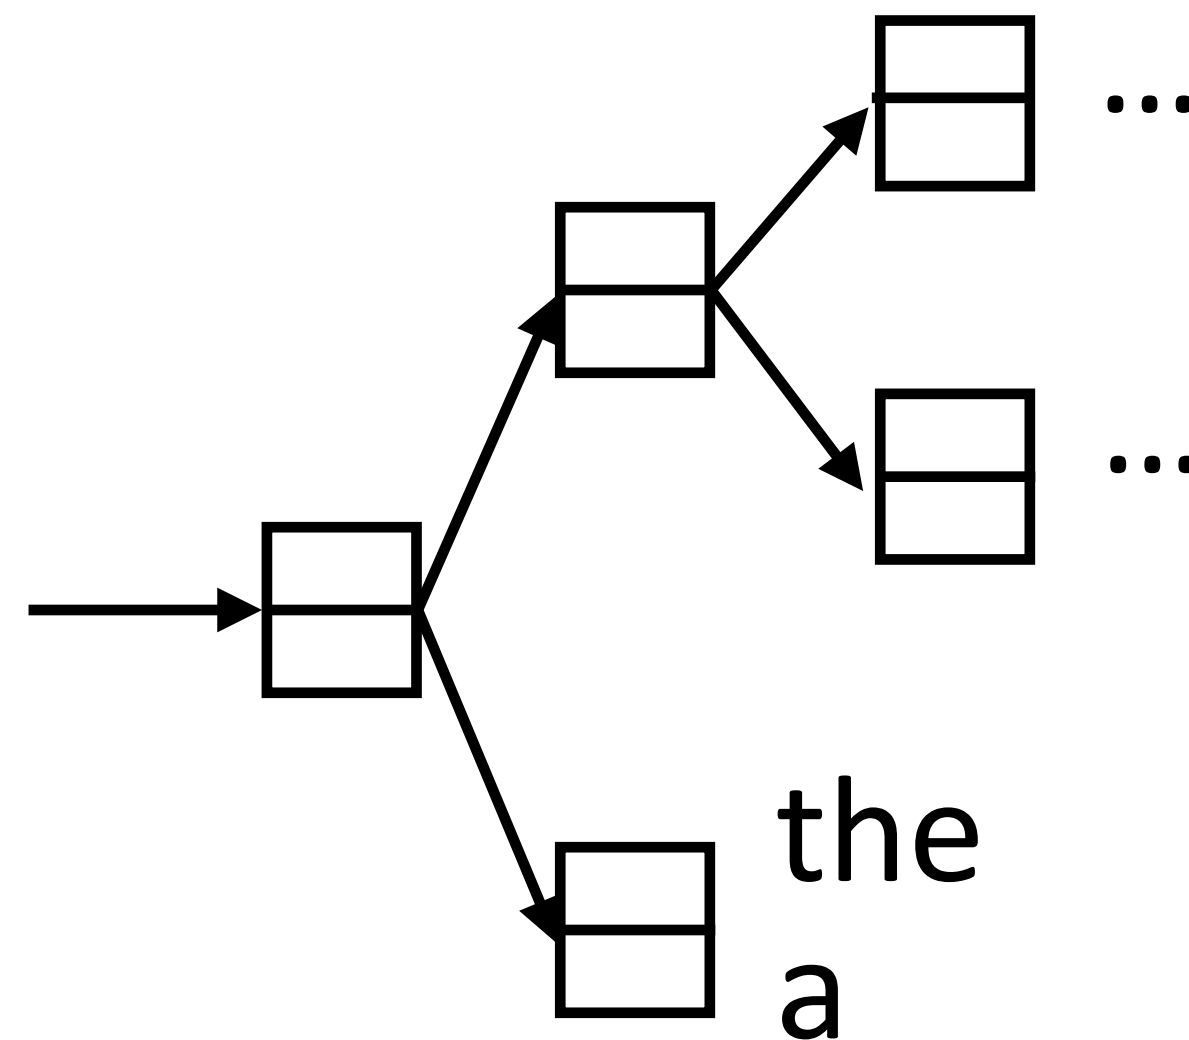
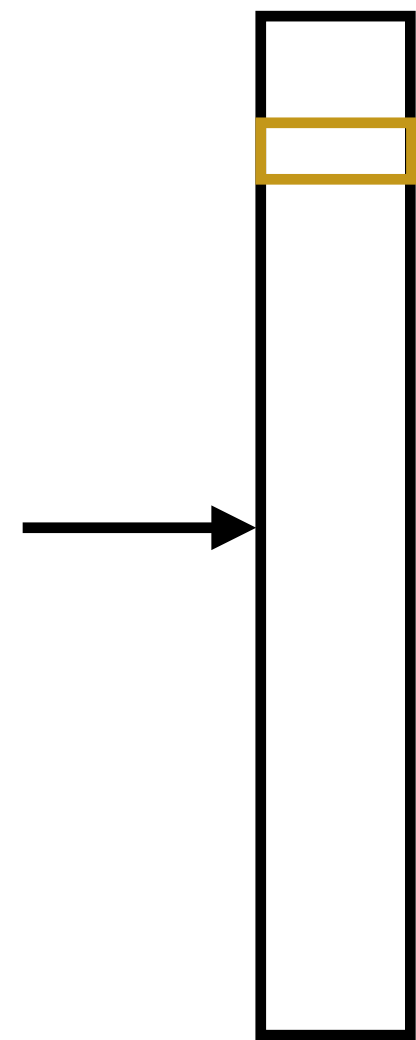
- ▶ Huffman encode vocabulary, use binary classifiers to decide which branch to take
- ▶  $\log(|V|)$  binary decisions



# Hierarchical Softmax

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG



- ▶ Huffman encode vocabulary, use binary classifiers to decide which branch to take
- ▶  $\log(|V|)$  binary decisions

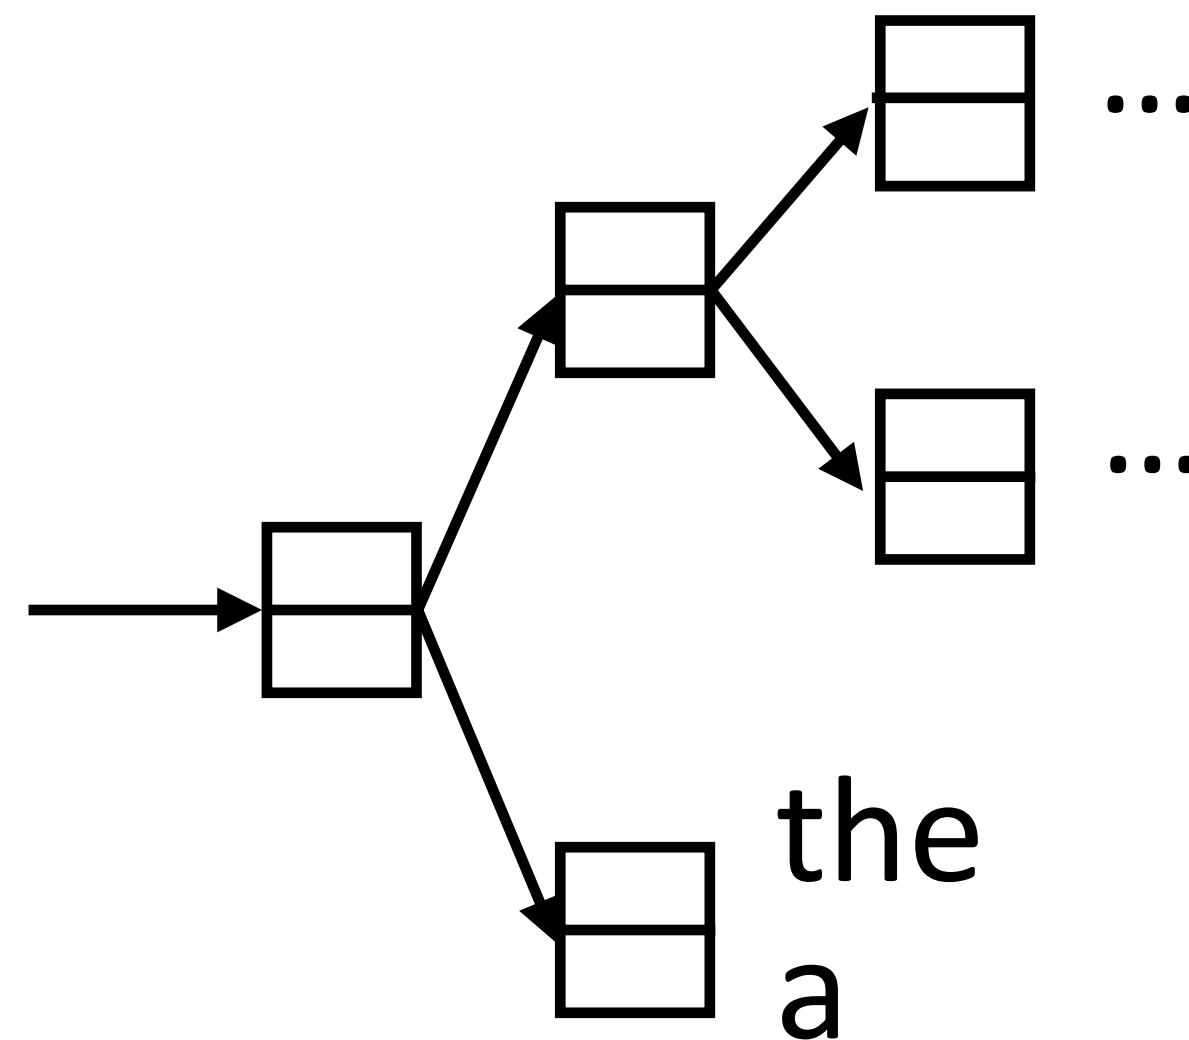
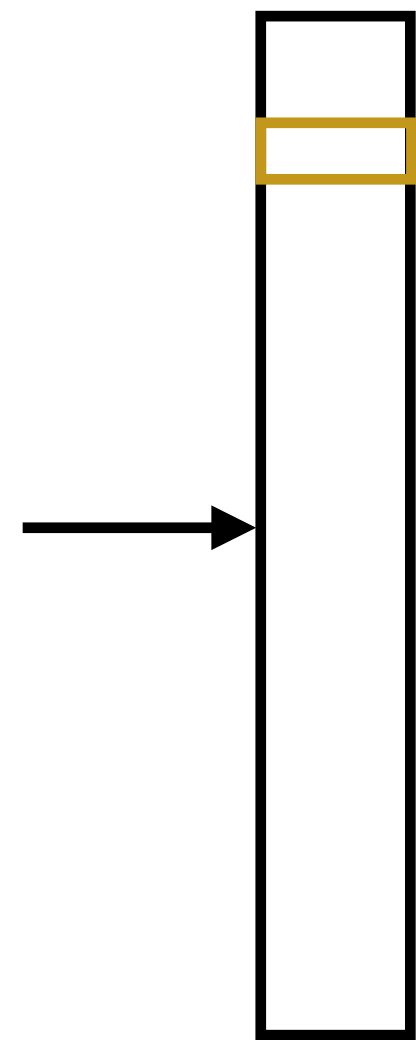
- ▶ Standard softmax:  
 $[|V| \times d] \times d$

- ▶ Hierarchical softmax:  
 $\log(|V|)$  dot products of size  $d$ ,

# Hierarchical Softmax

$$P(w|w_{-1}, w_{+1}) = \text{softmax}(W(c(w_{-1}) + c(w_{+1}))) \quad P(w'|w) = \text{softmax}(We(w))$$

- ▶ Matmul + softmax over  $|V|$  is very slow to compute for CBOW and SG



- ▶ Huffman encode vocabulary, use binary classifiers to decide which branch to take
- ▶  $\log(|V|)$  binary decisions

- ▶ Standard softmax:  
 $[|V| \times d] \times d$

- ▶ Hierarchical softmax:  
 $\log(|V|)$  dot products of size  $d$ ,  
 $|V| \times d$  parameters

Mikolov et al. (2013)

# Skip-Gram with Negative Sampling

---

# Skip-Gram with Negative Sampling

---

- ▶ Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

# Skip-Gram with Negative Sampling

---

- ▶ Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the)* => +1

# Skip-Gram with Negative Sampling

---

- ▶ Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the)* => +1

*(bit, cat)* => -1

# Skip-Gram with Negative Sampling

---

- ▶ Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the) => +1*

*(bit, cat) => -1*

*(bit, a) => -1*

*(bit, fish) => -1*

# Skip-Gram with Negative Sampling

---

- ▶ Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the) => +1*

*(bit, cat) => -1*

*(bit, a) => -1*

*(bit, fish) => -1*

$$P(y = 1|w, c) = \frac{e^{w \cdot c}}{e^{w \cdot c} + 1}$$



# Skip-Gram with Negative Sampling

---

- Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the)* => +1

*(bit, cat)* => -1

*(bit, a)* => -1

*(bit, fish)* => -1

$$P(y = 1|w, c) = \frac{e^{w \cdot c}}{e^{w \cdot c} + 1}$$

words in similar contexts select for similar  $c$  vectors

# Skip-Gram with Negative Sampling

---

- Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the)* => +1

*(bit, cat)* => -1

*(bit, a)* => -1

*(bit, fish)* => -1

$$P(y = 1|w, c) = \frac{e^{w \cdot c}}{e^{w \cdot c} + 1}$$

words in similar contexts select for similar  $c$  vectors

- $d \times |V|$  vectors,  $d \times |V|$  context vectors (same # of params as before)

# Skip-Gram with Negative Sampling

- Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the)* => +1

*(bit, cat)* => -1

*(bit, a)* => -1

*(bit, fish)* => -1

$$P(y = 1|w, c) = \frac{e^{w \cdot c}}{e^{w \cdot c} + 1}$$

words in similar contexts select for similar  $c$  vectors

- $d \times |V|$  vectors,  $d \times |V|$  context vectors (same # of params as before)

- Objective =  $\log P(y = 1|w, c) - \frac{1}{k} \sum_{i=1}^n \log P(y = 0|w_i, c)$

Mikolov et al. (2013)

# Skip-Gram with Negative Sampling

- Take (word, context) pairs and classify them as “real” or not. Create random negative examples by sampling from unigram distribution

*(bit, the)* => +1

*(bit, cat)* => -1

*(bit, a)* => -1

*(bit, fish)* => -1

$$P(y = 1|w, c) = \frac{e^{w \cdot c}}{e^{w \cdot c} + 1}$$

words in similar contexts select for similar  $c$  vectors

- $d \times |V|$  vectors,  $d \times |V|$  context vectors (same # of params as before)

- Objective =  $\log P(y = 1|w, c) - \frac{1}{k} \sum_{i=1}^n \log P(y = 0|w_i, c)$  sampled

Mikolov et al. (2013)

# Connections with Matrix Factorization

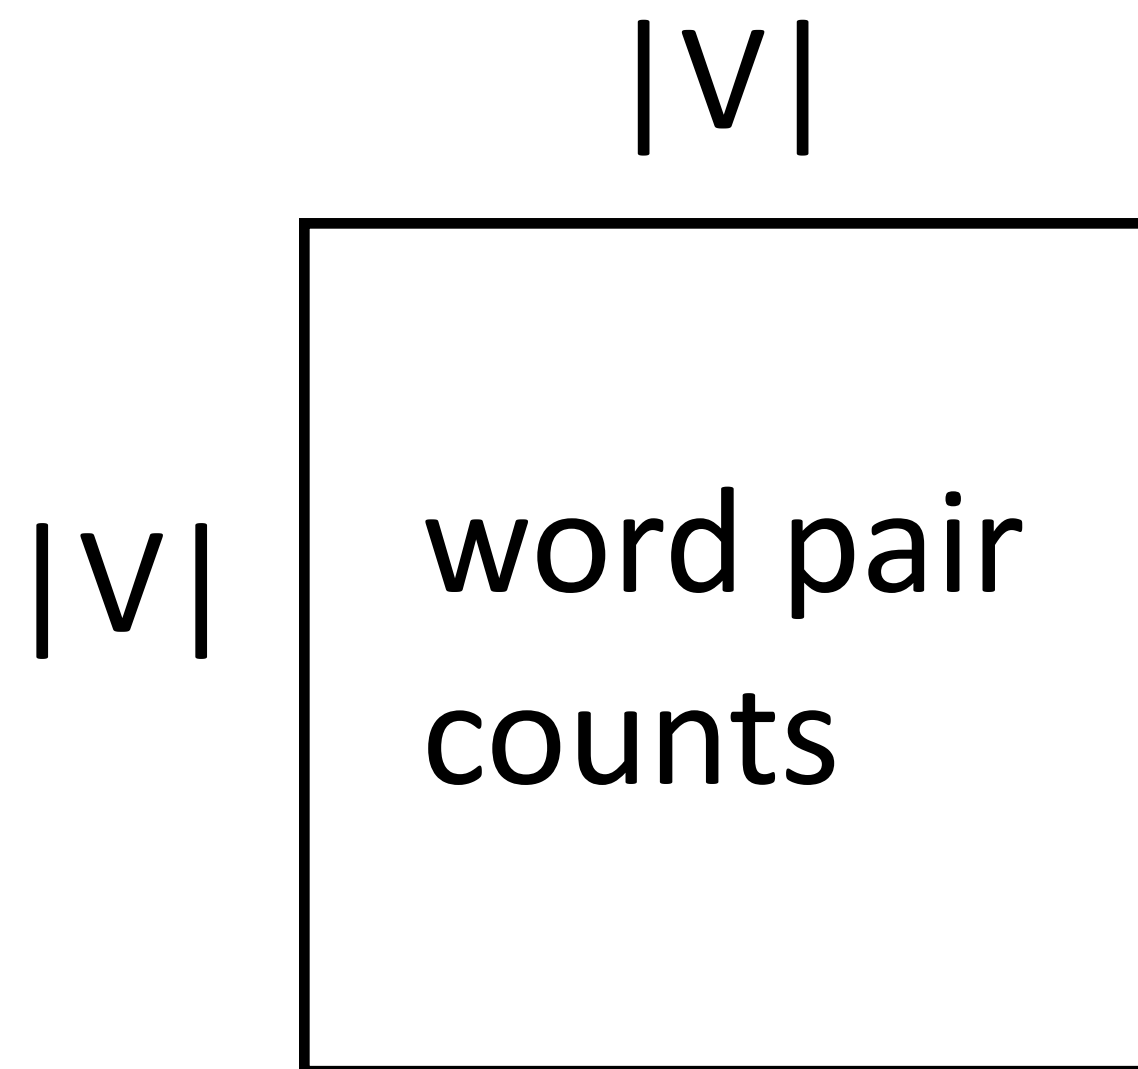
---

- ▶ Skip-gram model looks at word-word co-occurrences and produces two types of vectors

# Connections with Matrix Factorization

---


- ▶ Skip-gram model looks at word-word co-occurrences and produces two types of vectors



# Connections with Matrix Factorization

---

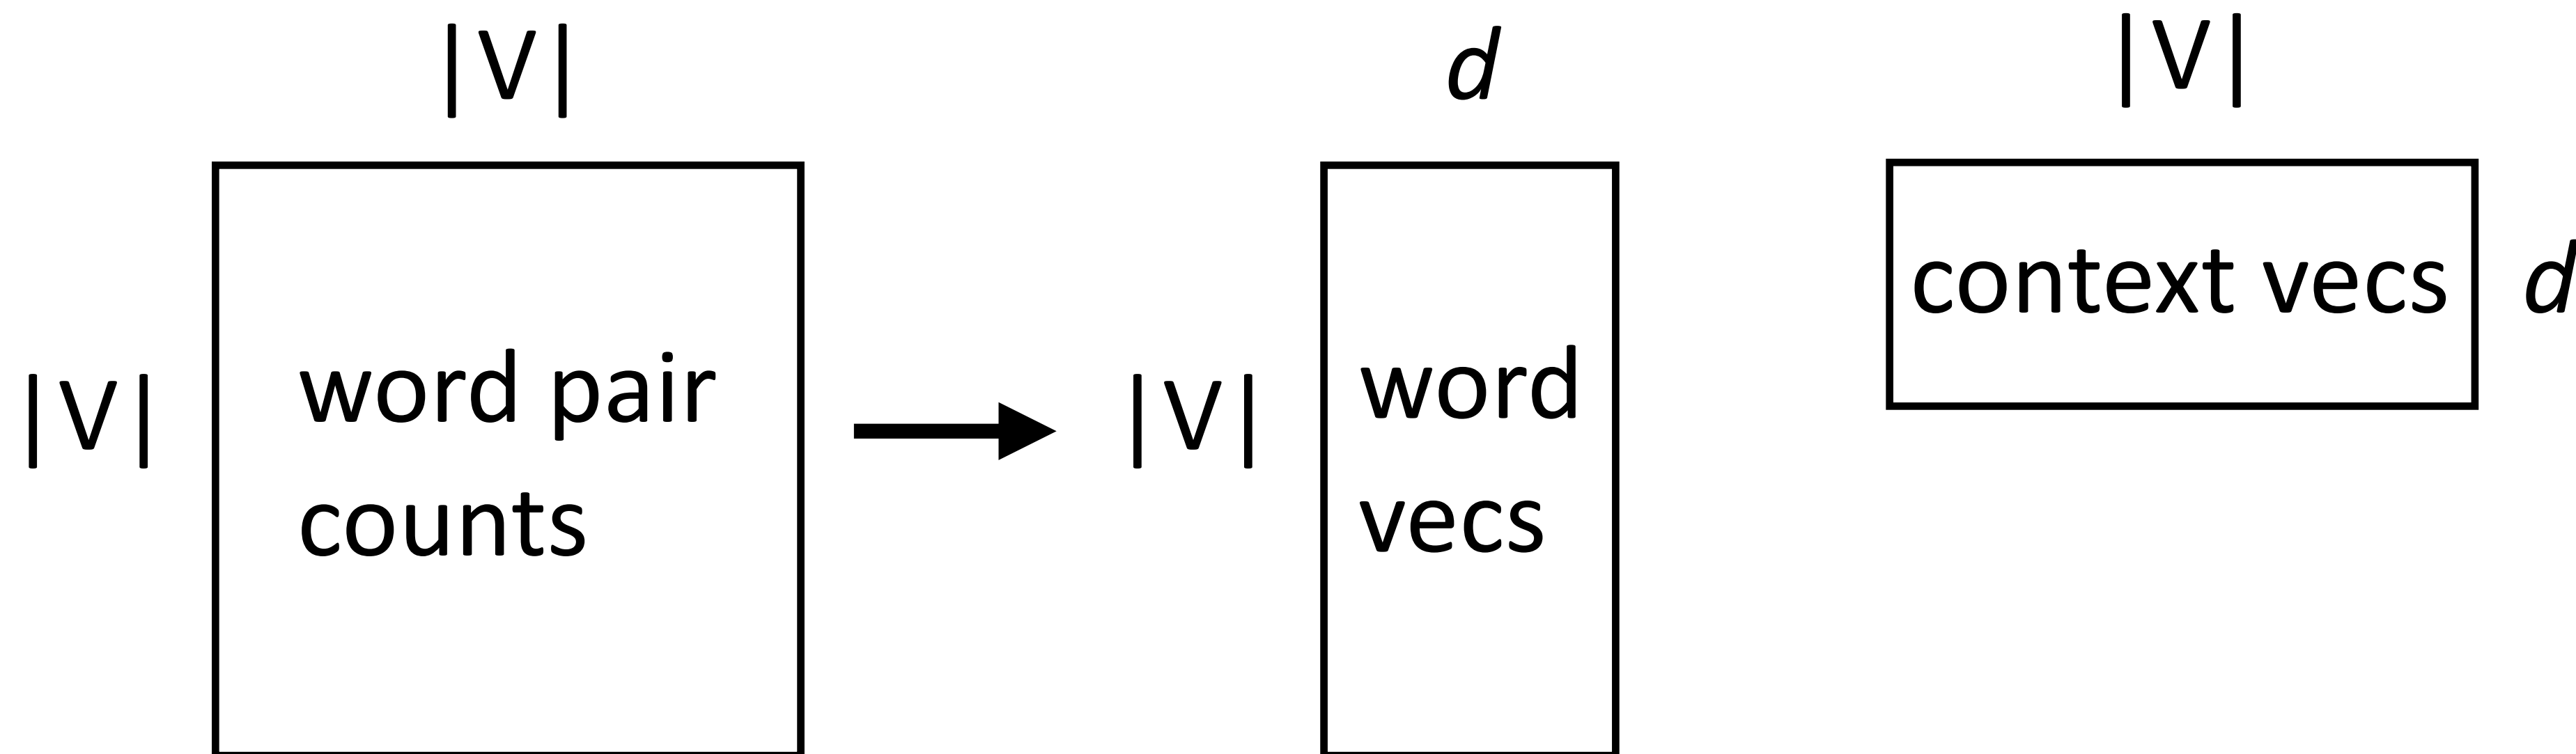
- ▶ Skip-gram model looks at word-word co-occurrences and produces two types of vectors



# Connections with Matrix Factorization

---

- ▶ Skip-gram model looks at word-word co-occurrences and produces two types of vectors

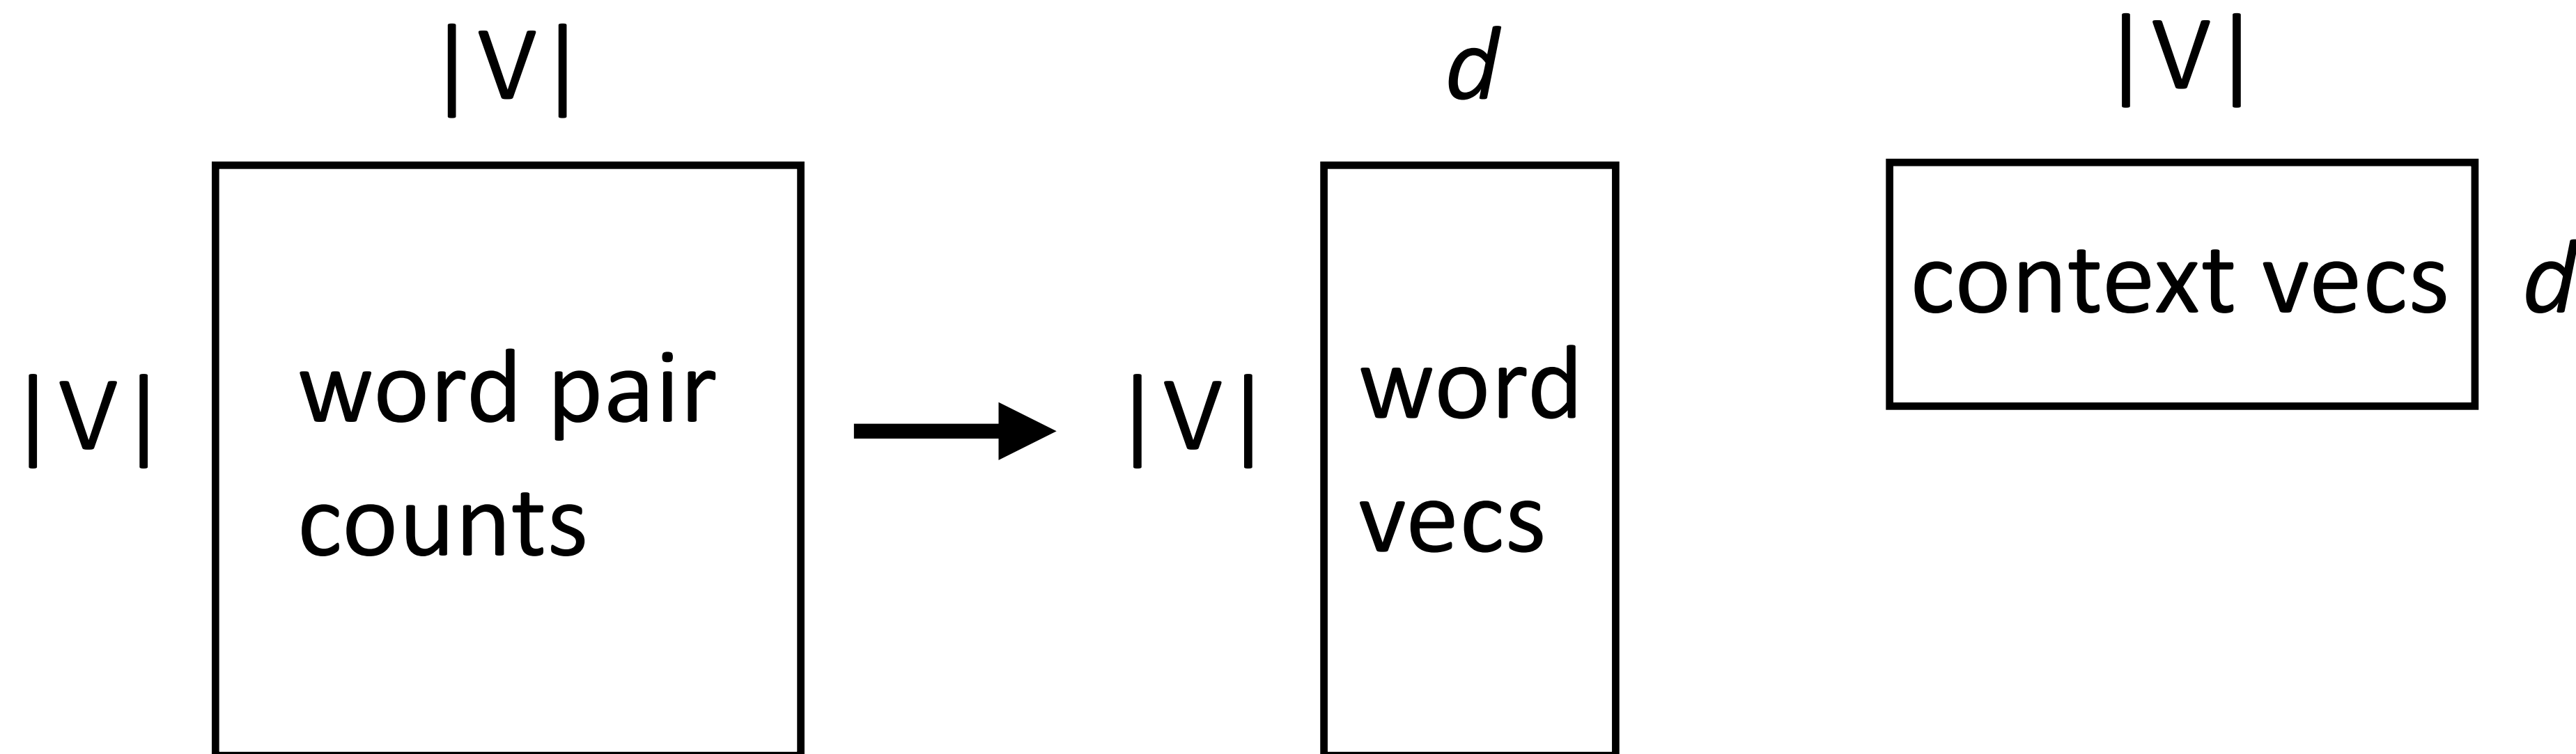




# Connections with Matrix Factorization

---

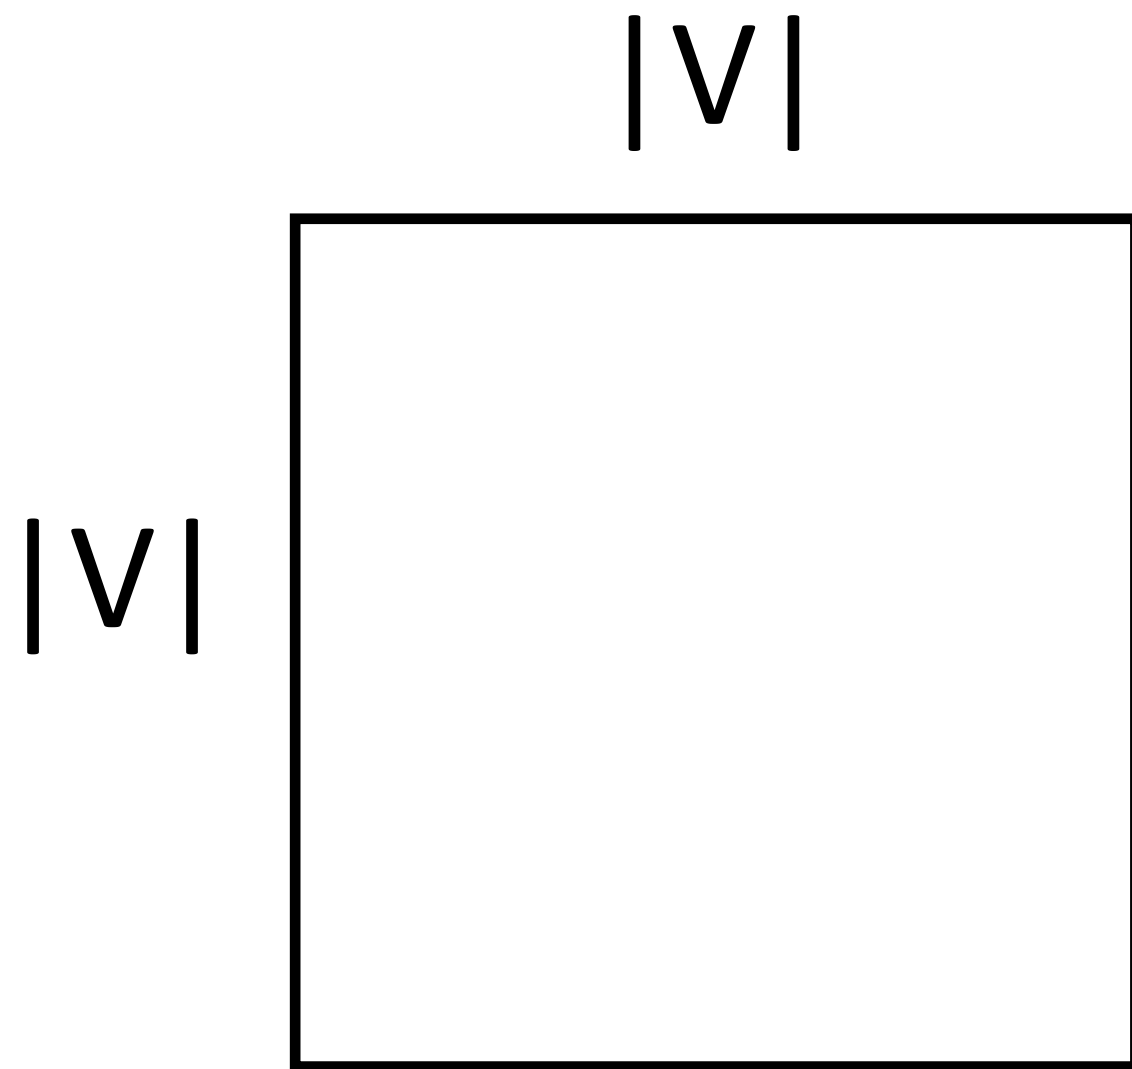
- ▶ Skip-gram model looks at word-word co-occurrences and produces two types of vectors



- ▶ Looks almost like a matrix factorization...can we interpret it this way?


# Skip-Gram as Matrix Factorization

---



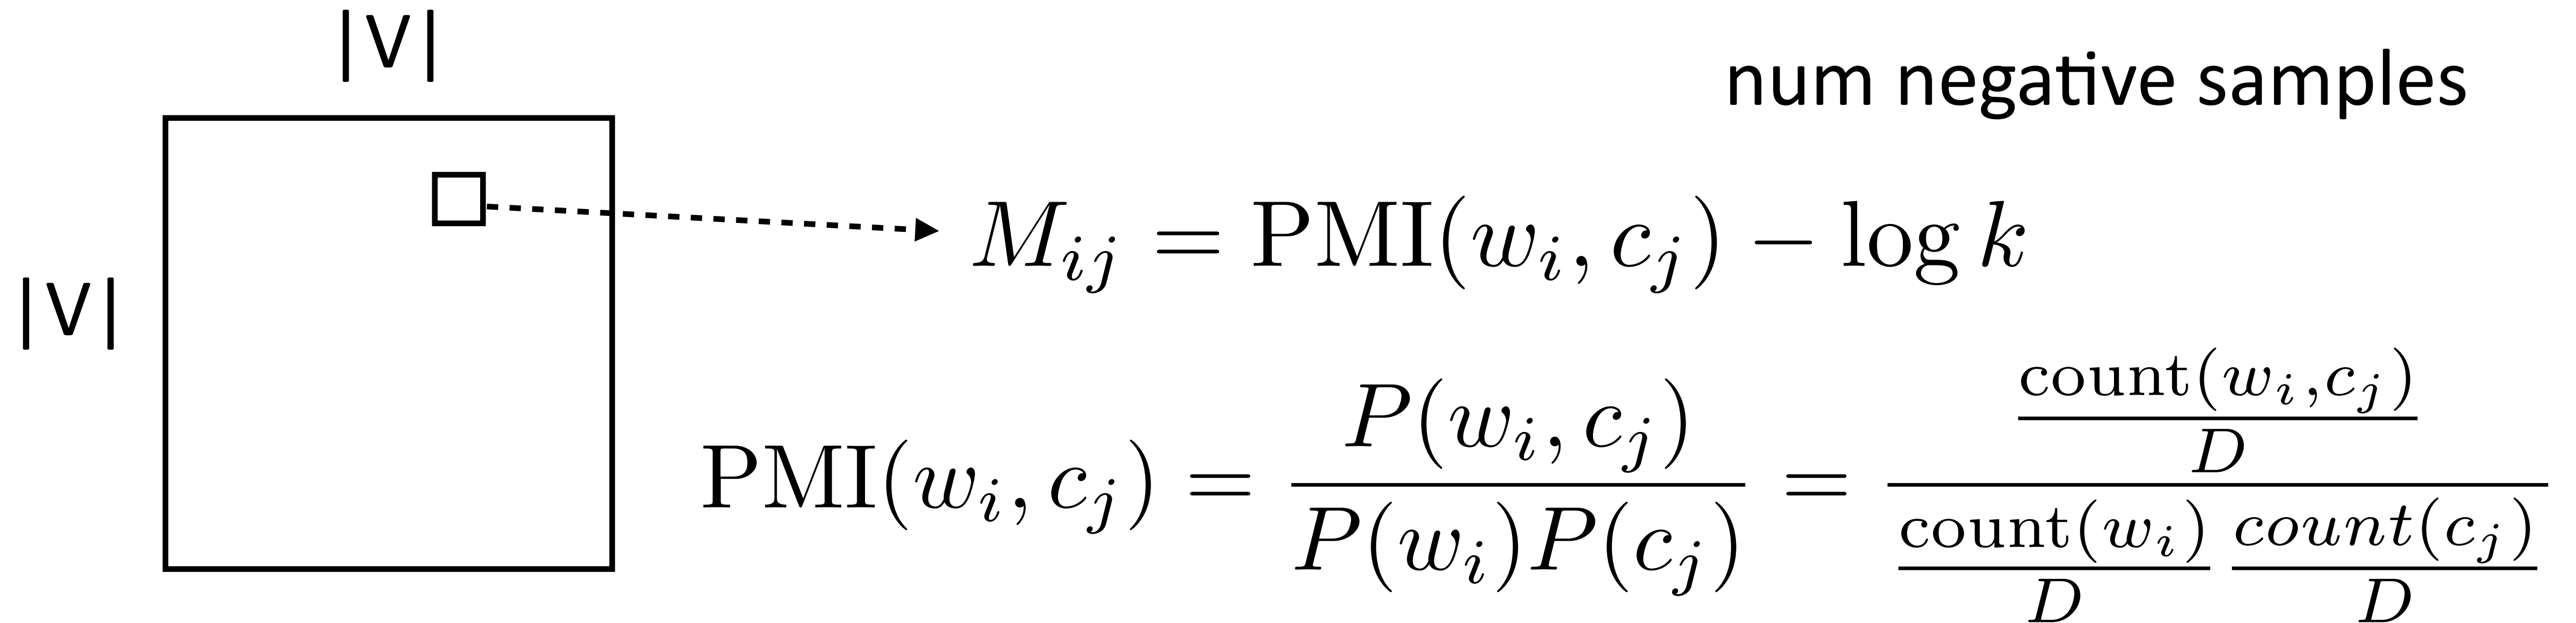
# Skip-Gram as Matrix Factorization

---




# Skip-Gram as Matrix Factorization

---



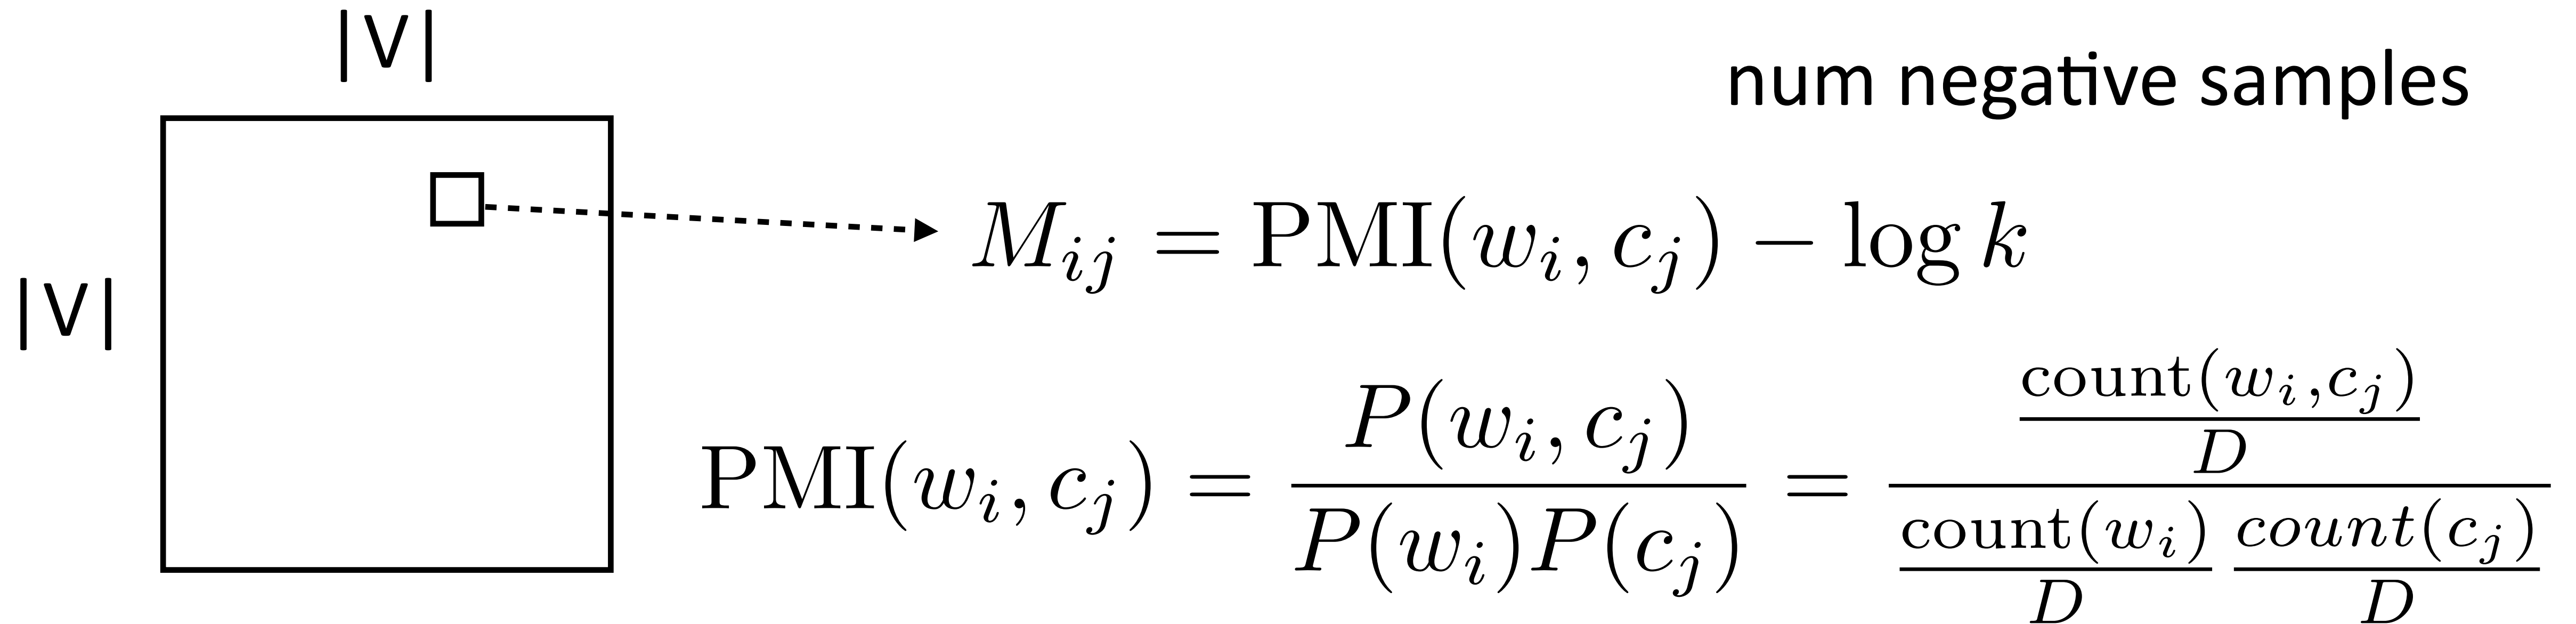
# Skip-Gram as Matrix Factorization

---



Skip-gram objective *exactly* corresponds to factoring this matrix:

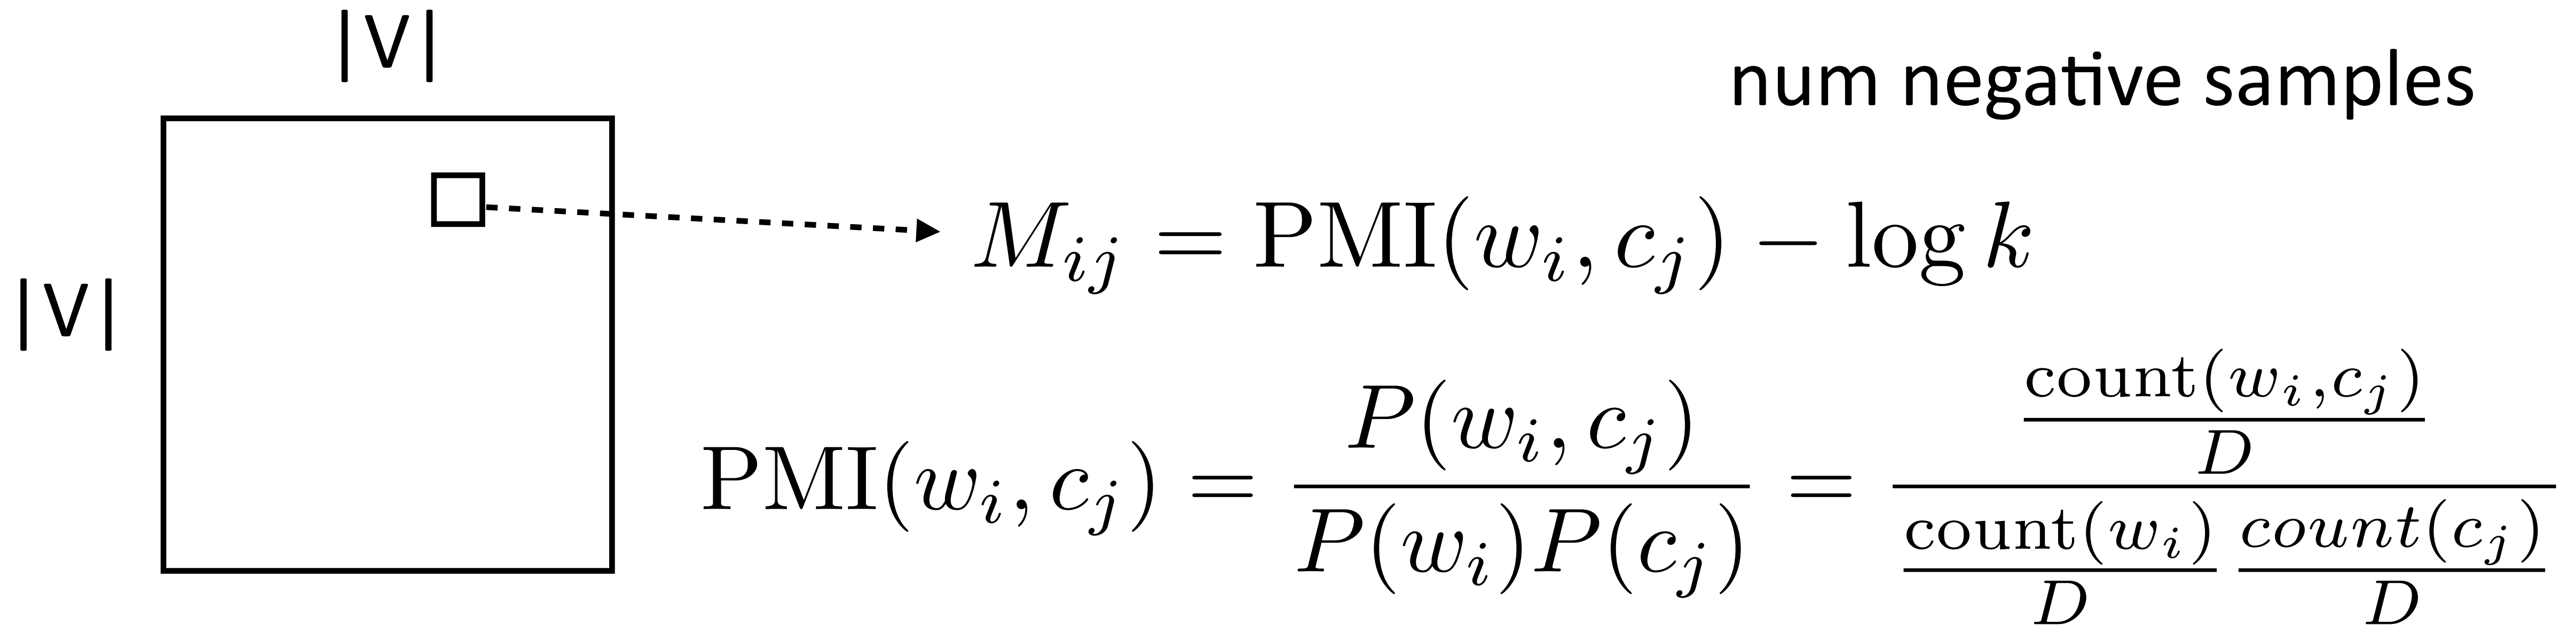
# Skip-Gram as Matrix Factorization



Skip-gram objective *exactly* corresponds to factoring this matrix:

- ▶ If we sample negative examples from the uniform distribution over words

# Skip-Gram as Matrix Factorization



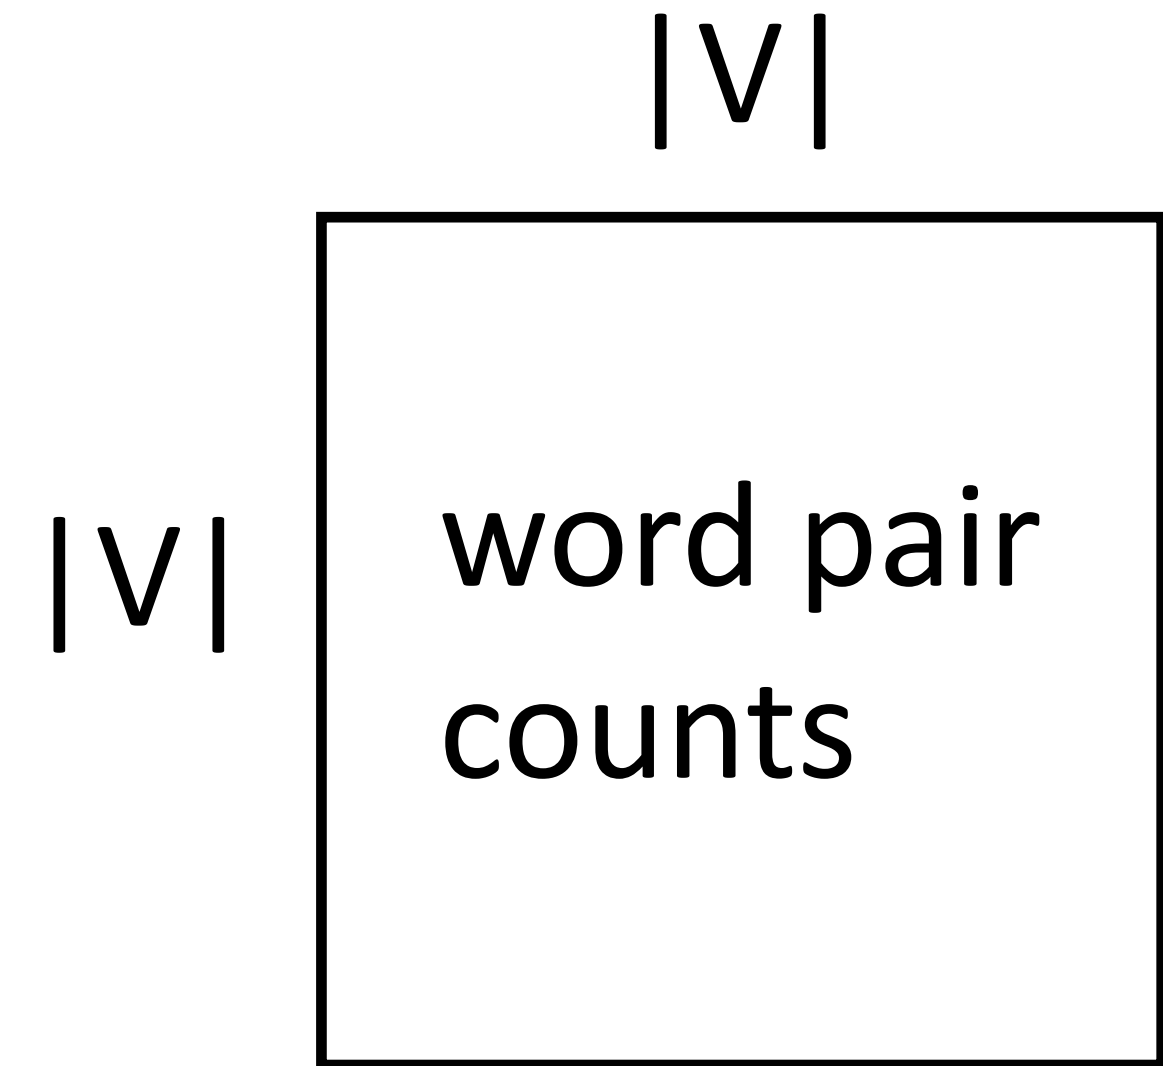
Skip-gram objective *exactly* corresponds to factoring this matrix:

- ▶ If we sample negative examples from the uniform distribution over words
- ▶ ...and it's a *weighted* factorization problem (weighted by word freq)

# GloVe (Global Vectors)

---

- ▶ Also operates on counts matrix, weighted regression on the log co-occurrence matrix

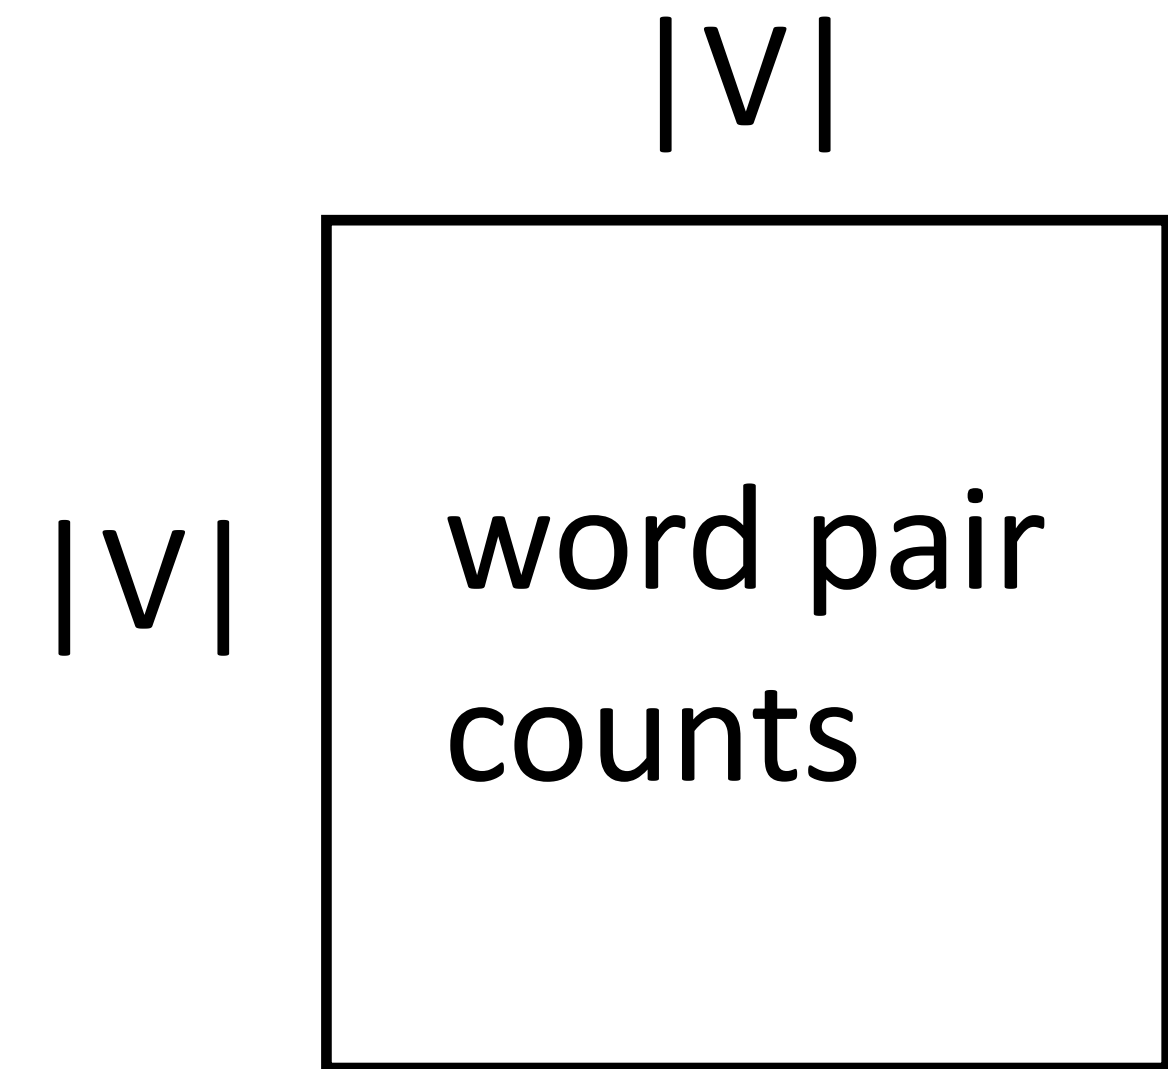




# GloVe (Global Vectors)

---

- ▶ Also operates on counts matrix, weighted regression on the log co-occurrence matrix

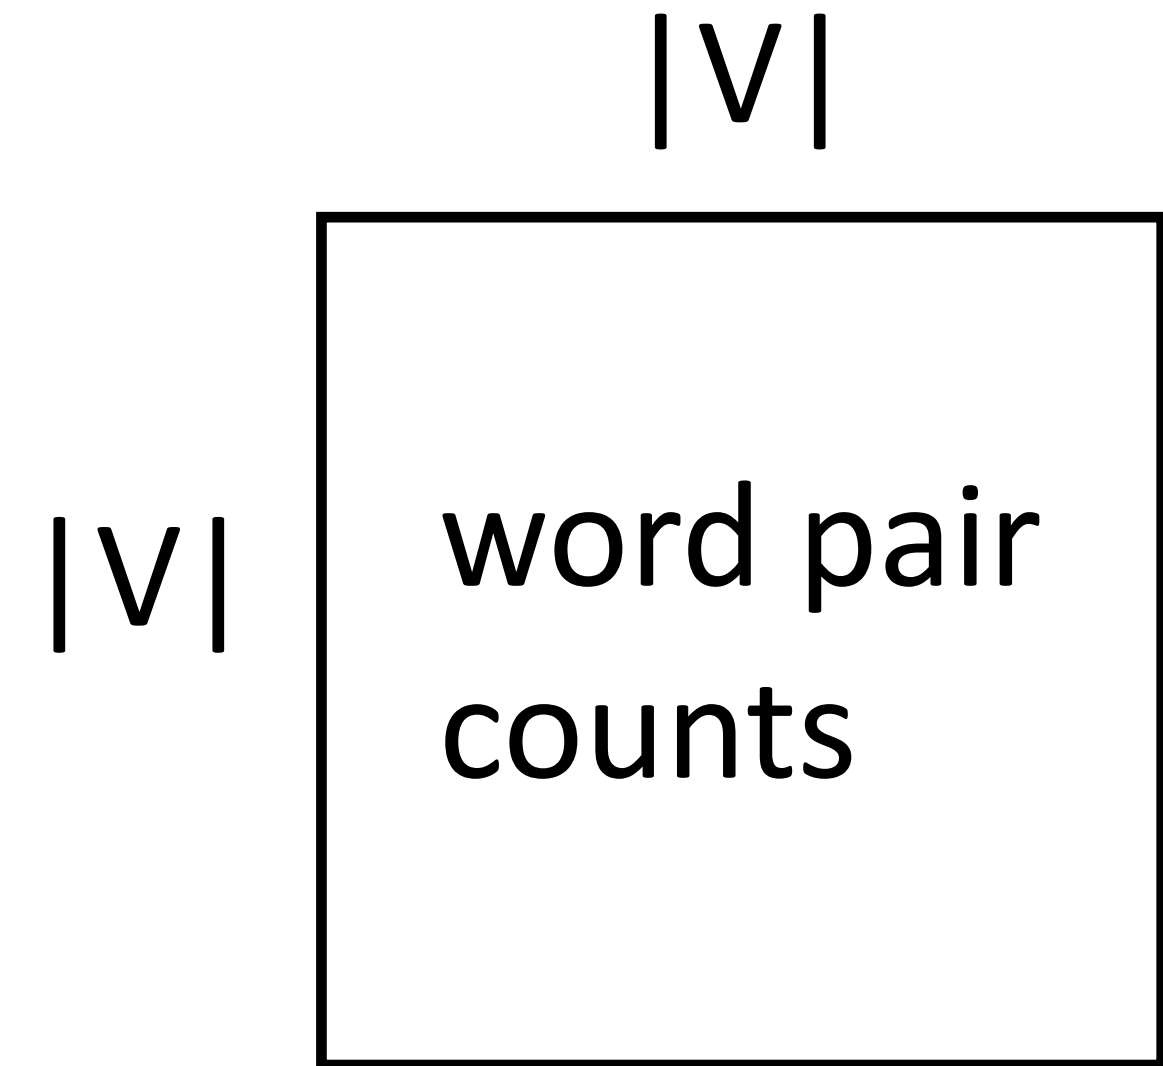


- ▶ Loss 
$$= \sum_{i,j} f(\text{count}(w_i, c_j)) \left( w_i^\top c_j + a_i + b_j - \log \text{count}(w_i, c_j) \right)^2$$

# GloVe (Global Vectors)

---

- ▶ Also operates on counts matrix, weighted regression on the log co-occurrence matrix

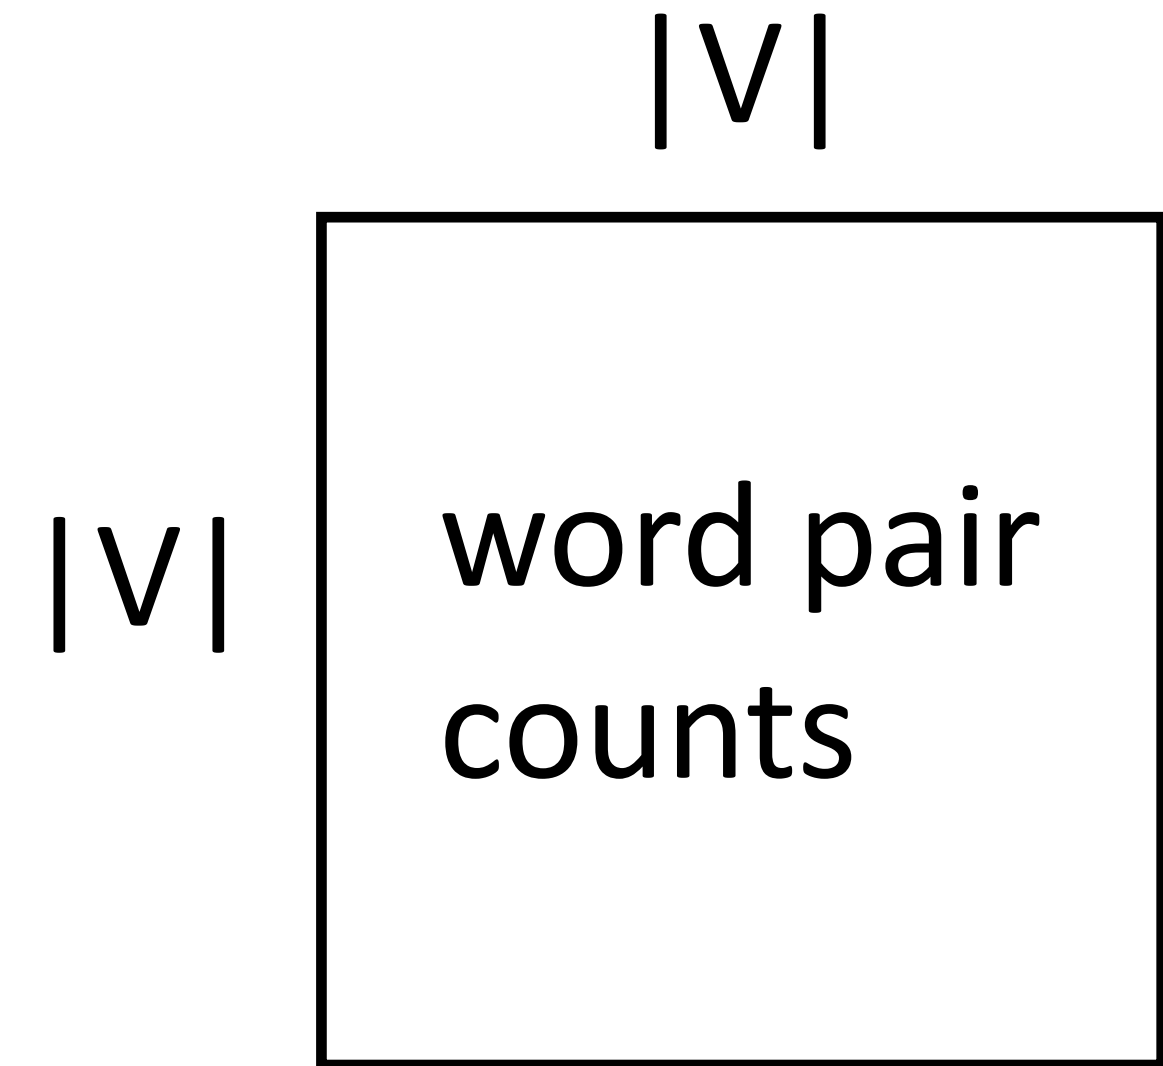


- ▶ Loss  $= \sum_{i,j} f(\text{count}(w_i, c_j)) (w_i^\top c_j + a_i + b_j - \log \text{count}(w_i, c_j))^2$
- ▶ Constant in the dataset size (just need counts), quadratic in voc size

# GloVe (Global Vectors)

---

- ▶ Also operates on counts matrix, weighted regression on the log co-occurrence matrix



- ▶ Loss  $= \sum_{i,j} f(\text{count}(w_i, c_j)) (w_i^\top c_j + a_i + b_j - \log \text{count}(w_i, c_j))^2$
- ▶ Constant in the dataset size (just need counts), quadratic in voc size
- ▶ By far the most common (uncontextualized) word vectors used today (5000+ citations)

Pennington et al. (2014)

# Preview: Context-dependent Embeddings

---

- ▶ How to handle different word senses? One vector for *balls*

they dance at balls

they hit the balls

# Preview: Context-dependent Embeddings

---

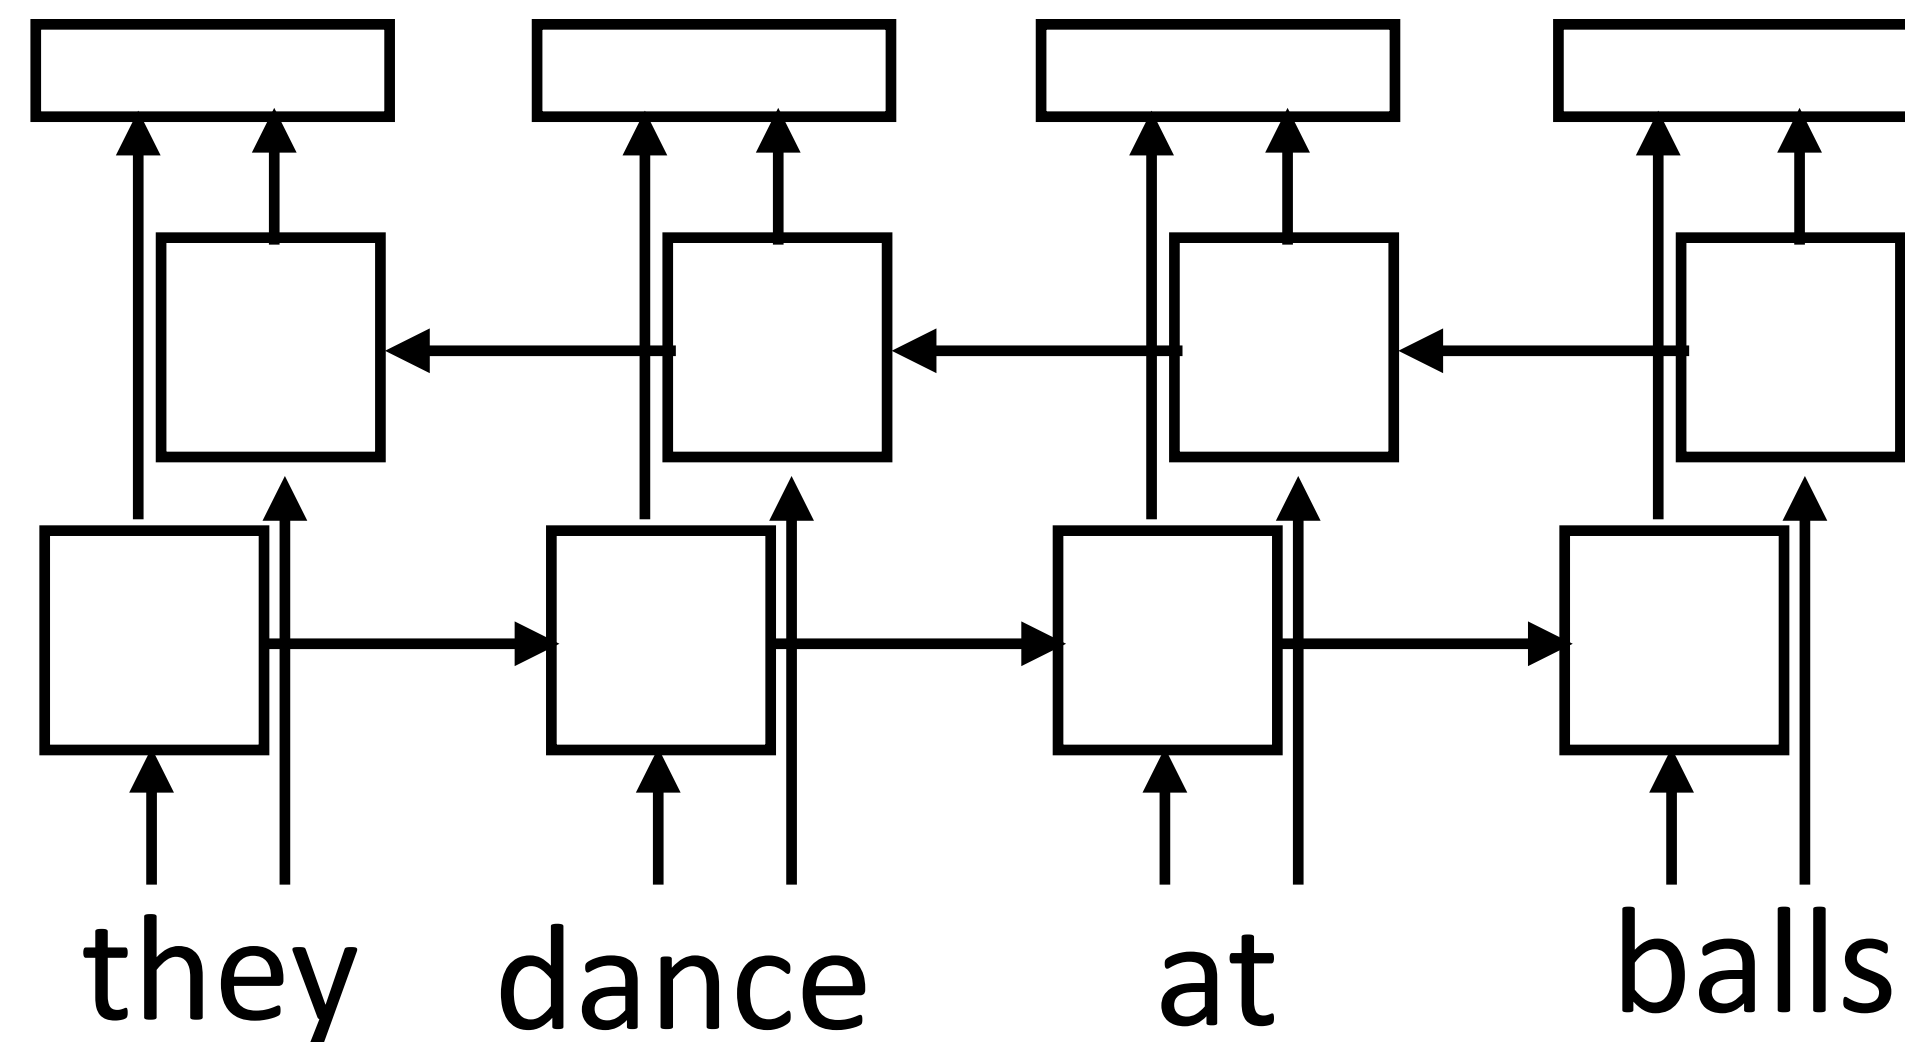
- ▶ How to handle different word senses? One vector for *balls*

they dance at balls                      they hit the balls

- ▶ Train a neural language model to predict the next word given previous words in the sentence, use its internal representations as word vectors

# Preview: Context-dependent Embeddings

- ▶ How to handle different word senses? One vector for *balls*

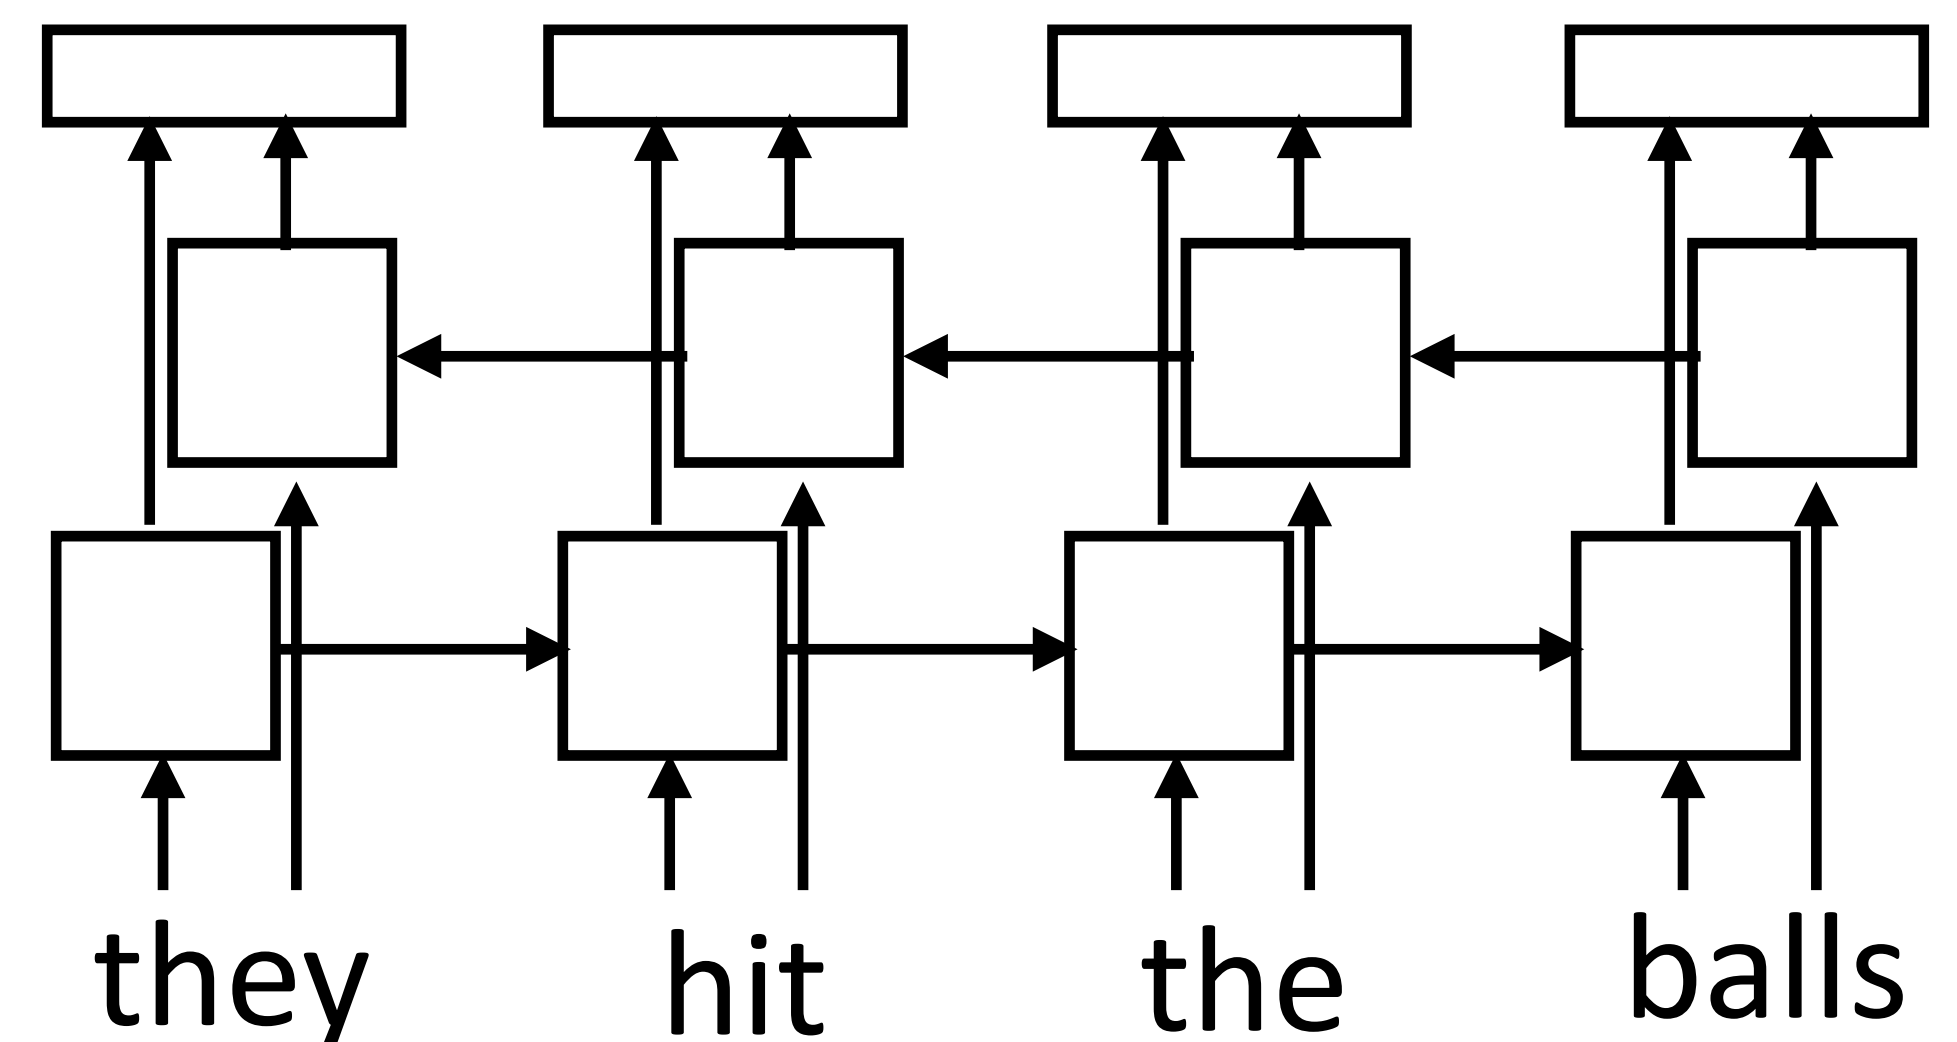
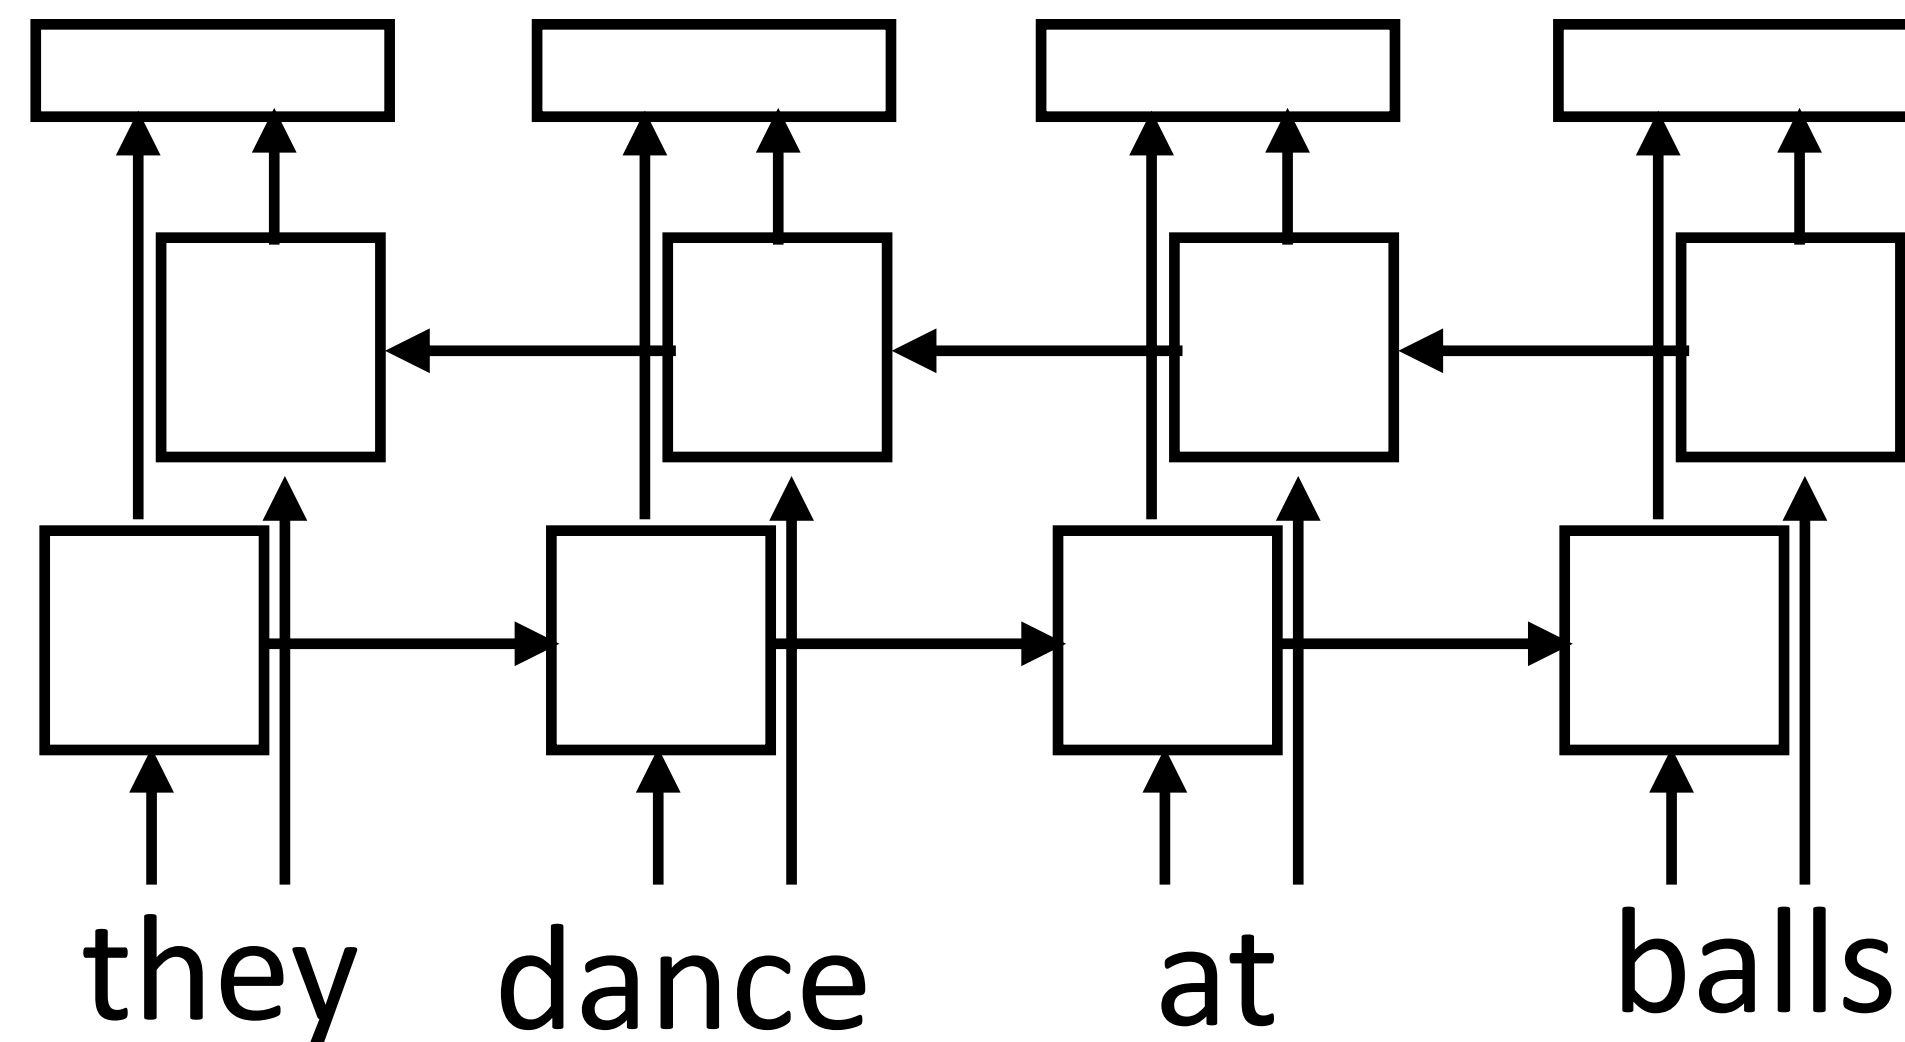


they hit the balls

- ▶ Train a neural language model to predict the next word given previous words in the sentence, use its internal representations as word vectors

# Preview: Context-dependent Embeddings

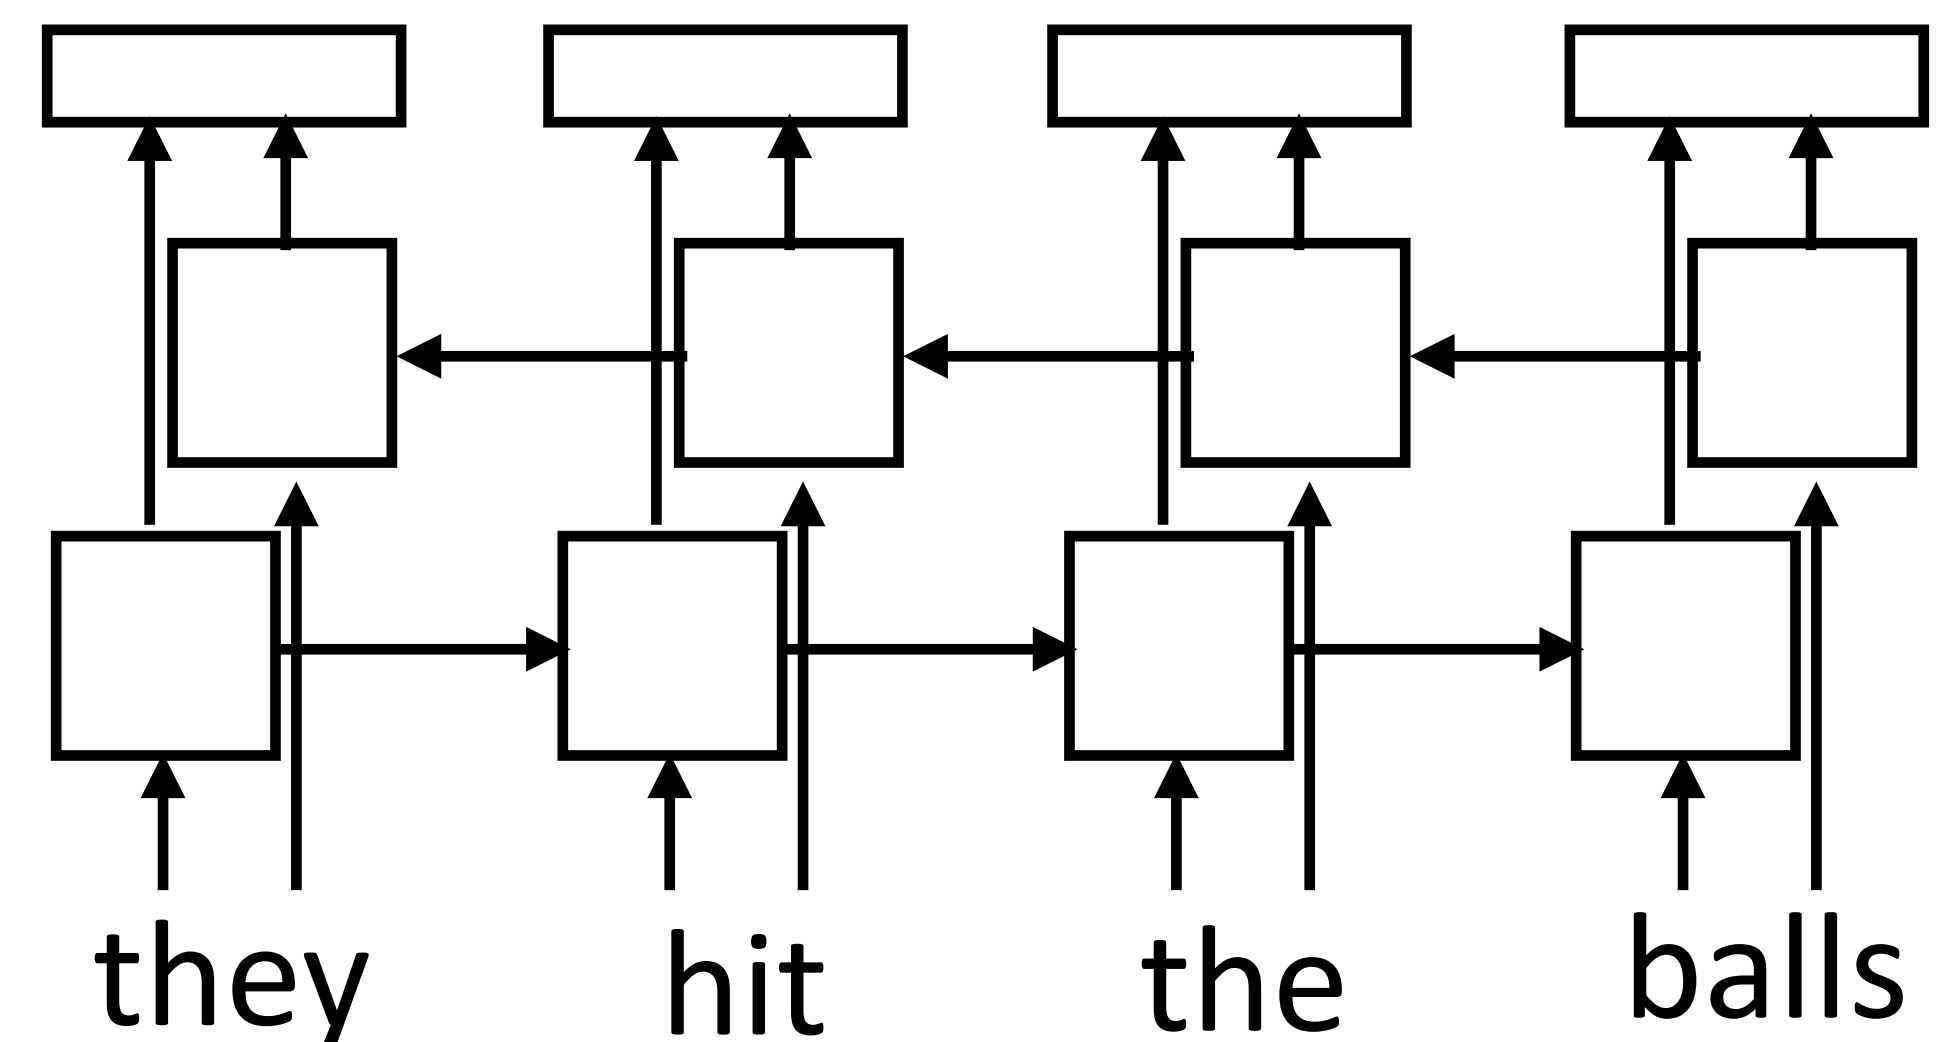
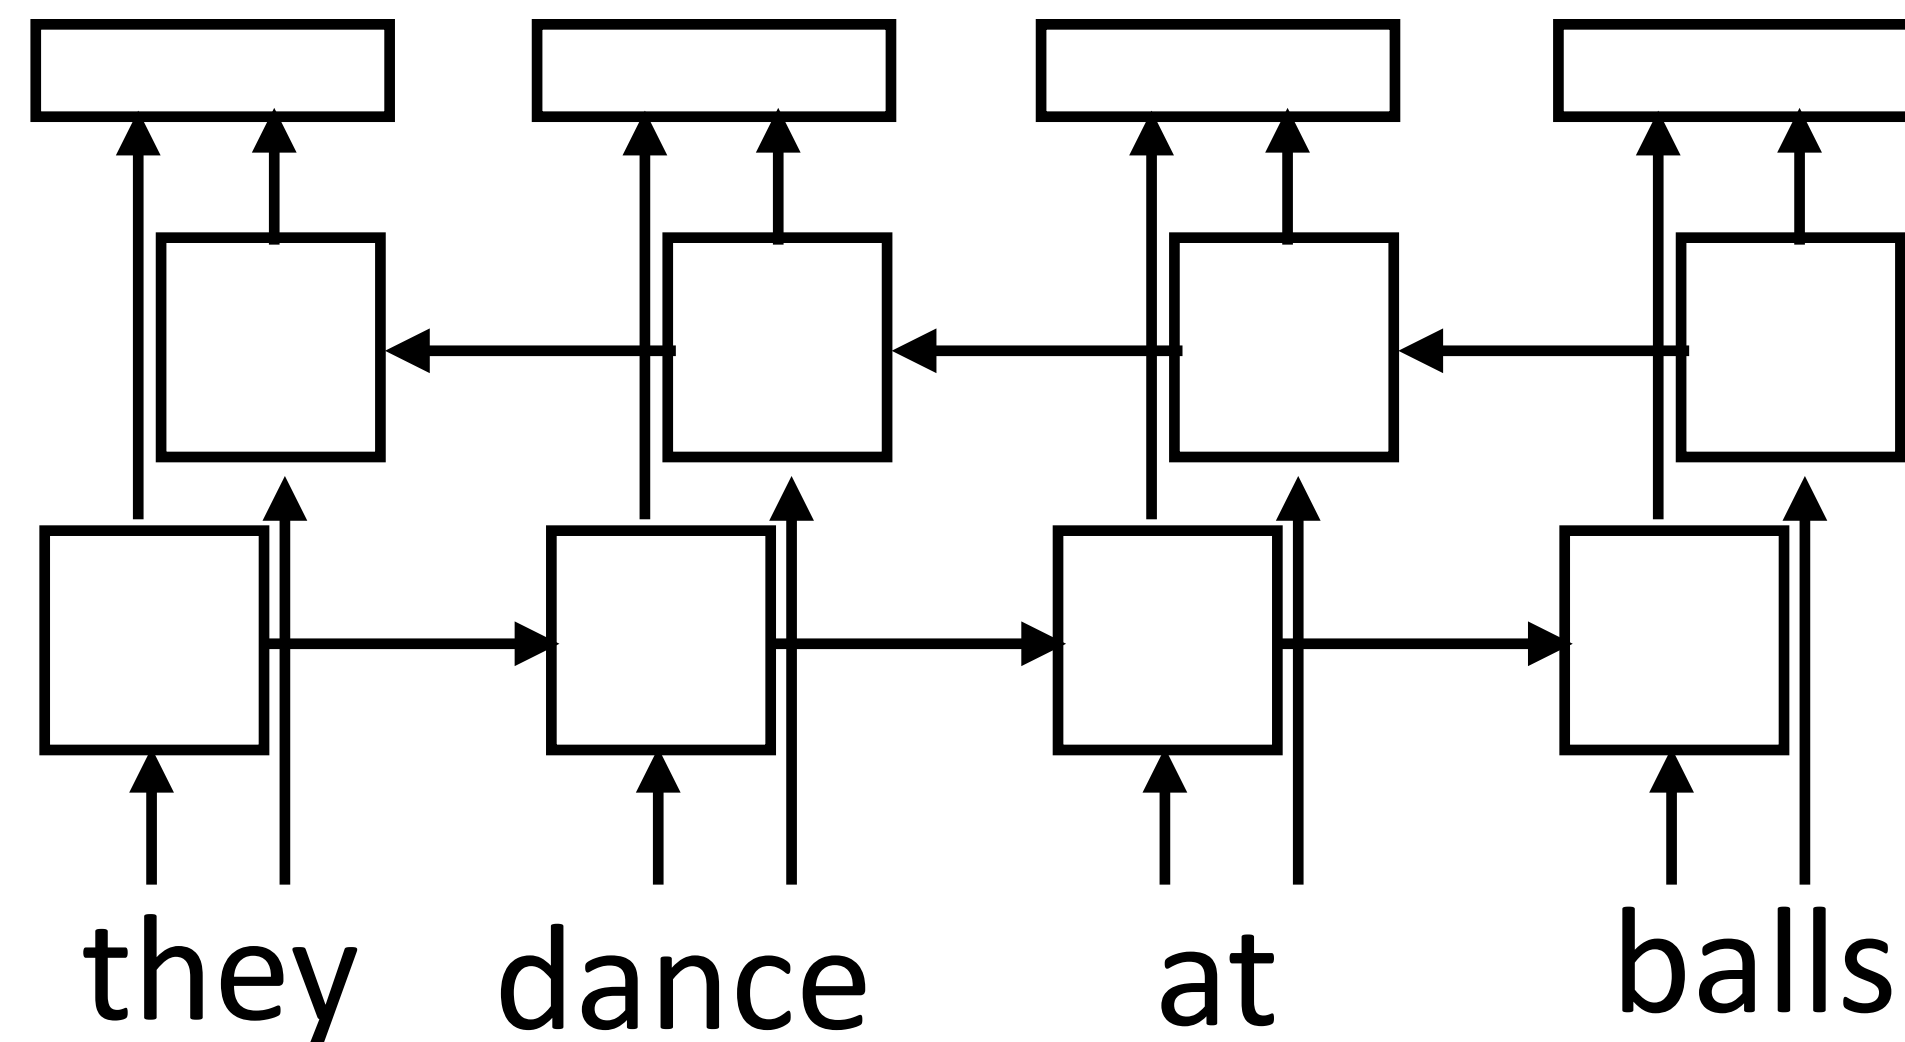
- How to handle different word senses? One vector for *balls*



- Train a neural language model to predict the next word given previous words in the sentence, use its internal representations as word vectors

# Preview: Context-dependent Embeddings

- ▶ How to handle different word senses? One vector for *balls*

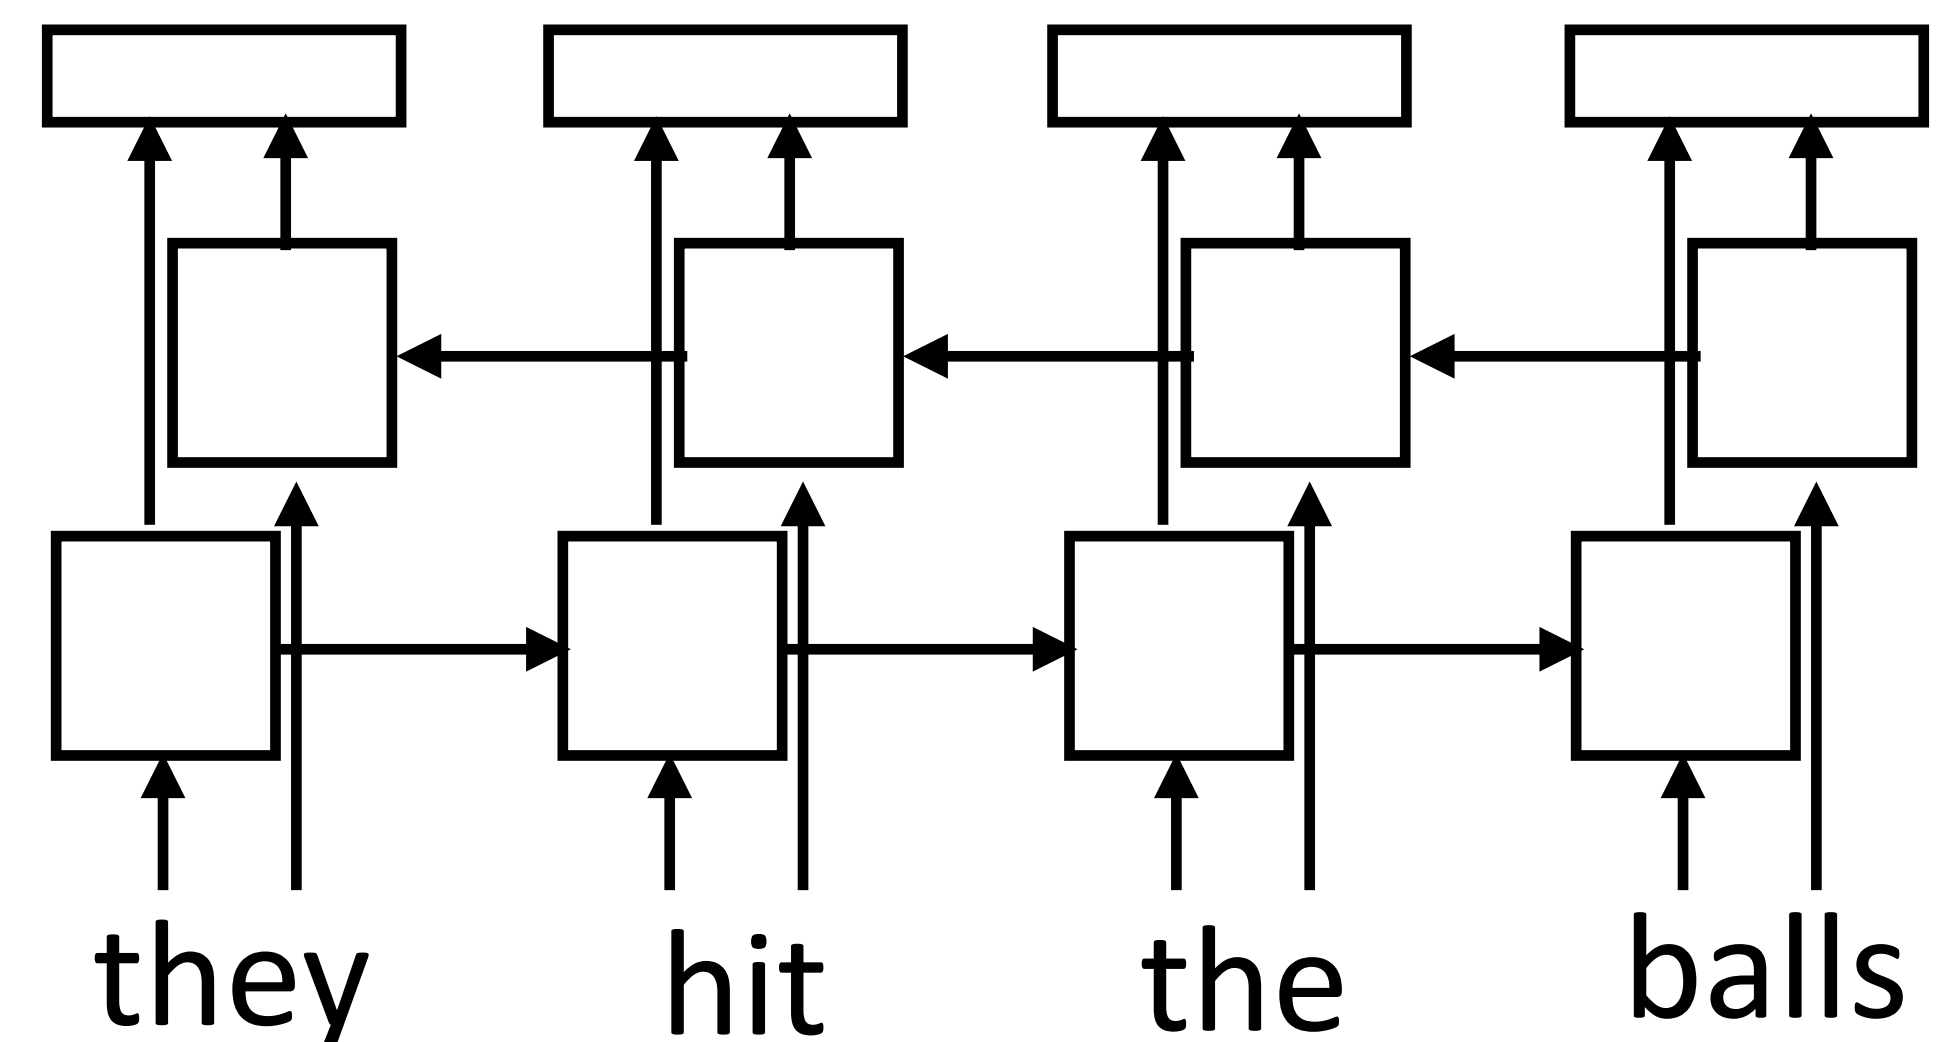
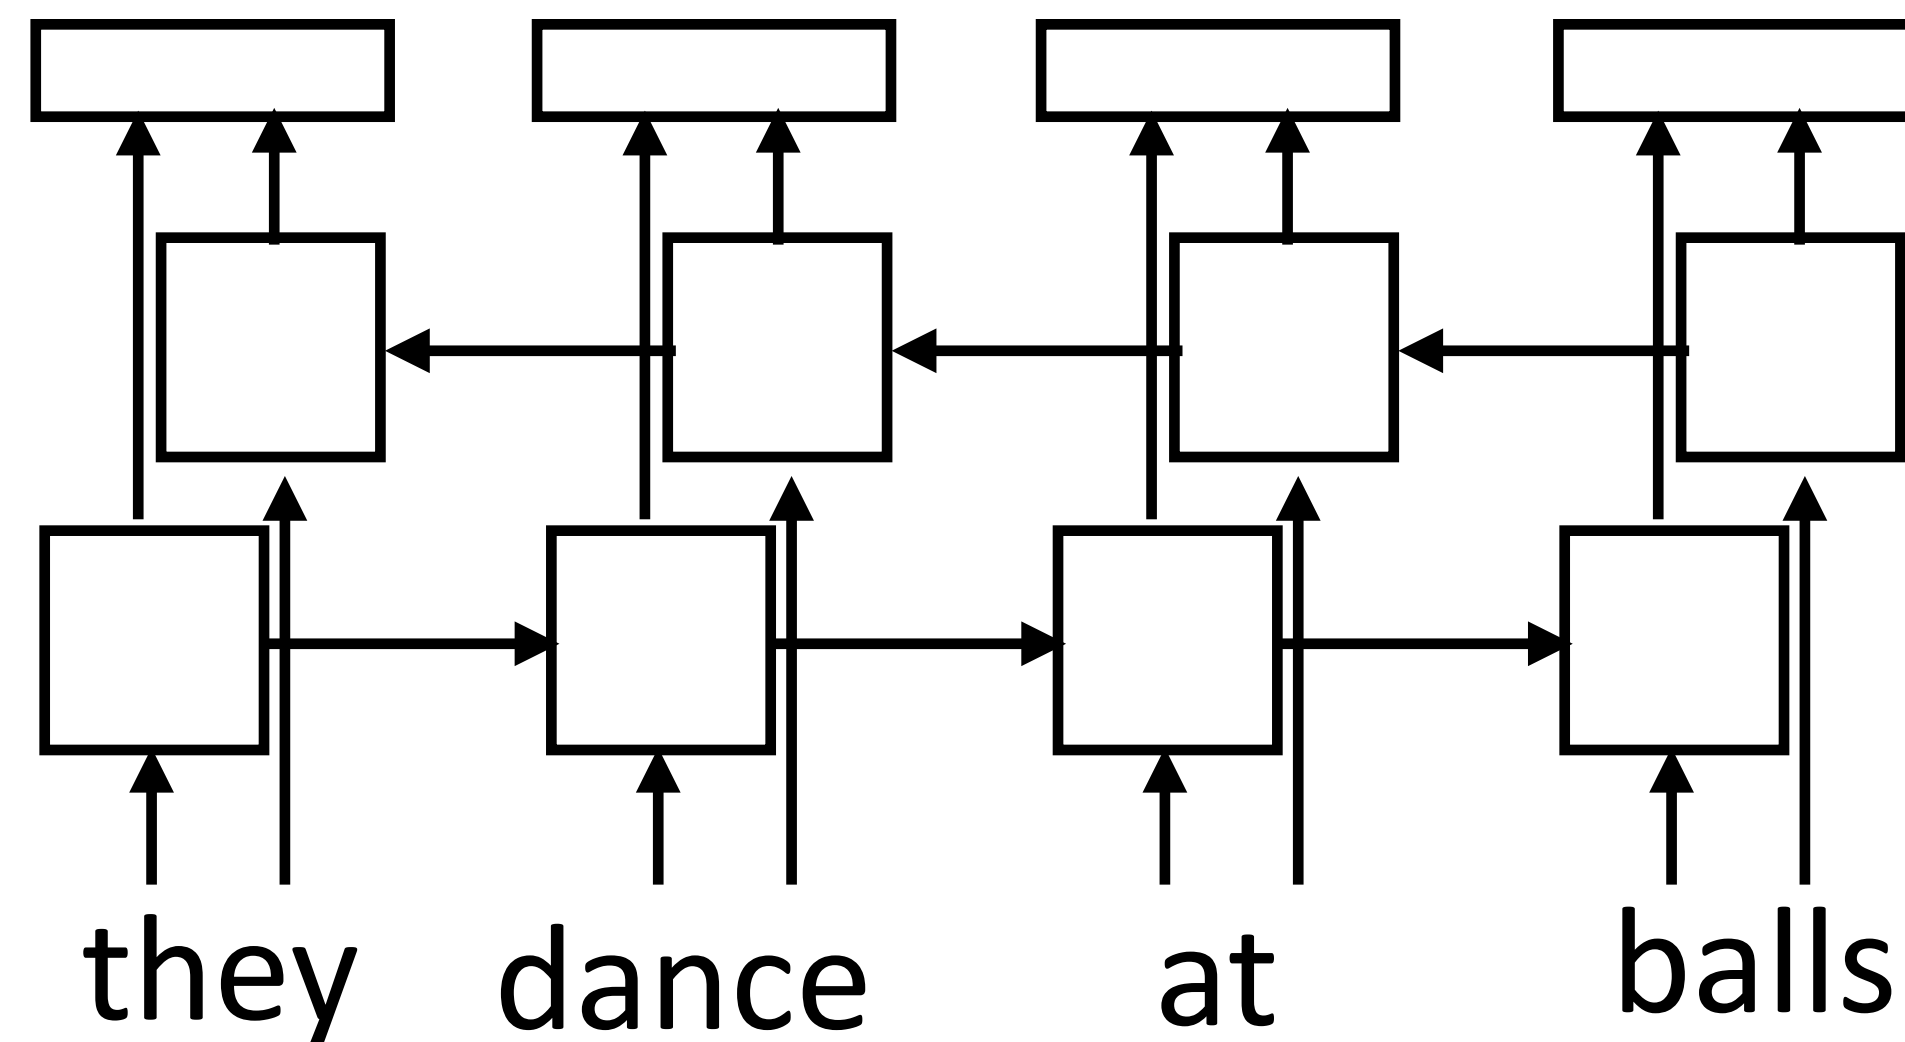


- ▶ Train a neural language model to predict the next word given previous words in the sentence, use its internal representations as word vectors
- ▶ *Context-sensitive* word embeddings: depend on rest of the sentence



# Preview: Context-dependent Embeddings

- ▶ How to handle different word senses? One vector for *balls*



- ▶ Train a neural language model to predict the next word given previous words in the sentence, use its internal representations as word vectors
- ▶ *Context-sensitive* word embeddings: depend on rest of the sentence
- ▶ *Huge* improvements across nearly all NLP tasks over GloVe

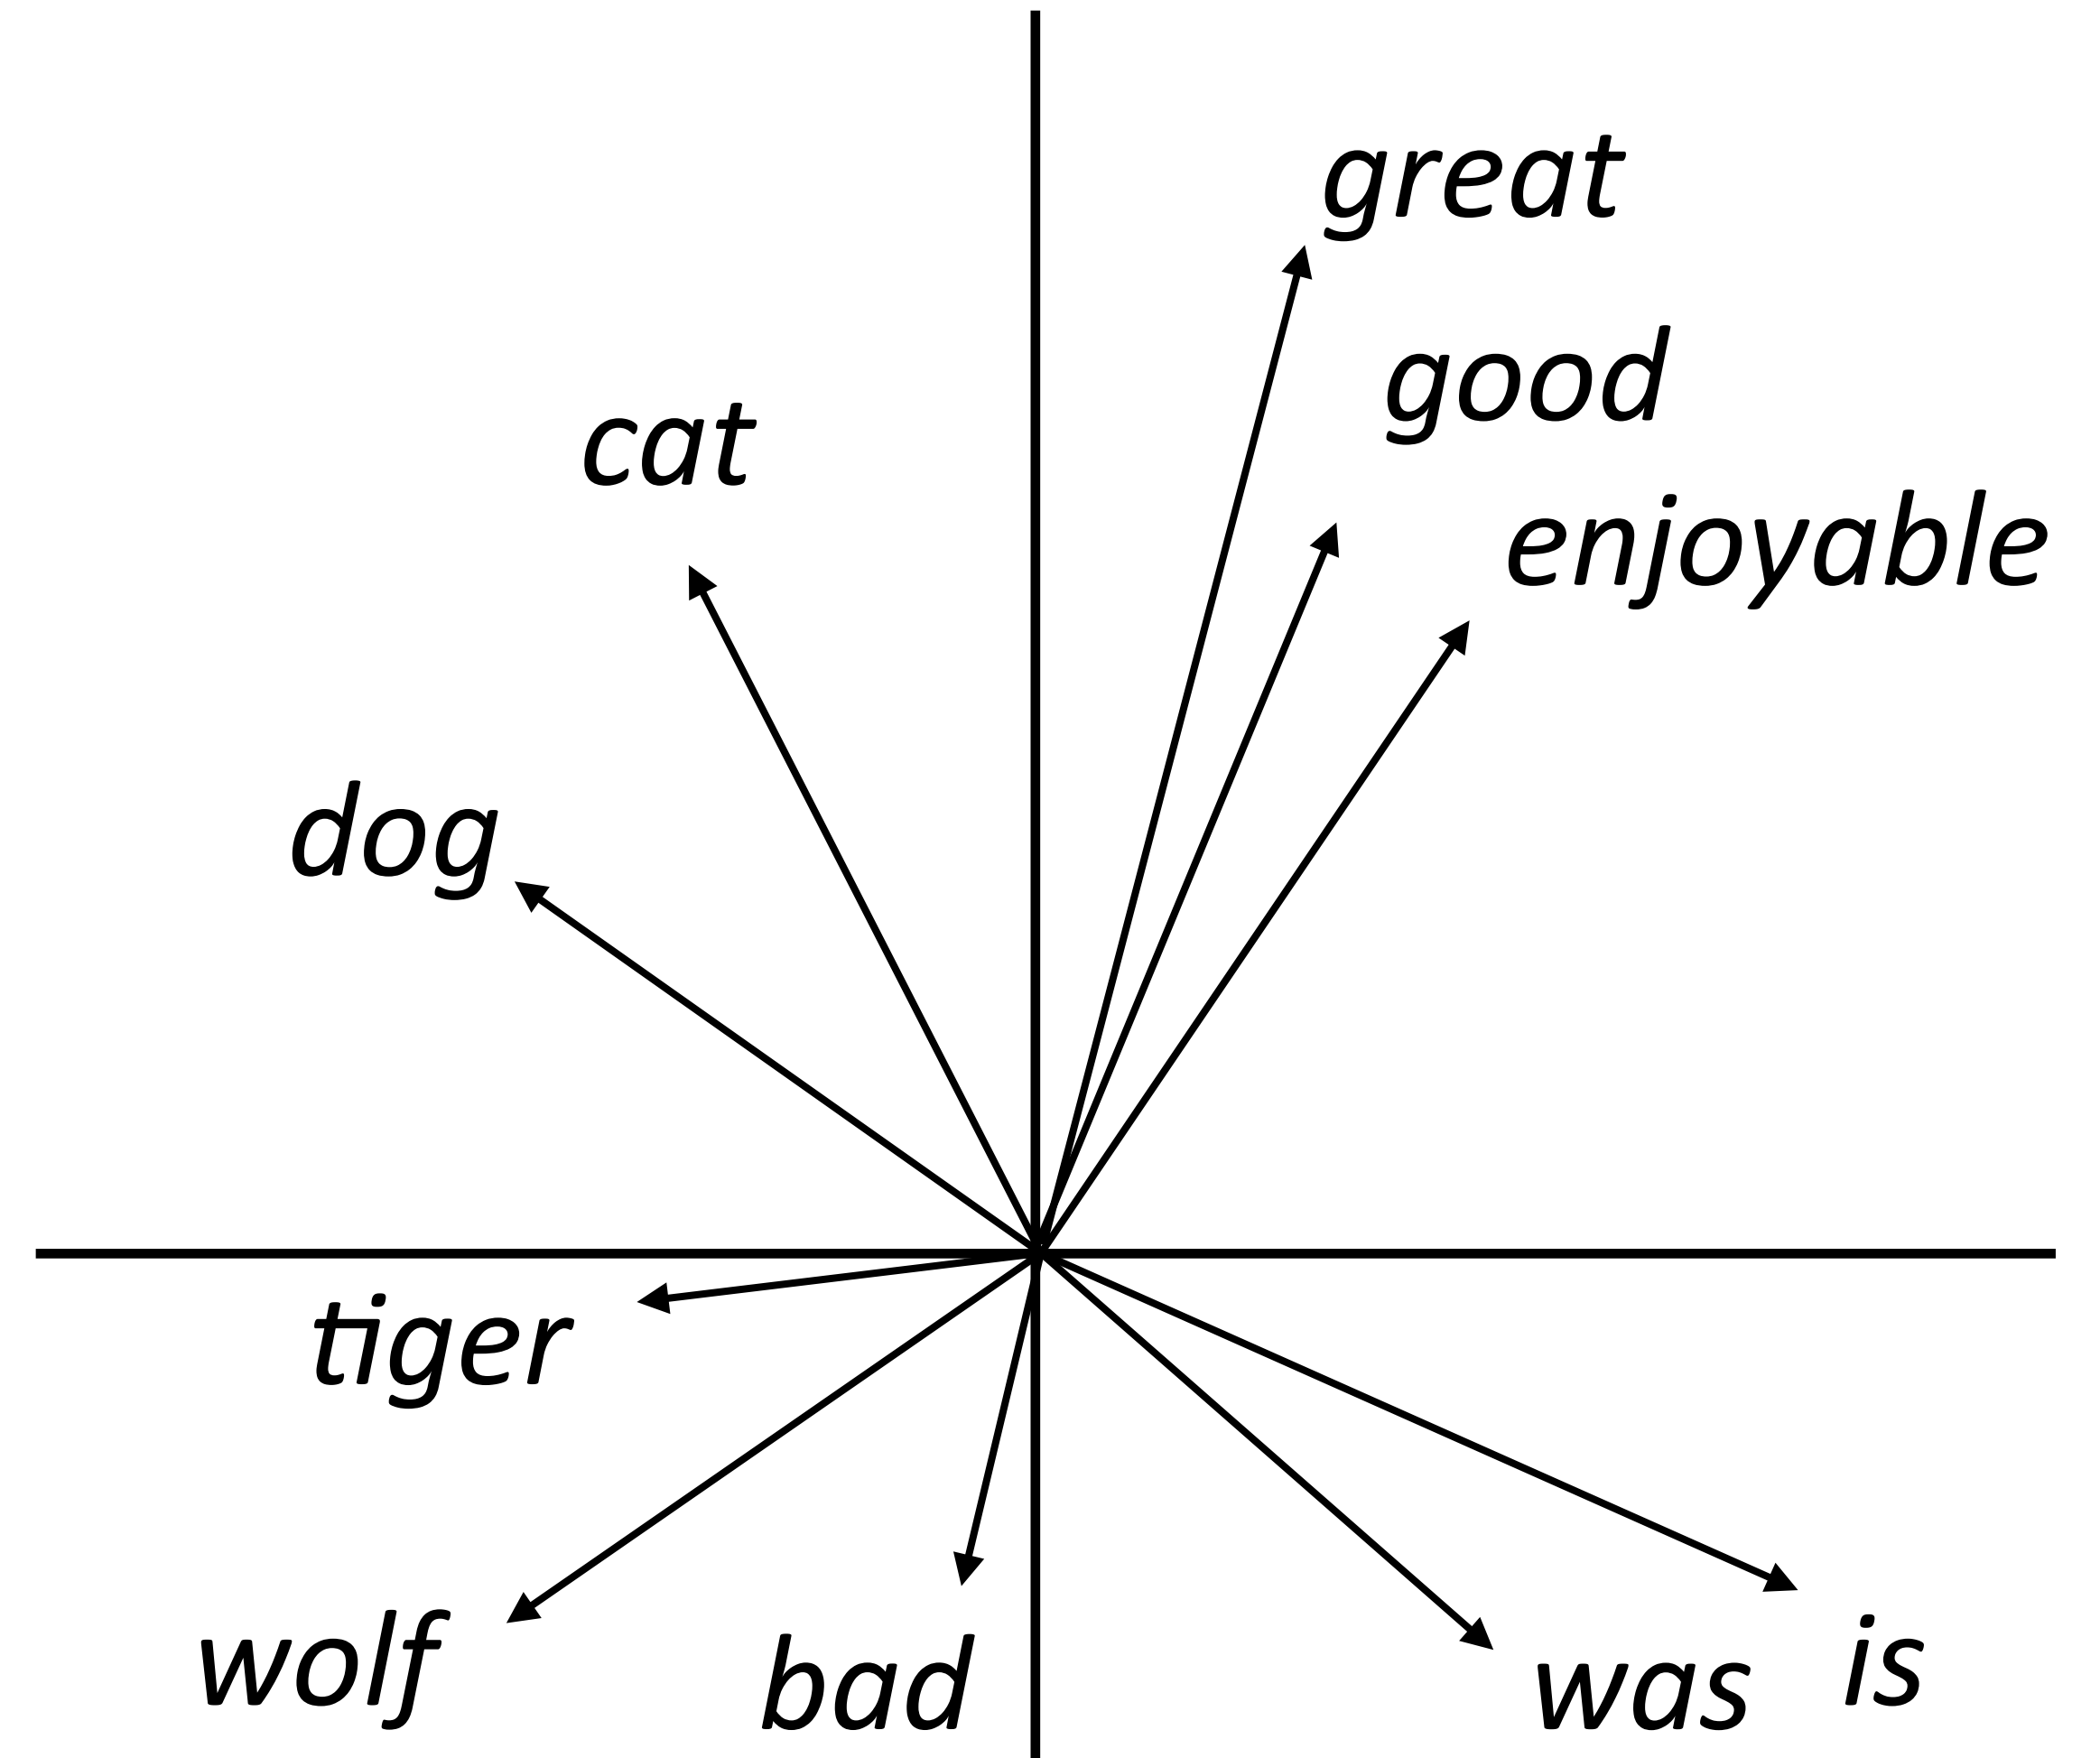
Peters et al. (2018)

# Evaluation

# Evaluating Word Embeddings

---

- What properties of language should word embeddings capture?




# Evaluating Word Embeddings

---

- ▶ What properties of language should word embeddings capture?

- ▶ Similarity: similar words are close to each other



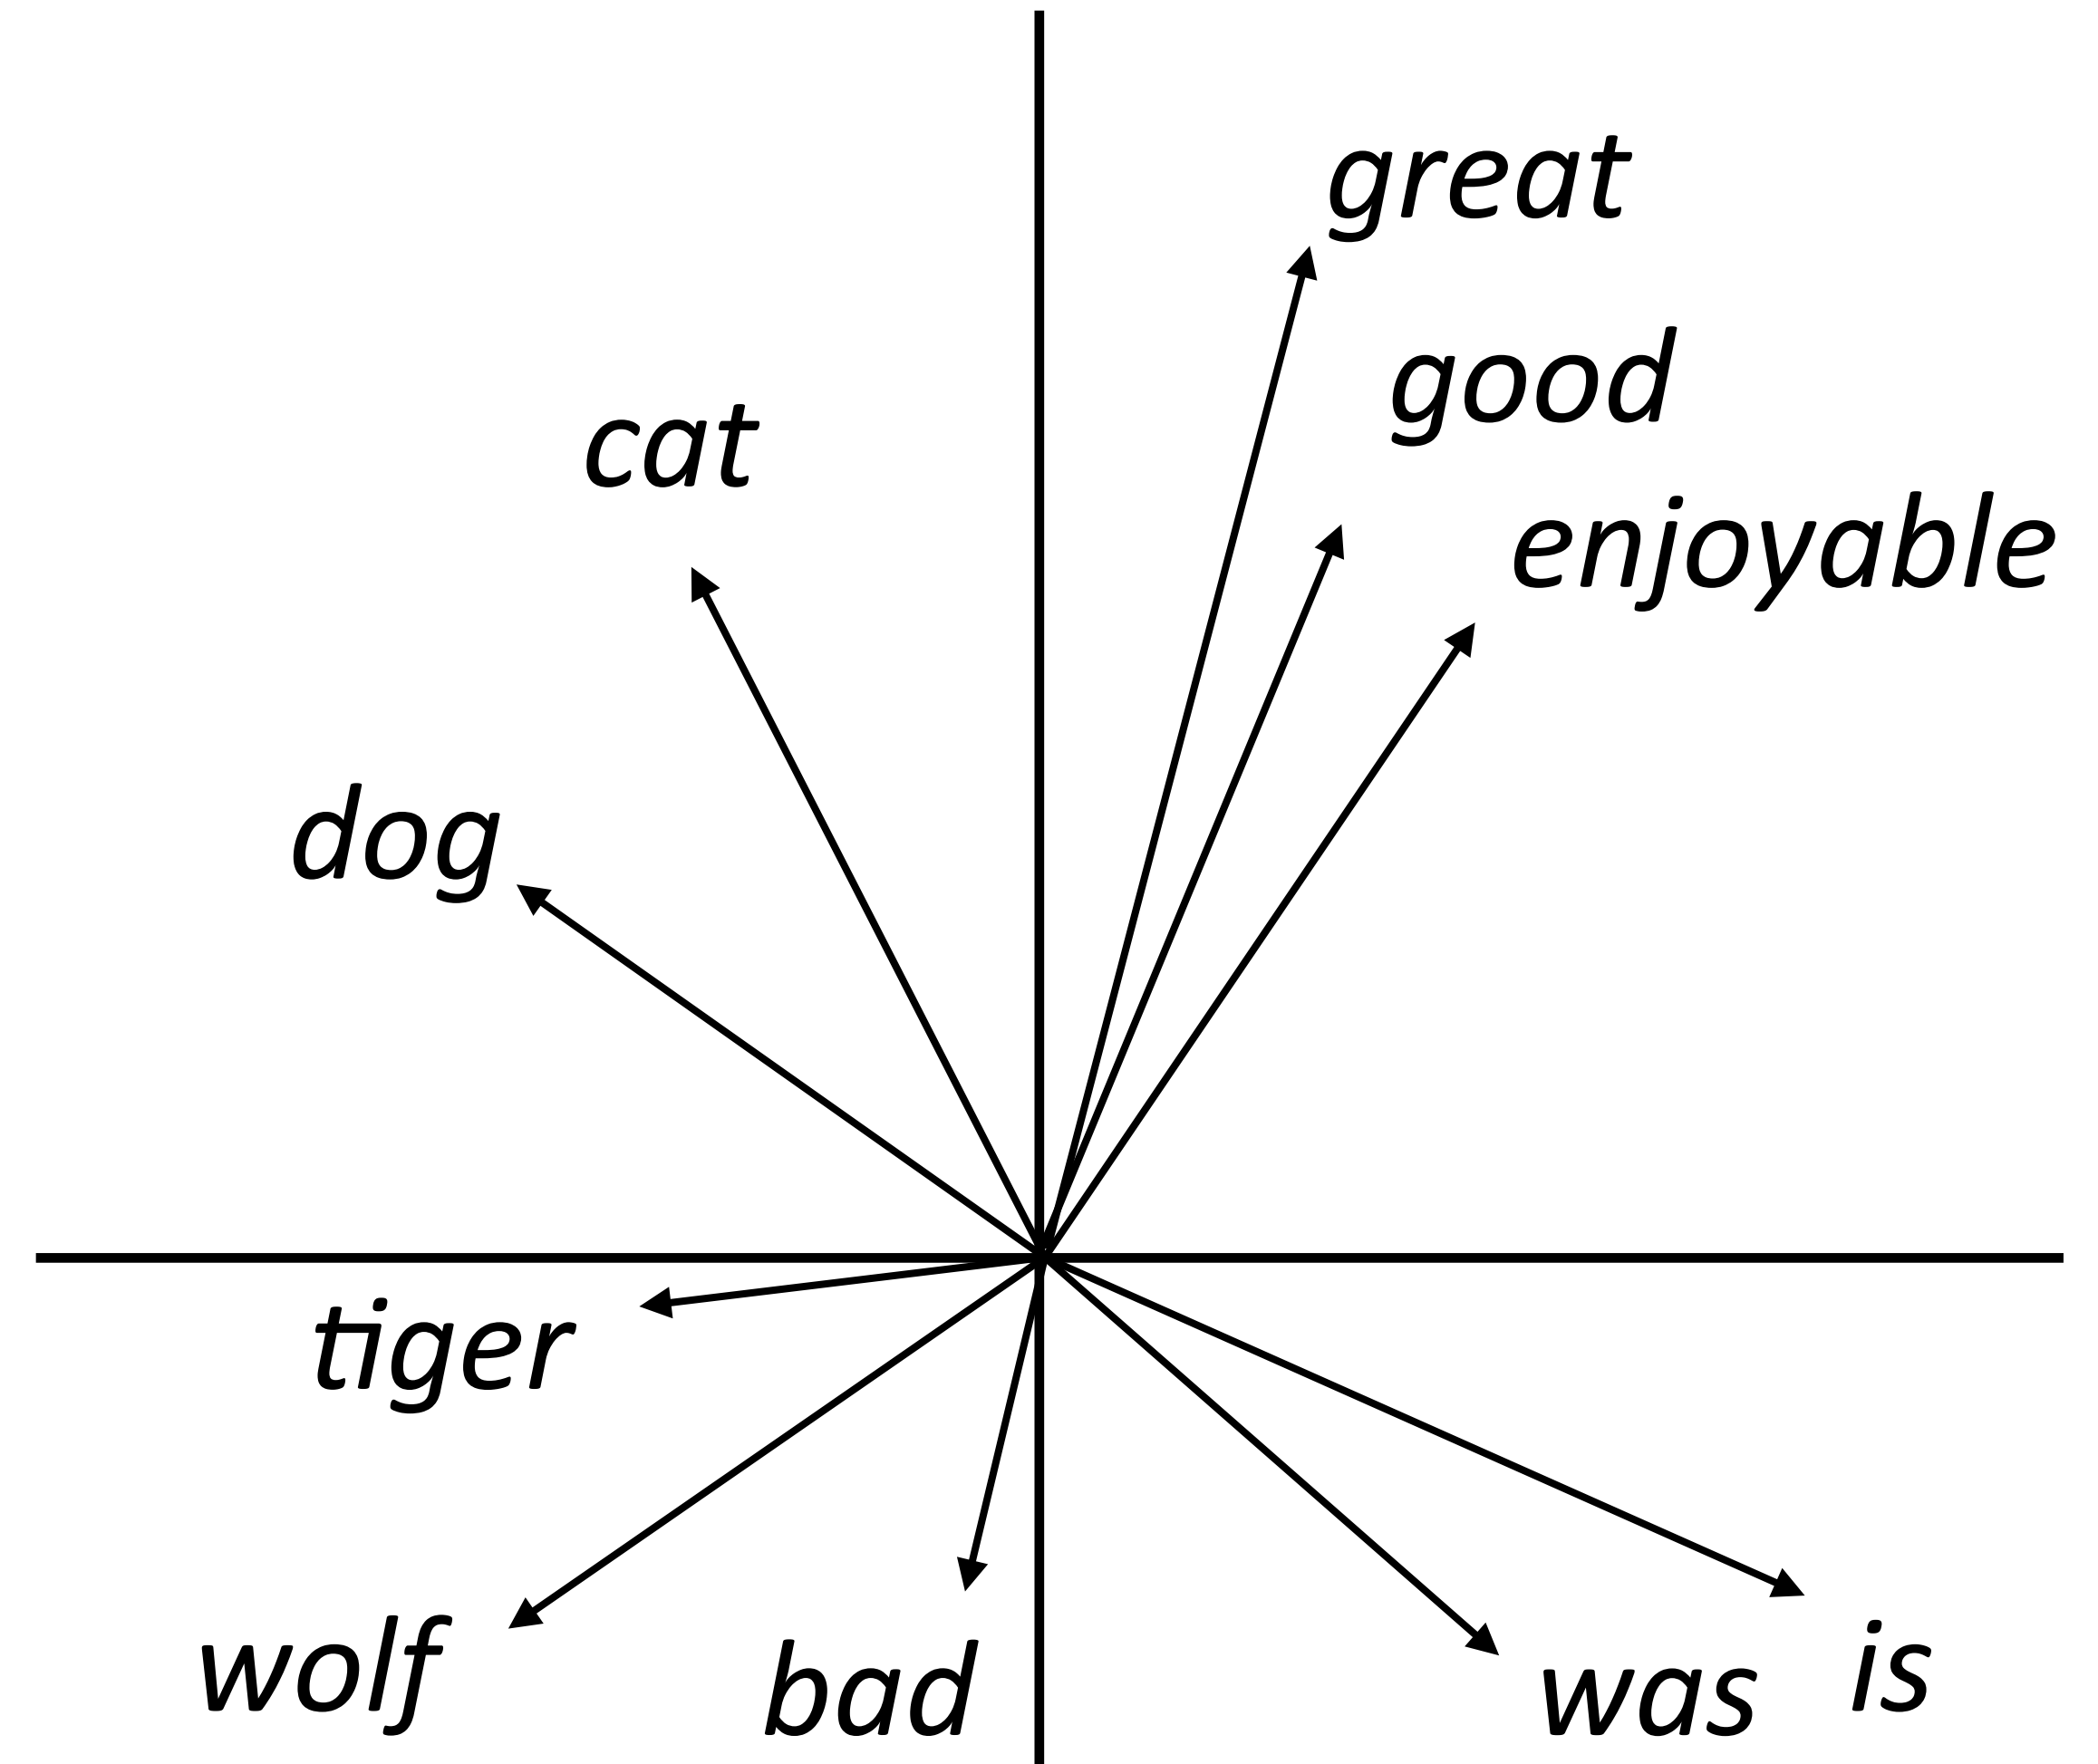
# Evaluating Word Embeddings

---

- ▶ What properties of language should word embeddings capture?

- ▶ Similarity: similar words are close to each other

- ▶ Analogy:



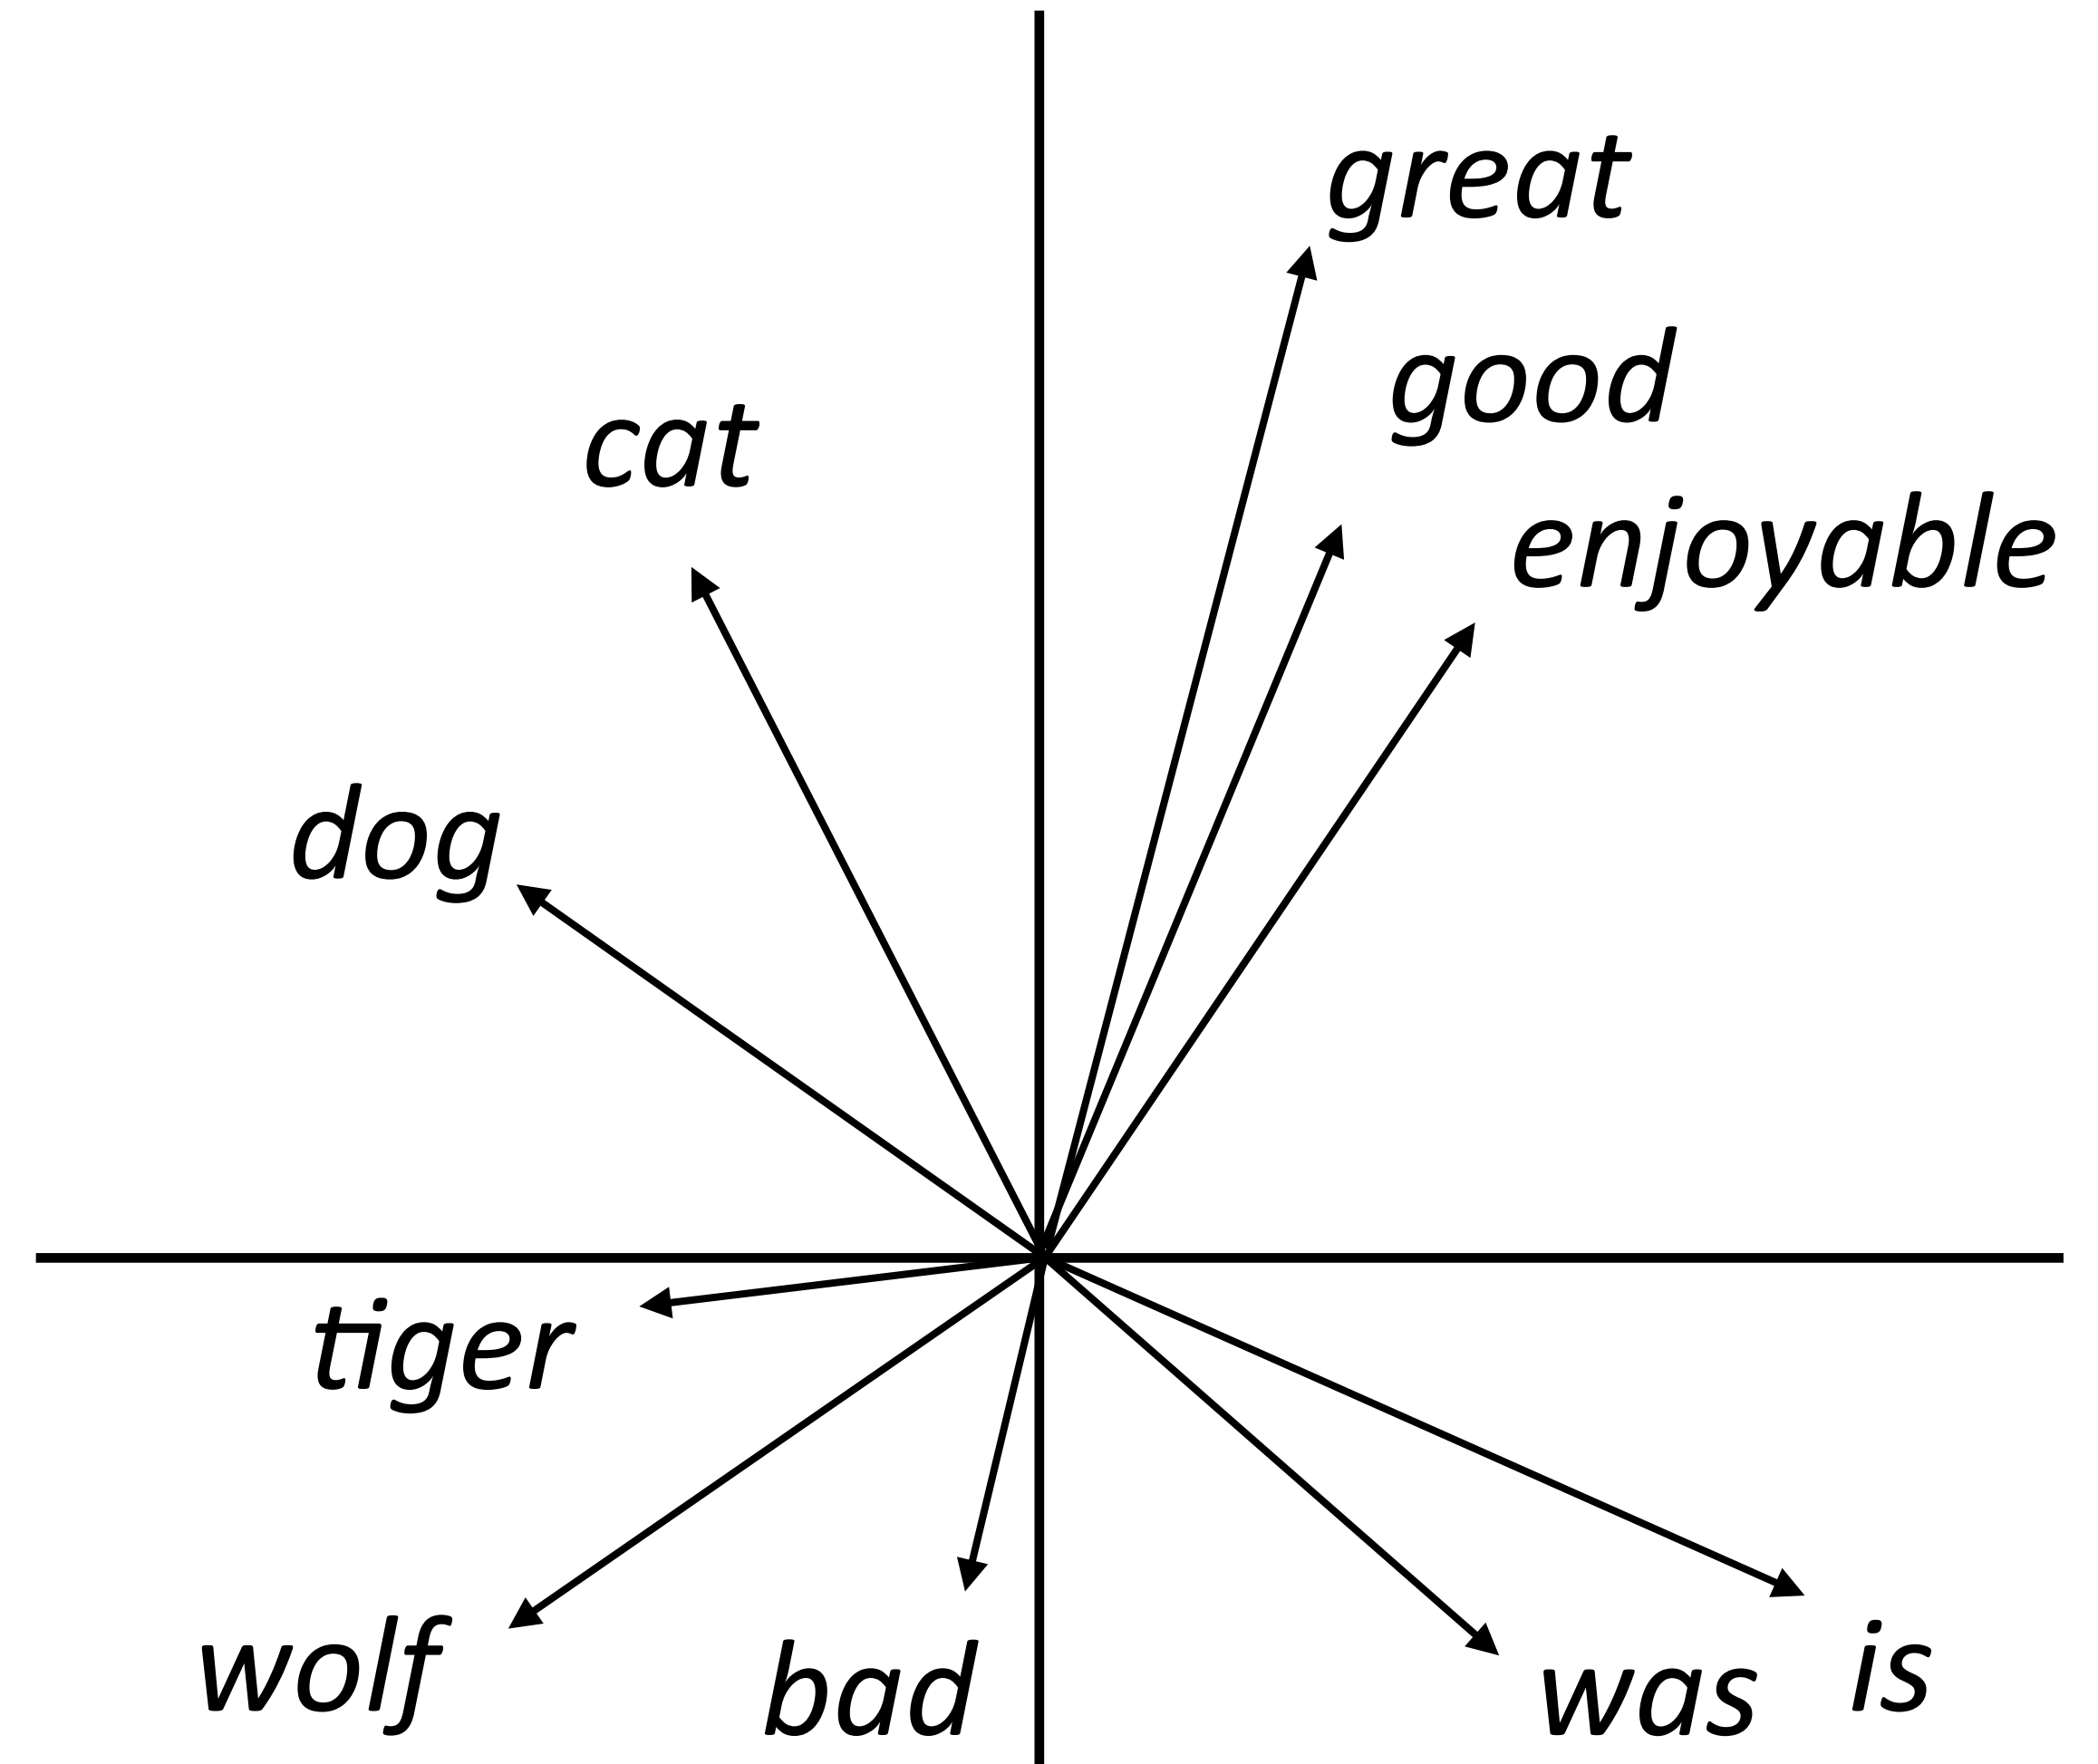
# Evaluating Word Embeddings

---

- ▶ What properties of language should word embeddings capture?

- ▶ Similarity: similar words are close to each other

- ▶ Analogy:  
good is to best as smart is to ???



# Evaluating Word Embeddings

---


- ▶ What properties of language should word embeddings capture?

- ▶ Similarity: similar words are close to each other

- ▶ Analogy:

good is to best as smart is to ???

Paris is to France as Tokyo is to ???



# Similarity

---

Method	WordSim Similarity	WordSim Relatedness	Bruni et al. MEN	Radinsky et al. M. Turk	Luong et al. Rare Words	Hill et al. SimLex
PPMI	.755	<b>.697</b>	.745	.686	.462	.393
SVD	<b>.793</b>	.691	<b>.778</b>	.666	<b>.514</b>	.432
SGNS	<b>.793</b>	.685	.774	<b>.693</b>	.470	<b>.438</b>
GloVe	.725	.604	.729	.632	.403	.398

- ▶ SVD = singular value decomposition on PMI matrix



# Similarity

Method	WordSim Similarity	WordSim Relatedness	Bruni et al. MEN	Radinsky et al. M. Turk	Luong et al. Rare Words	Hill et al. SimLex
PPMI	.755	<b>.697</b>	.745	.686	.462	.393
SVD	<b>.793</b>	.691	<b>.778</b>	.666	<b>.514</b>	.432
SGNS	<b>.793</b>	.685	.774	<b>.693</b>	.470	<b>.438</b>
GloVe	.725	.604	.729	.632	.403	.398

- ▶ SVD = singular value decomposition on PMI matrix
- ▶ GloVe does not appear to be the best when experiments are carefully controlled, but it depends on hyperparameters + these distinctions don't matter in practice

# Hypernymy Detection

---

- ▶ Hypernyms: detective *is a* person, dog *is a* animal

# Hypernymy Detection

---

- ▶ Hypernyms: detective *is a* person, dog *is a* animal
- ▶ Do word vectors encode these relationships?

# Hypernymy Detection

---

- ▶ Hypernyms: detective *is a* person, dog *is a* animal
- ▶ Do word vectors encode these relationships?

Dataset	TM14	Kotlerman 2010	HypeNet	WordNet	Avg (10 datasets)
Random	52.0	30.8	24.5	55.2	23.2
Word2Vec + C	52.1	<b>39.5</b>	20.7	<b>63.0</b>	25.3
GE + C	53.9	36.0	21.6	58.2	26.1
GE + KL	52.0	39.4	23.7	54.4	25.9
DIVE + C· $\Delta$ S	<b>57.2</b>	36.6	<b>32.0</b>	60.9	<b>32.7</b>

# Hypernymy Detection

---

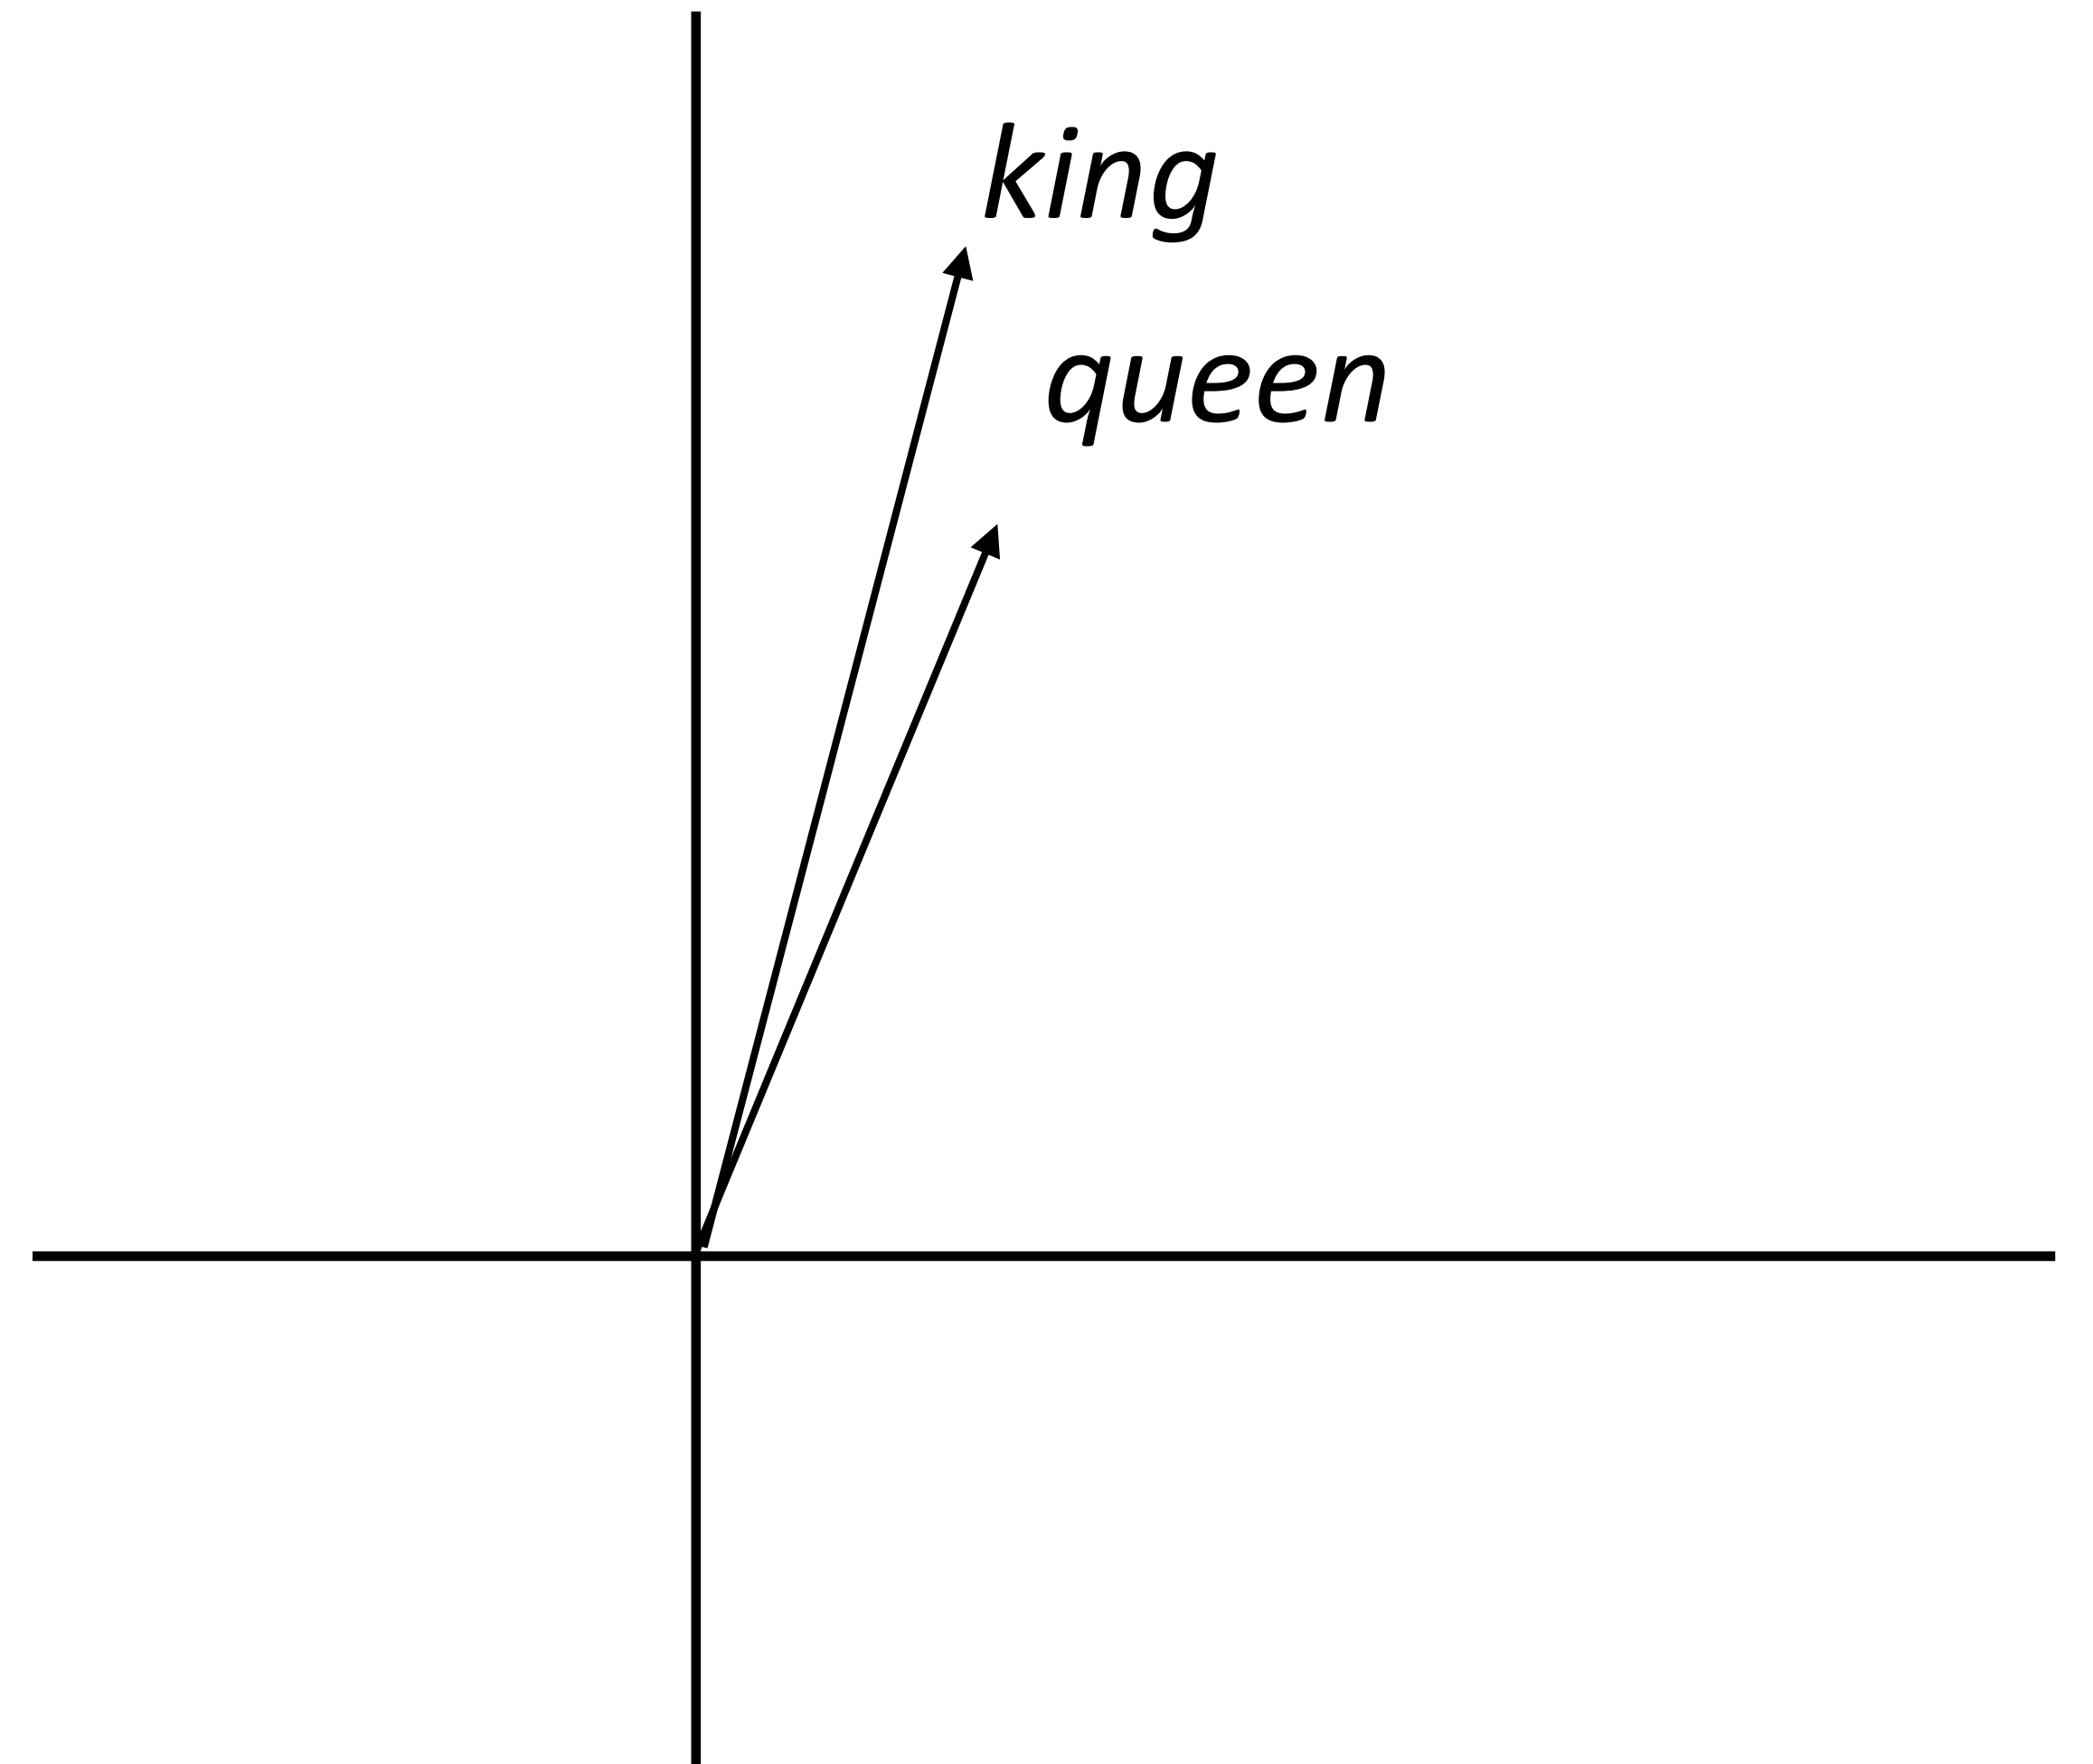
- ▶ Hypernyms: detective *is a* person, dog *is a* animal
- ▶ Do word vectors encode these relationships?

Dataset	TM14	Kotlerman 2010	HypeNet	WordNet	Avg (10 datasets)
Random	52.0	30.8	24.5	55.2	23.2
Word2Vec + C	52.1	<b>39.5</b>	20.7	<b>63.0</b>	25.3
GE + C	53.9	36.0	21.6	58.2	26.1
GE + KL	52.0	39.4	23.7	54.4	25.9
DIVE + C· $\Delta$ S	<b>57.2</b>	36.6	<b>32.0</b>	60.9	<b>32.7</b>

- ▶ word2vec (SGNS) works barely better than random guessing here

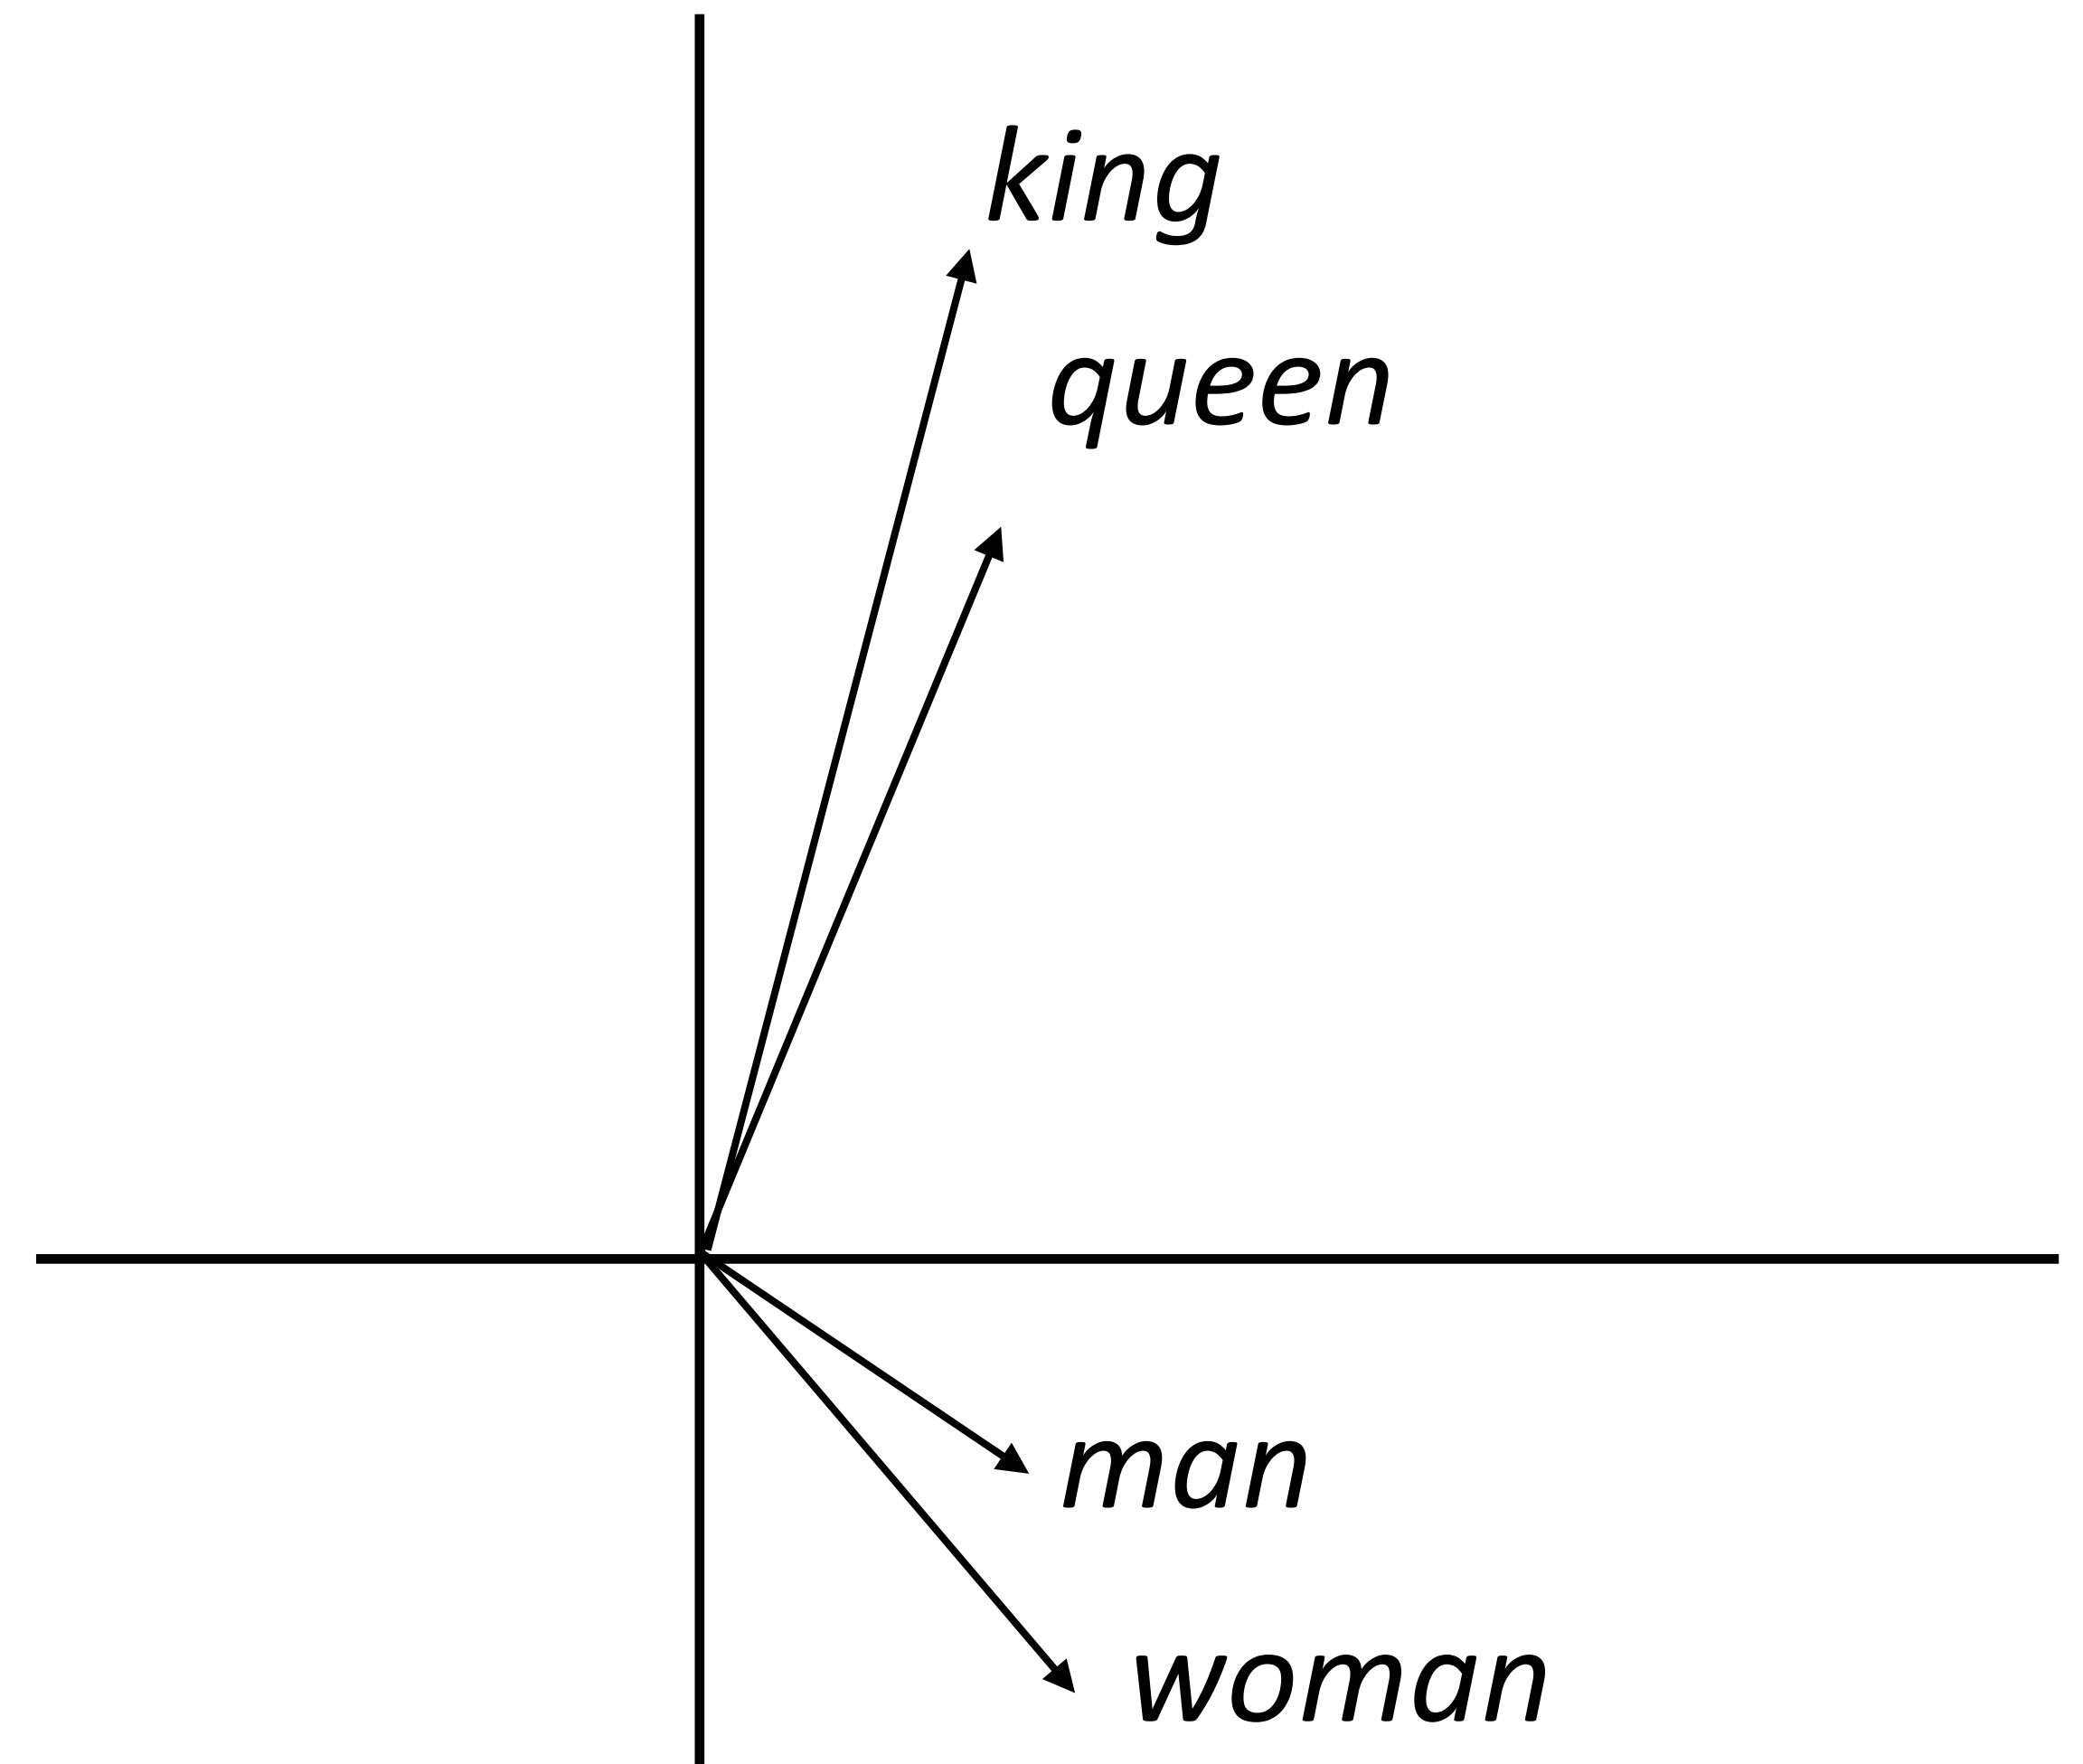
# Analogies

---



# Analogies

---






# Analogies

---

*(king - man) + woman = queen*

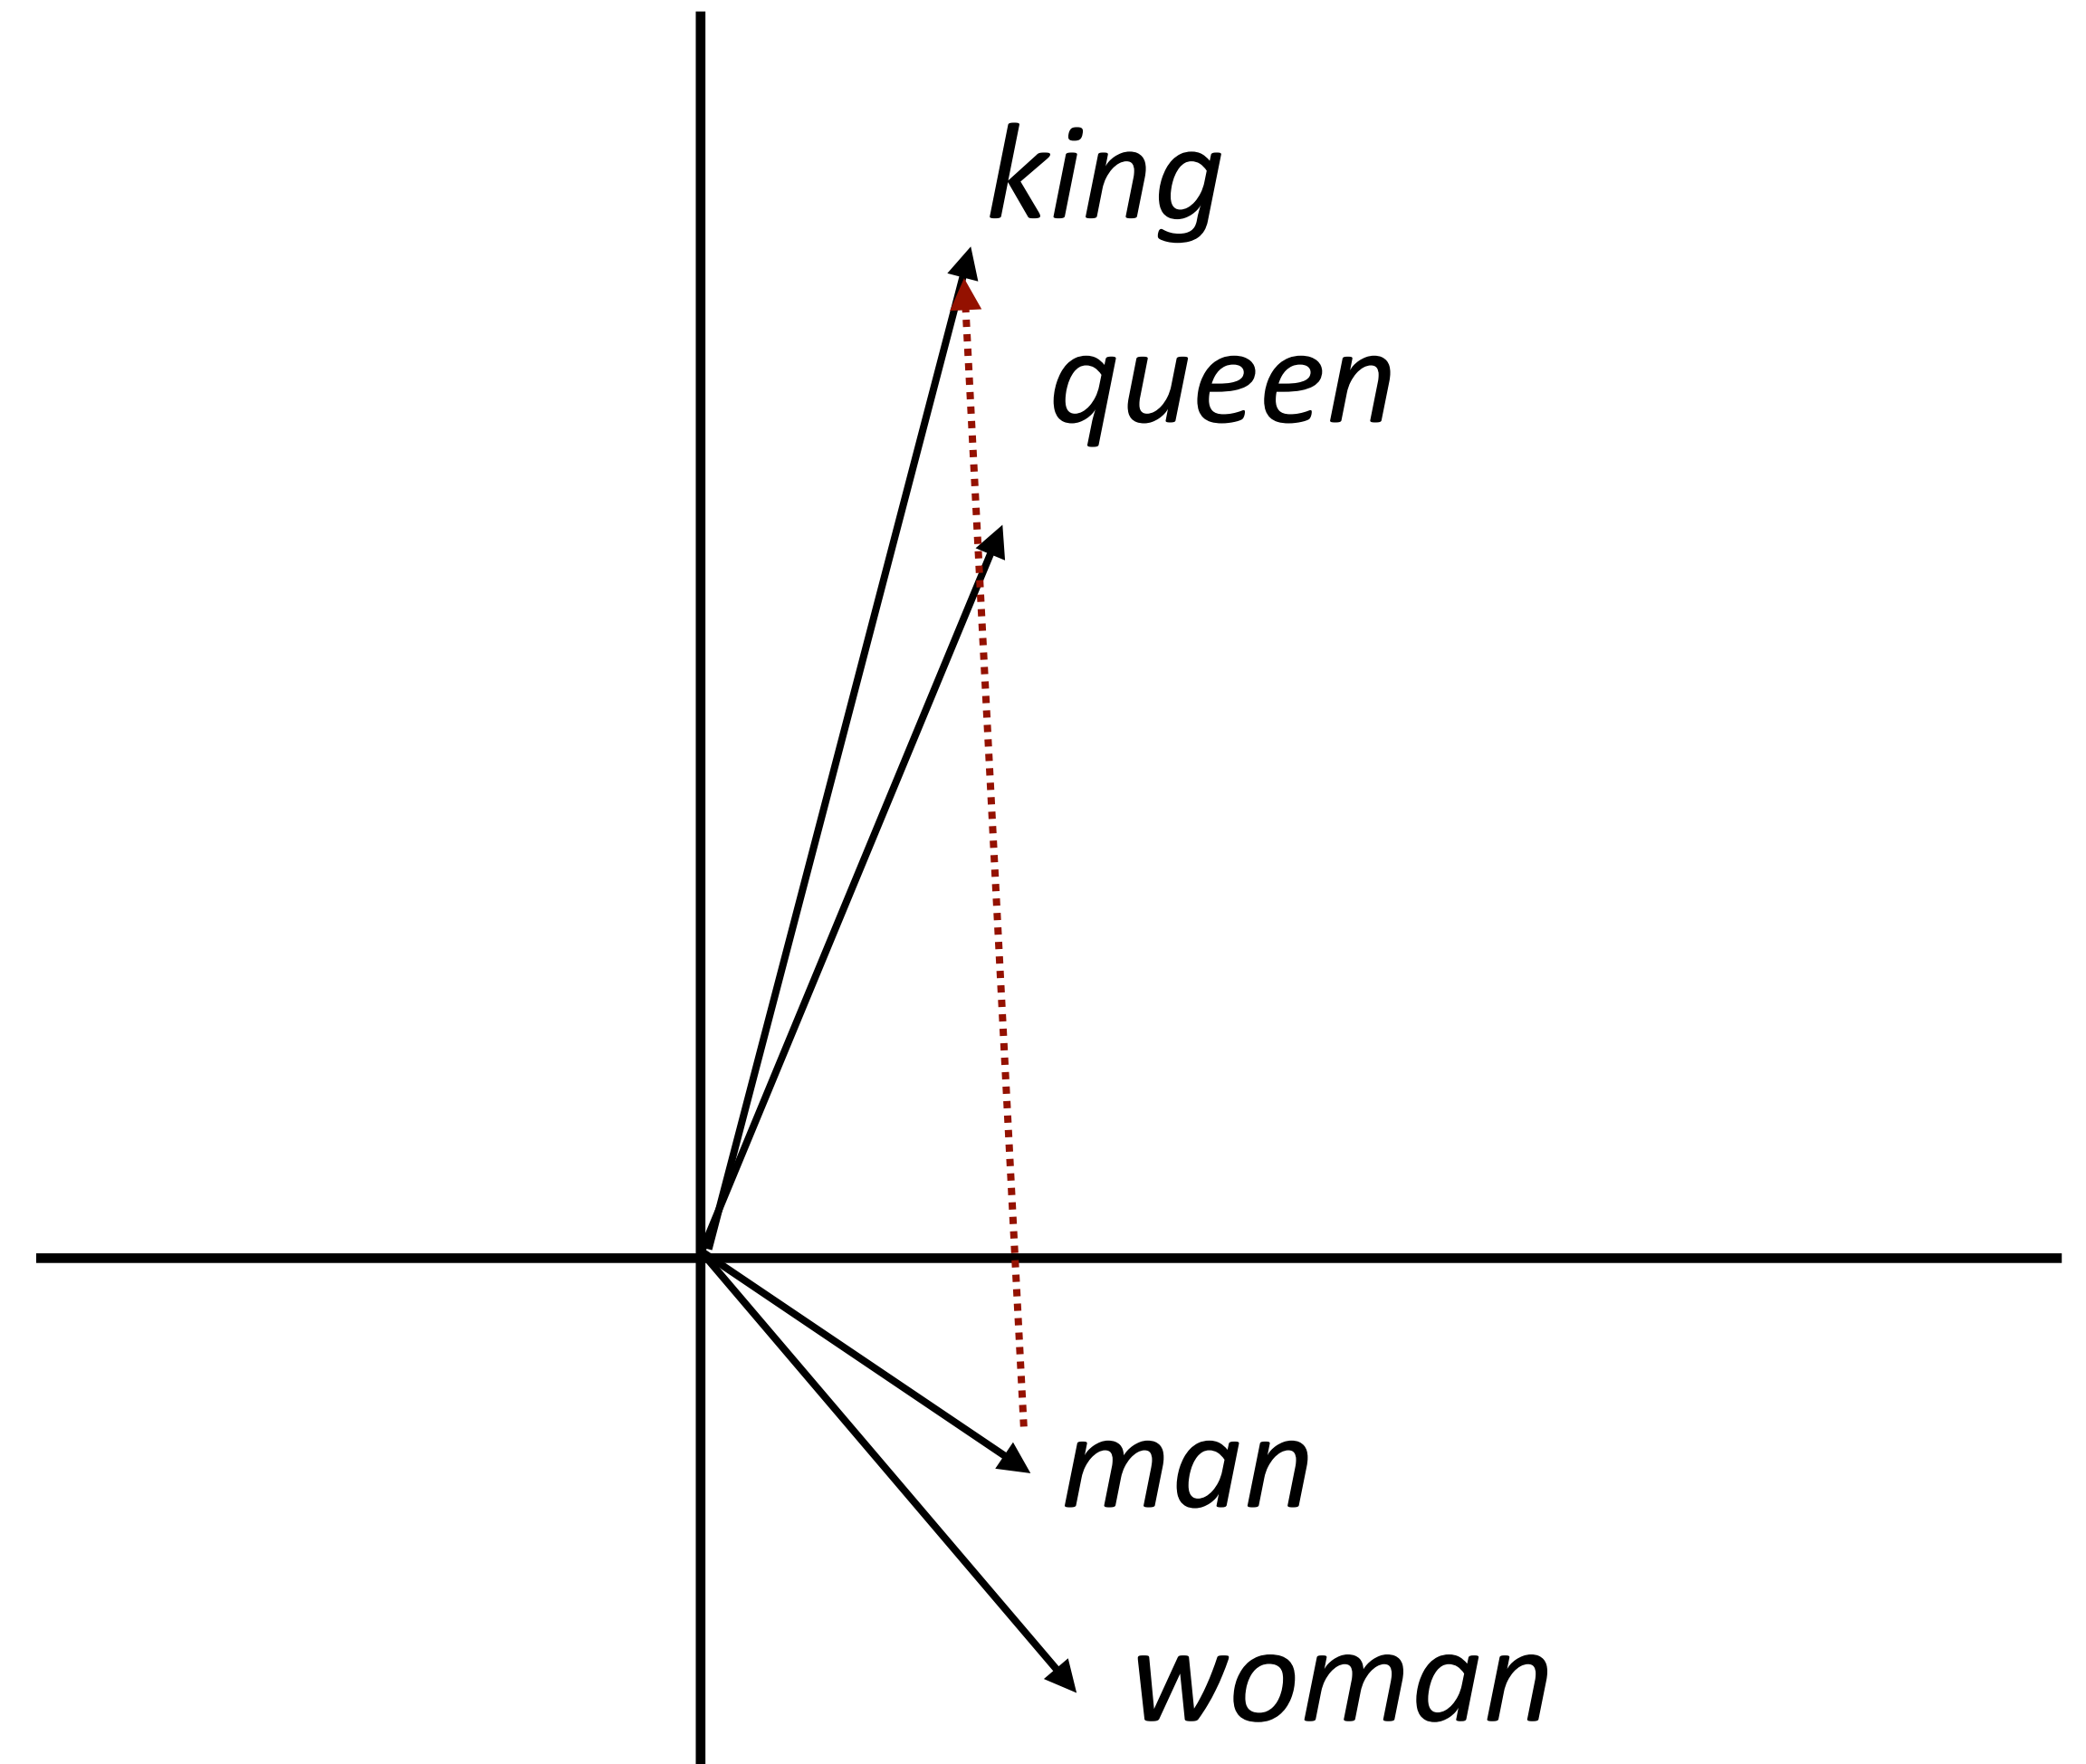




# Analogies

---

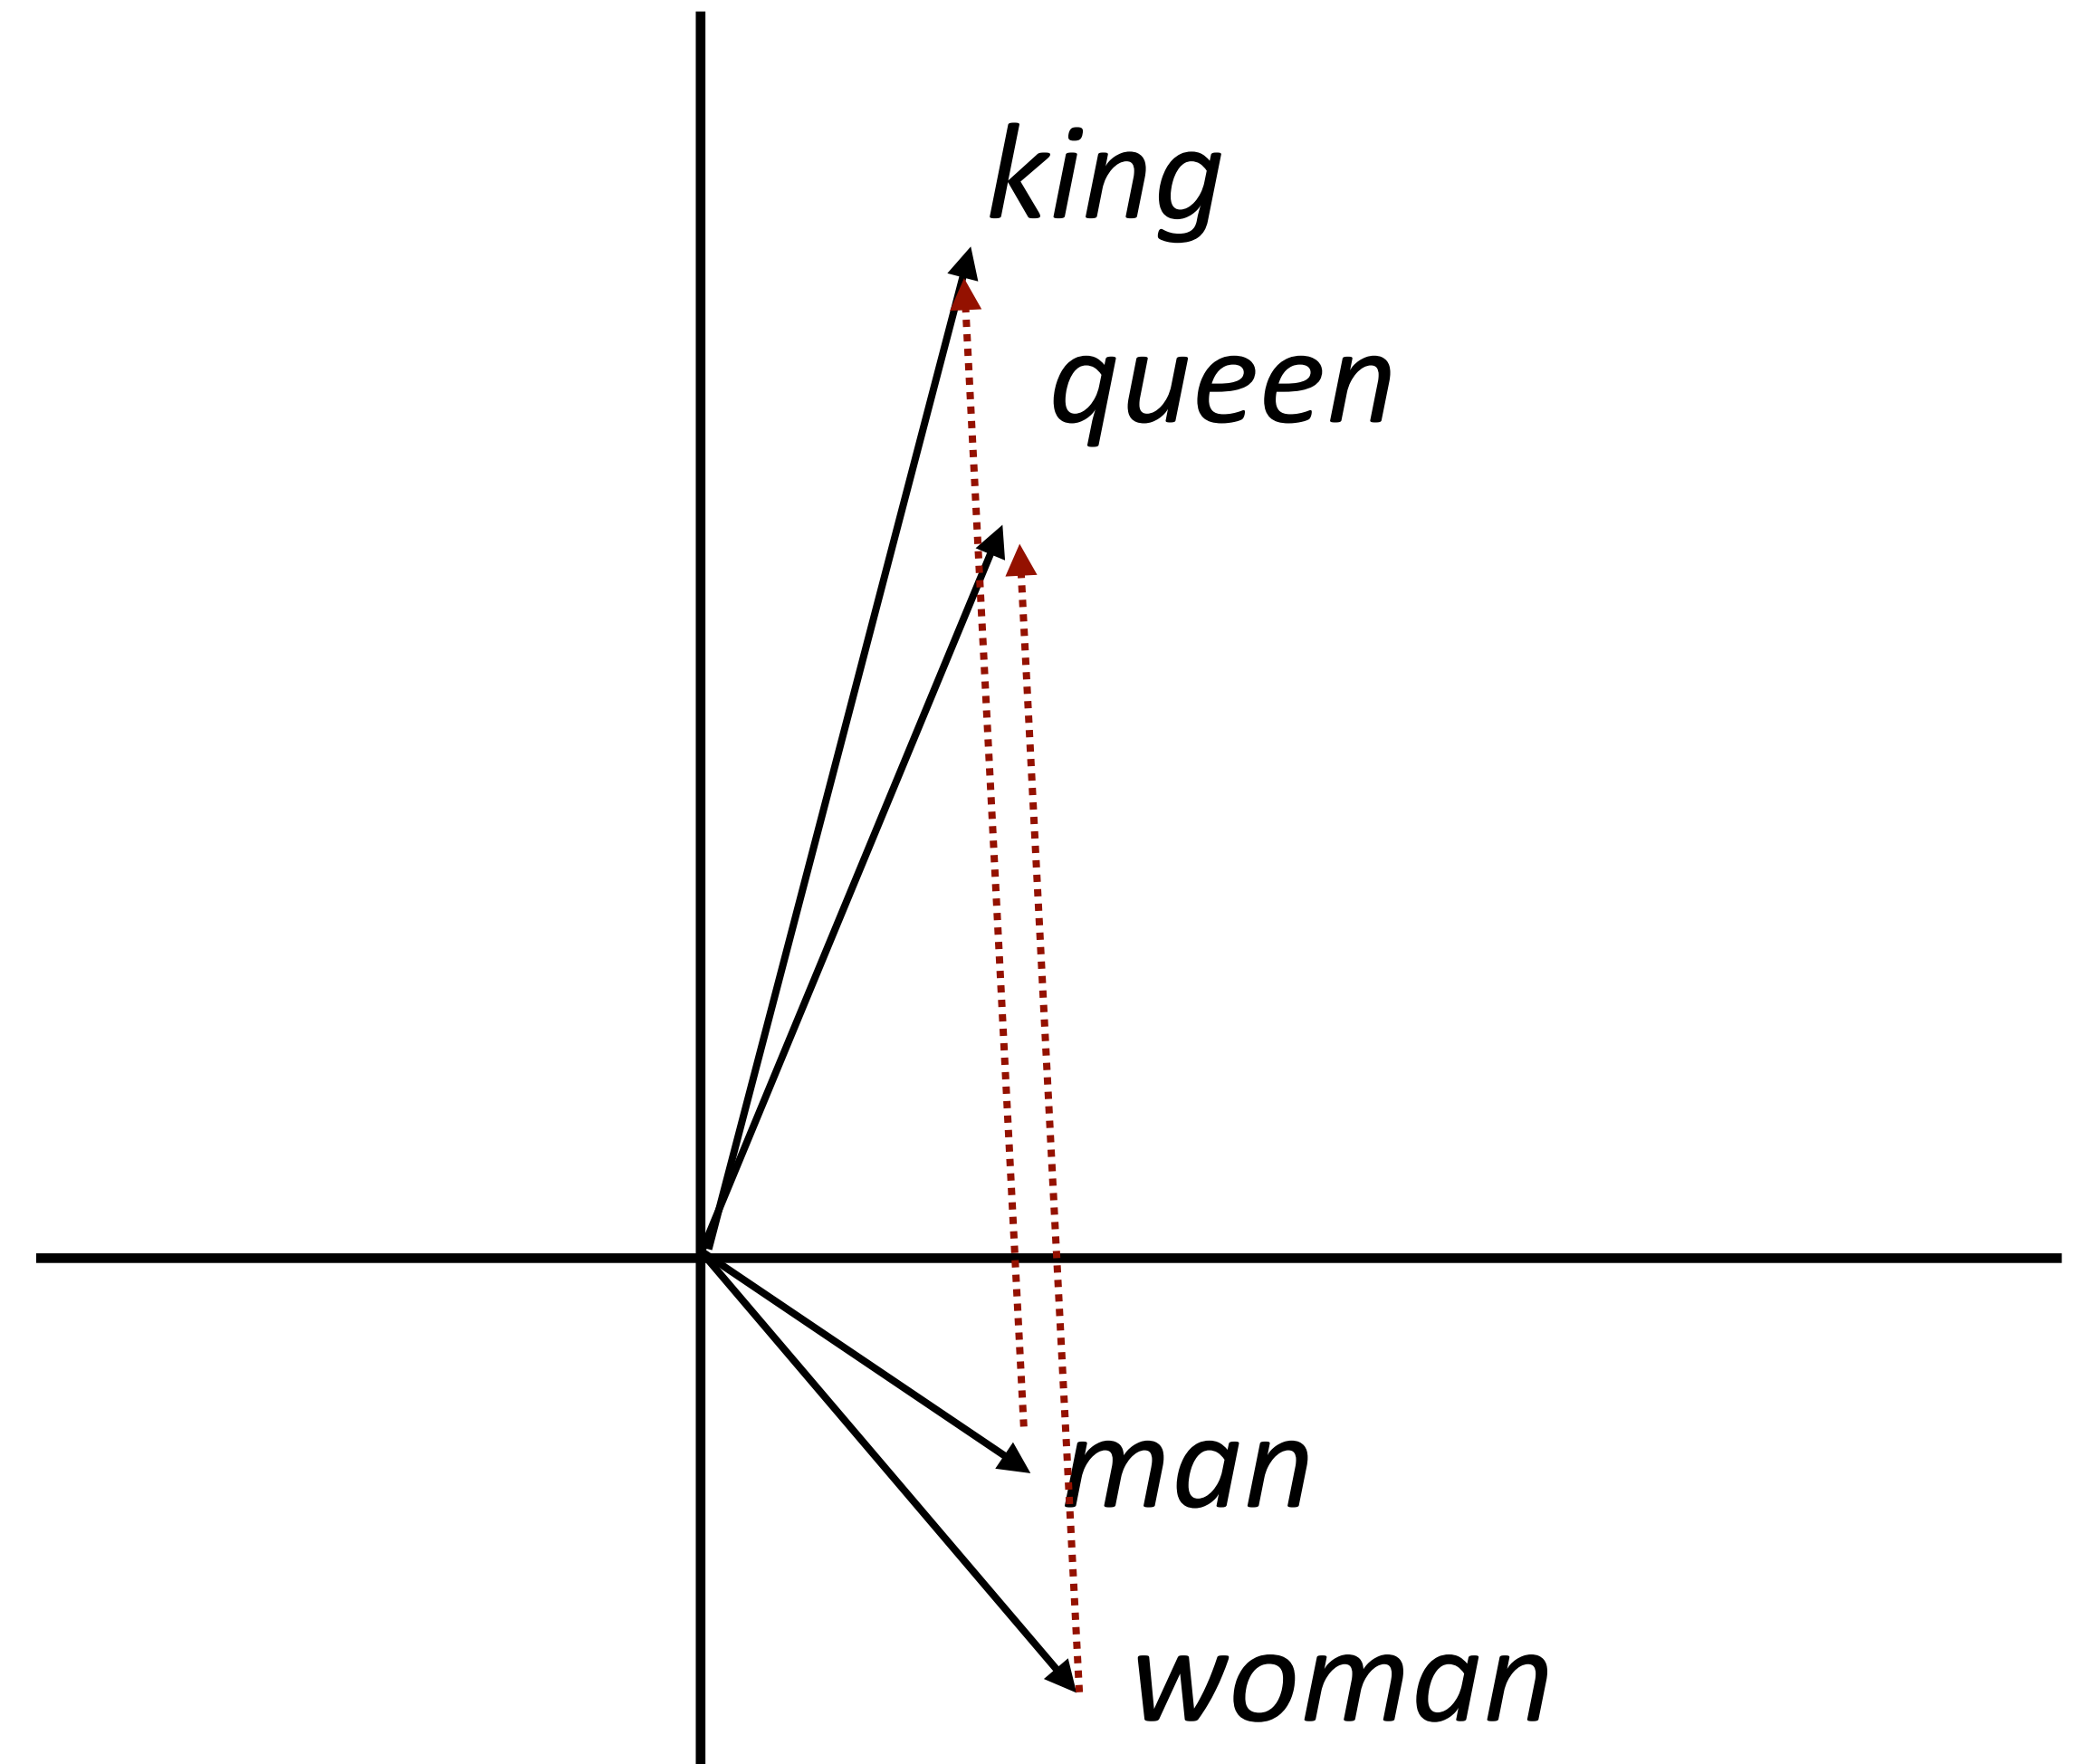
*(king - man) + woman = queen*



# Analogies

---

*(king - man) + woman = queen*

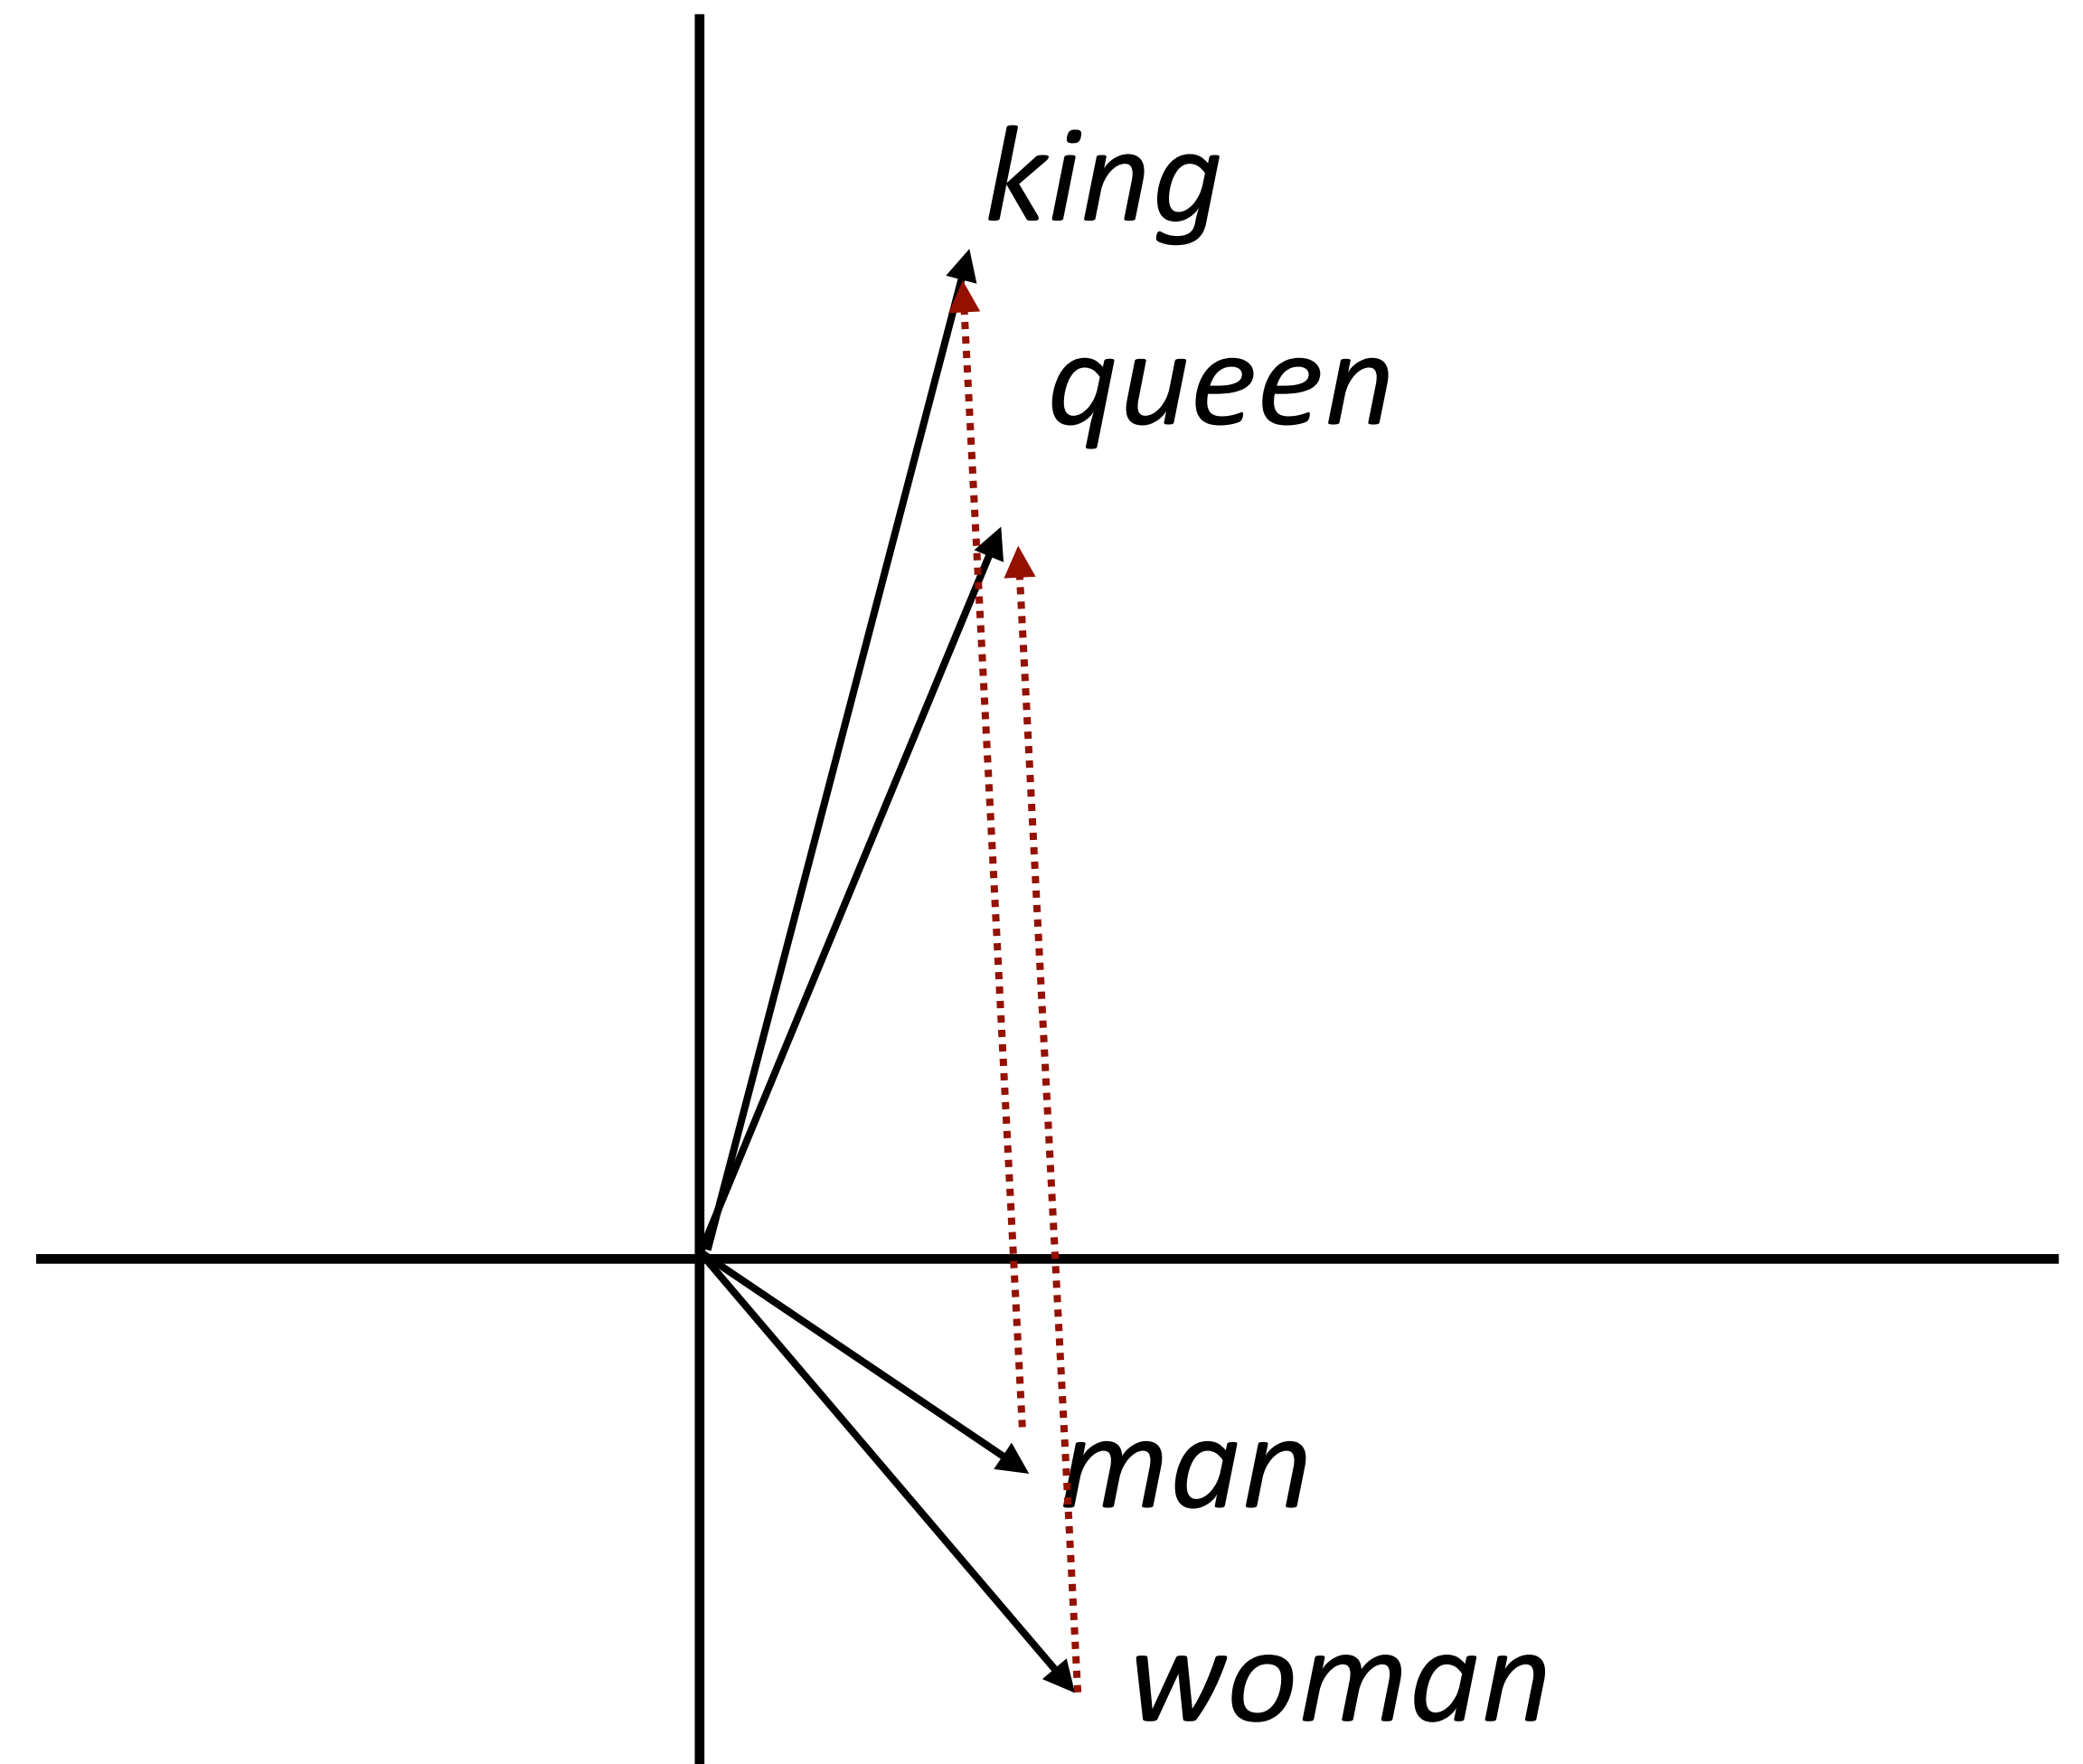


# Analogies

---

*(king - man) + woman = queen*

*king + (woman - man) = queen*



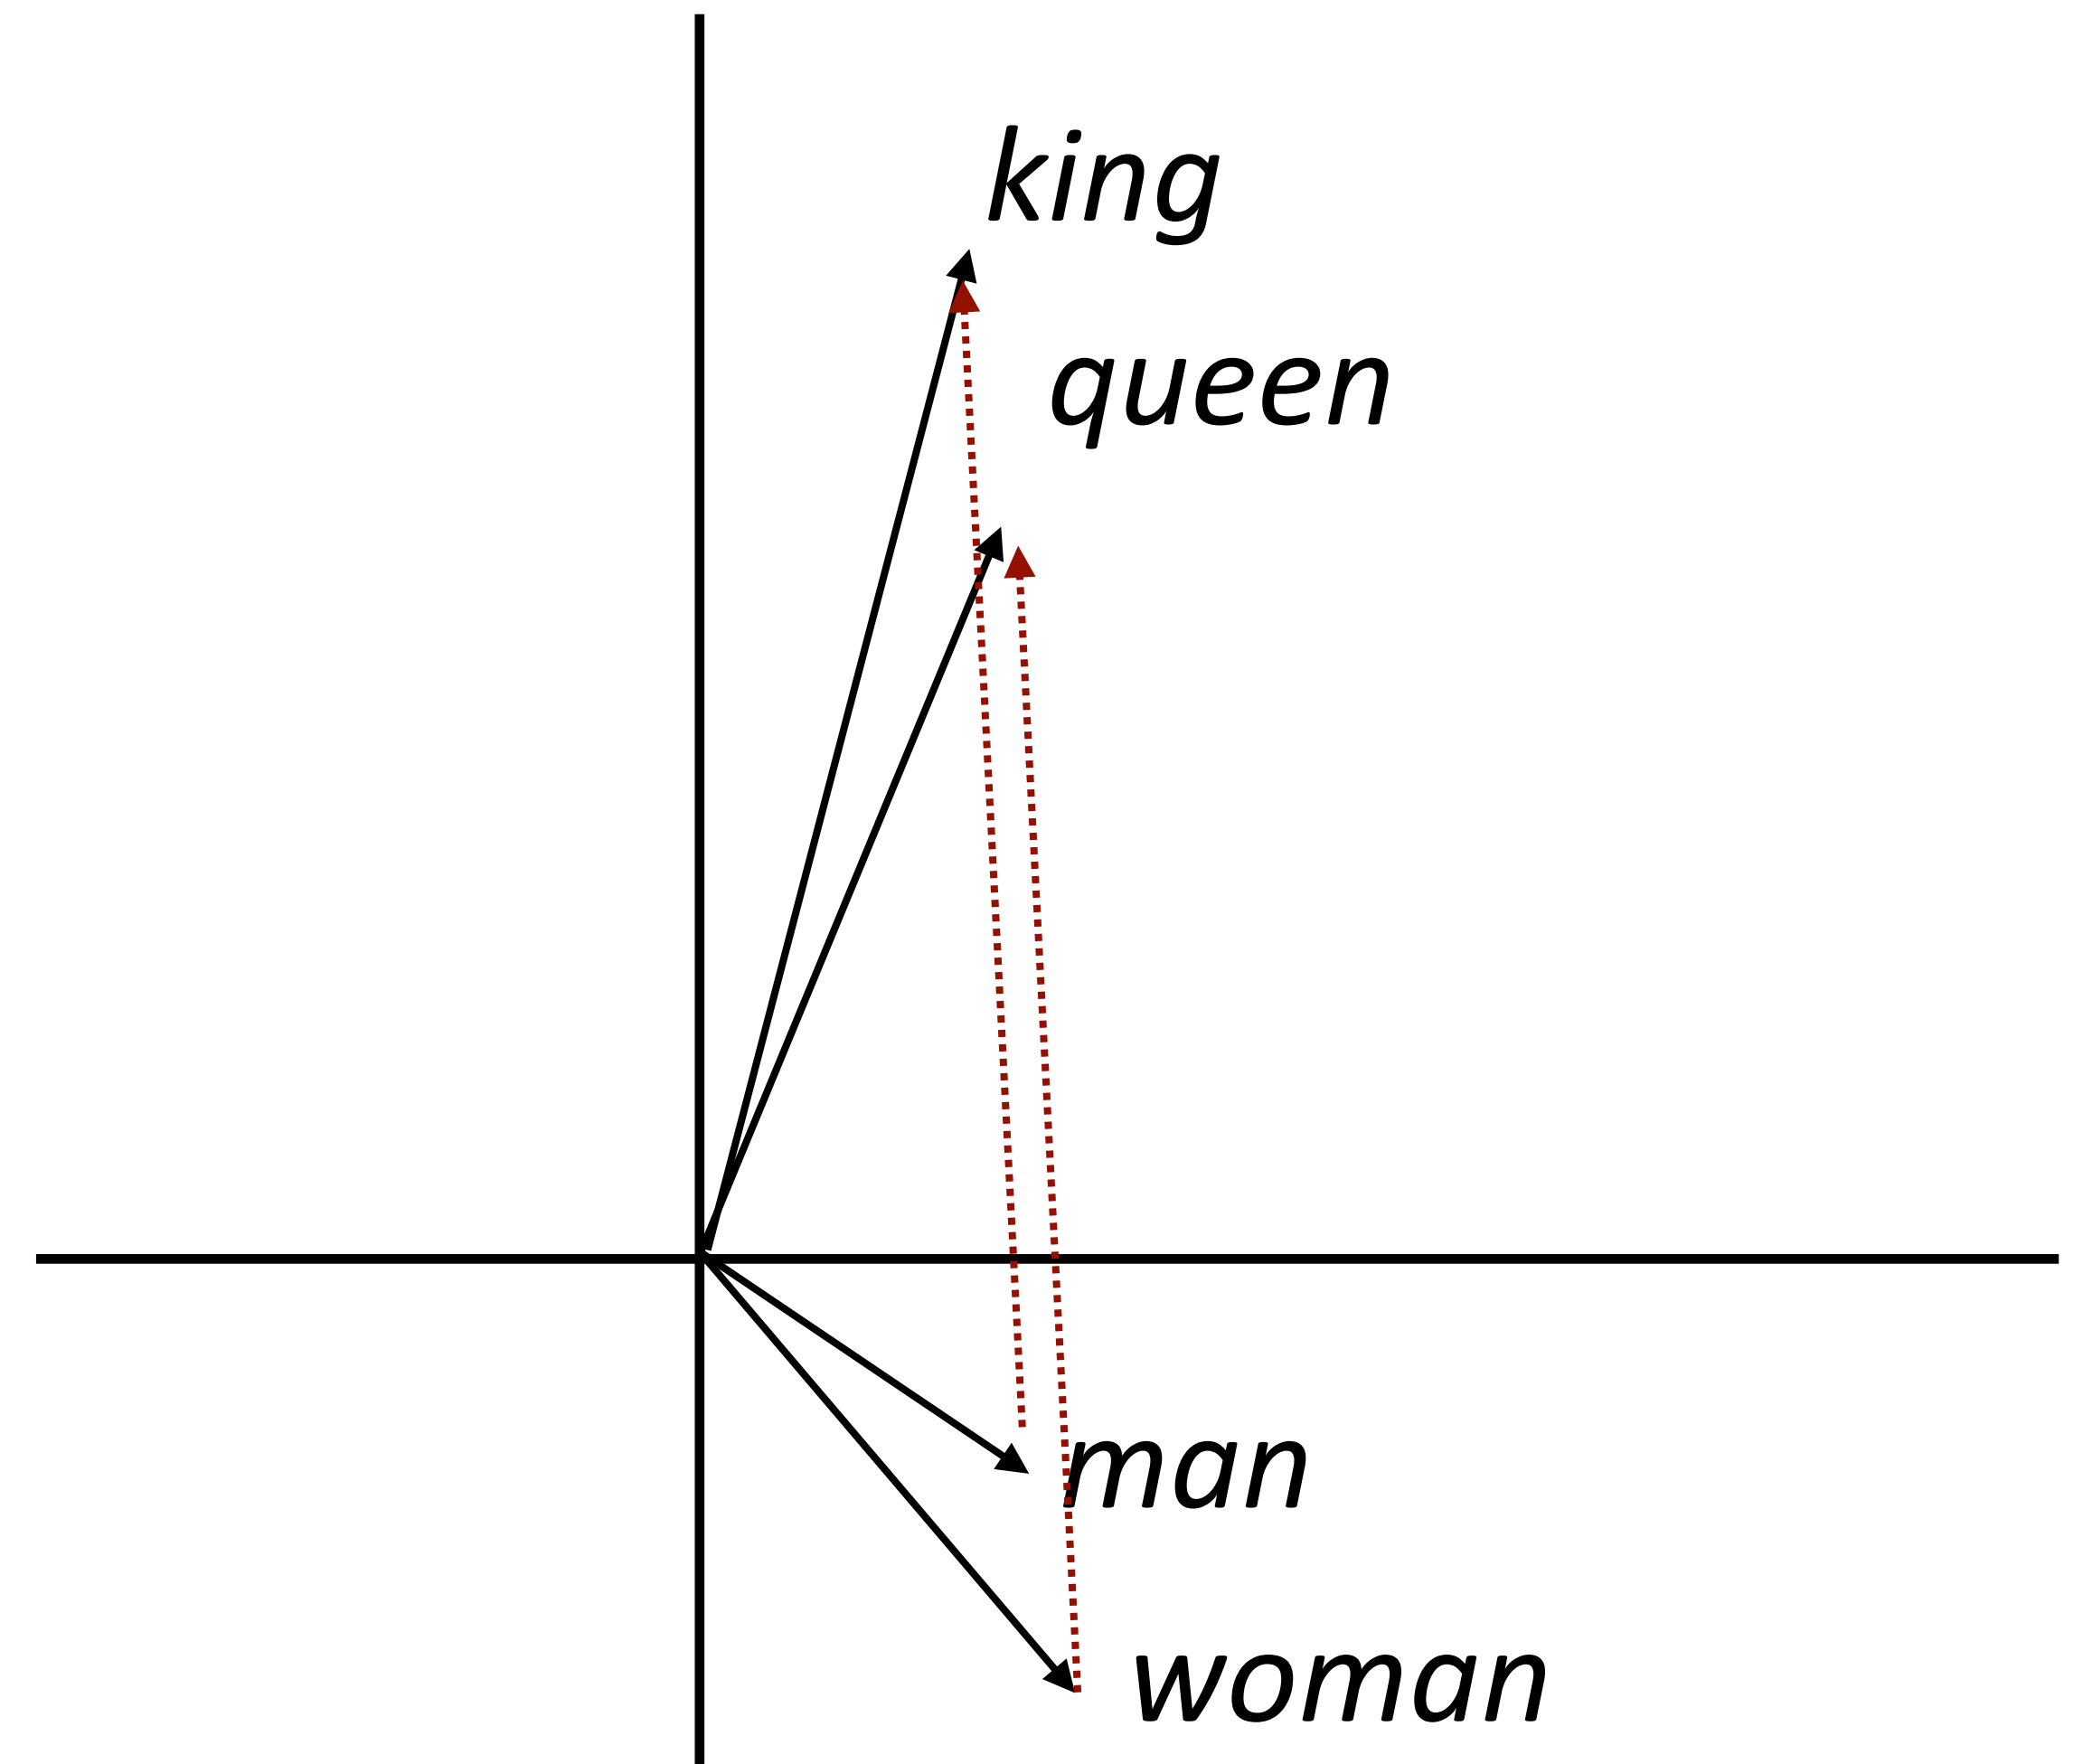
# Analogies

---

*(king - man) + woman = queen*

*king + (woman - man) = queen*

- Why would this be?



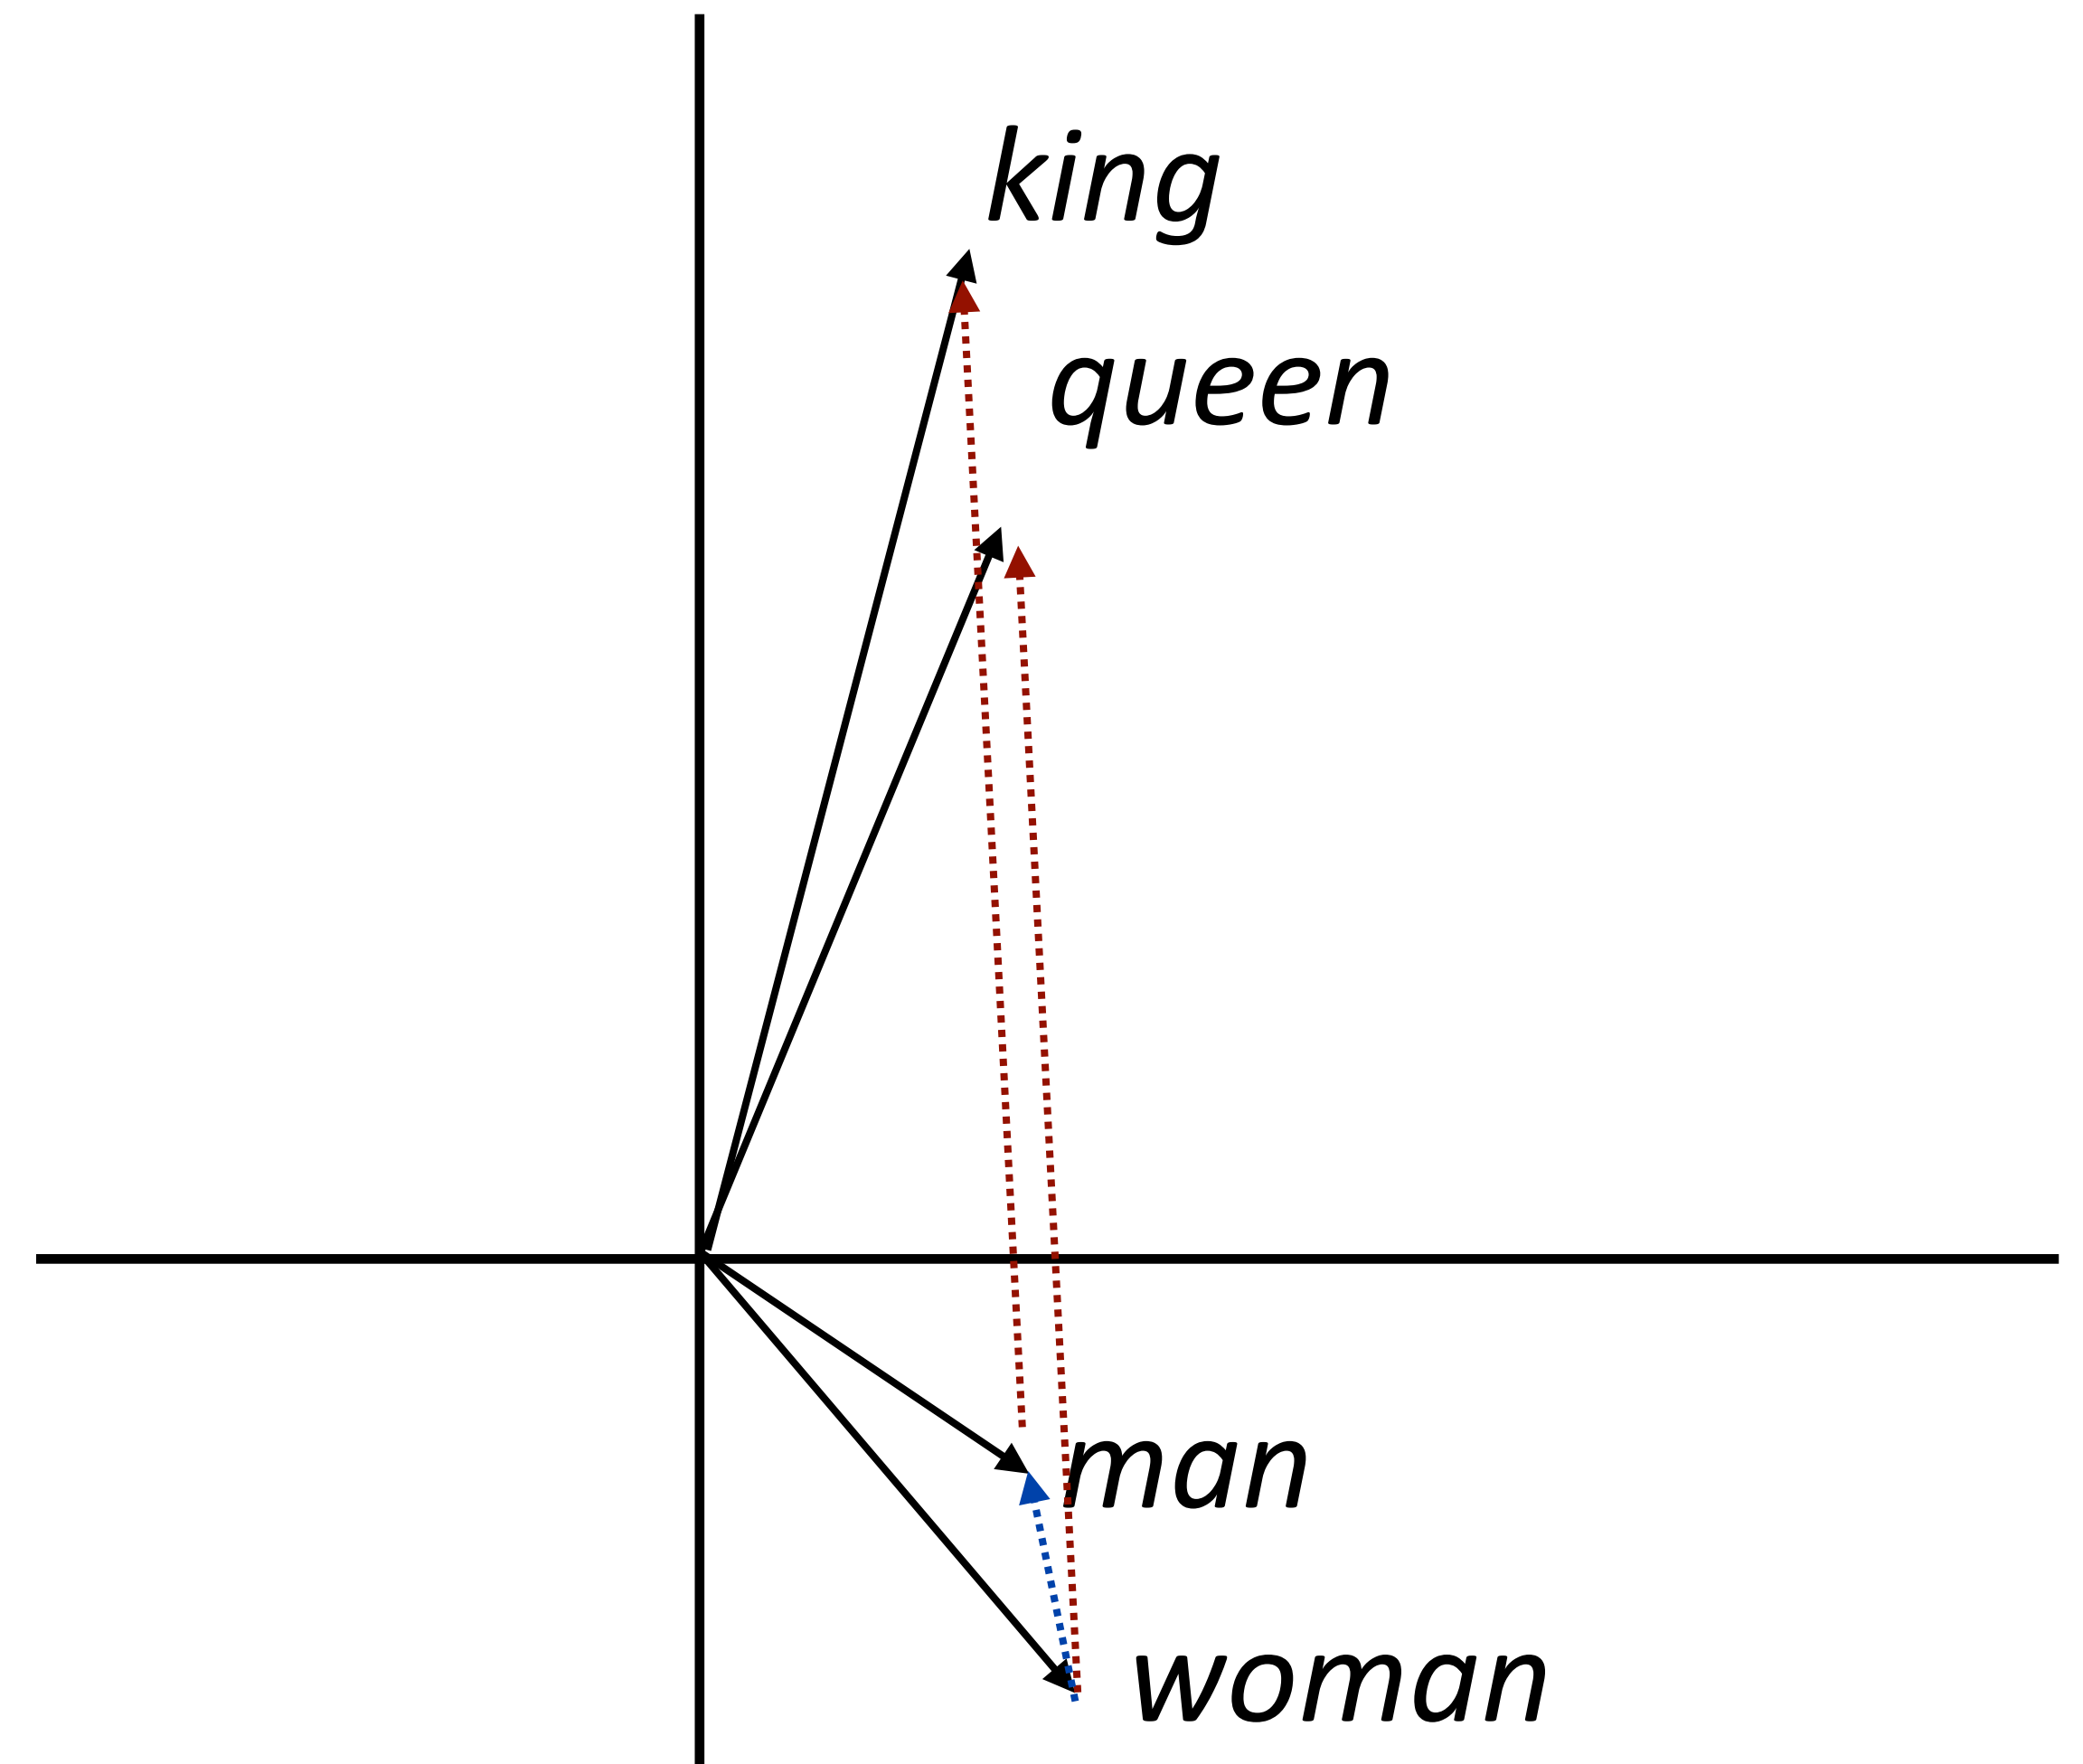
# Analogy

---

$(king - man) + woman = queen$

$king + (woman - man) = queen$

- Why would this be?



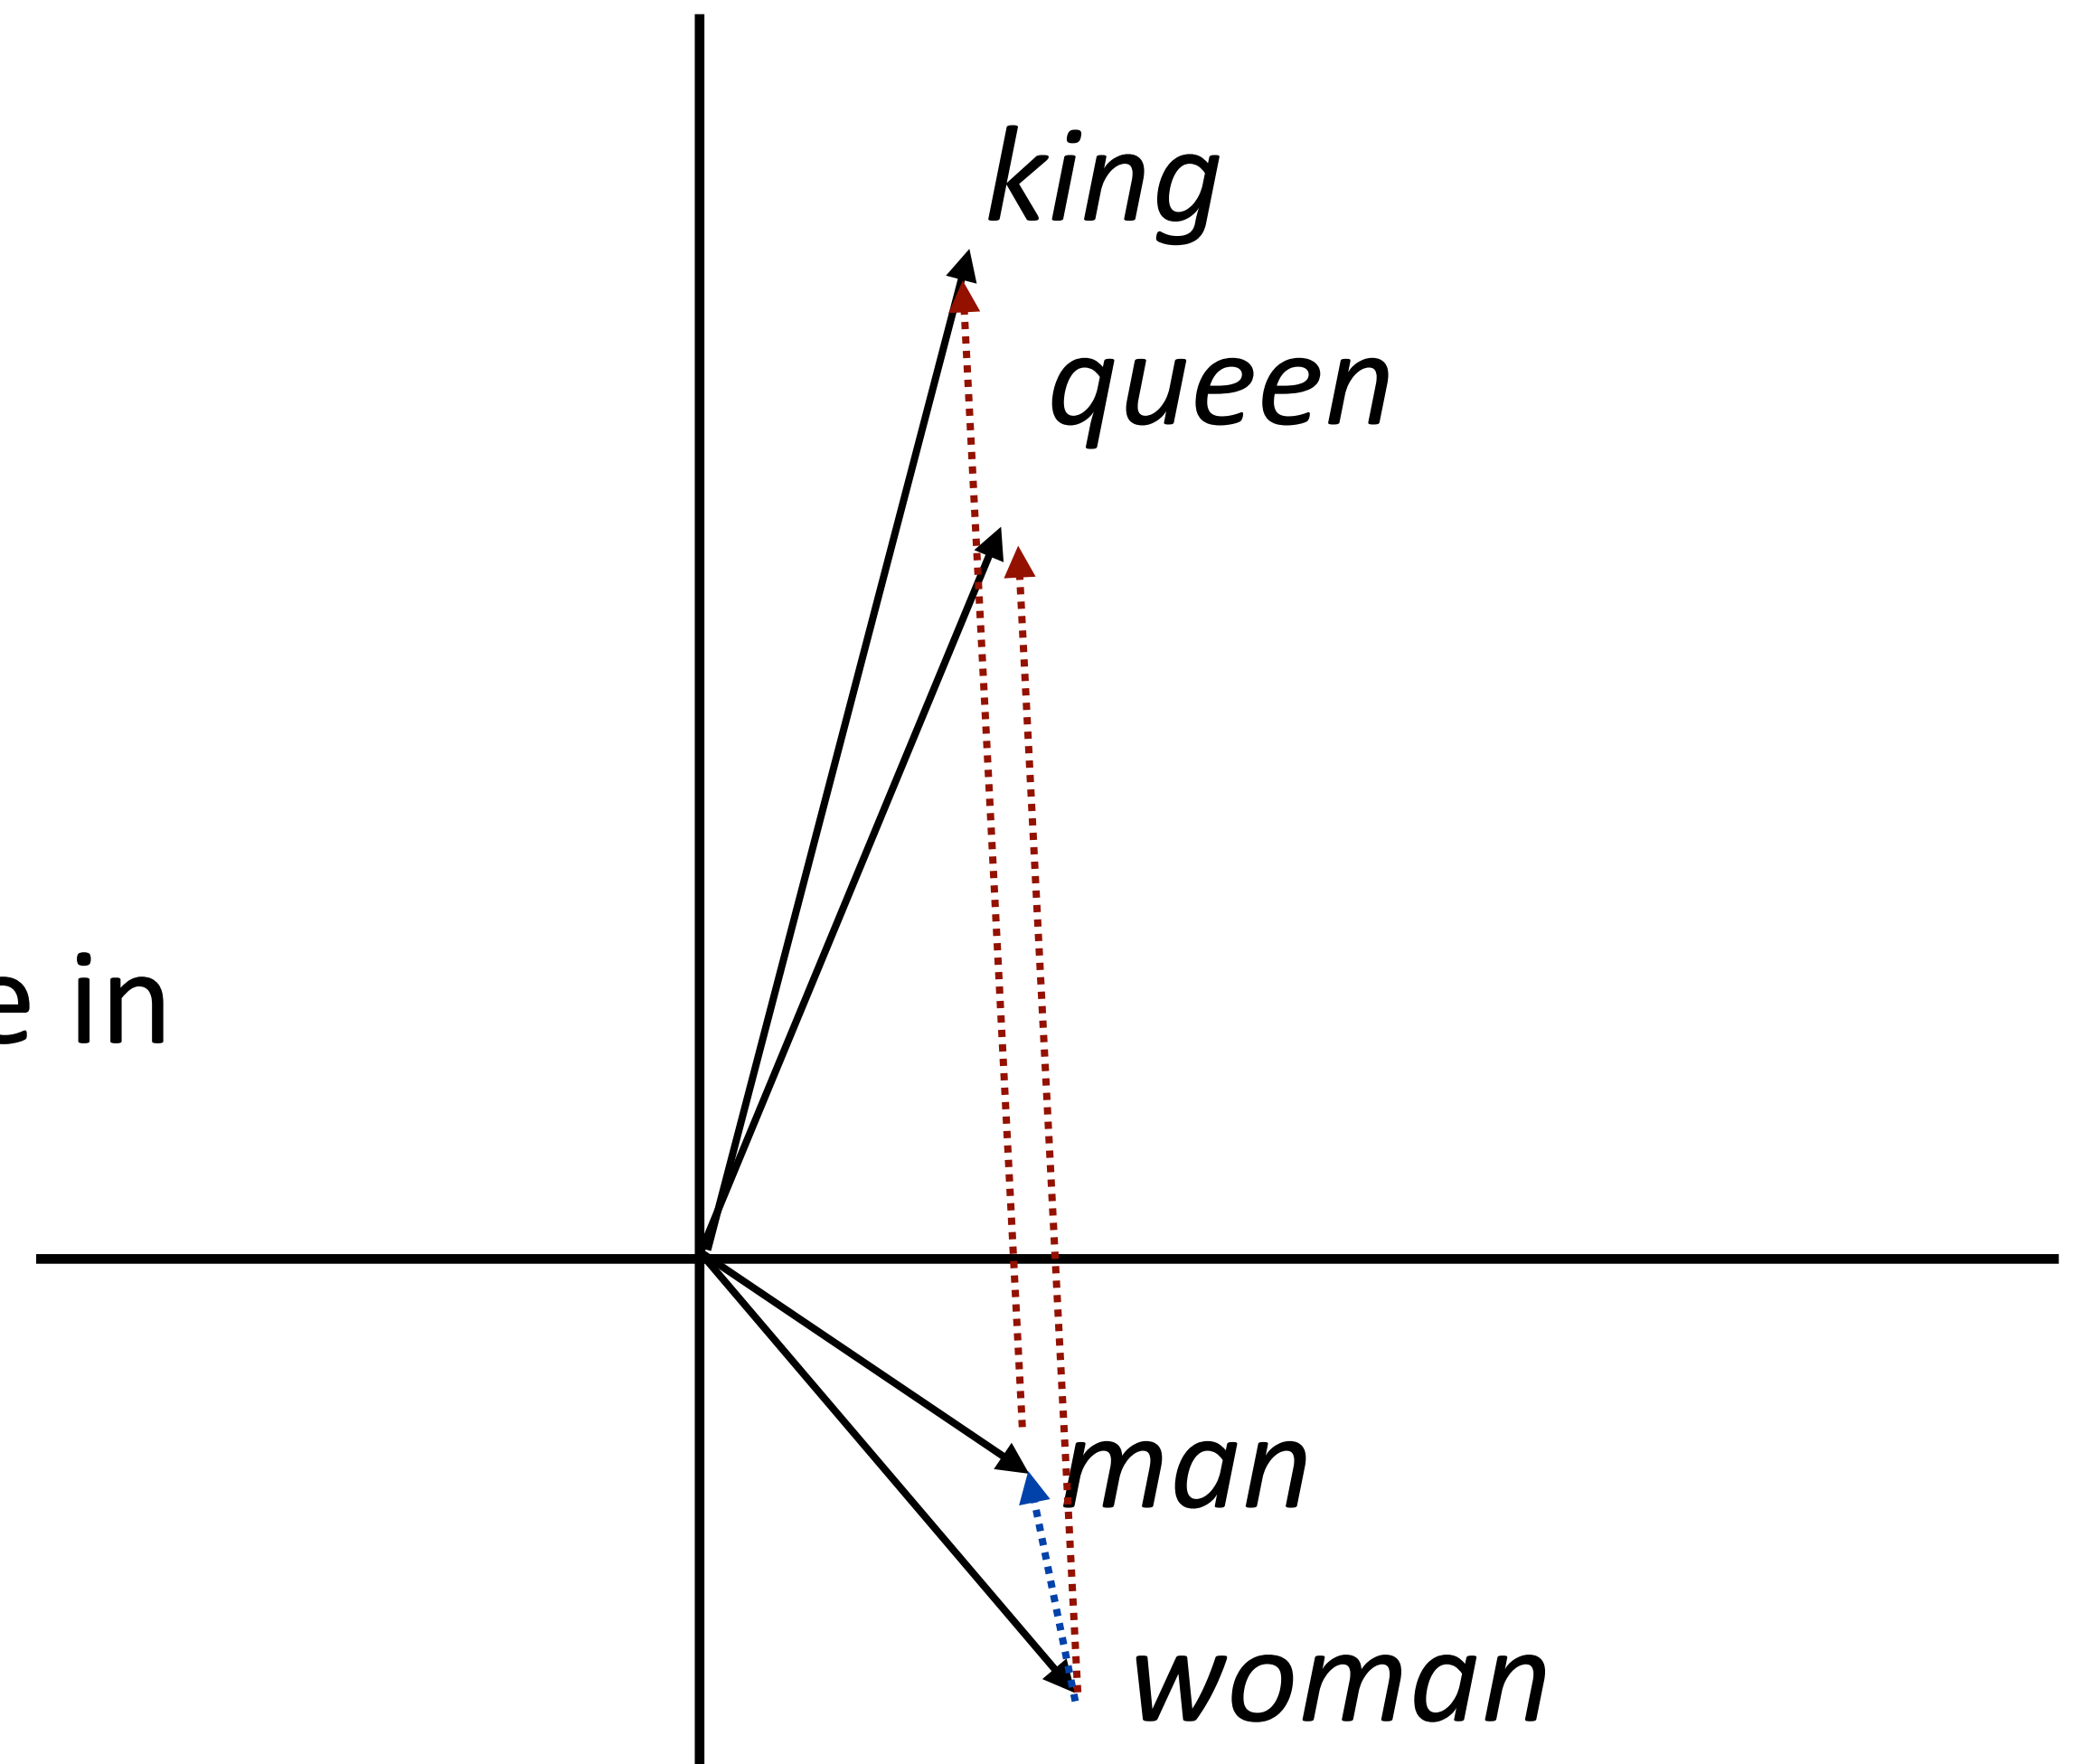
# Analogies

---

$(king - man) + woman = queen$

$king + (woman - man) = queen$

- ▶ Why would this be?
- ▶ woman - man captures the difference in the contexts that these occur in

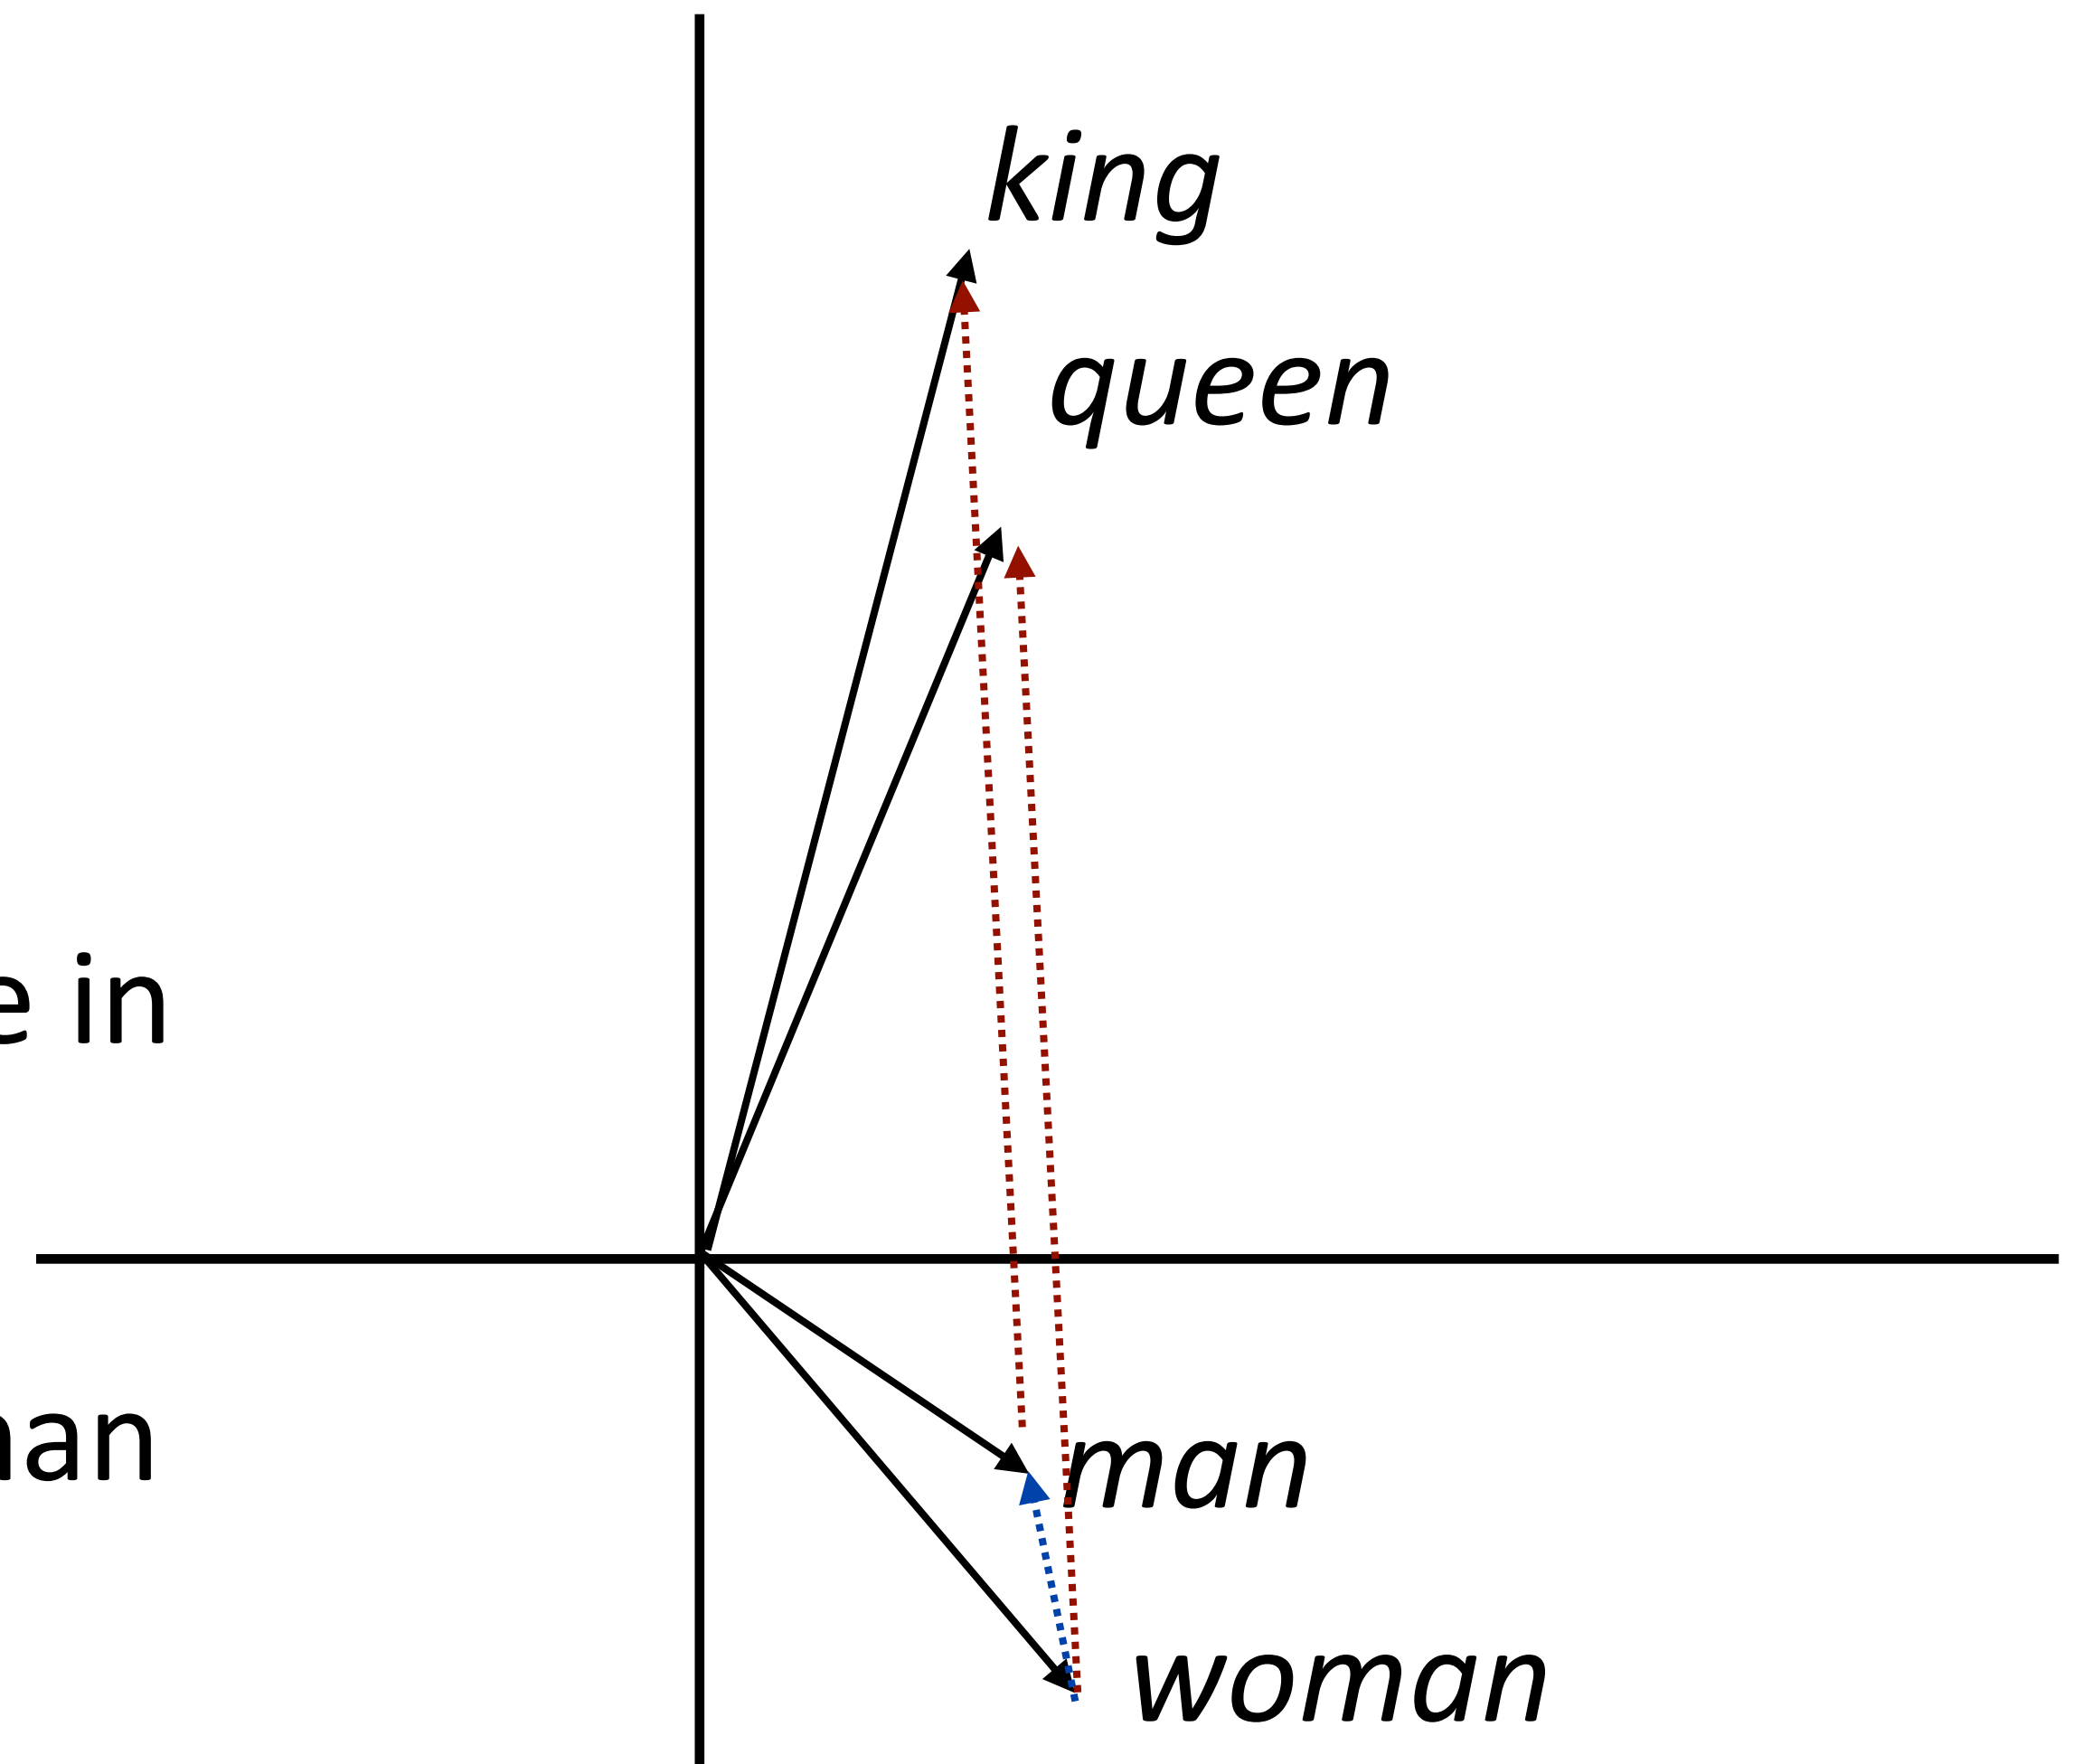


# Analogies

*(king - man) + woman = queen*

*king + (woman - man) = queen*

- ▶ Why would this be?
- ▶ woman - man captures the difference in the contexts that these occur in
- ▶ Dominant change: more “he” with man and “she” with woman — similar to difference between king and queen



# Analogies

---

Method	Google	MSR
	Add / Mul	Add / Mul
PPMI	.553 / .679	.306 / .535
SVD	.554 / .591	.408 / .468
SGNS	.676 / <b>.688</b>	.618 / <b>.645</b>
GloVe	.569 / .596	.533 / .580



# Analogies

---

Method	Google	MSR
	Add / Mul	Add / Mul
PPMI	.553 / .679	.306 / .535
SVD	.554 / .591	.408 / .468
SGNS	.676 / <b>.688</b>	.618 / <b>.645</b>
GloVe	.569 / .596	.533 / .580

- ▶ These methods can perform well on analogies on two different datasets using two different methods

# Analogies

---

Method	Google	MSR
	Add / Mul	Add / Mul
PPMI	.553 / .679	.306 / .535
SVD	.554 / .591	.408 / .468
SGNS	.676 / <b>.688</b>	.618 / <b>.645</b>
GloVe	.569 / .596	.533 / .580

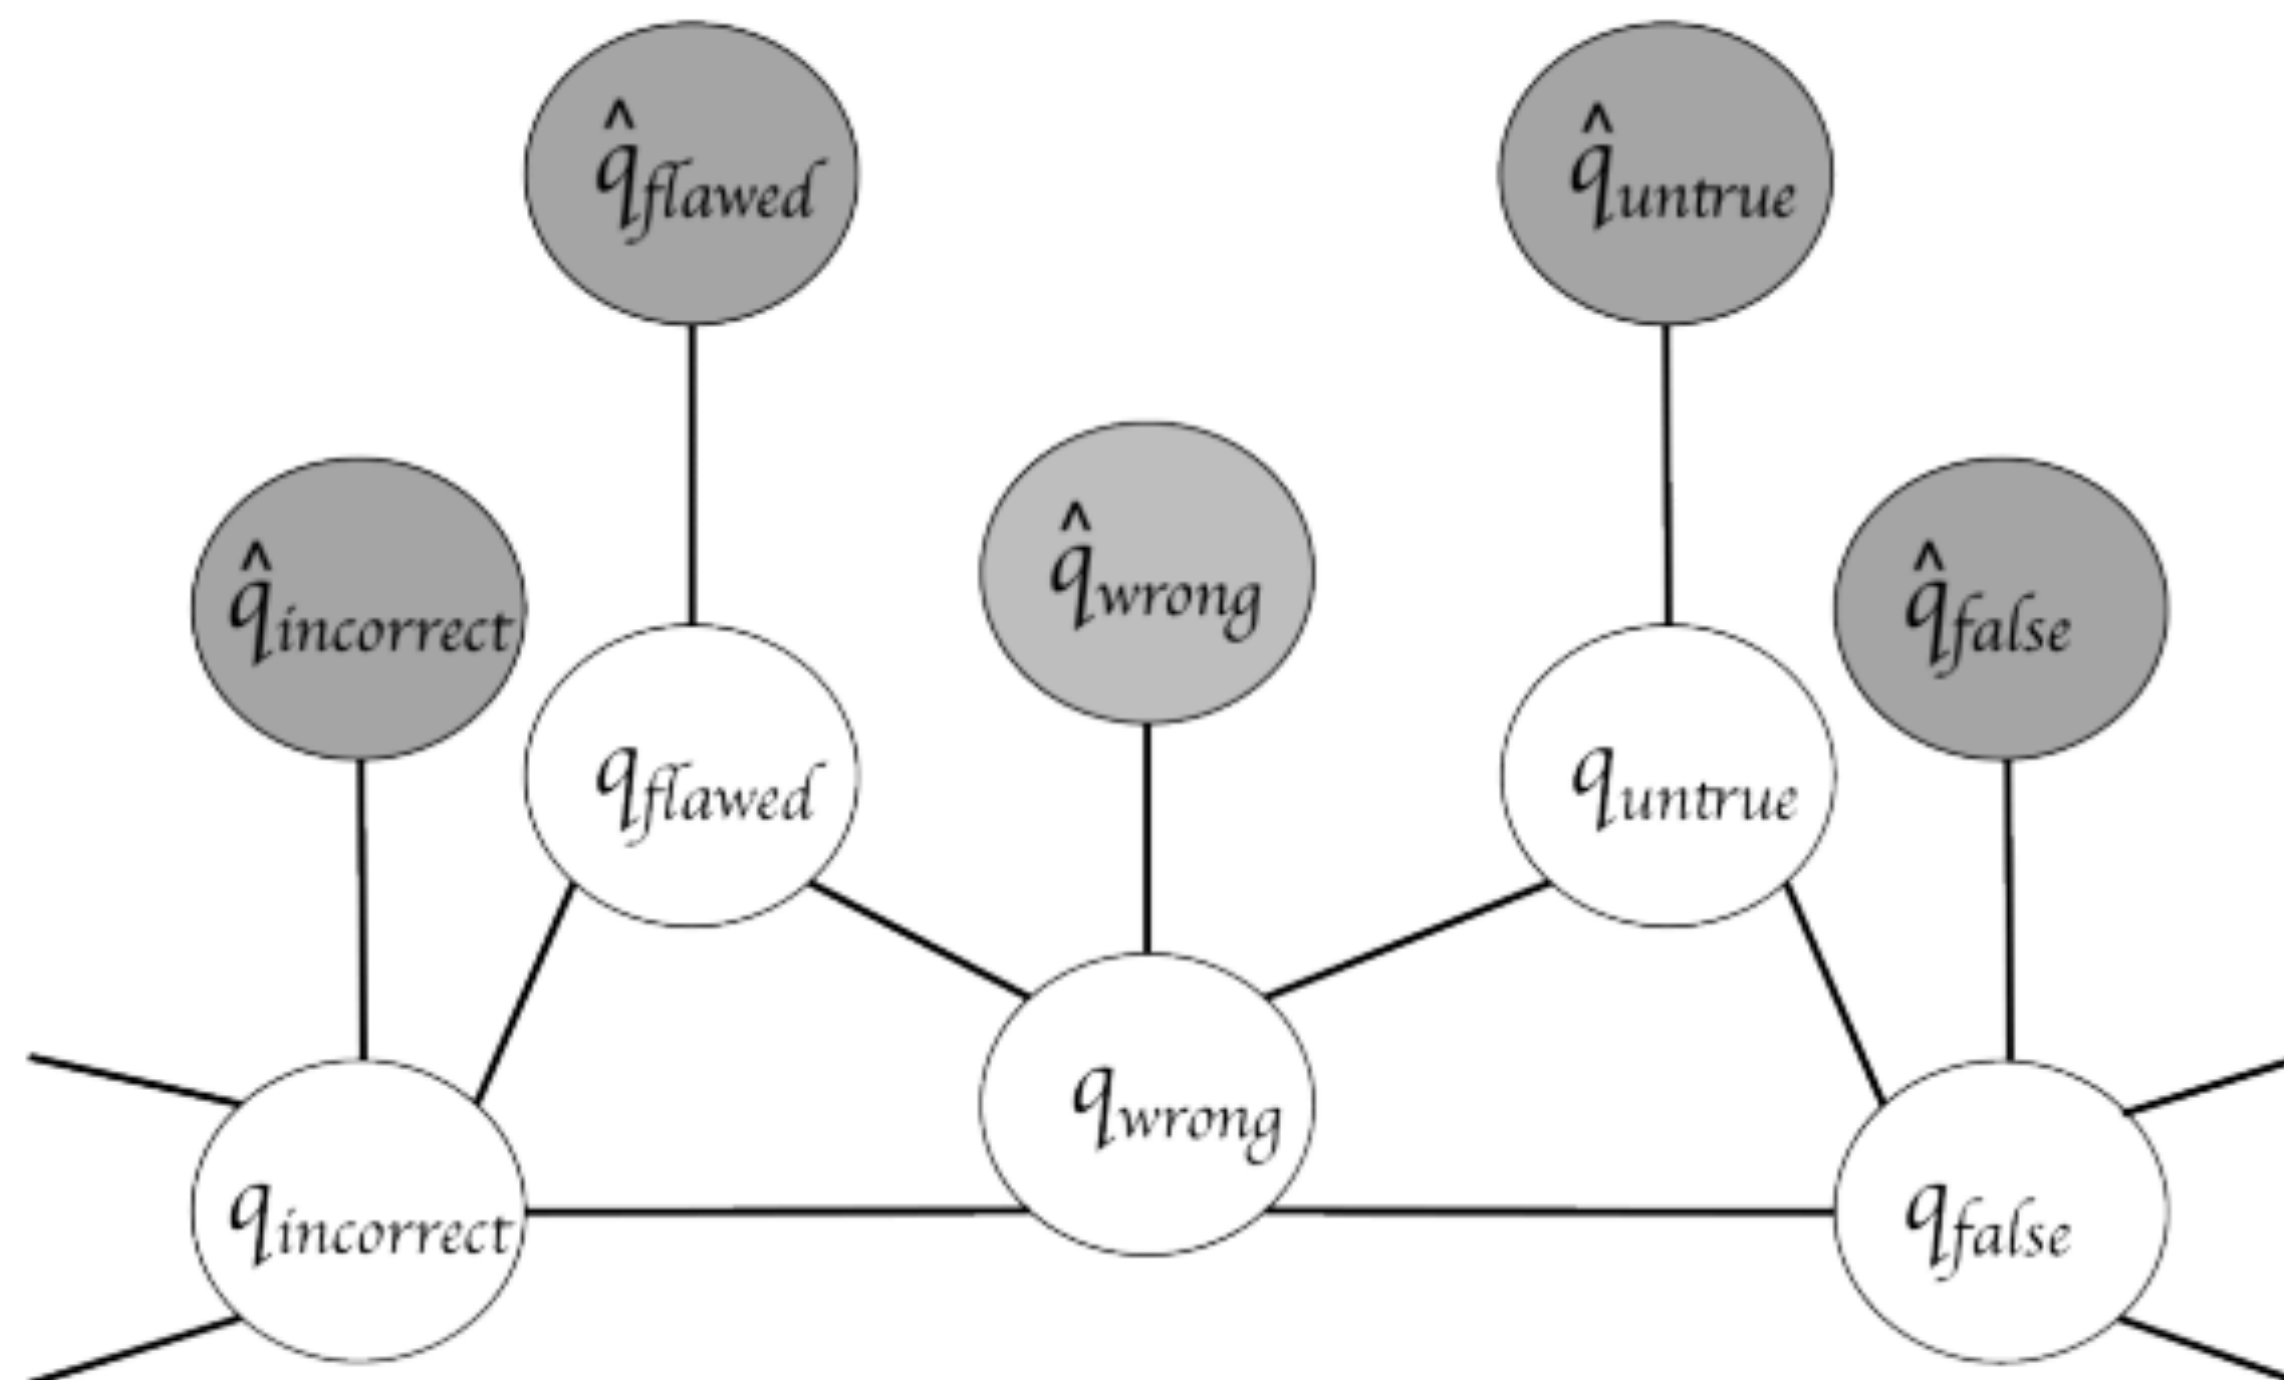
- ▶ These methods can perform well on analogies on two different datasets using two different methods

$$\text{Maximizing for } b: \text{Add} = \cos(b, a_2 - a_1 + b_1) \quad \text{Mul} = \frac{\cos(b_2, a_2) \cos(b_2, b_1)}{\cos(b_2, a_1) + \epsilon}$$

Levy et al. (2015)

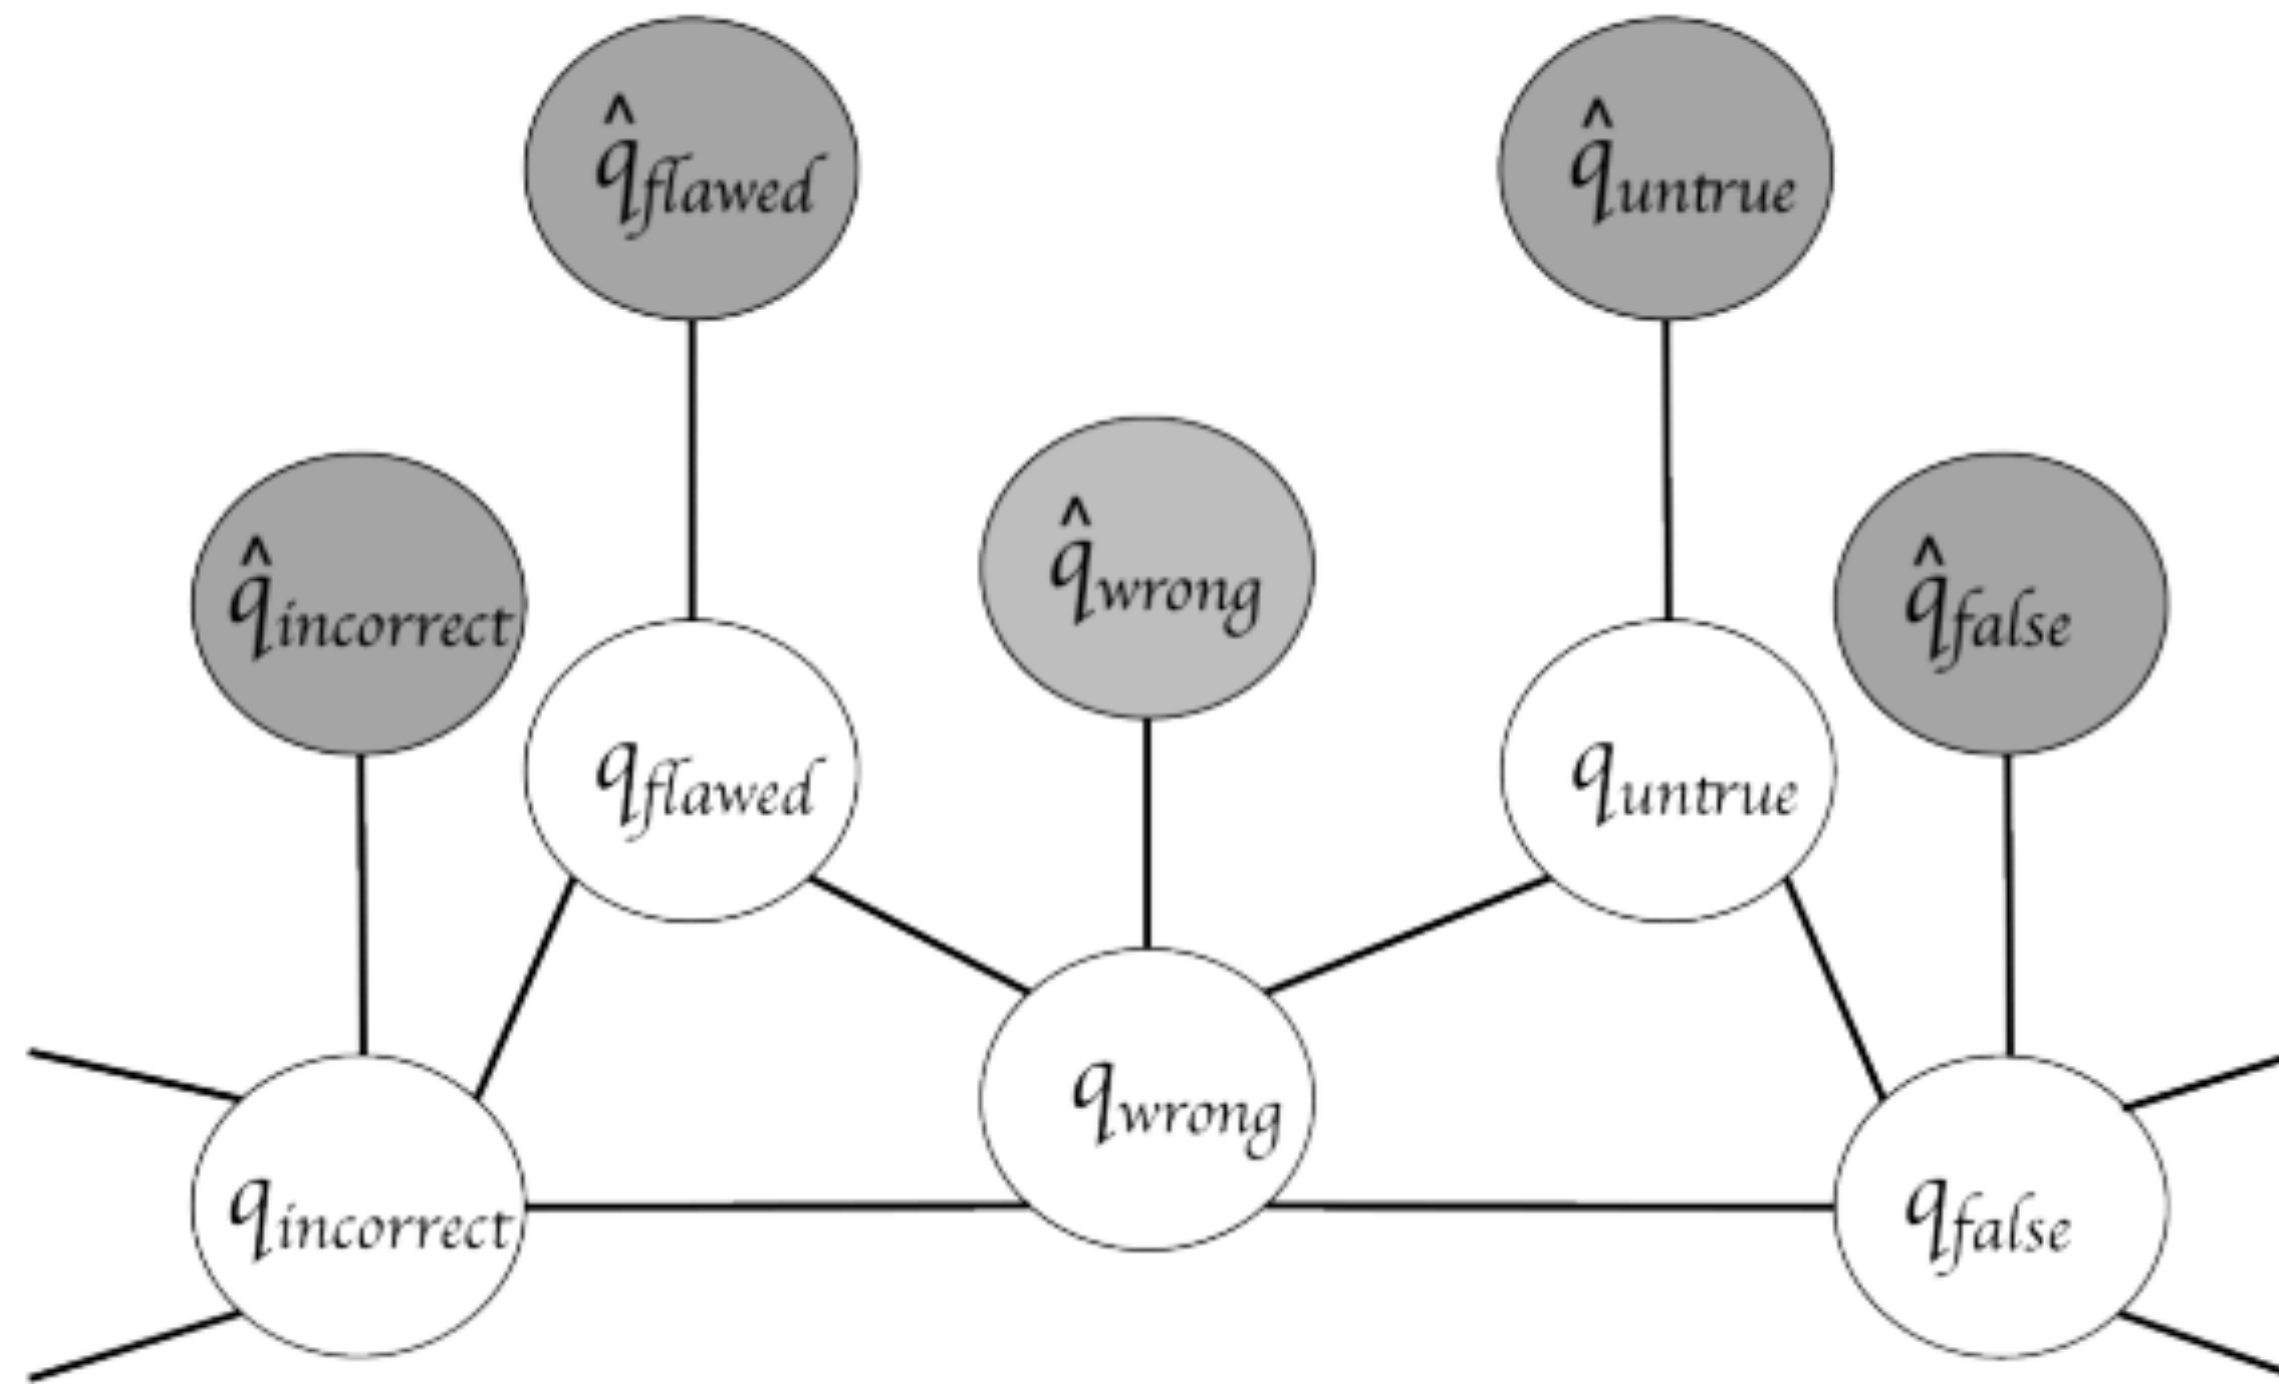
# Using Semantic Knowledge

---




# Using Semantic Knowledge

---



- Structure derived from a resource like WordNet

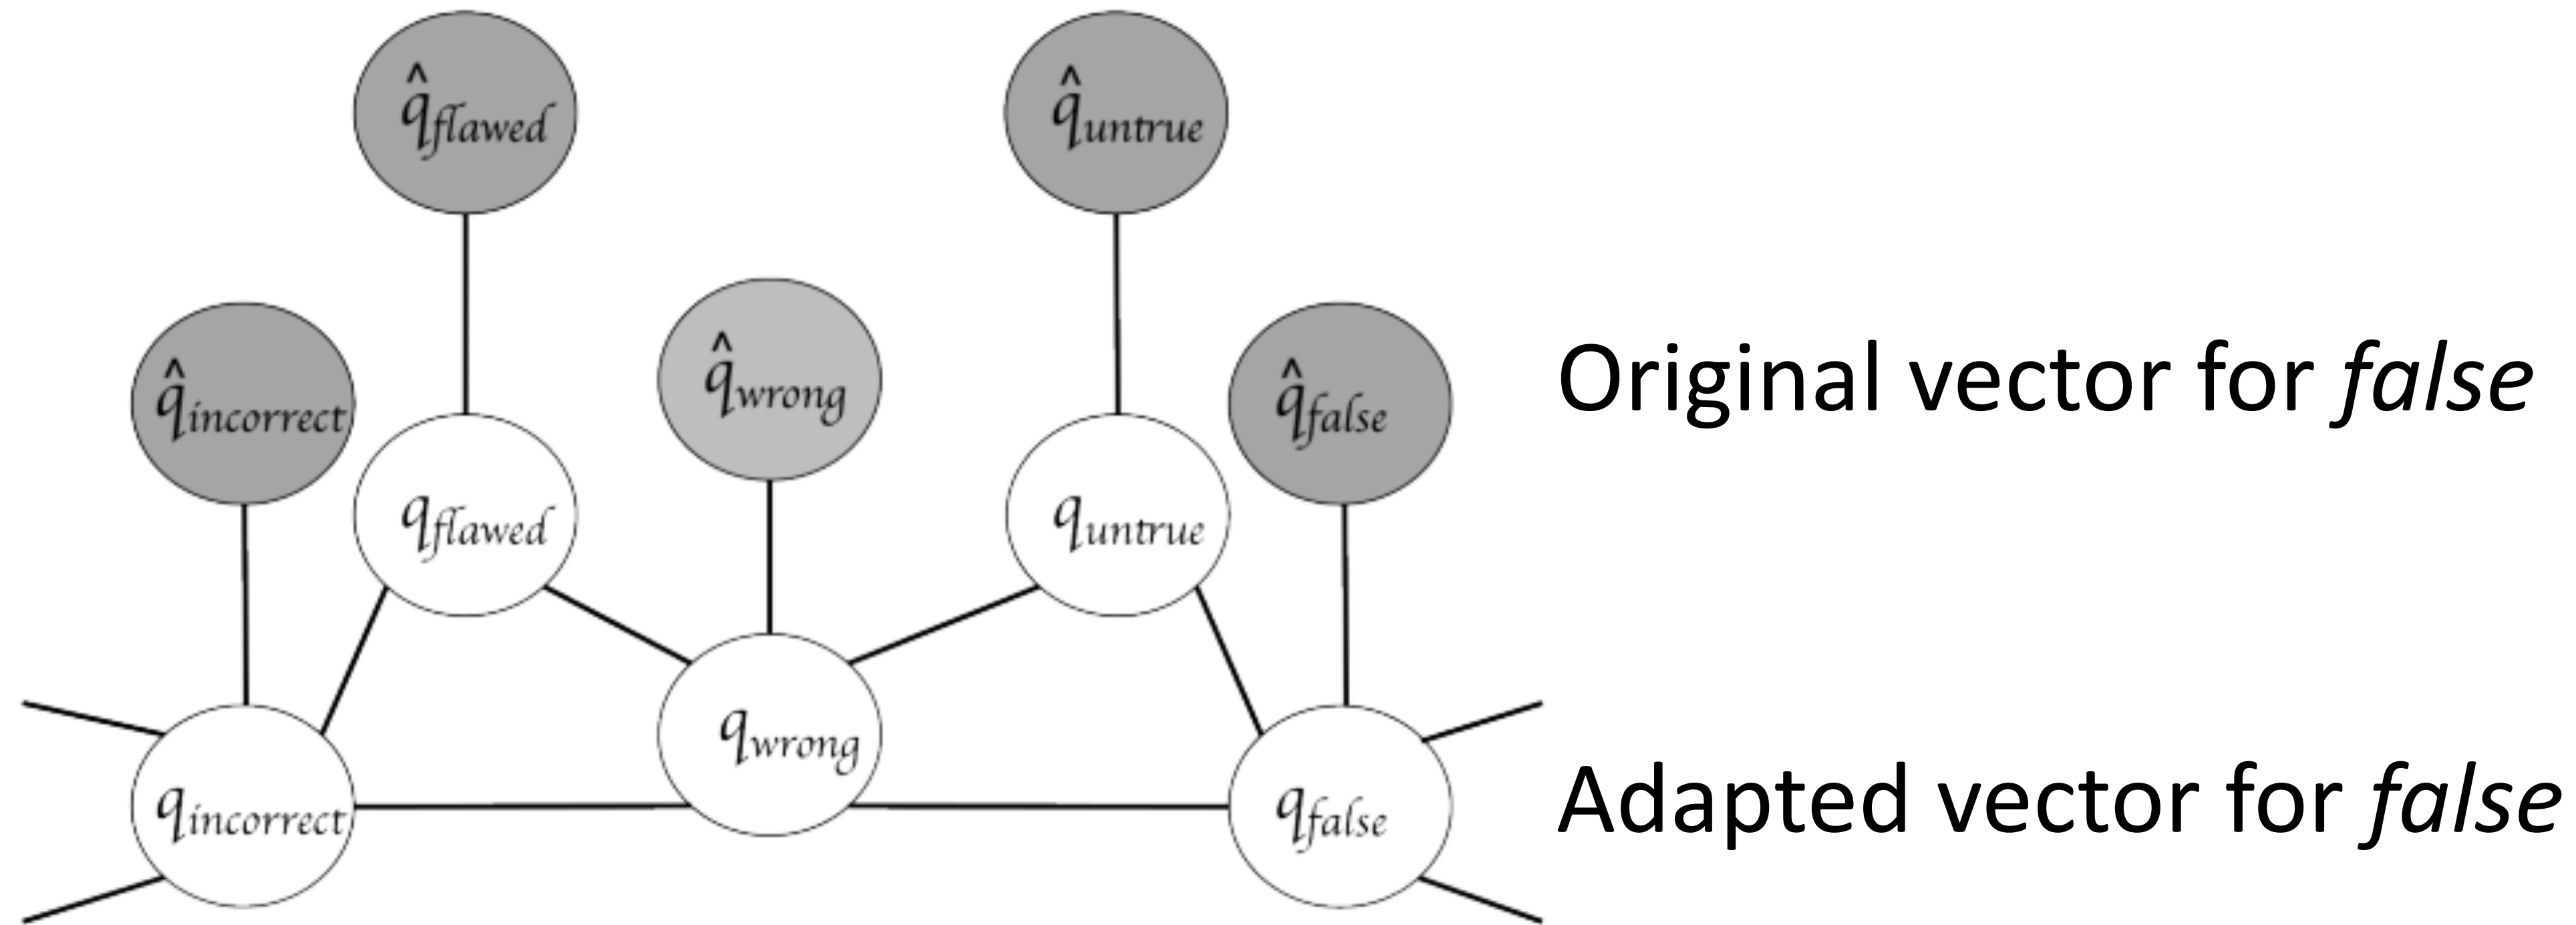
# Using Semantic Knowledge



- Structure derived from a resource like WordNet



# Using Semantic Knowledge



- ▶ Structure derived from a resource like WordNet
- ▶ Doesn't help most problems

# Using Word Embeddings

---

# Using Word Embeddings

---

- ▶ Approach 1: learn embeddings as parameters from your data
  - ▶ Often works pretty well



# Using Word Embeddings

---

- ▶ Approach 1: learn embeddings as parameters from your data
  - ▶ Often works pretty well
- ▶ Approach 2: initialize using GloVe/ELMo, keep fixed
  - ▶ Faster because no need to update these parameters

# Using Word Embeddings

---

- ▶ Approach 1: learn embeddings as parameters from your data
  - ▶ Often works pretty well
- ▶ Approach 2: initialize using GloVe/ELMo, keep fixed
  - ▶ Faster because no need to update these parameters
- ▶ Approach 3: initialize using GloVe, fine-tune
  - ▶ Works best for some tasks, but not used for ELMo

# Compositional Semantics

---

# Compositional Semantics

---

- ▶ What if we want embedding representations for whole sentences?

# Compositional Semantics

---

- ▶ What if we want embedding representations for whole sentences?
- ▶ Skip-*thought* vectors (Kiros et al., 2015), similar to skip-gram generalized to a sentence level (more later)

# Compositional Semantics

---

- ▶ What if we want embedding representations for whole sentences?
- ▶ Skip-*thought* vectors (Kiros et al., 2015), similar to skip-gram generalized to a sentence level (more later)
- ▶ Is there a way we can compose vectors to make sentence representations? Summing?

# Compositional Semantics

---

- ▶ What if we want embedding representations for whole sentences?
- ▶ Skip-*thought* vectors (Kiros et al., 2015), similar to skip-gram generalized to a sentence level (more later)
- ▶ Is there a way we can compose vectors to make sentence representations? Summing?
- ▶ Will return to this in a few weeks as we move on to syntax and semantics

# Takeaways

---



# Takeaways

---

- ▶ Lots to tune with neural networks

# Takeaways

---

- ▶ Lots to tune with neural networks
  - ▶ Training: optimizer, initializer, regularization (dropout), ...

# Takeaways

---

- ▶ Lots to tune with neural networks
  - ▶ Training: optimizer, initializer, regularization (dropout), ...
  - ▶ Hyperparameters: dimensionality of word embeddings, layers, ...

# Takeaways

---

- ▶ Lots to tune with neural networks
  - ▶ Training: optimizer, initializer, regularization (dropout), ...
  - ▶ Hyperparameters: dimensionality of word embeddings, layers, ...
- ▶ Word vectors: learning word  $\rightarrow$  context mappings has given way to matrix factorization approaches (constant in dataset size)

# Takeaways

---

- ▶ Lots to tune with neural networks
  - ▶ Training: optimizer, initializer, regularization (dropout), ...
  - ▶ Hyperparameters: dimensionality of word embeddings, layers, ...
- ▶ Word vectors: learning word  $\rightarrow$  context mappings has given way to matrix factorization approaches (constant in dataset size)
- ▶ Lots of pretrained embeddings work well in practice, they capture some desirable properties

# Takeaways

---

- ▶ Lots to tune with neural networks
  - ▶ Training: optimizer, initializer, regularization (dropout), ...
  - ▶ Hyperparameters: dimensionality of word embeddings, layers, ...
- ▶ Word vectors: learning word  $\rightarrow$  context mappings has given way to matrix factorization approaches (constant in dataset size)
- ▶ Lots of pretrained embeddings work well in practice, they capture some desirable properties
- ▶ Even better: context-sensitive word embeddings (ELMo)

# Takeaways

---

- ▶ Lots to tune with neural networks
  - ▶ Training: optimizer, initializer, regularization (dropout), ...
  - ▶ Hyperparameters: dimensionality of word embeddings, layers, ...
- ▶ Word vectors: learning word  $\rightarrow$  context mappings has given way to matrix factorization approaches (constant in dataset size)
- ▶ Lots of pretrained embeddings work well in practice, they capture some desirable properties
- ▶ Even better: context-sensitive word embeddings (ELMo)
- ▶ Next time: RNNs and CNNs